

# H2 Database Engine

Version 1.4.198 (2019-02-22)

## Table of Contents

Quickstart.....	21
Embedding H2 in an Application.....	21
The H2 Console Application.....	21
Step-by-Step.....	21
Installation.....	21
Start the Console.....	22
Login.....	22
Sample.....	23
Execute.....	24
Disconnect.....	25
End.....	25
Installation.....	26
Requirements.....	26
Database Engine.....	26
H2 Console.....	26
Supported Platforms.....	26
Installing the Software.....	26
Directory Structure.....	26
Tutorial.....	28
Starting and Using the H2 Console.....	28
Firewall.....	29
Testing Java.....	30
Error Message 'Port may be in use'.....	30
Using another Port.....	30
Connecting to the Server using a Browser.....	30
Multiple Concurrent Sessions.....	30
Login.....	31
Error Messages.....	31
Adding Database Drivers.....	31
Using the H2 Console.....	31
Inserting Table Names or Column Names.....	31
Disconnecting and Stopping the Application.....	32
Special H2 Console Syntax.....	32
Settings of the H2 Console.....	33
Connecting to a Database using JDBC.....	34
Creating New Databases.....	34
Using the Server.....	35
Starting the Server Tool from Command Line.....	35
Connecting to the TCP Server.....	35
Starting the TCP Server within an Application.....	35
Stopping a TCP Server from Another Process.....	36
Using Hibernate.....	36

Using TopLink and Glassfish.....	36
Using EclipseLink.....	37
Using Apache ActiveMQ.....	37
Using H2 within NetBeans.....	37
Using H2 with jOOQ.....	38
Using Databases in Web Applications.....	39
Embedded Mode.....	39
Server Mode.....	39
Using a Servlet Listener to Start and Stop a Database.....	39
Using the H2 Console Servlet.....	40
Android.....	41
CSV (Comma Separated Values) Support.....	42
Reading a CSV File from Within a Database.....	42
Importing Data from a CSV File.....	43
Writing a CSV File from Within a Database.....	43
Writing a CSV File from a Java Application.....	43
Reading a CSV File from a Java Application.....	43
Upgrade, Backup, and Restore.....	44
Database Upgrade.....	44
Backup using the Script Tool.....	44
Restore from a Script.....	45
Online Backup.....	45
Command Line Tools.....	45
The Shell Tool.....	46
Using OpenOffice Base.....	47
Java Web Start / JNLP.....	48
Using a Connection Pool.....	48
Fulltext Search.....	48
Using the Native Fulltext Search.....	49
Using the Apache Lucene Fulltext Search.....	49
User-Defined Variables.....	51
Date and Time.....	51
Using Spring.....	52
Using the TCP Server.....	52
OSGi.....	52
Java Management Extension (JMX).....	53
Features.....	54
Feature List.....	54
Main Features.....	54
Additional Features.....	55
SQL Support.....	55
Security Features.....	55
Other Features and Tools.....	56
H2 in Use.....	56
Connection Modes.....	56

Embedded Mode.....	57
Server Mode.....	57
Mixed Mode.....	58
Database URL Overview.....	58
Connecting to an Embedded (Local) Database.....	60
In-Memory Databases.....	60
Database Files Encryption.....	61
Creating a New Database with File Encryption.....	61
Connecting to an Encrypted Database.....	61
Encrypting or Decrypting a Database.....	61
Database File Locking.....	62
Opening a Database Only if it Already Exists.....	62
Closing a Database.....	63
Delayed Database Closing.....	63
Don't Close a Database when the VM Exits.....	63
Execute SQL on Connection.....	63
Ignore Unknown Settings.....	64
Changing Other Settings when Opening a Connection.....	64
Custom File Access Mode.....	64
Multiple Connections.....	65
Opening Multiple Databases at the Same Time.....	65
Multiple Connections to the Same Database: Client/Server.....	65
Multithreading Support.....	65
Locking, Lock-Timeout, Deadlocks.....	65
Avoiding Deadlocks.....	66
Database File Layout.....	67
Moving and Renaming Database Files.....	67
Backup.....	68
Logging and Recovery.....	68
Compatibility.....	68
Compatibility Modes.....	68
DB2 Compatibility Mode.....	68
Derby Compatibility Mode.....	69
HSQLDB Compatibility Mode.....	69
MS SQL Server Compatibility Mode.....	69
MySQL Compatibility Mode.....	70
Oracle Compatibility Mode.....	70
PostgreSQL Compatibility Mode.....	71
Ignite Compatibility Mode.....	71
Auto-Reconnect.....	72
Automatic Mixed Mode.....	72
Page Size.....	73
Using the Trace Options.....	73
Trace Options.....	73
Setting the Maximum Size of the Trace File.....	74

Java Code Generation.....	74
Using Other Logging APIs.....	75
Read Only Databases.....	75
Read Only Databases in Zip or Jar File.....	75
Opening a Corrupted Database.....	76
Computed Columns / Function Based Index.....	76
Multi-Dimensional Indexes.....	77
User-Defined Functions and Stored Procedures.....	77
Referencing a Compiled Method.....	78
Declaring Functions as Source Code.....	78
Method Overloading.....	79
Function Data Type Mapping.....	79
Functions That Require a Connection.....	80
Functions Throwing an Exception.....	80
Functions Returning a Result Set.....	80
Using SimpleResultSet.....	80
Using a Function as a Table.....	81
Pluggable or User-Defined Tables.....	81
Triggers.....	82
Compacting a Database.....	84
Cache Settings.....	84
External authentication (Experimental).....	85
Performance.....	87
Performance Comparison.....	87
Embedded.....	87
Client-Server.....	88
Benchmark Results and Comments.....	89
H2.....	89
HSQLDB.....	89
Derby.....	89
PostgreSQL.....	90
MySQL.....	90
Firebird.....	90
Why Oracle / MS SQL Server / DB2 are Not Listed.....	90
About this Benchmark.....	90
How to Run.....	90
Separate Process per Database.....	91
Number of Connections.....	91
Real-World Tests.....	91
Comparing Embedded with Server Databases.....	91
Test Platform.....	91
Multiple Runs.....	91
Memory Usage.....	91
Delayed Operations.....	92
Transaction Commit / Durability.....	92

Using Prepared Statements.....	92
Currently Not Tested: Startup Time.....	92
PolePosition Benchmark.....	92
Database Performance Tuning.....	94
Keep Connections Open or Use a Connection Pool.....	94
Use a Modern JVM.....	94
Virus Scanners.....	94
Using the Trace Options.....	94
Index Usage.....	95
Index Hints.....	95
How Data is Stored Internally.....	95
Optimizer.....	96
Expression Optimization.....	96
COUNT(*) Optimization.....	96
Updating Optimizer Statistics / Column Selectivity.....	96
In-Memory (Hash) Indexes.....	97
Use Prepared Statements.....	97
Prepared Statements and IN(... ).....	97
Optimization Examples.....	98
Cache Size and Type.....	98
Data Types.....	98
Sorted Insert Optimization.....	98
Using the Built-In Profiler.....	98
Application Profiling.....	99
Analyze First.....	99
Database Profiling.....	100
Statement Execution Plans.....	100
Displaying the Scan Count.....	101
Special Optimizations.....	102
How Data is Stored and How Indexes Work.....	102
Indexes.....	103
Using Multiple Indexes.....	105
Fast Database Import.....	105
Advanced.....	107
Result Sets.....	108
Statements that Return a Result Set.....	108
Limiting the Number of Rows.....	108
Large Result Sets and External Sorting.....	108
Large Objects.....	108
Storing and Reading Large Objects.....	108
When to use CLOB/BLOB.....	108
Large Object Compression.....	109
Linked Tables.....	109
Updatable Views.....	110
Transaction Isolation.....	110

Multi-Version Concurrency Control (MVCC).....	111
Table Level Locking (PageStore engine).....	112
Lock Timeout.....	112
Clustering / High Availability.....	112
Using the CreateCluster Tool.....	113
Detect Which Cluster Instances are Running.....	114
Clustering Algorithm and Limitations.....	114
Two Phase Commit.....	115
Compatibility.....	115
Transaction Commit when Autocommit is On.....	115
Keywords / Reserved Words.....	116
Standards Compliance.....	116
Supported Character Sets, Character Encoding, and Unicode.....	116
Run as Windows Service.....	116
Install the Service.....	117
Start the Service.....	117
Connect to the H2 Console.....	117
Stop the Service.....	117
Uninstall the Service.....	117
Additional JDBC drivers.....	117
ODBC Driver.....	118
ODBC Installation.....	118
Starting the Server.....	118
ODBC Configuration.....	119
PG Protocol Support Limitations.....	119
Security Considerations.....	120
Using Microsoft Access.....	120
Using H2 in Microsoft .NET.....	120
Using the ADO.NET API on .NET.....	120
Using the JDBC API on .NET.....	120
ACID.....	121
Atomicity.....	121
Consistency.....	121
Isolation.....	122
Durability.....	122
Durability Problems.....	122
Ways to (Not) Achieve Durability.....	122
Running the Durability Test.....	123
Using the Recover Tool.....	124
File Locking Protocols.....	124
File Locking Method 'File'.....	125
File Locking Method 'Socket'.....	126
File Locking Method 'FS'.....	126
Using Passwords.....	127
Using Secure Passwords.....	127

Passwords: Using Char Arrays instead of Strings.....	127
Passing the User Name and/or Password in the URL.....	128
Password Hash.....	128
Protection against SQL Injection.....	129
What is SQL Injection.....	129
Disabling Literals.....	129
Using Constants.....	130
Using the ZERO() Function.....	130
Protection against Remote Access.....	130
Restricting Class Loading and Usage.....	131
Security Protocols.....	131
User Password Encryption.....	131
File Encryption.....	132
Wrong Password / User Name Delay.....	133
HTTPS Connections.....	133
TLS Connections.....	134
Universally Unique Identifiers (UUID).....	134
Spatial Features.....	135
Recursive Queries.....	136
Settings Read from System Properties.....	137
Setting the Server Bind Address.....	137
Pluggable File System.....	137
Split File System.....	138
Database Upgrade.....	139
Java Objects Serialization.....	139
Custom Data Types Handler API.....	140
Limits and Limitations.....	140
Glossary and Links.....	141
SQL Grammar.....	143
Index.....	143
Commands (Data Manipulation).....	143
Commands (Data Definition).....	143
Commands (Other).....	144
Datetime fields.....	146
Other Grammar.....	146
Commands (Data Manipulation).....	148
SELECT.....	148
INSERT.....	150
UPDATE.....	150
DELETE.....	151
BACKUP.....	151
CALL.....	151
EXPLAIN.....	151
MERGE.....	152
MERGE USING.....	152



RUNSCRIPT.....	153
SCRIPT.....	153
SHOW.....	154
Explicit table.....	154
Table value.....	155
WITH.....	155
Commands (Data Definition).....	156
ALTER INDEX RENAME.....	156
ALTER SCHEMA RENAME.....	156
ALTER SEQUENCE.....	157
ALTER TABLE ADD.....	157
ALTER TABLE ADD CONSTRAINT.....	157
ALTER TABLE RENAME CONSTRAINT.....	158
ALTER TABLE ALTER COLUMN.....	158
ALTER TABLE DROP COLUMN.....	159
ALTER TABLE DROP CONSTRAINT.....	160
ALTER TABLE SET.....	160
ALTER TABLE RENAME.....	160
ALTER USER ADMIN.....	160
ALTER USER RENAME.....	161
ALTER USER SET PASSWORD.....	161
ALTER VIEW RECOMPILE.....	161
ALTER VIEW RENAME.....	162
ANALYZE.....	162
COMMENT.....	162
CREATE AGGREGATE.....	163
CREATE ALIAS.....	163
CREATE CONSTANT.....	164
CREATE DOMAIN.....	165
CREATE INDEX.....	165
CREATE LINKED TABLE.....	165
CREATE ROLE.....	166
CREATE SCHEMA.....	167
CREATE SEQUENCE.....	167
CREATE TABLE.....	167
CREATE TRIGGER.....	168
CREATE USER.....	170
CREATE VIEW.....	170
DROP AGGREGATE.....	171
DROP ALIAS.....	171
DROP ALL OBJECTS.....	171
DROP CONSTANT.....	172
DROP DOMAIN.....	172
DROP INDEX.....	172
DROP ROLE.....	173

DROP SCHEMA.....	173
DROP SEQUENCE.....	173
DROP TABLE.....	173
DROP TRIGGER.....	174
DROP USER.....	174
DROP VIEW.....	174
TRUNCATE TABLE.....	174
Commands (Other).....	175
CHECKPOINT.....	175
CHECKPOINT SYNC.....	175
COMMIT.....	175
COMMIT TRANSACTION.....	176
GRANT RIGHT.....	176
GRANT ALTER ANY SCHEMA.....	176
GRANT ROLE.....	177
HELP.....	177
PREPARE COMMIT.....	177
REVOKE RIGHT.....	177
REVOKE ROLE.....	178
ROLLBACK.....	178
ROLLBACK TRANSACTION.....	178
SAVEPOINT.....	179
SET @.....	179
SET ALLOW_LITERALS.....	179
SET AUTOCOMMIT.....	180
SET CACHE_SIZE.....	180
SET CLUSTER.....	180
SET BINARY_COLLATION.....	181
SET UUID_COLLATION.....	181
SET BUILTIN_ALIAS_OVERRIDE.....	182
SET COLLATION.....	182
SET COMPRESS_LOB.....	183
SET DATABASE_EVENT_LISTENER.....	183
SET DB_CLOSE_DELAY.....	183
SET DEFAULT_LOCK_TIMEOUT.....	184
SET DEFAULT_TABLE_TYPE.....	184
SET EXCLUSIVE.....	185
SET IGNORECASE.....	185
SET JAVA_OBJECT_SERIALIZER.....	185
SET LAZY_QUERY_EXECUTION.....	186
SET LOG.....	186
SET LOCK_MODE.....	187
SET LOCK_TIMEOUT.....	188
SET MAX_LENGTH_INPLACE_LOB.....	188
SET MAX_LOG_SIZE.....	188

SET MAX_MEMORY_ROWS.....	189
SET MAX_MEMORY_UNDO.....	189
SET MAX_OPERATION_MEMORY.....	189
SET MODE.....	190
SET MULTI_THREADED.....	190
SET OPTIMIZE_REUSE_RESULTS.....	191
SET PASSWORD.....	191
SET QUERY_STATISTICS.....	191
SET QUERY_STATISTICS_MAX_ENTRIES.....	192
SET QUERY_TIMEOUT.....	192
SET REFERENTIAL_INTEGRITY.....	192
SET RETENTION_TIME.....	193
SET SALT_HASH.....	193
SET SCHEMA.....	194
SET SCHEMA_SEARCH_PATH.....	194
SET THROTTLE.....	194
SET TRACE_LEVEL.....	195
SET TRACE_MAX_FILE_SIZE.....	195
SET UNDO_LOG.....	195
SET WRITE_DELAY.....	196
SHUTDOWN.....	196
Datetime fields.....	197
Datetime field.....	197
Year field.....	197
Month field.....	197
Day of month field.....	198
Hour field.....	198
Minute field.....	198
Second field.....	198
Millisecond field.....	198
Microsecond field.....	199
Nanosecond field.....	199
Timezone hour field.....	199
Timezone minute field.....	199
Day of week field.....	200
ISO week year field.....	200
ISO day of week field.....	200
Week of year field.....	200
ISO week of year field.....	201
Quarter field.....	201
Day of year field.....	201
Epoch field.....	201
Other Grammar.....	202
Alias.....	202
And Condition.....	202

Array.....	202
Boolean.....	202
Bytes.....	203
Case.....	203
Case When.....	203
Cipher.....	203
Column Definition.....	204
Comments.....	204
Compare.....	205
Condition.....	205
Condition Right Hand Side.....	205
Constraint.....	206
Constraint Name Definition.....	207
Csv Options.....	207
Data Type.....	208
Date.....	208
Decimal.....	208
Digit.....	209
Dollar Quoted String.....	209
Expression.....	209
Factor.....	210
Hex.....	210
Hex Number.....	210
Index Column.....	210
Insert columns and source.....	211
Insert values.....	211
Int.....	211
Long.....	211
Merge when clause.....	212
Merge when matched clause.....	212
Merge when not matched clause.....	212
Name.....	212
Null.....	213
Number.....	213
Numeric.....	213
Operand.....	213
Order.....	214
Quoted Name.....	214
Referential Constraint.....	214
Referential Action.....	215
Script Compression Encryption.....	215
Row value expression.....	215
Select Expression.....	216
Sequence options.....	216
Set clause list.....	216

String.....	216
Summand.....	217
Table Expression.....	217
Within group specification.....	217
Wildcard expression.....	218
Window name or specification.....	218
Window specification.....	218
Window frame.....	219
Window frame preceding.....	220
Window frame bound.....	220
Term.....	220
Column.....	221
Time.....	221
Timestamp.....	221
Timestamp with time zone.....	222
Date and time.....	222
INTERVAL YEAR.....	222
INTERVAL MONTH.....	222
INTERVAL DAY.....	223
INTERVAL HOUR.....	223
INTERVAL MINUTE.....	223
INTERVAL SECOND.....	223
INTERVAL YEAR TO MONTH.....	223
INTERVAL DAY TO HOUR.....	224
INTERVAL DAY TO MINUTE.....	224
INTERVAL DAY TO SECOND.....	224
INTERVAL HOUR TO MINUTE.....	224
INTERVAL HOUR TO SECOND.....	225
INTERVAL MINUTE TO SECOND.....	225
Interval.....	225
Value.....	225
Functions.....	227
Index.....	227
Aggregate Functions.....	227
Numeric Functions.....	227
String Functions.....	228
Time and Date Functions.....	229
System Functions.....	230
Window Functions.....	231
Aggregate Functions.....	232
AVG.....	232
BIT_AND.....	232
BIT_OR.....	232
EVERY.....	232
ANY.....	233

COUNT.....	233
LISTAGG.....	233
ARRAY_AGG.....	234
MAX.....	234
MIN.....	235
SUM.....	235
SELECTIVITY.....	235
STDDEV_POP.....	235
STDDEV_SAMP.....	236
VAR_POP.....	236
VAR_SAMP.....	236
RANK aggregate.....	237
DENSE_RANK aggregate.....	237
PERCENT_RANK aggregate.....	237
CUME_DIST aggregate.....	238
PERCENTILE_CONT.....	238
PERCENTILE_DISC.....	238
MEDIAN.....	239
MODE.....	239
ENVELOPE.....	240
Numeric Functions.....	240
ABS.....	240
ACOS.....	240
ASIN.....	241
ATAN.....	241
COS.....	241
COSH.....	241
COT.....	242
SIN.....	242
SINH.....	242
TAN.....	242
TANH.....	243
ATAN2.....	243
BITAND.....	243
BITGET.....	243
BITOR.....	244
BITXOR.....	244
MOD.....	244
CEILING.....	244
DEGREES.....	244
EXP.....	245
FLOOR.....	245
LN.....	245
LOG10.....	245
ORA_HASH.....	246

RADIANS.....	246
SQRT.....	246
PI.....	246
POWER.....	247
RAND.....	247
RANDOM_UUID.....	247
ROUND.....	248
ROUNDMAGIC.....	248
SECURE_RAND.....	248
SIGN.....	248
ENCRYPT.....	249
DECRYPT.....	249
HASH.....	249
TRUNCATE.....	249
COMPRESS.....	250
EXPAND.....	250
ZERO.....	250
String Functions.....	251
ASCII.....	251
BIT_LENGTH.....	251
LENGTH.....	251
OCTET_LENGTH.....	251
CHAR.....	252
CONCAT.....	252
CONCAT_WS.....	252
DIFFERENCE.....	252
HEXTORAW.....	253
RAWTOHEX.....	253
INSTR.....	253
INSERT Function.....	254
LOWER.....	254
UPPER.....	254
LEFT.....	254
RIGHT.....	255
LOCATE.....	255
POSITION.....	255
LPAD.....	255
RPAD.....	256
LTRIM.....	256
RTRIM.....	256
TRIM.....	256
REGEXP_REPLACE.....	257
REGEXP_LIKE.....	257
REPEAT.....	258
REPLACE.....	258

SOUNDEX.....	258
SPACE.....	258
STRINGDECODE.....	259
STRINGENCODE.....	259
STRINGTOUTF8.....	259
SUBSTRING.....	259
UTF8TOSTRING.....	260
XMLATTR.....	260
XMLNODE.....	260
XMLCOMMENT.....	261
XMLCDATA.....	261
XMLSTARTDOC.....	261
XMLTEXT.....	261
TO_CHAR.....	262
TRANSLATE.....	262
Time and Date Functions.....	262
CURRENT_DATE.....	262
CURRENT_TIME.....	262
CURRENT_TIMESTAMP.....	263
LOCALTIME.....	263
LOCALTIMESTAMP.....	264
DATEADD.....	264
DATEDIFF.....	264
DAYNAME.....	265
DAY_OF_MONTH.....	265
DAY_OF_WEEK.....	265
ISO_DAY_OF_WEEK.....	265
DAY_OF_YEAR.....	266
EXTRACT.....	266
FORMATDATETIME.....	266
HOUR.....	266
MINUTE.....	267
MONTH.....	267
MONTHNAME.....	267
PARSEDATETIME.....	267
QUARTER.....	268
SECOND.....	268
WEEK.....	268
ISO_WEEK.....	268
YEAR.....	269
ISO_YEAR.....	269
System Functions.....	269
ARRAY_GET.....	269
ARRAY_LENGTH.....	269
ARRAY_CONTAINS.....	270



ARRAY_CAT.....	270
ARRAY_APPEND.....	270
ARRAY_SLICE.....	270
AUTOCOMMIT.....	271
CANCEL_SESSION.....	271
CASEWHEN Function.....	271
CAST.....	271
COALESCE.....	272
CONVERT.....	272
CURRVAL.....	272
CSVREAD.....	273
CSVWRITE.....	274
DATABASE.....	274
DATABASE_PATH.....	275
DECODE.....	275
DISK_SPACE_USED.....	275
SIGNAL.....	275
ESTIMATED_ENVELOPE.....	276
FILE_READ.....	276
FILE_WRITE.....	276
GREATEST.....	277
IDENTITY.....	277
IFNULL.....	277
LEAST.....	277
LOCK_MODE.....	278
LOCK_TIMEOUT.....	278
LINK_SCHEMA.....	278
MEMORY_FREE.....	278
MEMORY_USED.....	279
NEXTVAL.....	279
NULLIF.....	279
NVL2.....	279
READONLY.....	280
ROWNUM.....	280
SCHEMA.....	280
SCOPE_IDENTITY.....	281
SESSION_ID.....	281
SET.....	281
TABLE.....	281
TRANSACTION_ID.....	282
TRUNCATE_VALUE.....	282
UNNEST.....	282
USER.....	283
H2VERSION.....	283
Window Functions.....	283

ROW_NUMBER.....	283
RANK.....	284
DENSE_RANK.....	284
PERCENT_RANK.....	284
CUME_DIST.....	285
NTILE.....	285
LEAD.....	286
LAG.....	286
FIRST_VALUE.....	287
LAST_VALUE.....	287
NTH_VALUE.....	288
RATIO_TO_REPORT.....	288
Data Types.....	289
Index.....	289
INT.....	289
BOOLEAN.....	290
TINYINT.....	290
SMALLINT.....	290
BIGINT.....	290
IDENTITY.....	291
DECIMAL.....	291
DOUBLE.....	291
REAL.....	292
TIME.....	292
DATE.....	292
TIMESTAMP.....	293
TIMESTAMP WITH TIME ZONE.....	293
BINARY.....	294
OTHER.....	294
VARCHAR.....	294
VARCHAR_IGNORECASE.....	295
CHAR.....	295
BLOB.....	296
CLOB.....	296
UUID.....	297
ARRAY.....	297
ENUM.....	297
GEOMETRY.....	298
INTERVAL.....	299
Interval Data Types.....	299
INTERVAL YEAR.....	299
INTERVAL MONTH.....	300
INTERVAL DAY.....	300
INTERVAL HOUR.....	300
INTERVAL MINUTE.....	300

INTERVAL SECOND.....	301
INTERVAL YEAR TO MONTH.....	301
INTERVAL DAY TO HOUR.....	301
INTERVAL DAY TO MINUTE.....	302
INTERVAL DAY TO SECOND.....	302
INTERVAL HOUR TO MINUTE.....	302
INTERVAL HOUR TO SECOND.....	303
INTERVAL MINUTE TO SECOND.....	303
System Tables.....	304
Information Schema.....	304
CATALOGS.....	304
COLLATIONS.....	304
COLUMNS.....	304
COLUMN_PRIVILEGES.....	304
CONSTANTS.....	304
CONSTRAINTS.....	305
CROSS_REFERENCES.....	305
DOMAINS.....	305
FUNCTION_ALIASES.....	305
FUNCTION_COLUMNS.....	305
HELP.....	305
INDEXES.....	305
IN_DOUBT.....	306
KEY_COLUMN_USAGE.....	306
LOCKS.....	306
QUERY_STATISTICS.....	306
REFERENTIAL_CONSTRAINTS.....	306
RIGHTS.....	306
ROLES.....	306
SCHEMATA.....	306
SEQUENCES.....	307
SESSIONS.....	307
SESSION_STATE.....	307
SETTINGS.....	307
SYNONYMS.....	307
TABLES.....	307
TABLE_CONSTRAINTS.....	307
TABLE_PRIVILEGES.....	307
TABLE_TYPES.....	308
TRIGGERS.....	308
TYPE_INFO.....	308
USERS.....	308
VIEWS.....	308
Range Table.....	308
Build.....	310

Portability.....	310
Environment.....	310
Building the Software.....	310
Build Targets.....	311
Using Apache Lucene.....	311
Using Maven 2.....	311
Using a Central Repository.....	311
Maven Plugin to Start and Stop the TCP Server.....	312
Using Snapshot Version.....	312
Using Eclipse.....	312
Translating.....	313
Submitting Source Code Changes.....	313
Reporting Problems or Requests.....	314
Automated Build.....	315
Generating Railroad Diagrams.....	315
History and Roadmap.....	316
Change Log.....	316
Roadmap.....	316
History of this Database Engine.....	316
Why Java.....	316
Supporters.....	317
Frequently Asked Questions.....	320
I Have a Problem or Feature Request.....	320
Are there Known Bugs? When is the Next Release?.....	320
Is this Database Engine Open Source?.....	321
Is Commercial Support Available?.....	321
How to Create a New Database?.....	321
How to Connect to a Database?.....	321
Where are the Database Files Stored?.....	321
What is the Size Limit (Maximum Size) of a Database?.....	322
Is it Reliable?.....	322
Why is Opening my Database Slow?.....	323
My Query is Slow.....	323
H2 is Very Slow.....	323
Column Names are Incorrect?.....	324
Float is Double?.....	324
How to Translate this Project?.....	324
How to Contribute to this Project?.....	324

# Quickstart

Embedding H2 in an Application  
The H2 Console Application


## Embedding H2 in an Application

This database can be used in embedded mode, or in server mode. To use it in embedded mode, you need to:

- Add the h2\*.jar to the classpath (H2 does not have any dependencies)
- Use the JDBC driver class: org.h2.Driver
- The database URL jdbc:h2:~/test opens the database test in your user home directory
- A new database is automatically created

## The H2 Console Application

The Console lets you access a SQL database using a browser interface.



If you don't have Windows XP, or if something does not work as expected, please see the detailed description in the [Tutorial](#).

## Step-by-Step

### Installation


Install the software using the Windows Installer (if you did not yet do that).

## Start the Console

Click [Start], [All Programs], [H2], and [H2 Console (Command Line)]:




A new console window appears:



Also, a new browser page should open with the URL <http://localhost:8082>. You may get a security warning from the firewall. If you don't want other computers in the network to access the database on your machine, you can let the firewall block these connections. Only local connections are required at this time.

## Login

Select [Generic H2] and click [Connect]:



You are now logged in.

## Sample

Click on the [Sample SQL Script]:

The screenshot shows a database management tool interface. At the top, there is a toolbar with icons for connection, execution, and help. Below the toolbar, the connection details are shown as 'jdbc:h2:test'. A tree view on the left shows the database structure with 'INFORMATION\_SCHEMA' and 'Users' folders. A 'Run' button and a 'Clear' button are next to the 'SQL statement:' label. Below this is a large text area for the SQL statement. A section titled 'Important Commands' contains a table with icons and their actions. Another section titled 'Sample SQL Script' contains a table with operations and their corresponding SQL statements. The 'Sample SQL Script' table is highlighted with a red border.

Icon	Action
?	Displays this Help Page
≡	Shows the Command History
▶	Executes the current SQL statement
🔌	Disconnects from the database

Operations	SQL statements
Delete the table if it exists	DROP TABLE IF EXISTS TEST;
Create a new table with ID and NAME columns	CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255));
Add a new row	INSERT INTO TEST VALUES(1, 'Hello');
Add another row	INSERT INTO TEST VALUES(2, 'World');
Query the table	SELECT * FROM TEST ORDER BY ID;
Change data in a row	UPDATE TEST SET NAME='Hi' WHERE ID=1;
Remove a row	DELETE FROM TEST WHERE ID=2;

The SQL commands appear in the command area.

## Execute

Click [Run]

The screenshot shows a database management tool interface. At the top, there is a toolbar with icons for connection, execution, and help. Below the toolbar, on the left, is a tree view showing the database structure: 'jdbc:h2:test' (database icon), 'INFORMATION\_SCHEMA' (folder icon), and 'Users' (users icon). To the right of the tree view, there are two buttons: 'Run' (highlighted with a red box) and 'Clear'. Next to these buttons is a text area labeled 'SQL statement:' containing the following SQL code:

```
DROP TABLE IF EXISTS TEST;  
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255));  
INSERT INTO TEST VALUES(1, 'Hello');  
INSERT INTO TEST VALUES(2, 'World');  
SELECT * FROM TEST ORDER BY ID;  
UPDATE TEST SET NAME='Hi' WHERE ID=1;  
DELETE FROM TEST WHERE ID=2;
```

Below the SQL statement area, there is a section titled 'Important Commands' with a table listing various actions and their corresponding icons:

Icon	Action
	Displays this Help Page
	Shows the Command History
	Executes the current SQL statement
	Disconnects from the database


Below the 'Important Commands' section, there is a section titled 'Sample SQL Script' with a table listing various operations and their corresponding SQL statements:

Operations	SQL statements
Delete the table if it exists	DROP TABLE IF EXISTS TEST;
Create a new table with ID and NAME columns	CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255));
Add a new row	INSERT INTO TEST VALUES(1, 'Hello');
Add another row	INSERT INTO TEST VALUES(2, 'World');
Query the table	SELECT * FROM TEST ORDER BY ID;
Change data in a row	UPDATE TEST SET NAME='Hi' WHERE ID=1;
Remove a row	DELETE FROM TEST WHERE ID=2;

On the left side, a new entry TEST is added below the database icon. The



operations and results of the statements are shown below the script.



The screenshot shows a SQL IDE interface with a toolbar at the top containing icons for connection, execution, and settings. Below the toolbar, on the left, is a tree view showing the database structure: 'jdbc:h2:test' with a sub-entry 'TEST' highlighted. The main area is divided into two panes. The top pane, titled 'SQL statement:', contains the following SQL script:

```
DROP TABLE IF EXISTS TEST;  
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255));  
INSERT INTO TEST VALUES(1, 'Hello');  
INSERT INTO TEST VALUES(2, 'World');  
SELECT * FROM TEST ORDER BY ID;  
UPDATE TEST SET NAME='Hi' WHERE ID=1;  
DELETE FROM TEST WHERE ID=2;
```

The bottom pane shows the execution results for each statement, including the update count and execution time:

```
DROP TABLE IF EXISTS TEST  
Update Count: 0  
(0 ms)  
  
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255))  
Update Count: 0  
(0 ms)  
  
INSERT INTO TEST VALUES(1, 'Hello')  
Update Count: 1  
(0 ms)  
  
INSERT INTO TEST VALUES(2, 'World')  
Update Count: 1  
(0 ms)  
  
SELECT * FROM TEST ORDER BY ID  


| ID | NAME  |
|----|-------|
| 1  | Hello |
| 2  | World |

  
(2 rows, 0 ms)  
  
UPDATE TEST SET NAME='Hi' WHERE ID=1  
Update Count: 1
```

## Disconnect

Click on [Disconnect]:



to close the connection.

## End

Close the console window. For more information, see the [Tutorial](#).

# Installation

[Requirements](#)

[Supported Platforms](#)

[Installing the Software](#)

[Directory Structure](#)

## Requirements

To run this database, the following software stack is known to work. Other software most likely also works, but is not tested as much.

### Database Engine

- Windows XP or Vista, Mac OS X, or Linux
- Oracle Java 7 or newer
- Recommended Windows file system: NTFS (FAT32 only supports files up to 4 GB)

### H2 Console

- Mozilla Firefox

## Supported Platforms

As this database is written in Java, it can run on many different platforms. It is tested with Java 7. Currently, the database is developed and tested on Windows 8 and Mac OS X using Java 7, but it also works in many other operating systems and using other Java runtime environments. All major operating systems (Windows XP, Windows Vista, Windows 7, Mac OS, Ubuntu,...) are supported.

## Installing the Software

To install the software, run the installer or unzip it to a directory of your choice.

## Directory Structure

After installing, you should get the following directory structure:

Directory	Contents
-----------	----------


bin	JAR and batch files
docs	Documentation
docs/html	HTML pages
docs/javadoc	Javadoc files
ext	External dependencies (downloaded when building)
service	Tools to run the database as a Windows Service
src	Source files
src/docsrc	Documentation sources
src/installer	Installer, shell, and release build script
src/main	Database engine source code
src/test	Test source code
src/tools	Tools and database adapters source code

# Tutorial

- Starting and Using the H2 Console
- Special H2 Console Syntax
- Settings of the H2 Console
- Connecting to a Database using JDBC
- Creating New Databases
- Using the Server
- Using Hibernate
- Using TopLink and Glassfish
- Using EclipseLink
- Using Apache ActiveMQ
- Using H2 within NetBeans
- Using H2 with jOOQ
- Using Databases in Web Applications
- Android
- CSV (Comma Separated Values) Support
- Upgrade, Backup, and Restore
- Command Line Tools
- The Shell Tool
- Using OpenOffice Base
- Java Web Start / JNLP
- Using a Connection Pool
- Fulltext Search
- User-Defined Variables
- Date and Time
- Using Spring
- OSGi
- Java Management Extension (JMX)


## Starting and Using the H2 Console

The H2 Console application lets you access a database using a browser. This can be a H2 database, or another database that supports the JDBC API.



This is a client/server application, so both a server and a client (a browser) are required to run it.

Depending on your platform and environment, there are multiple ways to start the H2 Console:

OS	Start
Windows	Click [Start], [All Programs], [H2], and [H2 Console (Command Line)] An icon will be added to the system tray:  If you don't get the window and the system tray icon, then maybe Java is not installed correctly (in this case, try another way to start the application). A browser window should open and point to the login page at <a href="http://localhost:8082">http://localhost:8082</a> .
Windows	Open a file browser, navigate to h2/bin, and double click on h2.bat. A console window appears. If there is a problem, you will see an error message in this window. A browser window will open and point to the login page (URL: <a href="http://localhost:8082">http://localhost:8082</a> ).
Any	Double click on the h2*.jar file. This only works if the .jar suffix is associated with Java.
Any	Open a console window, navigate to the directory h2/bin, and type: <pre>java -jar h2*.jar</pre>

## Firewall

If you start the server, you may get a security warning from the firewall (if you have installed one). If you don't want other computers in the network to access the application on your machine, you can let the firewall block those connections. The connection from the local machine will still work. Only if you want other computers to access the database on this computer, you need allow remote connections in the firewall.

It has been reported that when using Kaspersky 7.0 with firewall, the H2 Console is very slow when connecting over the IP address. A workaround is to connect using 'localhost'.

A small firewall is already built into the server: other computers may not connect to the server by default. To change this, go to 'Preferences' and select 'Allow connections from other computers'.

## Testing Java

To find out which version of Java is installed, open a command prompt and type:

```
java -version
```

If you get an error message, you may need to add the Java binary directory to the path environment variable.

## Error Message 'Port may be in use'

You can only start one instance of the H2 Console, otherwise you will get the following error message: "The Web server could not be started. Possible cause: another server is already running...". It is possible to start multiple console applications on the same computer (using different ports), but this is usually not required as the console supports multiple concurrent connections.

## Using another Port

If the default port of the H2 Console is already in use by another application, then a different port needs to be configured. The settings are stored in a properties file. For details, see [Settings of the H2 Console](#). The relevant entry is webPort.

If no port is specified for the TCP and PG servers, each service will try to listen on its default port. If the default port is already in use, a random port is used.

## Connecting to the Server using a Browser

If the server started successfully, you can connect to it using a web browser. Javascript needs to be enabled. If you started the server on the same computer as the browser, open the URL `http://localhost:8082`. If you want to connect to the application from another computer, you need to provide the IP address of the server, for example: `http://192.168.0.2:8082`. If you enabled TLS on the server side, the URL needs to start with `https://`.

## Multiple Concurrent Sessions

Multiple concurrent browser sessions are supported. As that the database objects reside on the server, the amount of concurrent work is limited by the memory available to the server application.

## Login

At the login page, you need to provide connection information to connect to a database. Set the JDBC driver class of your database, the JDBC URL, user name, and password. If you are done, click [Connect].

You can save and reuse previously saved settings. The settings are stored in a properties file (see [Settings of the H2 Console](#)).

## Error Messages

Error messages in are shown in red. You can show/hide the stack trace of the exception by clicking on the message.

## Adding Database Drivers

To register additional JDBC drivers (MySQL, PostgreSQL, HSQLDB,...), add the jar file names to the environment variables H2DRIVERS or CLASSPATH. Example (Windows): to add the HSQLDB JDBC driver C:\Programs\hsqldb\lib\hsqldb.jar, set the environment variable H2DRIVERS to C:\Programs\hsqldb\lib\hsqldb.jar.

Multiple drivers can be set; entries need to be separated by ; (Windows) or : (other operating systems). Spaces in the path names are supported. The settings must not be quoted.

## Using the H2 Console

The H2 Console application has three main panels: the toolbar on top, the tree on the left, and the query/result panel on the right. The database objects (for example, tables) are listed on the left. Type a SQL command in the query panel and click [Run]. The result appears just below the command.

## Inserting Table Names or Column Names

To insert table and column names into the script, click on the item in the tree. If you click on a table while the query is empty, then `SELECT * FROM ...` is added. While typing a query, the table that was used is expanded in the tree. For example if you type `SELECT * FROM TEST T WHERE T.` then the table TEST is expanded.

## Disconnecting and Stopping the Application

To log out of the database, click [Disconnect] in the toolbar panel. However, the server is still running and ready to accept new sessions.

To stop the server, right click on the system tray icon and select [Exit]. If you don't have the system tray icon, navigate to [Preferences] and click [Shutdown], press [Ctrl]+[C] in the console where the server was started (Windows), or close the console window.

## Special H2 Console Syntax

The H2 Console supports a few built-in commands. Those are interpreted within the H2 Console, so they work with any database. Built-in commands need to be at the beginning of a statement (before any remarks), otherwise they are not parsed correctly. If in doubt, add ; before the command.

Command(s)	Description
@autocommit_true; @autocommit_false;	Enable or disable autocommit.
@cancel;	Cancel the currently running statement.
@columns null null TEST; @index_info null null TEST; @tables; @tables null null TEST;	Call the corresponding DatabaseMetaData.get method. Patterns are case sensitive (usually identifiers are uppercase). For information about the parameters, see the Javadoc documentation. Missing parameters at the end of the line are set to null. The complete list of metadata commands is: @attributes, @best_row_identifier, @catalogs, @columns, @column_privileges, @cross_references, @exported_keys, @imported_keys, @index_info, @primary_keys, @procedures, @procedure_columns, @schemas, @super_tables, @super_types, @tables, @table_privileges, @table_types, @type_info, @udts, @version_columns
@edit select * from test;	Use an updatable result set.
@generated insert into test( ) values();	Show the result of Statement.getGeneratedKeys().
@history;	List the command history.
@info;	Display the result of various Connection and DatabaseMetaData methods.



@list select * from test;	Show the result set in list format (each column on its own line, with row numbers).
@loop 1000 select ?, ?/*rnd*/; @loop 1000 @statement select ?;	Run the statement this many times. Parameters (?) are set using a loop from 0 up to x - 1. Random values are used for each ?/*rnd*/. A Statement object is used instead of a PreparedStatement if @statement is used. Result sets are read until ResultSet.next() returns false. Timing information is printed.
@maxrows 20;	Set the maximum number of rows to display.
@memory;	Show the used and free memory. This will call System.gc().
@meta select 1;	List the ResultSetMetaData after running the query.
@parameter_meta select ?;	Show the result of the PreparedStatement.getParameterMetaData() calls. The statement is not executed.
@prof_start; call hash('SHA256', ", 10000 00); @prof_stop;	Start/stop the built-in profiling tool. The top 3 stack traces of the statement(s) between start and stop are listed (if there are 3).
@prof_start; @sleep 10; @prof_stop;	Sleep for a number of seconds. Used to profile a long running query or operation that is running in another session (but in the same process).
@transaction_isolation; @transaction_isolation 2;	Display (without parameters) or change (with parameters 1, 2, 4, 8) the transaction isolation level.

## Settings of the H2 Console

The settings of the H2 Console are stored in a configuration file called .h2.server.properties in you user home directory. For Windows installations, the user home directory is usually C:\Documents and Settings\[username] or C:\Users\[username]. The configuration file contains the settings of the application and is automatically created when the H2 Console is first started. Supported settings are:

- webAllowOthers: allow other computers to connect.
- webPort: the port of the H2 Console
- webSSL: use encrypted TLS (HTTPS) connections.

- webAdminPassword: password to access preferences and tools of H2 Console.

In addition to those settings, the properties of the last recently used connection are listed in the form <number>=<name>|<driver>|<url>|<user> using the escape character \. Example: 1=Generic H2 (Embedded)|org.h2.Driver|jdbc\:h2\:~/test|sa

## Connecting to a Database using JDBC

To connect to a database, a Java application first needs to load the database driver, and then get a connection. A simple way to do that is using the following code:

```
import java.sql.*;
public class Test {
    public static void main(String[] a)
        throws Exception {
        Connection conn = DriverManager.
            getConnection("jdbc:h2:~/test", "sa", "");
        // add application code here
        conn.close();
    }
}
```

This code opens a connection (using `DriverManager.getConnection()`). The driver name is "org.h2.Driver". The database URL always needs to start with `jdbc:h2:` to be recognized by this database. The second parameter in the `getConnection()` call is the user name (sa for System Administrator in this example). The third parameter is the password. In this database, user names are not case sensitive, but passwords are.

## Creating New Databases

By default, if the database specified in the [embedded](#) URL does not yet exist, a new (empty) database is created automatically. The user that created the database automatically becomes the administrator of this database.

Auto-creation of databases can be disabled, see [Opening a Database Only if it Already Exists](#).

H2 Console does not allow creation of databases by default. If H2 Console is launched with icon in the system tray its context menu can be used to create a new database.

## Using the Server

H2 currently supports three server: a web server (for the H2 Console), a TCP server (for client/server connections) and an PG server (for PostgreSQL clients). Please note that only the web server supports browser connections. The servers can be started in different ways, one is using the Server tool. Starting the server doesn't open a database - databases are opened as soon as a client connects.

### Starting the Server Tool from Command Line

To start the Server tool from the command line with the default settings, run:

```
java -cp h2*.jar org.h2.tools.Server
```

This will start the tool with the default options. To get the list of options and default values, run:

```
java -cp h2*.jar org.h2.tools.Server -?
```

There are options available to use other ports, and start or not start parts.

### Connecting to the TCP Server

To remotely connect to a database using the TCP server, use the following driver and database URL:

- JDBC driver class: org.h2.Driver
- Database URL: jdbc:h2:tcp://localhost/~/test

For details about the database URL, see also in Features. Please note that you can't connection with a web browser to this URL. You can only connect using a H2 client (over JDBC).

### Starting the TCP Server within an Application

Servers can also be started and stopped from within an application. Sample code:

```
import org.h2.tools.Server;
...
// start the TCP Server
Server server = Server.createTcpServer(args).start();
...
// stop the TCP Server
```

```
server.stop();
```

## Stopping a TCP Server from Another Process

The TCP server can be stopped from another process. To stop the server from the command line, run:

```
java org.h2.tools.Server -tcpShutdown tcp://localhost:9092
```

To stop the server from a user application, use the following code:

```
org.h2.tools.Server.shutdownTcpServer("tcp://localhost:9094");
```

This function will only stop the TCP server. If other server were started in the same process, they will continue to run. To avoid recovery when the databases are opened the next time, all connections to the databases should be closed before calling this method. To stop a remote server, remote connections must be enabled on the server. Shutting down a TCP server can be protected using the option `-tcpPassword` (the same password must be used to start and stop the TCP server).

## Using Hibernate

This database supports Hibernate version 3.1 and newer. You can use the HSQLDB Dialect, or the native H2 Dialect.

When using Hibernate, try to use the H2Dialect if possible. When using the H2Dialect, compatibility modes such as `MODE=MySQL` are not supported. When using such a compatibility mode, use the Hibernate dialect for the corresponding database instead of the H2Dialect; but please note H2 does not support all features of all databases.

## Using TopLink and Glassfish

To use H2 with Glassfish (or Sun AS), set the Datasource Classname to `org.h2.jdbcx.JdbcDataSource`. You can set this in the GUI at Application Server - Resources - JDBC - Connection Pools, or by editing the file `sun-resources.xml`: at element `jdbc-connection-pool`, set the attribute `datasource-classname` to `org.h2.jdbcx.JdbcDataSource`.

The H2 database is compatible with HSQLDB and PostgreSQL. To take advantage of H2 specific features, use the H2Platform. The source code of this platform is included in H2 at

src/tools/oracle/toplink/essentials/platform/database/DatabasePlatform.java.txt. You will need to copy this file to your application, and rename it to .java. To enable it, change the following setting in persistence.xml:

```
<property
  name="toplink.target-database"
  value="oracle.toplink.essentials.platform.database.H2Platform"/>
```

In old versions of Glassfish, the property name is `toplink.platform.class.name`.

To use H2 within Glassfish, copy the `h2*.jar` to the directory `glassfish/glassfish/lib`.

## Using EclipseLink

To use H2 in EclipseLink, use the platform class `org.eclipse.persistence.platform.database.H2Platform`. If this platform is not available in your version of EclipseLink, you can use the `OraclePlatform` instead in many case. See also [H2Platform](#).

## Using Apache ActiveMQ

When using H2 as the backend database for Apache ActiveMQ, please use the `TransactDatabaseLocker` instead of the default locking mechanism. Otherwise the database file will grow without bounds. The problem is that the default locking mechanism uses an uncommitted `UPDATE` transaction, which keeps the transaction log from shrinking (causes the database file to grow). Instead of using an `UPDATE` statement, the `TransactDatabaseLocker` uses `SELECT ... FOR UPDATE` which is not problematic. To use it, change the ApacheMQ configuration element `<jdbcPersistenceAdapter>` element, property `databaseLocker="org.apache.activemq.store.jdbc.adapter.TransactDatabaseLocker"`. However, using the MVCC mode will again result in the same problem. Therefore, please do not use the MVCC mode in this case. Another (more dangerous) solution is to set `useDatabaseLock` to `false`.

## Using H2 within NetBeans

There is a known issue when using the Netbeans SQL Execution Window: before executing a query, another query in the form `SELECT COUNT(*) FROM <query>` is run. This is a problem for queries that modify state, such as `SELECT SEQ.NEXTVAL`. In this case, two sequence values are allocated instead of just one.

## Using H2 with jOOQ

jOOQ adds a thin layer on top of JDBC, allowing for type-safe SQL construction, including advanced SQL, stored procedures and advanced data types. jOOQ takes your database schema as a base for code generation. If this is your example schema:

```
CREATE TABLE USER (ID INT, NAME VARCHAR(50));
```

then run the jOOQ code generator on the command line using this command:

```
java -cp jooq.jar;jooq-meta.jar;jooq-codegen.jar;h2-1.3.158.jar;.  
org.jooq.util.GenerationTool /codegen.xml
```

...where codegen.xml is on the classpath and contains this information

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>  
<configuration xmlns="http://www.jooq.org/xsd/jooq-codegen-2.3.0.xsd">  
  <jdbc>  
    <driver>org.h2.Driver</driver>  
    <url>jdbc:h2:~/test</url>  
    <user>sa</user>  
    <password></password>  
  </jdbc>  
  <generator>  
    <name>org.jooq.util.DefaultGenerator</name>  
    <database>  
      <name>org.jooq.util.h2.H2Database</name>  
      <includes>.*</includes>  
      <excludes></excludes>  
      <inputSchema>PUBLIC</inputSchema>  
    </database>  
    <generate></generate>  
    <target>  
      <packageName>org.jooq.h2.generated</packageName>  
      <directory>./src</directory>  
    </target>  
  </generator>  
</configuration>
```

Using the generated source, you can query the database as follows:

```
Factory create = new H2Factory(connection);  
Result<UserRecord> result =  
create.selectFrom(USER)
```

```
.where(NAME.like("Johnny%"))  
.orderBy(ID)  
.fetch();
```

See more details on [jOOQ Homepage](#) and in the [jOOQ Tutorial](#)

## Using Databases in Web Applications

There are multiple ways to access a database from within web applications. Here are some examples if you use Tomcat or JBoss.

### Embedded Mode

The (currently) simplest solution is to use the database in the embedded mode, that means open a connection in your application when it starts (a good solution is using a Servlet Listener, see below), or when a session starts. A database can be accessed from multiple sessions and applications at the same time, as long as they run in the same process. Most Servlet Containers (for example Tomcat) are just using one process, so this is not a problem (unless you run Tomcat in clustered mode). Tomcat uses multiple threads and multiple classloaders. If multiple applications access the same database at the same time, you need to put the database jar in the shared/lib or server/lib directory. It is a good idea to open the database when the web application starts, and close it when the web application stops. If using multiple applications, only one (any) of them needs to do that. In the application, an idea is to use one connection per Session, or even one connection per request (action). Those connections should be closed after use if possible (but it's not that bad if they don't get closed).

### Server Mode

The server mode is similar, but it allows you to run the server in another process.

## Using a Servlet Listener to Start and Stop a Database

Add the h2\*.jar file to your web application, and add the following snippet to your web.xml file (between the context-param and the filter section):

```
<listener>  
  <listener-class>org.h2.server.web.DbStarter</listener-class>  
</listener>
```

For details on how to access the database, see the file DbStarter.java. By default this tool opens an embedded connection using the database URL jdbc:h2:~/test,

user name sa, and password sa. If you want to use this connection within your servlet, you can access as follows:

```
Connection conn = getServletContext().getAttribute("connection");
```

DbStarter can also start the TCP server, however this is disabled by default. To enable it, use the parameter db.tcpServer in the file web.xml. Here is the complete list of options. These options need to be placed between the description tag and the listener / filter tags:

```
<context-param>
  <param-name>db.url</param-name>
  <param-value>jdbc:h2:~/test</param-value>
</context-param>
<context-param>
  <param-name>db.user</param-name>
  <param-value>sa</param-value>
</context-param>
<context-param>
  <param-name>db.password</param-name>
  <param-value>sa</param-value>
</context-param>
<context-param>
  <param-name>db.tcpServer</param-name>
  <param-value>-tcpAllowOthers</param-value>
</context-param>
```

When the web application is stopped, the database connection will be closed automatically. If the TCP server is started within the DbStarter, it will also be stopped automatically.

## Using the H2 Console Servlet

The H2 Console is a standalone application and includes its own web server, but it can be used as a servlet as well. To do that, include the h2\*.jar file in your application, and add the following configuration to your web.xml:

```
<servlet>
  <servlet-name>H2Console</servlet-name>
  <servlet-class>org.h2.server.web.WebServlet</servlet-class>
  <!--
  <init-param>
    <param-name>webAllowOthers</param-name>
    <param-value></param-value>
```



```
</init-param>
<init-param>
  <param-name>trace</param-name>
  <param-value></param-value>
</init-param>
-->
<load-on-startup>1</load-on-startup>
</servlet>
<servlet-mapping>
  <servlet-name>H2Console</servlet-name>
  <url-pattern>/console/*</url-pattern>
</servlet-mapping>
```

For details, see also `src/tools/WEB-INF/web.xml`.

To create a web application with just the H2 Console, run the following command:

```
build warConsole
```

## Android

You can use this database on an Android device (using the Dalvik VM) instead of or in addition to SQLite. So far, only very few tests and benchmarks were run, but it seems that performance is similar to SQLite, except for opening and closing a database, which is not yet optimized in H2 (H2 takes about 0.2 seconds, and SQLite about 0.02 seconds). Read operations seem to be a bit faster than SQLite, and write operations seem to be slower. So far, only very few tests have been run, and everything seems to work as expected. Fulltext search was not yet tested, however the native fulltext search should work.

Reasons to use H2 instead of SQLite are:

- Full Unicode support including UPPER() and LOWER().
- Streaming API for BLOB and CLOB data.
- Fulltext search.
- Multiple connections.
- User defined functions and triggers.
- Database file encryption.
- Reading and writing CSV files (this feature can be used outside the database as well).
- Referential integrity and check constraints.
- Better data type and SQL support.
- In-memory databases, read-only databases, linked tables.

- Better compatibility with other databases which simplifies porting applications.
- Possibly better performance (so far for read operations).
- Server mode (accessing a database on a different machine over TCP/IP).

Currently only the JDBC API is supported (it is planned to support the Android database API in future releases). Both the regular H2 jar file and the smaller h2small-\*.jar can be used. To create the smaller jar file, run the command ./build.sh jarSmall (Linux / Mac OS) or build.bat jarSmall (Windows).

The database files needs to be stored in a place that is accessible for the application. Example:

```
String url = "jdbc:h2:/data/data/" +
    "com.example.hello" +
    "/data/hello" +
    ";FILE_LOCK=FS" +
    ";PAGE_SIZE=1024" +
    ";CACHE_SIZE=8192";
conn = DriverManager.getConnection(url);
...
```

Limitations: Using a connection pool is currently not supported, because the required javax.sql. classes are not available on Android.

## CSV (Comma Separated Values) Support

The CSV file support can be used inside the database using the functions CSVREAD and CSVWRITE, or it can be used outside the database as a standalone tool.

### Reading a CSV File from Within a Database

A CSV file can be read using the function CSVREAD. Example:

```
SELECT * FROM CSVREAD('test.csv');
```

Please note for performance reason, CSVREAD should not be used inside a join. Instead, import the data first (possibly into a temporary table), create the required indexes if necessary, and then query this table.

## Importing Data from a CSV File

A fast way to load or import data (sometimes called 'bulk load') from a CSV file is to combine table creation with import. Optionally, the column names and data types can be set when creating the table. Another option is to use INSERT INTO ... SELECT.

```
CREATE TABLE TEST AS SELECT * FROM CSVREAD('test.csv');
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255))
  AS SELECT * FROM CSVREAD('test.csv');
```

## Writing a CSV File from Within a Database

The built-in function CSVWRITE can be used to create a CSV file from a query. Example:

```
CREATE TABLE TEST(ID INT, NAME VARCHAR);
INSERT INTO TEST VALUES(1, 'Hello'), (2, 'World');
CALL CSVWRITE('test.csv', 'SELECT * FROM TEST');
```

## Writing a CSV File from a Java Application

The Csv tool can be used in a Java application even when not using a database at all. Example:

```
import java.sql.*;
import org.h2.tools.Csv;
import org.h2.tools.SimpleResultSet;
public class TestCsv {
    public static void main(String[] args) throws Exception {
        SimpleResultSet rs = new SimpleResultSet();
        rs.addColumn("NAME", Types.VARCHAR, 255, 0);
        rs.addColumn("EMAIL", Types.VARCHAR, 255, 0);
        rs.addRow("Bob Meier", "bob.meier@abcde.abc");
        rs.addRow("John Jones", "john.jones@abcde.abc");
        new Csv().write("data/test.csv", rs, null);
    }
}
```

## Reading a CSV File from a Java Application

It is possible to read a CSV file without opening a database. Example:

```

import java.sql.*;
import org.h2.tools.Csv;
public class TestCsv {
    public static void main(String[] args) throws Exception {
        ResultSet rs = new Csv().read("data/test.csv", null, null);
        ResultSetMetaData meta = rs.getMetaData();
        while (rs.next()) {
            for (int i = 0; i < meta.getColumnCount(); i++) {
                System.out.println(
                    meta.getColumnLabel(i + 1) + ": " +
                    rs.getString(i + 1));
            }
            System.out.println();
        }
        rs.close();
    }
}

```

## Upgrade, Backup, and Restore

### Database Upgrade

The recommended way to upgrade from one version of the database engine to the next version is to create a backup of the database (in the form of a SQL script) using the old engine, and then execute the SQL script using the new engine.

### Backup using the Script Tool

The recommended way to backup a database is to create a compressed SQL script file. This will result in a small, human readable, and database version independent backup. Creating the script will also verify the checksums of the database file. The Script tool is ran as follows:

```

java org.h2.tools.Script -url jdbc:h2:~/test -user sa -script test.zip -options
compression zip

```

It is also possible to use the SQL command SCRIPT to create the backup of the database. For more information about the options, see the SQL command SCRIPT. The backup can be done remotely, however the file will be created on the server side. The built in FTP server could be used to retrieve the file from the server.

## Restore from a Script

To restore a database from a SQL script file, you can use the RunScript tool:

```
java org.h2.tools.RunScript -url jdbc:h2:~/test -user sa -script test.zip -options  
compression zip
```

For more information about the options, see the SQL command RUNSCRIPT. The restore can be done remotely, however the file needs to be on the server side. The built in FTP server could be used to copy the file to the server. It is also possible to use the SQL command RUNSCRIPT to execute a SQL script. SQL script files may contain references to other script files, in the form of RUNSCRIPT commands. However, when using the server mode, the references script files need to be available on the server side.

## Online Backup

The BACKUP SQL statement and the Backup tool both create a zip file with the database file. However, the contents of this file are not human readable.

The resulting backup is transactionally consistent, meaning the consistency and atomicity rules apply.

```
BACKUP TO 'backup.zip'
```

The Backup tool (org.h2.tools.Backup) can not be used to create a online backup; the database must not be in use while running this program.

Creating a backup by copying the database files while the database is running is not supported, except if the file systems support creating snapshots. With other file systems, it can't be guaranteed that the data is copied in the right order.

## Command Line Tools

This database comes with a number of command line tools. To get more information about a tool, start it with the parameter '-?', for example:

```
java -cp h2*.jar org.h2.tools.Backup -?
```

The command line tools are:

- Backup creates a backup of a database.
- ChangeFileEncryption allows changing the file encryption password or algorithm of a database.

- Console starts the browser based H2 Console.
- ConvertTraceFile converts a .trace.db file to a Java application and SQL script.
- CreateCluster creates a cluster from a standalone database.
- DeleteDbFiles deletes all files belonging to a database.
- Recover helps recovering a corrupted database.
- Restore restores a backup of a database.
- RunScript runs a SQL script against a database.
- Script allows converting a database to a SQL script for backup or migration.
- Server is used in the server mode to start a H2 server.
- Shell is a command line database tool.

The tools can also be called from an application by calling the main or another public method. For details, see the Javadoc documentation.

## The Shell Tool

The Shell tool is a simple interactive command line tool. To start it, type:

```
java -cp h2*.jar org.h2.tools.Shell
```

You will be asked for a database URL, JDBC driver, user name, and password. The connection setting can also be set as command line parameters. After connecting, you will get the list of options. The built-in commands don't need to end with a semicolon, but SQL statements are only executed if the line ends with a semicolon ;. This allows to enter multi-line statements:

```
sql> select * from test
...> where id = 0;
```

By default, results are printed as a table. For results with many column, consider using the list mode:

```
sql> list
Result list mode is now on
sql> select * from test;
ID : 1
NAME: Hello

ID : 2
NAME: World
(2 rows, 0 ms)
```

# Using OpenOffice Base

OpenOffice.org Base supports database access over the JDBC API. To connect to a H2 database using OpenOffice Base, you first need to add the JDBC driver to OpenOffice. The steps to connect to a H2 database are:

- Start OpenOffice Writer, go to [Tools], [Options]
- Make sure you have selected a Java runtime environment in OpenOffice.org / Java
- Click [Class Path...], [Add Archive...]
- Select your h2 jar file (location is up to you, could be wherever you choose)
- Click [OK] (as much as needed), stop OpenOffice (including the Quickstarter)
- Start OpenOffice Base
- Connect to an existing database; select [JDBC]; [Next]
- Example datasource URL: jdbc:h2:~/test
- JDBC driver class: org.h2.Driver

Now you can access the database stored in the current users home directory.

To use H2 in NeoOffice (OpenOffice without X11):

- In NeoOffice, go to [NeoOffice], [Preferences]
- Look for the page under [NeoOffice], [Java]
- Click [Class Path], [Add Archive...]
- Select your h2 jar file (location is up to you, could be wherever you choose)
- Click [OK] (as much as needed), restart NeoOffice.

Now, when creating a new database using the "Database Wizard" :

- Click [File], [New], [Database].
- Select [Connect to existing database] and the select [JDBC]. Click next.
- Example datasource URL: jdbc:h2:~/test
- JDBC driver class: org.h2.Driver

Another solution to use H2 in NeoOffice is:

- Package the h2 jar within an extension package
- Install it as a Java extension in NeoOffice

This can be done by create it using the NetBeans OpenOffice plugin. See also [Extensions Development](#).

## Java Web Start / JNLP

When using Java Web Start / JNLP (Java Network Launch Protocol), permissions tags must be set in the .jnlp file, and the application .jar file must be signed. Otherwise, when trying to write to the file system, the following exception will occur: java.security.AccessControlException: access denied (java.io.FilePermission ... read). Example permission tags:

```
<security>
  <all-permissions/>
</security>
```

## Using a Connection Pool

For H2, opening a connection is fast if the database is already open. Still, using a connection pool improves performance if you open and close connections a lot. A simple connection pool is included in H2. It is based on the [Mini Connection Pool Manager](#) from Christian d'Heureuse. There are other, more complex, open source connection pools available, for example the [Apache Commons DBCP](#). For H2, it is about twice as faster to get a connection from the built-in connection pool than to get one using DriverManager.getConnection(). The build-in connection pool is used as follows:

```
import java.sql.*;
import org.h2.jdbcx.JdbcConnectionPool;
public class Test {
    public static void main(String[] args) throws Exception {
        JdbcConnectionPool cp = JdbcConnectionPool.create(
            "jdbc:h2:~/test", "sa", "sa");
        for (int i = 0; i < args.length; i++) {
            Connection conn = cp.getConnection();
            conn.createStatement().execute(args[i]);
            conn.close();
        }
        cp.dispose();
    }
}
```

## Fulltext Search

H2 includes two fulltext search implementations. One is using Apache Lucene, and the other (the native implementation) stores the index data in special tables in the database.



## Using the Native Fulltext Search

To initialize, call:

```
CREATE ALIAS IF NOT EXISTS FT_INIT FOR "org.h2.fulltext.FullText.init";  
CALL FT_INIT();
```

You need to initialize it in each database where you want to use it. Afterwards, you can create a fulltext index for a table using:

```
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR);  
INSERT INTO TEST VALUES(1, 'Hello World');  
CALL FT_CREATE_INDEX('PUBLIC', 'TEST', NULL);
```

PUBLIC is the schema name, TEST is the table name. The list of column names (comma separated) is optional, in this case all columns are indexed. The index is updated in realtime. To search the index, use the following query:

```
SELECT * FROM FT_SEARCH('Hello', 0, 0);
```

This will produce a result set that contains the query needed to retrieve the data:

```
QUERY: "PUBLIC"."TEST" WHERE "ID"=1
```

To drop an index on a table:

```
CALL FT_DROP_INDEX('PUBLIC', 'TEST');
```

To get the raw data, use `FT_SEARCH_DATA('Hello', 0, 0);`. The result contains the columns `SCHEMA` (the schema name), `TABLE` (the table name), `COLUMNS` (an array of column names), and `KEYS` (an array of objects). To join a table, use a join as in: `SELECT T.* FROM FT_SEARCH_DATA('Hello', 0, 0) FT, TEST T WHERE FT.TABLE='TEST' AND T.ID=FT.KEYS[0];`

You can also call the index from within a Java application:

```
org.h2.fulltext.FullText.search(conn, text, limit, offset);  
org.h2.fulltext.FullText.searchData(conn, text, limit, offset);
```

## Using the Apache Lucene Fulltext Search

To use the Apache Lucene full text search, you need the Lucene library in the classpath. Apache Lucene 5.5.5 or later version up to 7.6.0 is required. Newer versions may also work, but were not tested. How to do that depends on the

application; if you use the H2 Console, you can add the Lucene jar file to the environment variables H2DRIVERS or CLASSPATH. To initialize the Lucene fulltext search in a database, call:

```
CREATE ALIAS IF NOT EXISTS FTL_INIT FOR
"org.h2.fulltext.FullTextLucene.init";
CALL FTL_INIT();
```

You need to initialize it in each database where you want to use it. Afterwards, you can create a full text index for a table using:

```
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR);
INSERT INTO TEST VALUES(1, 'Hello World');
CALL FTL_CREATE_INDEX('PUBLIC', 'TEST', NULL);
```

PUBLIC is the schema name, TEST is the table name. The list of column names (comma separated) is optional, in this case all columns are indexed. The index is updated in realtime. To search the index, use the following query:

```
SELECT * FROM FTL_SEARCH('Hello', 0, 0);
```

This will produce a result set that contains the query needed to retrieve the data:

```
QUERY: "PUBLIC"."TEST" WHERE "ID"=1
```

To drop an index on a table (be warned that this will re-index all of the full-text indices for the entire database):

```
CALL FTL_DROP_INDEX('PUBLIC', 'TEST');
```

To get the raw data, use `FTL_SEARCH_DATA('Hello', 0, 0);`. The result contains the columns SCHEMA (the schema name), TABLE (the table name), COLUMNS (an array of column names), and KEYS (an array of objects). To join a table, use a join as in: `SELECT T.* FROM FTL_SEARCH_DATA('Hello', 0, 0) FT, TEST T WHERE FT.TABLE='TEST' AND T.ID=FT.KEYS[0];`

You can also call the index from within a Java application:

```
org.h2.fulltext.FullTextLucene.search(conn, text, limit, offset);
org.h2.fulltext.FullTextLucene.searchData(conn, text, limit, offset);
```

The Lucene fulltext search supports searching in specific column only. Column names must be uppercase (except if the original columns are double quoted). For column names starting with an underscore (\_), another underscore needs to be added. Example:

```
CREATE ALIAS IF NOT EXISTS FTL_INIT FOR
"org.h2.fulltext.FullTextLucene.init";
CALL FTL_INIT();
DROP TABLE IF EXISTS TEST;
CREATE TABLE TEST(ID INT PRIMARY KEY, FIRST_NAME VARCHAR,
LAST_NAME VARCHAR);
CALL FTL_CREATE_INDEX('PUBLIC', 'TEST', NULL);
INSERT INTO TEST VALUES(1, 'John', 'Wayne');
INSERT INTO TEST VALUES(2, 'Elton', 'John');
SELECT * FROM FTL_SEARCH_DATA('John', 0, 0);
SELECT * FROM FTL_SEARCH_DATA('LAST_NAME:John', 0, 0);
CALL FTL_DROP_ALL();
```

## User-Defined Variables

This database supports user-defined variables. Variables start with @ and can be used wherever expressions or parameters are allowed. Variables are not persisted and session scoped, that means only visible from within the session in which they are defined. A value is usually assigned using the SET command:

```
SET @USER = 'Joe';
```

The value can also be changed using the SET() method. This is useful in queries:

```
SET @TOTAL = NULL;
SELECT X, SET(@TOTAL, IFNULL(@TOTAL, 1.) * X) F FROM SYSTEM_RANGE(1,
50);
```

Variables that are not set evaluate to NULL. The data type of a user-defined variable is the data type of the value assigned to it, that means it is not necessary (or possible) to declare variable names before using them. There are no restrictions on the assigned values; large objects (LOBs) are supported as well. Rolling back a transaction does not affect the value of a user-defined variable.

## Date and Time

Date, time and timestamp values support ISO 8601 formatting, including time zone:

```
CALL TIMESTAMP '2008-01-01 12:00:00+01:00';
```

If the time zone is not set, the value is parsed using the current time zone setting of the system. Date and time information is stored in H2 database files with or without time zone information depending on used data type.

- With `TIMESTAMP` data type if the database is opened using another system time zone, the date and time will be the same. That means if you store the value `'2000-01-01 12:00:00'` in one time zone, then close the database and open the database again in a different time zone, you will also get `'2000-01-01 12:00:00'`. Please note that changing the time zone after the H2 driver is loaded is not supported.
- With `TIMESTAMP WITH TIME ZONE` data type time zone offset is stored and if you store the value `'2008-01-01 12:00:00+01:00'` it remains the same even if you close and reopen the database with a different time zone. If you store the value with specified time zone name like `'2008-01-01 12:00:00 Europe/Berlin'` this name will be converted to time zone offset. Names of time zones are not stored.

## Using Spring

### Using the TCP Server

Use the following configuration to start and stop the H2 TCP server using the Spring Framework:

```
<bean id = "org.h2.tools.Server"
      class="org.h2.tools.Server"
      factory-method="createTcpServer"
      init-method="start"
      destroy-method="stop">
  <constructor-arg value="-tcp,-tcpAllowOthers,-tcpPort,8043" />
</bean>
```

The `destroy-method` will help prevent exceptions on hot-redeployment or when restarting the server.

## OSGi

The standard H2 jar can be dropped in as a bundle in an OSGi container. H2 implements the JDBC Service defined in OSGi Service Platform Release 4 Version 4.2 Enterprise Specification. The H2 Data Source Factory service is registered with the following properties: `OSGI_JDBC_DRIVER_CLASS=org.h2.Driver` and `OSGI_JDBC_DRIVER_NAME=H2 JDBC Driver`. The `OSGI_JDBC_DRIVER_VERSION` property reflects the version of the driver as is.

The following standard configuration properties are supported: JDBC\_USER, JDBC\_PASSWORD, JDBC\_DESCRIPTION, JDBC\_DATASOURCE\_NAME, JDBC\_NETWORK\_PROTOCOL, JDBC\_URL, JDBC\_SERVER\_NAME, JDBC\_PORT\_NUMBER. Any other standard property will be rejected. Non-standard properties will be passed on to H2 in the connection URL.

## Java Management Extension (JMX)

Management over JMX is supported, but not enabled by default. To enable JMX, append ;JMX=TRUE to the database URL when opening the database. Various tools support JMX, one such tool is the jconsole. When opening the jconsole, connect to the process where the database is open (when using the server mode, you need to connect to the server process). Then go to the MBeans section. Under org.h2 you will find one entry per database. The object name of the entry is the database short name, plus the path (each colon is replaced with an underscore character).

The following attributes and operations are supported:

- CacheSize: the cache size currently in use in KB.
- CacheSizeMax (read/write): the maximum cache size in KB.
- Exclusive: whether this database is open in exclusive mode or not.
- FileReadCount: the number of file read operations since the database was opened.
- FileSize: the file size in KB.
- FileWriteCount: the number of file write operations since the database was opened.
- FileWriteCountTotal: the number of file write operations since the database was created.
- LogMode (read/write): the current transaction log mode. See SET LOG for details.
- Mode: the compatibility mode (REGULAR if no compatibility mode is used).
- MultiThreaded: true if multi-threaded is enabled.
- Mvcc: true if MVCC is enabled.
- ReadOnly: true if the database is read-only.
- TraceLevel (read/write): the file trace level.
- Version: the database version in use.
- listSettings: list the database settings.
- listSessions: list the open sessions, including currently executing statement (if any) and locked tables (if any).

To enable JMX, you may need to set the system properties com.sun.management.jmxremote and com.sun.management.jmxremote.port as required by the JVM.

# Features

Feature List  
H2 in Use  
Connection Modes  
Database URL Overview  
Connecting to an Embedded (Local) Database  
In-Memory Databases  
Database Files Encryption  
Database File Locking  
Opening a Database Only if it Already Exists  
Closing a Database  
Ignore Unknown Settings  
Changing Other Settings when Opening a Connection  
Custom File Access Mode  
Multiple Connections  
Database File Layout  
Logging and Recovery  
Compatibility  
Auto-Reconnect  
Automatic Mixed Mode  
Page Size  
Using the Trace Options  
Using Other Logging APIs  
Read Only Databases  
Read Only Databases in Zip or Jar File  
Computed Columns / Function Based Index  
Multi-Dimensional Indexes  
User-Defined Functions and Stored Procedures  
Pluggable or User-Defined Tables  
Triggers  
Compacting a Database  
Cache Settings  
External Authentication (Experimental)

## Feature List

### Main Features

- Very fast database engine
- Open source
- Written in Java

- Supports standard SQL, JDBC API
- Embedded and Server mode, Clustering support
- Strong security features
- The PostgreSQL ODBC driver can be used
- Multi version concurrency

## **Additional Features**

- Disk based or in-memory databases and tables, read-only database support, temporary tables
- Transaction support (read committed), 2-phase-commit
- Multiple connections, table level locking
- Cost based optimizer, using a genetic algorithm for complex queries, zero-administration
- Scrollable and updatable result set support, large result set, external result sorting, functions can return a result set
- Encrypted database (AES), SHA-256 password encryption, encryption functions, SSL

## **SQL Support**

- Support for multiple schemas, information schema
- Referential integrity / foreign key constraints with cascade, check constraints
- Inner and outer joins, subqueries, read only views and inline views
- Triggers and Java functions / stored procedures
- Many built-in functions, including XML and lossless data compression
- Wide range of data types including large objects (BLOB/CLOB) and arrays
- Sequence and autoincrement columns, computed columns (can be used for function based indexes)
- ORDER BY, GROUP BY, HAVING, UNION, OFFSET / FETCH (including PERCENT and WITH TIES), LIMIT, TOP, DISTINCT / DISTINCT ON (...)
- Collation support, including support for the ICU4J library
- Support for users and roles
- Compatibility modes for IBM DB2, Apache Derby, HSQLDB, MS SQL Server, MySQL, Oracle, and PostgreSQL.

## **Security Features**

- Includes a solution for the SQL injection problem
- User password authentication uses SHA-256 and salt
- For server mode connections, user passwords are never transmitted in plain text over the network (even when using insecure connections; this only

- applies to the TCP server and not to the H2 Console however; it also doesn't apply if you set the password in the database URL)
- All database files (including script files that can be used to backup data) can be encrypted using the AES-128 encryption algorithm
- The remote JDBC driver supports TCP/IP connections over TLS
- The built-in web server supports connections over TLS
- Passwords can be sent to the database using char arrays instead of Strings

## Other Features and Tools

- Small footprint (around 2 MB), low memory requirements
- Multiple index types (b-tree, tree, hash)
- Support for multi-dimensional indexes
- CSV (comma separated values) file support
- Support for linked tables, and a built-in virtual 'range' table
- Supports the EXPLAIN PLAN statement; sophisticated trace options
- Database closing can be delayed or disabled to improve the performance
- Web-based Console application (translated to many languages) with autocomplete
- The database can generate SQL script files
- Contains a recovery tool that can dump the contents of the database
- Support for variables (for example to calculate running totals)
- Automatic re-compilation of prepared statements
- Uses a small number of database files
- Uses a checksum for each record and log entry for data integrity
- Well tested (high code coverage, randomized stress tests)

## H2 in Use

For a list of applications that work with or use H2, see: [Links](#).

## Connection Modes


The following connection modes are supported:

- Embedded mode (local connections using JDBC)
- Server mode (remote connections using JDBC or ODBC over TCP/IP)
- Mixed mode (local and remote connections at the same time)



## Embedded Mode


In embedded mode, an application opens a database from within the same JVM using JDBC. This is the fastest and easiest connection mode. The disadvantage is that a database may only be open in one virtual machine (and class loader) at any time. As in all modes, both persistent and in-memory databases are supported. There is no limit on the number of database open concurrently, or on the number of open connections.



## Server Mode

When using the server mode (sometimes called remote mode or client/server mode), an application opens a database remotely using the JDBC or ODBC API. A server needs to be started within the same or another virtual machine, or on another computer. Many applications can connect to the same database at the same time, by connecting to this server. Internally, the server process opens the database(s) in embedded mode.


The server mode is slower than the embedded mode, because all data is transferred over TCP/IP. As in all modes, both persistent and in-memory databases are supported. There is no limit on the number of database open concurrently per server, or on the number of open connections.



## Mixed Mode

The mixed mode is a combination of the embedded and the server mode. The first application that connects to a database does that in embedded mode, but also starts a server so that other applications (running in different processes or virtual machines) can concurrently access the same data. The local connections are as fast as if the database is used in just the embedded mode, while the remote connections are a bit slower.

The server can be started and stopped from within the application (using the server API), or automatically (automatic mixed mode). When using the [automatic mixed mode](#), all clients that want to connect to the database (no matter if it's an local or remote connection) can do so using the exact same database URL.



## Database URL Overview

This database supports multiple connection modes and connection settings. This is achieved using different database URLs. Settings in the URLs are not case sensitive.

Topic	URL Format and Examples
<a href="#">Embedded (local) connection</a>	<code>jdbc:h2:[file:][&lt;path&gt;]&lt;databaseName&gt;</code> <code>jdbc:h2:~/test</code> <code>jdbc:h2:file:/data/sample</code> <code>jdbc:h2:file:C:/data/sample (Windows only)</code>
<a href="#">In-memory (private)</a>	<code>jdbc:h2:mem:</code>
<a href="#">In-memory (named)</a>	<code>jdbc:h2:mem:&lt;databaseName&gt;</code> <code>jdbc:h2:mem:test_mem</code>
<a href="#">Server mode (remote connections) using TCP/IP</a>	<code>jdbc:h2:tcp://&lt;server&gt;[:&lt;port&gt;]/[&lt;path&gt;]&lt;databaseName&gt;</code> <code>jdbc:h2:tcp://localhost/~/test</code>

	jdbc:h2:tcp://dbserve:8084/~/sample jdbc:h2:tcp://localhost/mem:test
Server mode (remote connections) using TLS	jdbc:h2:ssl://<server>[:<port>]/ [<path>]<databaseName> jdbc:h2:ssl://localhost:8085/~/sample;
Using encrypted files	jdbc:h2:<url>;CIPHER=AES jdbc:h2:ssl://localhost/~/test;CIPHER=AES jdbc:h2:file:~/secure;CIPHER=AES
File locking methods	jdbc:h2:<url>;FILE_LOCK={FILE SOCKET NO} jdbc:h2:file:~/private;CIPHER=AES;FILE_LOCK=SOCKET
Only open if it already exists	jdbc:h2:<url>;IFEXISTS=TRUE jdbc:h2:file:~/sample;IFEXISTS=TRUE
Don't close the database when the VM exits	jdbc:h2:<url>;DB_CLOSE_ON_EXIT=FALSE
Execute SQL on connection	jdbc:h2:<url>;INIT=RUNSCRIPT FROM '~/create.sql' jdbc:h2:file:~/sample;INIT=RUNSCRIPT FROM '~/create.sql';RUNSCRIPT FROM '~/populate.sql'
User name and/or password	jdbc:h2:<url>[;USER=<username>] [;PASSWORD=<value>] jdbc:h2:file:~/sample;USER=sa;PASSWORD=123
Debug trace settings	jdbc:h2:<url>;TRACE_LEVEL_FILE=<level 0..3> jdbc:h2:file:~/sample;TRACE_LEVEL_FILE=3
Ignore unknown settings	jdbc:h2:<url>;IGNORE_UNKNOWN_SETTINGS=TRUE
Custom file access mode	jdbc:h2:<url>;ACCESS_MODE_DATA=rws
Database in a zip file	jdbc:h2:zip:<zipFileName>!/<databaseName> jdbc:h2:zip:~/db.zip!/test
Compatibility mode	jdbc:h2:<url>;MODE=<databaseType> jdbc:h2:~/test;MODE=MYSQL
Auto-reconnect	jdbc:h2:<url>;AUTO_RECONNECT=TRUE jdbc:h2:tcp://localhost/~/test;AUTO_RECONNECT=TRUE
Automatic mixed mode	jdbc:h2:<url>;AUTO_SERVER=TRUE jdbc:h2:~/test;AUTO_SERVER=TRUE
Page size	jdbc:h2:<url>;PAGE_SIZE=512

Changing other settings	<pre>jdbc:h2:&lt;url&gt;;&lt;setting&gt;=&lt;value&gt;[;&lt;setting&gt;=&lt;value&gt;...] jdbc:h2:file:~/sample;TRACE_LEVEL_SYSTEM_OUT=3</pre>
-------------------------	--

## Connecting to an Embedded (Local) Database

The database URL for connecting to a local database is `jdbc:h2:[file:] [<path>]<databaseName>`. The prefix `file:` is optional. If no or only a relative path is used, then the current working directory is used as a starting point. The case sensitivity of the path and database name depend on the operating system, however it is recommended to use lowercase letters only. The database name must be at least three characters long (a limitation of `File.createTempFile`). The database name must not contain a semicolon. To point to the user home directory, use `~/`, as in: `jdbc:h2:~/test`.

## In-Memory Databases

For certain use cases (for example: rapid prototyping, testing, high performance operations, read-only databases), it may not be required to persist data, or persist changes to the data. This database supports the in-memory mode, where the data is not persisted.

In some cases, only one connection to a in-memory database is required. This means the database to be opened is private. In this case, the database URL is `jdbc:h2:mem:`. Opening two connections within the same virtual machine means opening two different (private) databases.

Sometimes multiple connections to the same in-memory database are required. In this case, the database URL must include a name. Example: `jdbc:h2:mem:db1`. Accessing the same database using this URL only works within the same virtual machine and class loader environment.

To access an in-memory database from another process or from another computer, you need to start a TCP server in the same process as the in-memory database was created. The other processes then need to access the database over TCP/IP or TLS, using a database URL such as:  
`jdbc:h2:tcp://localhost/mem:db1`.

By default, closing the last connection to a database closes the database. For an in-memory database, this means the content is lost. To keep the database open, add `;DB_CLOSE_DELAY=-1` to the database URL. To keep the content of an in-memory database as long as the virtual machine is alive, use  
`jdbc:h2:mem:test;DB_CLOSE_DELAY=-1`.

## Database Files Encryption

The database files can be encrypted. Three encryption algorithms are supported:

- "AES" - also known as Rijndael, only AES-128 is implemented.
- "XTEA" - the 32 round version.
- "FOG" - pseudo-encryption only useful for hiding data from a text editor.

To use file encryption, you need to specify the encryption algorithm (the 'cipher') and the file password (in addition to the user password) when connecting to the database.

### Creating a New Database with File Encryption

By default, a new database is automatically created if it does not exist yet. To create an encrypted database, connect to it as it would already exist.

### Connecting to an Encrypted Database

The encryption algorithm is set in the database URL, and the file password is specified in the password field, before the user password. A single space separates the file password and the user password; the file password itself may not contain spaces. File passwords and user passwords are case sensitive. Here is an example to connect to a password-encrypted database:

```
String url = "jdbc:h2:~/test;CIPHER=AES";  
String user = "sa";  
String pwds = "filepwd userpwd";  
conn = DriverManager.  
    getConnection(url, user, pwds);
```

### Encrypting or Decrypting a Database

To encrypt an existing database, use the ChangeFileEncryption tool. This tool can also decrypt an encrypted database, or change the file encryption key. The tool is available from within the H2 Console in the tools section, or you can run it from the command line. The following command line will encrypt the database test in the user home directory with the file password filepwd and the encryption algorithm AES:

```
java -cp h2*.jar org.h2.tools.ChangeFileEncryption -dir ~ -db test -cipher AES  
-encrypt filepwd
```

## Database File Locking

Whenever a database is opened, a lock file is created to signal other processes that the database is in use. If database is closed, or if the process that opened the database terminates, this lock file is deleted.

The following file locking methods are implemented:

- The default method is FILE and uses a watchdog thread to protect the database file. The watchdog reads the lock file each second.
- The second method is SOCKET and opens a server socket. The socket method does not require reading the lock file every second. The socket method should only be used if the database files are only accessed by one (and always the same) computer.
- The third method is FS. This will use native file locking using `FileChannel.lock`.
- It is also possible to open the database without file locking; in this case it is up to the application to protect the database files. Failing to do so will result in a corrupted database. Using the method NO forces the database to not create a lock file at all. Please note that this is unsafe as another process is able to open the same database, possibly leading to data corruption.

To open the database with a different file locking method, use the parameter `FILE_LOCK`. The following code opens the database with the 'socket' locking method:

```
String url = "jdbc:h2:~/test;FILE_LOCK=SOCKET";
```

For more information about the algorithms, see [Advanced / File Locking Protocols](#).

## Opening a Database Only if it Already Exists

By default, when an application calls `DriverManager.getConnection(url, ...)` with [embedded](#) URL and the database specified in the URL does not yet exist, a new (empty) database is created. In some situations, it is better to restrict creating new databases, and only allow to open existing databases. To do this, add `;IFEXISTS=TRUE` to the database URL. In this case, if the database does not already exist, an exception is thrown when trying to connect. The connection only succeeds when the database already exists. The complete URL may look like this:

```
String url = "jdbc:h2:/data/sample;IFEXISTS=TRUE";
```

# Closing a Database

## Delayed Database Closing

Usually, a database is closed when the last connection to it is closed. In some situations this slows down the application, for example when it is not possible to keep at least one connection open. The automatic closing of a database can be delayed or disabled with the SQL statement `SET DB_CLOSE_DELAY <seconds>`. The parameter `<seconds>` specifies the number of seconds to keep a database open after the last connection to it was closed. The following statement will keep a database open for 10 seconds after the last connection was closed:

```
SET DB_CLOSE_DELAY 10
```

The value `-1` means the database is not closed automatically. The value `0` is the default and means the database is closed when the last connection is closed. This setting is persistent and can be set by an administrator only. It is possible to set the value in the database URL: `jdbc:h2:~/test;DB_CLOSE_DELAY=10`.

## Don't Close a Database when the VM Exits

By default, a database is closed when the last connection is closed. However, if it is never closed, the database is closed when the virtual machine exits normally, using a shutdown hook. In some situations, the database should not be closed in this case, for example because the database is still used at virtual machine shutdown (to store the shutdown process in the database for example). For those cases, the automatic closing of the database can be disabled in the database URL. The first connection (the one that is opening the database) needs to set the option in the database URL (it is not possible to change the setting afterwards). The database URL to disable database closing on exit is:

```
String url = "jdbc:h2:~/test;DB_CLOSE_ON_EXIT=FALSE";
```

## Execute SQL on Connection

Sometimes, particularly for in-memory databases, it is useful to be able to execute DDL or DML commands automatically when a client connects to a database. This functionality is enabled via the `INIT` property. Note that multiple commands may be passed to `INIT`, but the semicolon delimiter must be escaped, as in the example below.

```
String url = "jdbc:h2:mem:test;INIT=runscript from '~/create.sql'\\;runscript from '~/init.sql'";
```



Please note the double backslash is only required in a Java or properties file. In a GUI, or in an XML file, only one backslash is required:

```
<property name="url" value=
"jdbc:h2:mem:test;INIT=create schema if not exists test\;runscript from
'~/sql/init.sql'"
/>
```

Backslashes within the init script (for example within a runscript statement, to specify the folder names in Windows) need to be escaped as well (using a second backslash). It might be simpler to avoid backslashes in folder names for this reason; use forward slashes instead.

## Ignore Unknown Settings

Some applications (for example OpenOffice.org Base) pass some additional parameters when connecting to the database. Why those parameters are passed is unknown. The parameters `PREFERDOSLIKELINEENDS` and `IGNOREDRIVERPRIVILEGES` are such examples; they are simply ignored to improve the compatibility with OpenOffice.org. If an application passes other parameters when connecting to the database, usually the database throws an exception saying the parameter is not supported. It is possible to ignore such parameters by adding `;IGNORE_UNKNOWN_SETTINGS=TRUE` to the database URL.

## Changing Other Settings when Opening a Connection

In addition to the settings already described, other database settings can be passed in the database URL. Adding `;setting=value` at the end of a database URL is the same as executing the statement `SET setting value` just after connecting. For a list of supported settings, see [SQL Grammar](#) or the [DbSettings](#) javadoc.

## Custom File Access Mode

Usually, the database opens the database file with the access mode `rw`, meaning read-write (except for read only databases, where the mode `r` is used). To open a database in read-only mode if the database file is not read-only, use `ACCESS_MODE_DATA=r`. Also supported are `rws` and `rwd`. This setting must be specified in the database URL:

```
String url = "jdbc:h2:~/test;ACCESS_MODE_DATA=rws";
```



For more information see [Durability Problems](#). On many operating systems the access mode rws does not guarantee that the data is written to the disk.

## Multiple Connections

### Opening Multiple Databases at the Same Time

An application can open multiple databases at the same time, including multiple connections to the same database. The number of open database is only limited by the memory available.

### Multiple Connections to the Same Database: Client/Server

If you want to access the same database at the same time from different processes or computers, you need to use the client / server mode. In this case, one process acts as the server, and the other processes (that could reside on other computers as well) connect to the server via TCP/IP (or TLS over TCP/IP for improved security).

### Multithreading Support

This database is multithreading-safe. If an application is multi-threaded, it does not need to worry about synchronizing access to the database. An application should normally use one connection per thread. This database synchronizes access to the same connection, but other databases may not do this. To get higher concurrency, you need to use multiple connections.

An application can use multiple threads that access the same database at the same time. With default MVStore engine threads that use different connections can use the database concurrently. With PageStore engine requests to the same database are synchronized, that means that if one thread executes a long running query, the other threads need to wait. Concurrent database usage may be enabled for PageStore or disabled for MVStore with MULTI\_THREADED setting. Note that multi-threaded mode for PageStore engine is not tested well and has some issues; it should be used with caution.

### Locking, Lock-Timeout, Deadlocks

Please note MVCC is enabled in version 1.4.x by default, when using the MVStore. In this case, table level locking is not used. If [multi-version concurrency](#) is not used, the database uses table level locks to give each connection a consistent state of the data. There are two kinds of locks: read locks (shared locks) and

write locks (exclusive locks). All locks are released when the transaction commits or rolls back. When using the default transaction isolation level 'read committed', read locks are already released after each statement.

If a connection wants to reads from a table, and there is no write lock on the table, then a read lock is added to the table. If there is a write lock, then this connection waits for the other connection to release the lock. If a connection cannot get a lock for a specified time, then a lock timeout exception is thrown.

Usually, SELECT statements will generate read locks. This includes subqueries. Statements that modify data use write locks. It is also possible to lock a table exclusively without modifying data, using the statement SELECT ... FOR UPDATE. The statements COMMIT and ROLLBACK releases all open locks. The commands SAVEPOINT and ROLLBACK TO SAVEPOINT don't affect locks. The locks are also released when the autocommit mode changes, and for connections with autocommit set to true (this is the default), locks are released after each statement. The following statements generate locks:

Type of Lock	SQL Statement
Read	SELECT * FROM TEST; CALL SELECT MAX(ID) FROM TEST; SCRIPT;
Write	SELECT * FROM TEST WHERE 1=0 FOR UPDATE;
Write	INSERT INTO TEST VALUES(1, 'Hello'); INSERT INTO TEST SELECT * FROM TEST; UPDATE TEST SET NAME='Hi'; DELETE FROM TEST;
Write	ALTER TABLE TEST ...; CREATE INDEX ... ON TEST ...; DROP INDEX ...;

The number of seconds until a lock timeout exception is thrown can be set separately for each connection using the SQL command SET LOCK\_TIMEOUT <milliseconds>. The initial lock timeout (that is the timeout used for new connections) can be set using the SQL command SET DEFAULT\_LOCK\_TIMEOUT <milliseconds>. The default lock timeout is persistent.

## Avoiding Deadlocks

To avoid deadlocks, ensure that all transactions lock the tables in the same order (for example in alphabetical order), and avoid upgrading read locks to write locks. Both can be achieved using explicitly locking tables using SELECT ... FOR UPDATE.

Note that delete, insert and update operations issue table level locks with PageStore engine, but does not issue them with default MVStore engine.

## Database File Layout

The following files are created for persistent databases:

File Name	Description	Number of Files
test.h2.db	Database file. Contains the transaction log, indexes, and data for all tables. Format: <database>.h2.db	1 per database
test.lock.db	Database lock file. Automatically (re-)created while the database is in use. Format: <database>.lock.db	1 per database (only if in use)
test.trace.db	Trace file (if the trace option is enabled). Contains trace information. Format: <database>.trace.db Renamed to <database>.trace.db.old is too big.	0 or 1 per database
test.lob.db/*	Directory containing one file for each BLOB or CLOB value larger than a certain size. Format: <id>.t<tableId>.lob.db	1 per large object
test.123.temp.db	Temporary file. Contains a temporary blob or a large result set. Format: <database>.<id>.temp.db	1 per object

## Moving and Renaming Database Files

Database name and location are not stored inside the database files.

While a database is closed, the files can be moved to another directory, and they can be renamed as well (as long as all files of the same database start with the same name and the respective extensions are unchanged).

As there is no platform specific data in the files, they can be moved to other operating systems without problems.

## Backup

When the database is closed, it is possible to backup the database files.

To backup data while the database is running, the SQL commands SCRIPT and BACKUP can be used.

## Logging and Recovery

Whenever data is modified in the database and those changes are committed, the changes are written to the transaction log (except for in-memory objects). The changes to the main data area itself are usually written later on, to optimize disk access. If there is a power failure, the main data area is not up-to-date, but because the changes are in the transaction log, the next time the database is opened, the changes are re-applied automatically.

## Compatibility

All database engines behave a little bit different. Where possible, H2 supports the ANSI SQL standard, and tries to be compatible to other databases. There are still a few differences however:

In MySQL text columns are case insensitive by default, while in H2 they are case sensitive. However H2 supports case insensitive columns as well. To create the tables with case insensitive texts, append IGNORECASE=TRUE to the database URL (example: jdbc:h2:~/test;IGNORECASE=TRUE).

## Compatibility Modes

For certain features, this database can emulate the behavior of specific databases. However, only a small subset of the differences between databases are implemented in this way. Here is the list of currently supported modes and the differences to the regular mode:

### DB2 Compatibility Mode

To use the IBM DB2 mode, use the database URL jdbc:h2:~/test;MODE=DB2 or the SQL statement SET MODE DB2.

- For aliased columns, ResultSetMetaData.getColumnNames() returns the alias name and getTableName() returns null.
- Concatenating NULL with another value results in the other value.
- Support the pseudo-table SYSIBM.SYSDUMMY1.

- Timestamps with dash between date and time are supported.
- Datetime value functions return the same value within a command.

## Derby Compatibility Mode

To use the Apache Derby mode, use the database URL `jdbc:h2:~/test;MODE=Derby` or the SQL statement `SET MODE Derby`.

- For aliased columns, `ResultSetMetaData.getColumnName()` returns the alias name and `getTableName()` returns null.
- For unique indexes, NULL is distinct. That means only one row with NULL in one of the columns is allowed.
- Concatenating NULL with another value results in the other value.
- Support the pseudo-table `SYSIBM.SYSDUMMY1`.
- Datetime value functions return the same value within a command.

## HSQldb Compatibility Mode

To use the HSQldb mode, use the database URL `jdbc:h2:~/test;MODE=HSQldb` or the SQL statement `SET MODE HSQldb`.

- For aliased columns, `ResultSetMetaData.getColumnName()` returns the alias name and `getTableName()` returns null.
- When converting the scale of decimal data, the number is only converted if the new scale is smaller than the current scale. Usually, the scale is converted and 0s are added if required.
- For unique indexes, NULL is distinct. That means only one row with NULL in one of the columns is allowed.
- Text can be concatenated using `'+'`.
- Datetime value functions return the same value within a command.

## MS SQL Server Compatibility Mode

To use the MS SQL Server mode, use the database URL `jdbc:h2:~/test;MODE=MSSQLServer` or the SQL statement `SET MODE MSSQLServer`.

- For aliased columns, `ResultSetMetaData.getColumnName()` returns the alias name and `getTableName()` returns null.
- Identifiers may be quoted using square brackets as in `[Test]`.
- For unique indexes, NULL is distinct. That means only one row with NULL in one of the columns is allowed.
- Concatenating NULL with another value results in the other value.
- Text can be concatenated using `'+'`.

- MONEY data type is treated like NUMERIC(19, 4) data type. SMALLMONEY data type is treated like NUMERIC(10, 4) data type.
- IDENTITY can be used for automatic id generation on column level.
- Table hints are discarded. Example: SELECT \* FROM table WITH (NOLOCK).
- Datetime value functions return the same value within a command.

## MySQL Compatibility Mode

To use the MySQL mode, use the database URL jdbc:h2:~/test;MODE=MySQL or the SQL statement SET MODE MySQL. Use this mode for compatibility with MariaDB too.

- When inserting data, if a column is defined to be NOT NULL and NULL is inserted, then a 0 (or empty string, or the current timestamp for timestamp columns) value is used. Usually, this operation is not allowed and an exception is thrown.
- Creating indexes in the CREATE TABLE statement is allowed using INDEX(..) or KEY(..). Example: create table test(id int primary key, name varchar(255), key idx\_name(name));
- Meta data calls return identifiers in lower case.
- When converting a floating point number to an integer, the fractional digits are not truncated, but the value is rounded.
- Concatenating NULL with another value results in the other value.
- ON DUPLICATE KEY UPDATE is supported in INSERT statements, due to this feature VALUES has special non-standard meaning in some contexts.
- INSERT IGNORE is partially supported and may be used to skip rows with duplicate keys if ON DUPLICATE KEY UPDATE is not specified.
- REGEXP\_REPLACE() uses \ for back-references for compatibility with MariaDB.
- Datetime value functions return the same value within a command.

Text comparison in MySQL is case insensitive by default, while in H2 it is case sensitive (as in most other databases). H2 does support case insensitive text comparison, but it needs to be set separately, using SET IGNORECASE TRUE. This affects comparison using =, LIKE, REGEXP.

## Oracle Compatibility Mode

To use the Oracle mode, use the database URL jdbc:h2:~/test;MODE=Oracle or the SQL statement SET MODE Oracle.

- For aliased columns, ResultSetMetaData.getColumnNames() returns the alias name and getTableName() returns null.

- When using unique indexes, multiple rows with NULL in all columns are allowed, however it is not allowed to have multiple rows with the same values otherwise.
- Concatenating NULL with another value results in the other value.
- Empty strings are treated like NULL values.
- REGEXP\_REPLACE() uses \ for back-references.
- DATE data type is treated like TIMESTAMP(0) data type.
- Datetime value functions return the same value within a command.

## PostgreSQL Compatibility Mode

To use the PostgreSQL mode, use the database URL `jdbc:h2:~/test;MODE=PostgreSQL` or the SQL statement `SET MODE PostgreSQL`.

- For aliased columns, `ResultSetMetaData.getColumnNames()` returns the alias name and `getTableName()` returns null.
- When converting a floating point number to an integer, the fractional digits are not be truncated, but the value is rounded.
- The system columns CTID and OID are supported.
- LOG(x) is base 10 in this mode.
- REGEXP\_REPLACE():
  - uses \ for back-references;
  - does not throw an exception when the flagsString parameter contains a 'g';
  - replaces only the first matched substring in the absence of the 'g' flag in the flagsString parameter.
- Fixed-width strings are padded with spaces.
- MONEY data type is treated like NUMERIC(19, 2) data type.
- Datetime value functions return the same value within a transaction.
- ARRAY\_SLICE() out of bounds parameters are silently corrected.
- EXTRACT function with DOW field returns (0-6), Sunday is 0.

## Ignite Compatibility Mode

To use the Ignite mode, use the database URL `jdbc:h2:~/test;MODE=Ignite` or the SQL statement `SET MODE Ignite`.

- Creating indexes in the CREATE TABLE statement is allowed using INDEX(..) or KEY(..). Example: `create table test(id int primary key, name varchar(255), key idx_name(name));`
- AFFINITY KEY and SHARD KEY keywords may be used in index definition.
- Datetime value functions return the same value within a transaction.

## Auto-Reconnect

The auto-reconnect feature causes the JDBC driver to reconnect to the database if the connection is lost. The automatic re-connect only occurs when auto-commit is enabled; if auto-commit is disabled, an exception is thrown. To enable this mode, append `;AUTO_RECONNECT=TRUE` to the database URL.

Re-connecting will open a new session. After an automatic re-connect, variables and local temporary tables definitions (excluding data) are re-created. The contents of the system table `INFORMATION_SCHEMA.SESSION_STATE` contains all client side state that is re-created.

If another connection uses the database in exclusive mode (enabled using `SET EXCLUSIVE 1` or `SET EXCLUSIVE 2`), then this connection will try to re-connect until the exclusive mode ends.

## Automatic Mixed Mode

Multiple processes can access the same database without having to start the server manually. To do that, append `;AUTO_SERVER=TRUE` to the database URL. You can use the same database URL independent of whether the database is already open or not. This feature doesn't work with in-memory databases. Example database URL:

```
jdbc:h2:/data/test;AUTO_SERVER=TRUE
```

Use the same URL for all connections to this database. Internally, when using this mode, the first connection to the database is made in embedded mode, and additionally a server is started internally (as a daemon thread). If the database is already open in another process, the server mode is used automatically. The IP address and port of the server are stored in the file `.lock.db`, that's why in-memory databases can't be supported.

The application that opens the first connection to the database uses the embedded mode, which is faster than the server mode. Therefore the main application should open the database first if possible. The first connection automatically starts a server on a random port. This server allows remote connections, however only to this database (to ensure that, the client reads `.lock.db` file and sends the random key that is stored there to the server). When the first connection is closed, the server stops. If other (remote) connections are still open, one of them will then start a server (auto-reconnect is enabled automatically).

All processes need to have access to the database files. If the first connection is closed (the connection that started the server), open transactions of other



connections will be rolled back (this may not be a problem if you don't disable autocommit). Explicit client/server connections (using jdbc:h2:tcp:// or ssl://) are not supported. This mode is not supported for in-memory databases.

Here is an example how to use this mode. Application 1 and 2 are not necessarily started on the same computer, but they need to have access to the database files. Application 1 and 2 are typically two different processes (however they could run within the same process).

```
// Application 1:  
DriverManager.getConnection("jdbc:h2:/data/test;AUTO_SERVER=TRUE");  
  
// Application 2:  
DriverManager.getConnection("jdbc:h2:/data/test;AUTO_SERVER=TRUE");
```

When using this feature, by default the server uses any free TCP port. The port can be set manually using AUTO\_SERVER\_PORT=9090.

## Page Size

The page size for new databases is 2 KB (2048), unless the page size is set explicitly in the database URL using PAGE\_SIZE= when the database is created. The page size of existing databases can not be changed, so this property needs to be set when the database is created.

## Using the Trace Options

To find problems in an application, it is sometimes good to see what database operations were executed. This database offers the following trace features:

- Trace to System.out and/or to a file
- Support for trace levels OFF, ERROR, INFO, DEBUG
- The maximum size of the trace file can be set
- It is possible to generate Java source code from the trace file
- Trace can be enabled at runtime by manually creating a file

## Trace Options

The simplest way to enable the trace option is setting it in the database URL. There are two settings, one for System.out (TRACE\_LEVEL\_SYSTEM\_OUT) tracing, and one for file tracing (TRACE\_LEVEL\_FILE). The trace levels are 0 for OFF, 1 for ERROR (the default), 2 for INFO, and 3 for DEBUG. A database URL with both levels set to DEBUG is:

```
jdbc:h2:~/test;TRACE_LEVEL_FILE=3;TRACE_LEVEL_SYSTEM_OUT=3
```

The trace level can be changed at runtime by executing the SQL command SET TRACE\_LEVEL\_SYSTEM\_OUT level (for System.out tracing) or SET TRACE\_LEVEL\_FILE level (for file tracing). Example:

```
SET TRACE_LEVEL_SYSTEM_OUT 3
```

## Setting the Maximum Size of the Trace File

When using a high trace level, the trace file can get very big quickly. The default size limit is 16 MB, if the trace file exceeds this limit, it is renamed to .old and a new file is created. If another such file exists, it is deleted. To limit the size to a certain number of megabytes, use SET TRACE\_MAX\_FILE\_SIZE mb. Example:

```
SET TRACE_MAX_FILE_SIZE 1
```

## Java Code Generation

When setting the trace level to INFO or DEBUG, Java source code is generated as well. This simplifies reproducing problems. The trace file looks like this:

```
...
12-20 20:58:09 jdbc[0]:
/**/dbMeta3.getURL();
12-20 20:58:09 jdbc[0]:
/**/dbMeta3.getTables(null, "", null, new String[]{"TABLE", "VIEW"});
...
```

To filter the Java source code, use the ConvertTraceFile tool as follows:

```
java -cp h2*.jar org.h2.tools.ConvertTraceFile
    -traceFile "~/test.trace.db" -javaClass "Test"
```

The generated file Test.java will contain the Java source code. The generated source code may be too large to compile (the size of a Java method is limited). If this is the case, the source code needs to be split in multiple methods. The password is not listed in the trace file and therefore not included in the source code.

## Using Other Logging APIs

By default, this database uses its own native 'trace' facility. This facility is called 'trace' and not 'log' within this database to avoid confusion with the transaction log. Trace messages can be written to both file and System.out. In most cases, this is sufficient, however sometimes it is better to use the same facility as the application, for example Log4j. To do that, this database support SLF4J.

[SLF4J](#) is a simple facade for various logging APIs and allows to plug in the desired implementation at deployment time. SLF4J supports implementations such as Logback, Log4j, Jakarta Commons Logging (JCL), Java logging, x4juli, and Simple Log.

To enable SLF4J, set the file trace level to 4 in the database URL:

```
jdbc:h2:~/test;TRACE_LEVEL_FILE=4
```

Changing the log mechanism is not possible after the database is open, that means executing the SQL statement SET TRACE\_LEVEL\_FILE 4 when the database is already open will not have the desired effect. To use SLF4J, all required jar files need to be in the classpath. The logger name is h2database. If it does not work, check the file <database>.trace.db for error messages.

## Read Only Databases

If the database files are read-only, then the database is read-only as well. It is not possible to create new tables, add or modify data in this database. Only SELECT and CALL statements are allowed. To create a read-only database, close the database. Then, make the database file read-only. When you open the database now, it is read-only. There are two ways an application can find out whether database is read-only: by calling Connection.isReadOnly() or by executing the SQL statement CALL READONLY().

Using the [Custom Access Mode](#) the database can also be opened in read-only mode, even if the database file is not read only.

## Read Only Databases in Zip or Jar File

To create a read-only database in a zip file, first create a regular persistent database, and then create a backup. The database must not have pending changes, that means you need to close all connections to the database first. To speed up opening the read-only database and running queries, the database should be closed using SHUTDOWN DEFRAG. If you are using a database named test, an easy way to create a zip file is using the Backup tool. You can start the

tool from the command line, or from within the H2 Console (Tools - Backup). Please note that the database must be closed when the backup is created. Therefore, the SQL statement BACKUP TO can not be used.

When the zip file is created, you can open the database in the zip file using the following database URL:

```
jdbc:h2:zip:~/data.zip!/test
```

Databases in zip files are read-only. The performance for some queries will be slower than when using a regular database, because random access in zip files is not supported (only streaming). How much this affects the performance depends on the queries and the data. The database is not read in memory; therefore large databases are supported as well. The same indexes are used as when using a regular database.

If the database is larger than a few megabytes, performance is much better if the database file is split into multiple smaller files, because random access in compressed files is not possible. See also the sample application [ReadOnlyDatabaseInZip](#).

## Opening a Corrupted Database

If a database cannot be opened because the boot info (the SQL script that is run at startup) is corrupted, then the database can be opened by specifying a database event listener. The exceptions are logged, but opening the database will continue.

## Computed Columns / Function Based Index

A computed column is a column whose value is calculated before storing. The formula is evaluated when the row is inserted, and re-evaluated every time the row is updated. One use case is to automatically update the last-modification time:

```
CREATE TABLE TEST(ID INT, NAME VARCHAR, LAST_MOD TIMESTAMP AS  
NOW());
```

Function indexes are not directly supported by this database, but they can be emulated by using computed columns. For example, if an index on the upper-case version of a column is required, create a computed column with the upper-case version of the original column, and create an index for this column:

```
CREATE TABLE ADDRESS(
```

```
ID INT PRIMARY KEY,  
NAME VARCHAR,  
UPPER_NAME VARCHAR AS UPPER(NAME)  
);  
CREATE INDEX IDX_U_NAME ON ADDRESS(UPPER_NAME);
```

When inserting data, it is not required (and not allowed) to specify a value for the upper-case version of the column, because the value is generated. But you can use the column when querying the table:

```
INSERT INTO ADDRESS(ID, NAME) VALUES(1, 'Miller');  
SELECT * FROM ADDRESS WHERE UPPER_NAME='MILLER';
```

## Multi-Dimensional Indexes

A tool is provided to execute efficient multi-dimension (spatial) range queries. This database does not support a specialized spatial index (R-Tree or similar). Instead, the B-Tree index is used. For each record, the multi-dimensional key is converted (mapped) to a single dimensional (scalar) value. This value specifies the location on a space-filling curve.

Currently, Z-order (also called N-order or Morton-order) is used; Hilbert curve could also be used, but the implementation is more complex. The algorithm to convert the multi-dimensional value is called bit-interleaving. The scalar value is indexed using a B-Tree index (usually using a computed column).

The method can result in a drastic performance improvement over just using an index on the first column. Depending on the data and number of dimensions, the improvement is usually higher than factor 5. The tool generates a SQL query from a specified multi-dimensional range. The method used is not database dependent, and the tool can easily be ported to other databases. For an example how to use the tool, please have a look at the sample code provided in `TestMultiDimension.java`.

## User-Defined Functions and Stored Procedures

In addition to the built-in functions, this database supports user-defined Java functions. In this database, Java functions can be used as stored procedures as well. A function must be declared (registered) before it can be used. A function can be defined using source code, or as a reference to a compiled class that is available in the classpath. By default, the function aliases are stored in the current schema.

## Referencing a Compiled Method

When referencing a method, the class must already be compiled and included in the classpath where the database is running. Only static Java methods are supported; both the class and the method must be public. Example Java class:

```
package acme;
import java.math.*;
public class Function {
    public static boolean isPrime(int value) {
        return new BigInteger(String.valueOf(value)).isProbablePrime(100);
    }
}
```

The Java function must be registered in the database by calling CREATE ALIAS ... FOR:

```
CREATE ALIAS IS_PRIME FOR "acme.Function.isPrime";
```

For a complete sample application, see `src/test/org/h2/samples/Function.java`.

## Declaring Functions as Source Code

When defining a function alias with source code, the database tries to compile the source code using the Sun Java compiler (the class `com.sun.tools.javac.Main`) if the `tools.jar` is in the classpath. If not, `javac` is run as a separate process. Only the source code is stored in the database; the class is compiled each time the database is re-opened. Source code is usually passed as dollar quoted text to avoid escaping problems, however single quotes can be used as well. Example:

```
CREATE ALIAS NEXT_PRIME AS $$
String nextPrime(String value) {
    return new BigInteger(value).nextProbablePrime().toString();
}
$$;
```

By default, the three packages `java.util`, `java.math`, `java.sql` are imported. The method name (`nextPrime` in the example above) is ignored. Method overloading is not supported when declaring functions as source code, that means only one method may be declared for an alias. If different import statements are required, they must be declared at the beginning and separated with the tag `@CODE`:

```
CREATE ALIAS IP_ADDRESS AS $$
import java.net.*;
```

```
@CODE
String ipAddress(String host) throws Exception {
    return InetAddress.getByName(host).getHostAddress();
}
$$;
```

The following template is used to create a complete Java class:

```
package org.h2.dynamic;
< import statements before the tag @CODE; if not set:
import java.util.*;
import java.math.*;
import java.sql.*;
>
public class <aliasName> {
    public static <sourceCode>
}
```

## Method Overloading

Multiple methods may be bound to a SQL function if the class is already compiled and included in the classpath. Each Java method must have a different number of arguments. Method overloading is not supported when declaring functions as source code.

## Function Data Type Mapping

Functions that accept non-nullable parameters such as `int` will not be called if one of those parameters is `NULL`. Instead, the result of the function is `NULL`. If the function should be called if a parameter is `NULL`, you need to use `java.lang.Integer` instead.

SQL types are mapped to Java classes and vice-versa as in the JDBC API. For details, see [Data Types](#). There are a few special cases: `java.lang.Object` is mapped to `OTHER` (a serialized object). Therefore, `java.lang.Object` can not be used to match all SQL types (matching all SQL types is not supported). The second special case is `Object[]`: arrays of any class are mapped to `ARRAY`. Objects of type `org.h2.value.Value` (the internal value class) are passed through without conversion.

## Functions That Require a Connection

If the first parameter of a Java function is a `java.sql.Connection`, then the connection to database is provided. This connection does not need to be closed before returning. When calling the method from within the SQL statement, this connection parameter does not need to be (can not be) specified.

## Functions Throwing an Exception

If a function throws an exception, then the current statement is rolled back and the exception is thrown to the application. `SQLException` are directly re-thrown to the calling application; all other exceptions are first converted to a `SQLException`.

## Functions Returning a Result Set

Functions may returns a result set. Such a function can be called with the `CALL` statement:

```
public static ResultSet query(Connection conn, String sql) throws SQLException {  
    return conn.createStatement().executeQuery(sql);  
}
```

```
CREATE ALIAS QUERY FOR "org.h2.samples.Function.query";  
CALL QUERY('SELECT * FROM TEST');
```

## Using SimpleResultSet

A function can create a result set using the `SimpleResultSet` tool:

```
import org.h2.tools.SimpleResultSet;  
...  
public static ResultSet simpleResultSet() throws SQLException {  
    SimpleResultSet rs = new SimpleResultSet();  
    rs.addColumn("ID", Types.INTEGER, 10, 0);  
    rs.addColumn("NAME", Types.VARCHAR, 255, 0);  
    rs.addRow(0, "Hello");  
    rs.addRow(1, "World");  
    return rs;  
}
```

```
CREATE ALIAS SIMPLE FOR "org.h2.samples.Function.simpleResultSet";  
CALL SIMPLE();
```



## Using a Function as a Table

A function that returns a result set can be used like a table. However, in this case the function is called at least twice: first while parsing the statement to collect the column names (with parameters set to null where not known at compile time). And then, while executing the statement to get the data (maybe multiple times if this is a join). If the function is called just to get the column list, the URL of the connection passed to the function is `jdbc:columnlist:connection`. Otherwise, the URL of the connection is `jdbc:default:connection`.

```
public static ResultSet getMatrix(Connection conn, Integer size)
    throws SQLException {
    SimpleResultSet rs = new SimpleResultSet();
    rs.addColumn("X", Types.INTEGER, 10, 0);
    rs.addColumn("Y", Types.INTEGER, 10, 0);
    String url = conn.getMetaData().getURL();
    if (url.equals("jdbc:columnlist:connection")) {
        return rs;
    }
    for (int s = size.intValue(), x = 0; x < s; x++) {
        for (int y = 0; y < s; y++) {
            rs.addRow(x, y);
        }
    }
    return rs;
}
```

```
CREATE ALIAS MATRIX FOR "org.h2.samples.Function.getMatrix";
SELECT * FROM MATRIX(4) ORDER BY X, Y;
```

## Pluggable or User-Defined Tables

For situations where you need to expose other data-sources to the SQL engine as a table, there are "pluggable tables". For some examples, have a look at the code in `org.h2.test.db.TestTableEngines`.

In order to create your own `TableEngine`, you need to implement the `org.h2.api.TableEngine` interface e.g. something like this:

```
package acme;
public static class MyTableEngine implements org.h2.api.TableEngine {

    private static class MyTable extends org.h2.table.TableBase {
        .. rather a lot of code here...
    }
}
```

```

    }

    public EndlessTable createTable(CreateTableData data) {
        return new EndlessTable(data);
    }
}

```

and then create the table from SQL like this:

```

CREATE TABLE TEST(ID INT, NAME VARCHAR)
ENGINE "acme.MyTableEngine";

```

It is also possible to pass in parameters to the table engine, like so:

```

CREATE TABLE TEST(ID INT, NAME VARCHAR) ENGINE "acme.MyTableEngine"
WITH "param1", "param2";

```

In which case the parameters are passed down in the `tableEngineParams` field of the `CreateTableData` object.

It is also possible to specify default table engine params on schema creation:

```

CREATE SCHEMA TEST_SCHEMA WITH "param1", "param2";

```

Params from the schema are used when `CREATE TABLE` issued on this schema does not have its own engine params specified.

## Triggers

This database supports Java triggers that are called before or after a row is updated, inserted or deleted. Triggers can be used for complex consistency checks, or to update related data in the database. It is also possible to use triggers to simulate materialized views. For a complete sample application, see `src/test/org/h2/samples/TriggerSample.java`. A Java trigger must implement the interface `org.h2.api.Trigger`. The trigger class must be available in the classpath of the database engine (when using the server mode, it must be in the classpath of the server).

```

import org.h2.api.Trigger;
...
public class TriggerSample implements Trigger {

    public void init(Connection conn, String schemaName, String triggerName,
        String tableName, boolean before, int type) {

```

```

    // initialize the trigger object is necessary
}

public void fire(Connection conn,
    Object[] oldRow, Object[] newRow)
    throws SQLException {
    // the trigger is fired
}

public void close() {
    // the database is closed
}

public void remove() {
    // the trigger was dropped
}
}

```

The connection can be used to query or update data in other tables. The trigger then needs to be defined in the database:

```

CREATE TRIGGER INV_INS AFTER INSERT ON INVOICE
FOR EACH ROW CALL "org.h2.samples.TriggerSample"

```

The trigger can be used to veto a change by throwing a SQLException.

As an alternative to implementing the Trigger interface, an application can extend the abstract class `org.h2.tools.TriggerAdapter`. This will allow to use the `ResultSet` interface within trigger implementations. In this case, only the `fire` method needs to be implemented:

```

import org.h2.tools.TriggerAdapter;
...
public class TriggerSample extends TriggerAdapter {

    public void fire(Connection conn, ResultSet oldRow, ResultSet newRow)
        throws SQLException {
        // the trigger is fired
    }

}

```

## Compacting a Database

Empty space in the database file re-used automatically. When closing the database, the database is automatically compacted for up to 200 milliseconds by default. To compact more, use the SQL statement SHUTDOWN COMPACT. However re-creating the database may further reduce the database size because this will re-build the indexes. Here is a sample function to do this:

```
public static void compact(String dir, String dbName,
    String user, String password) throws Exception {
    String url = "jdbc:h2:" + dir + "/" + dbName;
    String file = "data/test.sql";
    Script.execute(url, user, password, file);
    DeleteDbFiles.execute(dir, dbName, true);
    RunScript.execute(url, user, password, file, null, false);
}
```

See also the sample application `org.h2.samples.Compact`. The commands SCRIPT / RUNSCRIPT can be used as well to create a backup of a database and re-build the database from the script.

## Cache Settings

The database keeps most frequently used data in the main memory. The amount of memory used for caching can be changed using the setting `CACHE_SIZE`. This setting can be set in the database connection URL (`jdbc:h2:~/test;CACHE_SIZE=131072`), or it can be changed at runtime using `SET CACHE_SIZE` size. The size of the cache, as represented by `CACHE_SIZE` is measured in KB, with each KB being 1024 bytes. This setting has no effect for in-memory databases. For persistent databases, the setting is stored in the database and re-used when the database is opened the next time. However, when opening an existing database, the cache size is set to at most half the amount of memory available for the virtual machine (`Runtime.getRuntime().maxMemory()`), even if the cache size setting stored in the database is larger; however the setting stored in the database is kept. Setting the cache size in the database URL or explicitly using `SET CACHE_SIZE` overrides this value (even if larger than the physical memory). To get the current used maximum cache size, use the query `SELECT * FROM INFORMATION_SCHEMA.SETTINGS WHERE NAME = 'info.CACHE_MAX_SIZE'`

An experimental scan-resistant cache algorithm "Two Queue" (2Q) is available. To enable it, append `;CACHE_TYPE=TQ` to the database URL. The cache might not actually improve performance. If you plan to use it, please run your own test cases first.

Also included is an experimental second level soft reference cache. Rows in this cache are only garbage collected on low memory. By default the second level cache is disabled. To enable it, use the prefix `SOFT_`. Example: `jdbc:h2:~/test;CACHE_TYPE=SOFT_LRU`. The cache might not actually improve performance. If you plan to use it, please run your own test cases first.

To get information about page reads and writes, and the current caching algorithm in use, call `SELECT * FROM INFORMATION_SCHEMA.SETTINGS`. The number of pages read / written is listed.

## External authentication (Experimental)

External authentication allows to optionally validate user credentials externally (JAAS,LDAP,custom classes). Is also possible to temporary assign roles to externally authenticated users. **This feature is experimental and subject to change**

Master user cannot be externally authenticated

To enable external authentication on a database execute statement `SET AUTHENTICATOR TRUE`. This setting is persisted on the database.

To connect on a database by using external credentials client must append `AUTHREALM=H2` to the database URL. H2 is the identifier of the authentication realm (see later).

External authentication requires to send password to the server. For this reason it works only on local connection or remote over ssl

By default external authentication is performed through JAAS login interface (configuration name is h2). To configure JAAS add argument `-Djava.security.auth.login.config=jaas.conf` Here an example of [JAAS login configuration file](#) content:

```
h2 {
  com.sun.security.auth.module.LdapLoginModule REQUIRED \
    userProvider="ldap://127.0.0.1:10389"
  authIdentity="uid={USERNAME},ou=people,dc=example,dc=com" \
    debug=true useSSL=false ;
};
```

Is it possible to specify custom authentication settings by using JVM argument `-Dh2auth.configurationFile={urlOfH2Auth.xml}`. Here an example of `h2auth.xml` file content:

```

<h2Auth allowUserRegistration="false" createMissingRoles="true">

  <!-- realm: DUMMY authenticate users named DUMMY[0-9] with a static
password -->
  <realm name="DUMMY"
validatorClass="org.h2.security.auth.impl.FixedPasswordCredentialsValidator">
    <property name="userNamePattern" value="DUMMY[0-9]" />
    <property name="password" value="mock" />
  </realm>

  <!-- realm LDAPEXAMPLE:perform credentials validation on LDAP -->
  <realm name="LDAPEXAMPLE"
validatorClass="org.h2.security.auth.impl.LdapCredentialsValidator">
    <property name="bindDnPattern" value="uid=
%u,ou=people,dc=example,dc=com" />
    <property name="host" value="127.0.0.1" />
    <property name="port" value="10389" />
    <property name="secure" value="false" />
  </realm>

  <!-- realm JAAS: perform credentials validation by using JAAS api -->
  <realm name="JAAS"
validatorClass="org.h2.security.auth.impl.JaasCredentialsValidator">
    <property name="appName" value="H2" />
  </realm>

  <!--Assign to each user role @{REALM} -->
  <userToRolesMapper
class="org.h2.security.auth.impl.AssignRealmNameRole"/>

  <!--Assign to each user role REMOTEUSER -->
  <userToRolesMapper class="org.h2.security.auth.impl.StaticRolesMapper">
    <property name="roles" value="REMOTEUSER"/>
  </userToRolesMapper>
</h2Auth>

```

Custom credentials validators must implement the interface  
org.h2.api.CredentialsValidator

Custom criteria for role assignments must implement the interface  
org.h2.api.UserToRoleMapper

# Performance

[Performance Comparison](#)  
[PolePosition Benchmark](#)  
[Database Performance Tuning](#)  
[Using the Built-In Profiler](#)  
[Application Profiling](#)  
[Database Profiling](#)  
[Statement Execution Plans](#)  
[How Data is Stored and How Indexes Work](#)  
[Fast Database Import](#)

## Performance Comparison

In many cases H2 is faster than other (open source and not open source) database engines. Please note this is mostly a single connection benchmark run on one computer, with many very simple operations running against the database. This benchmark does not include very complex queries. The embedded mode of H2 is faster than the client-server mode because the per-statement overhead is greatly reduced.

### Embedded

Test Case	Unit	H2	HSQLDB	Derby
Simple: Init	ms	1019	1907	8280
Simple: Query (random)	ms	1304	873	1912
Simple: Query (sequential)	ms	835	1839	5415
Simple: Update (sequential)	ms	961	2333	21759
Simple: Delete (sequential)	ms	950	1922	32016
Simple: Memory Usage	MB	21	10	8
BenchA: Init	ms	919	2133	7528
BenchA: Transactions	ms	1219	2297	8541
BenchA: Memory Usage	MB	12	15	7
BenchB: Init	ms	905	1993	8049
BenchB: Transactions	ms	1091	583	1165
BenchB: Memory Usage	MB	17	11	8

BenchC: Init	ms	2491	4003	8064
BenchC: Transactions	ms	1979	803	2840
BenchC: Memory Usage	MB	19	22	9
Executed statements	#	1930995	1930995	1930995
Total time	ms	13673	20686	105569
Statements per second	#	141226	93347	18291

## Client-Server

Test Case	Unit	H2 (Server)	HSQLDB	Derby	PostgreSQL	MySQL
Simple: Init	ms	16338	17198	27860	30156	29409
Simple: Query (random)	ms	3399	2582	6190	3315	3342
Simple: Query (sequential)	ms	21841	18699	42347	30774	32611
Simple: Update (sequential)	ms	6913	7745	28576	32698	11350
Simple: Delete (sequential)	ms	8051	9751	42202	44480	16555
Simple: Memory Usage	MB	22	11	9	0	1
BenchA: Init	ms	12996	14720	24722	26375	26060
BenchA: Transactions	ms	10134	10250	18452	21453	15877
BenchA: Memory Usage	MB	13	15	9	0	1
BenchB: Init	ms	15264	16889	28546	31610	29747
BenchB: Transactions	ms	3017	3376	1842	2771	1433
BenchB: Memory Usage	MB	17	12	11	1	1
BenchC: Init	ms	14020	10407	17655	19520	17532
BenchC: Transactions	ms	5076	3160	6411	6063	4530



BenchC: Memory Usage	MB	19	21	11	1	1
Executed statements	#	1930995	1930995	1930995	1930995	1930995
Total time	ms	117049	114777	244803	249215	188446
Statements per second	#	16497	16823	7887	7748	10246

## Benchmark Results and Comments

### H2

Version 1.4.177 (2014-04-12) was used for the test. For most operations, the performance of H2 is about the same as for HSQLDB. One situation where H2 is slow is large result sets, because they are buffered to disk if more than a certain number of records are returned. The advantage of buffering is: there is no limit on the result set size.

### HSQLDB

Version 2.3.2 was used for the test. Cached tables are used in this test (hsqldb.default\_table\_type=cached), and the write delay is 1 second (SET WRITE\_DELAY 1).

### Derby

Version 10.10.1.1 was used for the test. Derby is clearly the slowest embedded database in this test. This seems to be a structural problem, because all operations are really slow. It will be hard for the developers of Derby to improve the performance to a reasonable level. A few problems have been identified: leaving autocommit on is a problem for Derby. If it is switched off during the whole test, the results are about 20% better for Derby. Derby calls `FileChannel.force(false)`, but only twice per log file (not on each commit). Disabling this call improves performance for Derby by about 2%. Unlike H2, Derby does not call `FileDescriptor.sync()` on each checkpoint. Derby supports a testing mode (system property `derby.system.durability=test`) where durability is disabled. According to the documentation, this setting should be used for testing only, as the database may not recover after a crash. Enabling this setting improves performance by a factor of 2.6 (embedded mode) or 1.4 (server mode). Even if enabled, Derby is still less than half as fast as H2 in default mode.

## PostgreSQL

Version 9.1.5 was used for the test. The following options were changed in postgresql.conf: fsync = off, commit\_delay = 1000. PostgreSQL is run in server mode. The memory usage number is incorrect, because only the memory usage of the JDBC driver is measured.

## MySQL

Version 5.1.65-log was used for the test. MySQL was run with the InnoDB backend. The setting innodb\_flush\_log\_at\_trx\_commit (found in the my.ini / my.cnf file) was set to 0. Otherwise (and by default), MySQL is slow (around 140 statements per second in this test) because it tries to flush the data to disk for each commit. For small transactions (when autocommit is on) this is really slow. But many use cases use small or relatively small transactions. Too bad this setting is not listed in the configuration wizard, and it always overwritten when using the wizard. You need to change this setting manually in the file my.ini / my.cnf, and then restart the service. The memory usage number is incorrect, because only the memory usage of the JDBC driver is measured.

## Firebird

Firebird 1.5 (default installation) was tested, but the results are not published currently. It is possible to run the performance test with the Firebird database, and any information on how to configure Firebird for higher performance are welcome.

## Why Oracle / MS SQL Server / DB2 are Not Listed

The license of these databases does not allow to publish benchmark results. This doesn't mean that they are fast. They are in fact quite slow, and need a lot of memory. But you will need to test this yourself. SQLite was not tested because the JDBC driver doesn't support transactions.

## About this Benchmark

### How to Run

This test was as follows:

```
build benchmark
```

## **Separate Process per Database**

For each database, a new process is started, to ensure the previous test does not impact the current test.

## **Number of Connections**

This is mostly a single-connection benchmark. BenchB uses multiple connections; the other tests use one connection.

## **Real-World Tests**

Good benchmarks emulate real-world use cases. This benchmark includes 4 test cases: BenchSimple uses one table and many small updates / deletes. BenchA is similar to the TPC-A test, but single connection / single threaded (see also: [www.tpc.org](http://www.tpc.org)). BenchB is similar to the TPC-B test, using multiple connections (one thread per connection). BenchC is similar to the TPC-C test, but single connection / single threaded.

## **Comparing Embedded with Server Databases**

This is mainly a benchmark for embedded databases (where the application runs in the same virtual machine as the database engine). However MySQL and PostgreSQL are not Java databases and cannot be embedded into a Java application. For the Java databases, both embedded and server modes are tested.

## **Test Platform**

This test is run on Mac OS X 10.6. No virus scanner was used, and disk indexing was disabled. The JVM used is Sun JDK 1.6.

## **Multiple Runs**

When a Java benchmark is run first, the code is not fully compiled and therefore runs slower than when running multiple times. A benchmark should always run the same test multiple times and ignore the first run(s). This benchmark runs three times, but only the last run is measured.

## **Memory Usage**

It is not enough to measure the time taken, the memory usage is important as well. Performance can be improved by using a bigger cache, but the amount of memory is limited. HSQLDB tables are kept fully in memory by default; this

benchmark uses 'disk based' tables for all databases. Unfortunately, it is not so easy to calculate the memory usage of PostgreSQL and MySQL, because they run in a different process than the test. This benchmark currently does not print memory usage of those databases.

## **Delayed Operations**

Some databases delay some operations (for example flushing the buffers) until after the benchmark is run. This benchmark waits between each database tested, and each database runs in a different process (sequentially).

## **Transaction Commit / Durability**

Durability means transaction committed to the database will not be lost. Some databases (for example MySQL) try to enforce this by default by calling `fsync()` to flush the buffers, but most hard drives don't actually flush all data. Calling the method slows down transaction commit a lot, but doesn't always make data durable. When comparing the results, it is important to think about the effect. Many database suggest to 'batch' operations when possible. This benchmark switches off autocommit when loading the data, and calls commit after each 1000 inserts. However many applications need 'short' transactions at runtime (a commit after each update). This benchmark commits after each update / delete in the simple benchmark, and after each business transaction in the other benchmarks. For databases that support delayed commits, a delay of one second is used.

## **Using Prepared Statements**

Wherever possible, the test cases use prepared statements.

## **Currently Not Tested: Startup Time**

The startup time of a database engine is important as well for embedded use. This time is not measured currently. Also, not tested is the time used to create a database and open an existing database. Here, one (wrapper) connection is opened at the start, and for each step a new connection is opened and then closed.

## **PolePosition Benchmark**

The PolePosition is an open source benchmark. The algorithms are all quite simple. It was developed / sponsored by db4o. This test was not run for a longer time, so please be aware that the results below are for older database versions (H2 version 1.1, HSQLDB 1.8, Java 1.4).

<b>Test Case</b>	<b>Unit</b>	<b>H2</b>	<b>HSQLDB</b>	<b>MySQL</b>
Melbourne write	ms	369	249	2022
Melbourne read	ms	47	49	93
Melbourne read_hot	ms	24	43	95
Melbourne delete	ms	147	133	176
Sepang write	ms	965	1201	3213
Sepang read	ms	765	948	3455
Sepang read_hot	ms	789	859	3563
Sepang delete	ms	1384	1596	6214
Bahrain write	ms	1186	1387	6904
Bahrain query_indexed_string	ms	336	170	693
Bahrain query_string	ms	18064	39703	41243
Bahrain query_indexed_int	ms	104	134	678
Bahrain update	ms	191	87	159
Bahrain delete	ms	1215	729	6812
Imola retrieve	ms	198	194	4036
Barcelona write	ms	413	832	3191
Barcelona read	ms	119	160	1177
Barcelona query	ms	20	5169	101
Barcelona delete	ms	388	319	3287
Total	ms	26724	53962	87112

There are a few problems with the PolePosition test:

- HSQLDB uses in-memory tables by default while H2 uses persistent tables. The HSQLDB version included in PolePosition does not support changing this, so you need to replace poleposition-0.20/lib/hsqldb.jar with a newer version (for example hsqldb-1.8.0.7.jar), and then use the setting `hsqldb.connecturl=jdbc:hsqldb:file:data/hsqldb/dbbench2;hsqldb.default_table_type=cached;sql.enforce_size=true` in the file `Jdbc.properties`.
- HSQLDB keeps the database open between tests, while H2 closes the database (losing all the cache). To change that, use the database URL `jdbc:h2:file:data/h2/dbbench;DB_CLOSE_DELAY=-1`

- The amount of cache memory is quite important, specially for the PolePosition test. Unfortunately, the PolePosition test does not take this into account.

## Database Performance Tuning

### Keep Connections Open or Use a Connection Pool

If your application opens and closes connections a lot (for example, for each request), you should consider using a [connection pool](#). Opening a connection using `DriverManager.getConnection` is specially slow if the database is closed. By default the database is closed if the last connection is closed.

If you open and close connections a lot but don't want to use a connection pool, consider keeping a 'sentinel' connection open for as long as the application runs, or use delayed database closing. See also [Closing a database](#).

### Use a Modern JVM

Newer JVMs are faster. Upgrading to the latest version of your JVM can provide a "free" boost to performance. Switching from the default Client JVM to the Server JVM using the `-server` command-line option improves performance at the cost of a slight increase in start-up time.

### Virus Scanners

Some virus scanners scan files every time they are accessed. It is very important for performance that database files are not scanned for viruses. The database engine never interprets the data stored in the files as programs, that means even if somebody would store a virus in a database file, this would be harmless (when the virus does not run, it cannot spread). Some virus scanners allow to exclude files by suffix. Ensure files ending with `.db` are not scanned.

### Using the Trace Options

If the performance hot spots are in the database engine, in many cases the performance can be optimized by creating additional indexes, or changing the schema. Sometimes the application does not directly generate the SQL statements, for example if an O/R mapping tool is used. To view the SQL statements and JDBC API calls, you can use the trace options. For more information, see [Using the Trace Options](#).

## Index Usage

This database uses indexes to improve the performance of SELECT, UPDATE, DELETE. If a column is used in the WHERE clause of a query, and if an index exists on this column, then the index can be used. Multi-column indexes are used if all or the first columns of the index are used. Both equality lookup and range scans are supported. Indexes are used to order result sets, but only if the condition uses the same index or no index at all. The results are sorted in memory if required. Indexes are created automatically for primary key and unique constraints. Indexes are also created for foreign key constraints, if required. For other columns, indexes need to be created manually using the CREATE INDEX statement.

## Index Hints

If you have determined that H2 is not using the optimal index for your query, you can use index hints to force H2 to use specific indexes.

```
SELECT * FROM TEST USE INDEX (index_name_1, index_name_2) WHERE X=1
```

Only indexes in the list will be used when choosing an index to use on the given table. There is no significance to order in this list.

It is possible that no index in the list is chosen, in which case a full table scan will be used.

An empty list of index names forces a full table scan to be performed.

Each index in the list must exist.

## How Data is Stored Internally

For persistent databases, if a table is created with a single column primary key of type BIGINT, INT, SMALLINT, TINYINT, then the data of the table is organized in this way. This is sometimes also called a "clustered index" or "index organized table".

H2 internally stores table data and indexes in the form of b-trees. Each b-tree stores entries as a list of unique keys (one or more columns) and data (zero or more columns). The table data is always organized in the form of a "data b-tree" with a single column key of type long. If a single column primary key of type BIGINT, INT, SMALLINT, TINYINT is specified when creating the table (or just after creating the table, but before inserting any rows), then this column is used as the key of the data b-tree. If no primary key has been specified, if the primary key column is of another data type, or if the primary key contains more than one

column, then a hidden auto-increment column of type BIGINT is added to the table, which is used as the key for the data b-tree. All other columns of the table are stored within the data area of this data b-tree (except for large BLOB, CLOB columns, which are stored externally).

For each additional index, one new "index b-tree" is created. The key of this b-tree consists of the indexed columns, plus the key of the data b-tree. If a primary key is created after the table has been created, or if the primary key contains multiple column, or if the primary key is not of the data types listed above, then the primary key is stored in a new index b-tree.

## **Optimizer**

This database uses a cost based optimizer. For simple and queries and queries with medium complexity (less than 7 tables in the join), the expected cost (running time) of all possible plans is calculated, and the plan with the lowest cost is used. For more complex queries, the algorithm first tries all possible combinations for the first few tables, and the remaining tables added using a greedy algorithm (this works well for most joins). Afterwards a genetic algorithm is used to test at most 2000 distinct plans. Only left-deep plans are evaluated.

## **Expression Optimization**

After the statement is parsed, all expressions are simplified automatically if possible. Operations are evaluated only once if all parameters are constant. Functions are also optimized, but only if the function is constant (always returns the same result for the same parameter values). If the WHERE clause is always false, then the table is not accessed at all.

## **COUNT(\*) Optimization**

If the query only counts all rows of a table, then the data is not accessed. However, this is only possible if no WHERE clause is used, that means it only works for queries of the form `SELECT COUNT(*) FROM table`.

## **Updating Optimizer Statistics / Column Selectivity**

When executing a query, at most one index per join can be used. If the same table is joined multiple times, for each join only one index is used (the same index could be used for both joins, or each join could use a different index). Example: for the query `SELECT * FROM TEST T1, TEST T2 WHERE T1.NAME='A' AND T2.ID=T1.ID`, two index can be used, in this case the index on NAME for T1 and the index on ID for T2.



If a table has multiple indexes, sometimes more than one index could be used. Example: if there is a table TEST(ID, NAME, FIRSTNAME) and an index on each column, then two indexes could be used for the query `SELECT * FROM TEST WHERE NAME='A' AND FIRSTNAME='B'`, the index on NAME or the index on FIRSTNAME. It is not possible to use both indexes at the same time. Which index is used depends on the selectivity of the column. The selectivity describes the 'uniqueness' of values in a column. A selectivity of 100 means each value appears only once, and a selectivity of 1 means the same value appears in many or most rows. For the query above, the index on NAME should be used if the table contains more distinct names than first names.

The SQL statement `ANALYZE` can be used to automatically estimate the selectivity of the columns in the tables. This command should be run from time to time to improve the query plans generated by the optimizer.

## In-Memory (Hash) Indexes

Using in-memory indexes, specially in-memory hash indexes, can speed up queries and data manipulation.

In-memory indexes are automatically used for in-memory databases, but can also be created for persistent databases using `CREATE MEMORY TABLE`. In many cases, the rows itself will also be kept in-memory. Please note this may cause memory problems for large tables.

In-memory hash indexes are backed by a hash table and are usually faster than regular indexes. However, hash indexes only supports direct lookup (`WHERE ID = ?`) but not range scan (`WHERE ID < ?`). To use hash indexes, use `HASH` as in: `CREATE UNIQUE HASH INDEX` and `CREATE TABLE ...(ID INT PRIMARY KEY HASH,...)`.

## Use Prepared Statements

If possible, use prepared statements with parameters.

## Prepared Statements and IN(...)

Avoid generating SQL statements with a variable size `IN(...)` list. Instead, use a prepared statement with arrays as in the following example:

```
PreparedStatement prep = conn.prepareStatement(
    "SELECT * FROM TEST WHERE ID = ANY(?)");
prep.setObject(1, new Object[] { "1", "2" });
ResultSet rs = prep.executeQuery();
```

## Optimization Examples

See `src/test/org/h2/samples/optimizations.sql` for a few examples of queries that benefit from special optimizations built into the database.

## Cache Size and Type

By default the cache size of H2 is quite small. Consider using a larger cache size, or enable the second level soft reference cache. See also [Cache Settings](#).

## Data Types

Each data type has different storage and performance characteristics:

- The DECIMAL/NUMERIC type is slower and requires more storage than the REAL and DOUBLE types.
- Text types are slower to read, write, and compare than numeric types and generally require more storage.
- See [Large Objects](#) for information on BINARY vs. BLOB and VARCHAR vs. CLOB performance.
- Parsing and formatting takes longer for the TIME, DATE, and TIMESTAMP types than the numeric types.
- SMALLINT/TINYINT/BOOLEAN are not significantly smaller or faster to work with than INTEGER in most modes.

## Sorted Insert Optimization

To reduce disk space usage and speed up table creation, an optimization for sorted inserts is available. When used, b-tree pages are split at the insertion point. To use this optimization, add SORTED before the SELECT statement:

```
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR) AS
  SORTED SELECT X, SPACE(100) FROM SYSTEM_RANGE(1, 100);
INSERT INTO TEST
  SORTED SELECT X, SPACE(100) FROM SYSTEM_RANGE(101, 200);
```

## Using the Built-In Profiler

A very simple Java profiler is built-in. To use it, use the following template:

```
import org.h2.util.Profiler;
Profiler prof = new Profiler();
prof.startCollecting();
```

```
// .... some long running process, at least a few seconds
prof.stopCollecting();
System.out.println(prof.getTop(3));
```

## Application Profiling

### Analyze First

Before trying to optimize performance, it is important to understand where the problem is (what part of the application is slow). Blind optimization or optimization based on guesses should be avoided, because usually it is not an efficient strategy. There are various ways to analyze an application. Sometimes two implementations can be compared using `System.currentTimeMillis()`. But this does not work for complex applications with many modules, and for memory problems.

A simple way to profile an application is to use the built-in profiling tool of java. Example:

```
java -Xrunhprof:cpu=samples,depth=16 com.acme.Test
```

Unfortunately, it is only possible to profile the application from start to end. Another solution is to create a number of full thread dumps. To do that, first run `jps -l` to get the process id, and then run `jstack <pid>` or `kill -QUIT <pid>` (Linux) or press `Ctrl+C` (Windows).

A simple profiling tool is included in H2. To use it, the application needs to be changed slightly. Example:

```
import org.h2.util;
...
Profiler profiler = new Profiler();
profiler.startCollecting();
// application code
System.out.println(profiler.getTop(3));
```

The profiler is built into the H2 Console tool, to analyze databases that open slowly. To use it, run the H2 Console, and then click on 'Test Connection'. Afterwards, click on "Test successful" and you get the most common stack traces, which helps to find out why it took so long to connect. You will only get the stack traces if opening the database took more than a few seconds.

## Database Profiling

The ConvertTraceFile tool generates SQL statement statistics at the end of the SQL script file. The format used is similar to the profiling data generated when using `java -Xrunhprof`. For this to work, the trace level needs to be 2 or higher (`TRACE_LEVEL_FILE=2`). The easiest way to set the trace level is to append the setting to the database URL, for example: `jdbc:h2:~/test;TRACE_LEVEL_FILE=2` or `jdbc:h2:tcp://localhost/~/test;TRACE_LEVEL_FILE=2`. As an example, execute the following script using the H2 Console:

```
SET TRACE_LEVEL_FILE 2;
DROP TABLE IF EXISTS TEST;
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255));
@LOOP 1000 INSERT INTO TEST VALUES(?, ?);
SET TRACE_LEVEL_FILE 0;
```

After running the test case, convert the `.trace.db` file using the ConvertTraceFile tool. The trace file is located in the same directory as the database file.

```
java -cp h2*.jar org.h2.tools.ConvertTraceFile
    -traceFile "~/test.trace.db" -script "~/test.sql"
```

The generated file `test.sql` will contain the SQL statements as well as the following profiling data (results vary):

```
-----
-- SQL Statement Statistics
-- time: total time in milliseconds (accumulated)
-- count: how many times the statement ran
-- result: total update count or row count
-----
-- self accu   time   count  result sql
-- 62% 62%    158   1000   1000 INSERT INTO TEST VALUES(?, ?);
-- 37% 100%    93    1      0 CREATE TABLE TEST(ID INT PRIMARY KEY...
-- 0% 100%     0     1      0 DROP TABLE IF EXISTS TEST;
-- 0% 100%     0     1      0 SET TRACE_LEVEL_FILE 3;
```

## Statement Execution Plans

The SQL statement `EXPLAIN` displays the indexes and optimizations the database uses for a statement. The following statements support `EXPLAIN`: `SELECT`, `UPDATE`, `DELETE`, `MERGE`, `INSERT`. The following query shows that the database uses the primary key index to search for rows:

```

EXPLAIN SELECT * FROM TEST WHERE ID=1;
SELECT
    TEST.ID,
    TEST.NAME
FROM PUBLIC.TEST
    /* PUBLIC.PRIMARY_KEY_2: ID = 1 */
WHERE ID = 1

```

For joins, the tables in the execution plan are sorted in the order they are processed. The following query shows the database first processes the table INVOICE (using the primary key). For each row, it will additionally check that the value of the column AMOUNT is larger than zero, and for those rows the database will search in the table CUSTOMER (using the primary key). The query plan contains some redundancy so it is a valid statement.

```

CREATE TABLE CUSTOMER(ID IDENTITY, NAME VARCHAR);
CREATE TABLE INVOICE(ID IDENTITY,
    CUSTOMER_ID INT REFERENCES CUSTOMER(ID),
    AMOUNT NUMBER);

EXPLAIN SELECT I.ID, C.NAME FROM CUSTOMER C, INVOICE I
WHERE I.ID=10 AND AMOUNT>0 AND C.ID=I.CUSTOMER_ID;

SELECT
    I.ID,
    C.NAME
FROM PUBLIC.INVOICE I
    /* PUBLIC.PRIMARY_KEY_9: ID = 10 */
    /* WHERE (I.ID = 10)
        AND (AMOUNT > 0)
    */
INNER JOIN PUBLIC.CUSTOMER C
    /* PUBLIC.PRIMARY_KEY_5: ID = I.CUSTOMER_ID */
    ON 1=1
WHERE (C.ID = I.CUSTOMER_ID)
    AND ((I.ID = 10)
        AND (AMOUNT > 0))

```

## Displaying the Scan Count

EXPLAIN ANALYZE additionally shows the scanned rows per table and pages read from disk per table or index. This will actually execute the query, unlike EXPLAIN which only prepares it. The following query scanned 1000 rows, and to do that had to read 85 pages from the data area of the table. Running the query twice

will not list the pages read from disk, because they are now in the cache. The tableScan means this query doesn't use an index.

```
EXPLAIN ANALYZE SELECT * FROM TEST;  
SELECT  
    TEST.ID,  
    TEST.NAME  
FROM PUBLIC.TEST  
    /* PUBLIC.TEST.tableScan */  
    /* scanCount: 1000 */  
/*  
total: 85  
TEST.TEST_DATA read: 85 (100%)  
*/
```

The cache will prevent the pages are read twice. H2 reads all columns of the row unless only the columns in the index are read. Except for large CLOB and BLOB, which are not store in the table.

## Special Optimizations

For certain queries, the database doesn't need to read all rows, or doesn't need to sort the result even if ORDER BY is used.

For queries of the form `SELECT COUNT(*), MIN(ID), MAX(ID) FROM TEST`, the query plan includes the line `/* direct lookup */` if the data can be read from an index.

For queries of the form `SELECT DISTINCT CUSTOMER_ID FROM INVOICE`, the query plan includes the line `/* distinct */` if there is an non-unique or multi-column index on this column, and if this column has a low selectivity.

For queries of the form `SELECT * FROM TEST ORDER BY ID`, the query plan includes the line `/* index sorted */` to indicate there is no separate sorting required.

For queries of the form `SELECT * FROM TEST GROUP BY ID ORDER BY ID`, the query plan includes the line `/* group sorted */` to indicate there is no separate sorting required.

## How Data is Stored and How Indexes Work

Internally, each row in a table is identified by a unique number, the row id. The rows of a table are stored with the row id as the key. The row id is a number of

type long. If a table has a single column primary key of type INT or BIGINT, then the value of this column is the row id, otherwise the database generates the row id automatically. There is a (non-standard) way to access the row id: using the `_ROWID_` pseudo-column:

```
CREATE TABLE ADDRESS(FIRST_NAME VARCHAR,  
    NAME VARCHAR, CITY VARCHAR, PHONE VARCHAR);  
INSERT INTO ADDRESS VALUES('John', 'Miller', 'Berne', '123 456 789');  
INSERT INTO ADDRESS VALUES('Philip', 'Jones', 'Berne', '123 012 345');  
SELECT _ROWID_, * FROM ADDRESS;
```

The data is stored in the database as follows:

<b><code>_ROWID_</code></b>	<b><code>FIRST_NAME</code></b>	<b><code>NAME</code></b>	<b><code>CITY</code></b>	<b><code>PHONE</code></b>
1	John	Miller	Berne	123 456 789
2	Philip	Jones	Berne	123 012 345

Access by row id is fast because the data is sorted by this key. Please note the row id is not available until after the row was added (that means, it can not be used in computed columns or constraints). If the query condition does not contain the row id (and if no other index can be used), then all rows of the table are scanned. A table scan iterates over all rows in the table, in the order of the row id. To find out what strategy the database uses to retrieve the data, use `EXPLAIN SELECT`:

```
SELECT * FROM ADDRESS WHERE NAME = 'Miller';  
  
EXPLAIN SELECT PHONE FROM ADDRESS WHERE NAME = 'Miller';  
SELECT  
    PHONE  
FROM PUBLIC.ADDRESS  
    /* PUBLIC.ADDRESS.tableScan */  
WHERE NAME = 'Miller';
```

## Indexes

An index internally is basically just a table that contains the indexed column(s), plus the row id:

```
CREATE INDEX INDEX_PLACE ON ADDRESS(CITY, NAME, FIRST_NAME);
```

In the index, the data is sorted by the indexed columns. So this index contains the following data:

<b><code>CITY</code></b>	<b><code>NAME</code></b>	<b><code>FIRST_NAME</code></b>	<b><code>_ROWID_</code></b>
--------------------------	--------------------------	--------------------------------	-----------------------------

Berne Jones Philip 2

Berne Miller John 1

When the database uses an index to query the data, it searches the index for the given data, and (if required) reads the remaining columns in the main data table (retrieved using the row id). An index on city, name, and first name (multi-column index) allows to quickly search for rows when the city, name, and first name are known. If only the city and name, or only the city is known, then this index is also used (so creating an additional index on just the city is not needed). This index is also used when reading all rows, sorted by the indexed columns. However, if only the first name is known, then this index is not used:

```
EXPLAIN SELECT PHONE FROM ADDRESS
  WHERE CITY = 'Berne' AND NAME = 'Miller'
  AND FIRST_NAME = 'John';
SELECT
  PHONE
FROM PUBLIC.ADDRESS
  /* PUBLIC.INDEX_PLACE: FIRST_NAME = 'John'
  AND CITY = 'Berne'
  AND NAME = 'Miller'
  */
WHERE (FIRST_NAME = 'John')
  AND ((CITY = 'Berne')
  AND (NAME = 'Miller'));
```

```
EXPLAIN SELECT PHONE FROM ADDRESS WHERE CITY = 'Berne';
SELECT
  PHONE
FROM PUBLIC.ADDRESS
  /* PUBLIC.INDEX_PLACE: CITY = 'Berne' */
WHERE CITY = 'Berne';
```

```
EXPLAIN SELECT * FROM ADDRESS ORDER BY CITY, NAME, FIRST_NAME;
SELECT
  ADDRESS.FIRST_NAME,
  ADDRESS.NAME,
  ADDRESS.CITY,
  ADDRESS.PHONE
FROM PUBLIC.ADDRESS
  /* PUBLIC.INDEX_PLACE */
ORDER BY 3, 2, 1
/* index sorted */;
```

```
EXPLAIN SELECT PHONE FROM ADDRESS WHERE FIRST_NAME = 'John';
```



```
SELECT
  PHONE
FROM PUBLIC.ADDRESS
  /* PUBLIC.ADDRESS.tableScan */
WHERE FIRST_NAME = 'John';
```

If your application often queries the table for a phone number, then it makes sense to create an additional index on it:

```
CREATE INDEX IDX_PHONE ON ADDRESS(PHONE);
```

This index contains the phone number, and the row id:

<b>PHONE</b>	<b>_ROWID_</b>
123 012 345	2
123 456 789	1

## Using Multiple Indexes

Within a query, only one index per logical table is used. Using the condition `PHONE = '123 567 789' OR CITY = 'Berne'` would use a table scan instead of first using the index on the phone number and then the index on the city. It makes sense to write two queries and combine them using `UNION`. In this case, each individual query uses a different index:

```
EXPLAIN SELECT NAME FROM ADDRESS WHERE PHONE = '123 567 789'
UNION SELECT NAME FROM ADDRESS WHERE CITY = 'Berne';
```

```
(SELECT
  NAME
FROM PUBLIC.ADDRESS
  /* PUBLIC.IDX_PHONE: PHONE = '123 567 789' */
WHERE PHONE = '123 567 789')
UNION
(SELECT
  NAME
FROM PUBLIC.ADDRESS
  /* PUBLIC.INDEX_PLACE: CITY = 'Berne' */
WHERE CITY = 'Berne')
```

## Fast Database Import

To speed up large imports, consider using the following options temporarily:

- SET LOG 0 (disabling the transaction log)
- SET CACHE\_SIZE (a large cache is faster)
- SET LOCK\_MODE 0 (disable locking)
- SET UNDO\_LOG 0 (disable the session undo log)

These options can be set in the database URL:

`jdbc:h2:~/test;LOG=0;CACHE_SIZE=65536;LOCK_MODE=0;UNDO_LOG=0`. Most of those options are not recommended for regular use, that means you need to reset them after use.

If you have to import a lot of rows, use a `PreparedStatement` or use CSV import. Please note that `CREATE TABLE(...) ... AS SELECT ...` is faster than `CREATE TABLE(...); INSERT INTO ... SELECT ....`

# Advanced

- Result Sets
- Large Objects
- Linked Tables
- Spatial Features
- Recursive Queries
- Updatable Views
- Transaction Isolation
- Multi-Version Concurrency Control (MVCC)
- Clustering / High Availability
- Two Phase Commit
- Compatibility
- Standards Compliance
- Run as Windows Service
- ODBC Driver
- Using H2 in Microsoft .NET
- ACID
- Durability Problems
- Using the Recover Tool
- File Locking Protocols
- Using Passwords
- Password Hash
- Protection against SQL Injection
- Protection against Remote Access
- Restricting Class Loading and Usage
- Security Protocols
- TLS Connections
- Universally Unique Identifiers (UUID)
- Settings Read from System Properties
- Setting the Server Bind Address
- Pluggable File System
- Split File System
- Database Upgrade
- Java Objects Serialization
- Custom Data Types Handler API
- Limits and Limitations
- Glossary and Links

# Result Sets

## Statements that Return a Result Set

The following statements return a result set: SELECT, EXPLAIN, CALL, SCRIPT, SHOW, HELP. All other statements return an update count.

## Limiting the Number of Rows

Before the result is returned to the application, all rows are read by the database. Server side cursors are not supported currently. If only the first few rows are interesting for the application, then the result set size should be limited to improve the performance. This can be done using FETCH in a query (example: SELECT \* FROM TEST FETCH FIRST 100 ROWS ONLY), or by using Statement.setMaxRows(max).

## Large Result Sets and External Sorting

For large result set, the result is buffered to disk. The threshold can be defined using the statement SET MAX\_MEMORY\_ROWS. If ORDER BY is used, the sorting is done using an external sort algorithm. In this case, each block of rows is sorted using quick sort, then written to disk; when reading the data, the blocks are merged together.

## Large Objects

### Storing and Reading Large Objects

If it is possible that the objects don't fit into memory, then the data type CLOB (for textual data) or BLOB (for binary data) should be used. For these data types, the objects are not fully read into memory, by using streams. To store a BLOB, use PreparedStatement.setBinaryStream. To store a CLOB, use PreparedStatement.setCharacterStream. To read a BLOB, use ResultSet.getBinaryStream, and to read a CLOB, use ResultSet.getCharacterStream. When using the client/server mode, large BLOB and CLOB data is stored in a temporary file on the client side.

### When to use CLOB/BLOB

By default, this database stores large LOB (CLOB and BLOB) objects separate from the main table data. Small LOB objects are stored in-place, the threshold can be set using [MAX\\_LENGTH\\_INPLACE\\_LOB](#), but there is still an overhead to use

CLOB/BLOB. Because of this, BLOB and CLOB should never be used for columns with a maximum size below about 200 bytes. The best threshold depends on the use case; reading in-place objects is faster than reading from separate files, but slows down the performance of operations that don't involve this column.

## Large Object Compression

The following feature is only available for the PageStore storage engine. For the MVStore engine (the default for H2 version 1.4.x), append `;COMPRESS=TRUE` to the database URL instead. CLOB and BLOB values can be compressed by using [SET COMPRESS\\_LOB](#). The LZF algorithm is faster but needs more disk space. By default compression is disabled, which usually speeds up write operations. If you store many large compressible values such as XML, HTML, text, and uncompressed binary files, then compressing can save a lot of disk space (sometimes more than 50%), and read operations may even be faster.

## Linked Tables

This database supports linked tables, which means tables that don't exist in the current database but are just links to another database. To create such a link, use the `CREATE LINKED TABLE` statement:

```
CREATE LINKED TABLE LINK('org.postgresql.Driver', 'jdbc:postgresql:test', 'sa', 'sa', 'TEST');
```

You can then access the table in the usual way. Whenever the linked table is accessed, the database issues specific queries over JDBC. Using the example above, if you issue the query `SELECT * FROM LINK WHERE ID=1`, then the following query is run against the PostgreSQL database: `SELECT * FROM TEST WHERE ID=1`. The same happens for insert and update statements. Only simple statements are executed against the target database, that means no joins (queries that contain joins are converted to simple queries). Prepared statements are used where possible.

To view the statements that are executed against the target table, set the trace level to 3.

If multiple linked tables point to the same database (using the same database URL), the connection is shared. To disable this, set the system property `h2.shareLinkedConnections=false`.

The statement [CREATE LINKED TABLE](#) supports an optional schema name parameter.

The following are not supported because they may result in a deadlock: creating a linked table to the same database, and creating a linked table to another database using the server mode if the other database is open in the same server (use the embedded mode instead).

Data types that are not supported in H2 are also not supported for linked tables, for example unsigned data types if the value is outside the range of the signed type. In such cases, the columns needs to be cast to a supported type.

## Updatable Views

By default, views are not updatable. To make a view updatable, use an "instead of" trigger as follows:

```
CREATE TRIGGER TRIGGER_NAME  
INSTEAD OF INSERT, UPDATE, DELETE  
ON VIEW_NAME  
FOR EACH ROW CALL "com.acme.TriggerClassName";
```

Update the base table(s) within the trigger as required. For details, see the sample application `org.h2.samples.UpdatableView`.

## Transaction Isolation

Please note that most data definition language (DDL) statements, such as "create table", commit the current transaction. See the [Grammar](#) for details.

Transaction isolation is provided for all data manipulation language (DML) statements.

Please note that with default MVStore storage engine table level locking is not used. Instead, rows are locked for update, and read committed is used in all cases except for explicitly selected read uncommitted transaction isolation level.

This database supports the following transaction isolation levels:

- **Read Committed**

This is the default level. Read locks are released immediately after executing the statement, but write locks are kept until the transaction commits. Higher concurrency is possible when using this level.

To enable, execute the SQL statement `SET LOCK_MODE 3` or append `;LOCK_MODE=3` to the database URL:  
`jdbc:h2:~/test;LOCK_MODE=3`

- **Serializable**

Both read locks and write locks are kept until the transaction commits. To enable, execute the SQL statement SET LOCK\_MODE 1 or append ;LOCK\_MODE=1 to the database URL:  
jdbc:h2:~/test;LOCK\_MODE=1

- **Read Uncommitted**

This level means that transaction isolation is disabled. This level is not supported by PageStore engine if multi-threaded mode is enabled. To enable, execute the SQL statement SET LOCK\_MODE 0 or append ;LOCK\_MODE=0 to the database URL:  
jdbc:h2:~/test;LOCK\_MODE=0

When using the isolation level 'serializable', dirty reads, non-repeatable reads, and phantom reads are prohibited.

- **Dirty Reads**

Means a connection can read uncommitted changes made by another connection.

Possible with: read uncommitted

- **Non-Repeatable Reads**

A connection reads a row, another connection changes a row and commits, and the first connection re-reads the same row and gets the new result.

Possible with: read uncommitted, read committed

- **Phantom Reads**

A connection reads a set of rows using a condition, another connection inserts a row that falls in this condition and commits, then the first connection re-reads using the same condition and gets the new row.

Possible with: read uncommitted, read committed

## **Multi-Version Concurrency Control (MVCC)**

With default MVStore engine delete, insert and update operations only issue a shared lock on the table. An exclusive lock is still used when adding or removing columns, when dropping the table, and when using SELECT ... FOR UPDATE. Connections only 'see' committed data, and own changes. That means, if connection A updates a row but doesn't commit this change yet, connection B will see the old value. Only when the change is committed, the new value is visible by other connections (read committed). If multiple connections concurrently try to update the same row, the database waits until it can apply the change, but at most until the lock timeout expires.

## Table Level Locking (PageStore engine)

With PageStore engine to make sure all connections only see consistent data, table level locking is used. This mechanism does not allow high concurrency, but is very fast. Shared locks and exclusive locks are supported. Before reading from a table, the database tries to add a shared lock to the table (this is only possible if there is no exclusive lock on the object by another connection). If the shared lock is added successfully, the table can be read. It is allowed that other connections also have a shared lock on the same object. If a connection wants to write to a table (update or delete a row), an exclusive lock is required. To get the exclusive lock, other connection must not have any locks on the object. After the connection commits, all locks are released. This database keeps all locks in memory. When a lock is released, and multiple connections are waiting for it, one of them is picked at random.

## Lock Timeout

If a connection cannot get a lock on an object, the connection waits for some amount of time (the lock timeout). During this time, hopefully the connection holding the lock commits and it is then possible to get the lock. If this is not possible because the other connection does not release the lock for some time, the unsuccessful connection will get a lock timeout exception. The lock timeout can be set individually for each connection.

## Clustering / High Availability

This database supports a simple clustering / high availability mechanism. The architecture is: two database servers run on two different computers, and on both computers is a copy of the same database. If both servers run, each database operation is executed on both computers. If one server fails (power, hardware or network failure), the other server can still continue to work. From this point on, the operations will be executed only on one server until the other server is back up.

Clustering can only be used in the server mode (the embedded mode does not support clustering). The cluster can be re-created using the CreateCluster tool without stopping the remaining server. Applications that are still connected are automatically disconnected, however when appending `;AUTO_RECONNECT=TRUE`, they will recover from that.

To initialize the cluster, use the following steps:

- Create a database



- Use the CreateCluster tool to copy the database to another location and initialize the clustering. Afterwards, you have two databases containing the same data.
- Start two servers (one for each copy of the database)
- You are now ready to connect to the databases with the client application(s)

## Using the CreateCluster Tool

To understand how clustering works, please try out the following example. In this example, the two databases reside on the same computer, but usually, the databases will be on different servers.

- Create two directories: server1, server2. Each directory will simulate a directory on a computer.
- Start a TCP server pointing to the first directory. You can do this using the command line:

```
java org.h2.tools.Server  
-tcp -tcpPort 9101  
-baseDir server1
```

- Start a second TCP server pointing to the second directory. This will simulate a server running on a second (redundant) computer. You can do this using the command line:

```
java org.h2.tools.Server  
-tcp -tcpPort 9102  
-baseDir server2
```

- Use the CreateCluster tool to initialize clustering. This will automatically create a new, empty database if it does not exist. Run the tool on the command line:

```
java org.h2.tools.CreateCluster  
-urlSource jdbc:h2:tcp://localhost:9101/~ /test  
-urlTarget jdbc:h2:tcp://localhost:9102/~ /test  
-user sa  
-serverList localhost:9101,localhost:9102
```

- You can now connect to the databases using an application or the H2 Console using the JDBC URL  
jdbc:h2:tcp://localhost:9101,localhost:9102/~ /test
- If you stop a server (by killing the process), you will notice that the other machine continues to work, and therefore the database is still accessible.

- To restore the cluster, you first need to delete the database that failed, then restart the server that was stopped, and re-run the CreateCluster tool.

## Detect Which Cluster Instances are Running

To find out which cluster nodes are currently running, execute the following SQL statement:

```
SELECT VALUE FROM INFORMATION_SCHEMA.SETTINGS WHERE  
NAME='CLUSTER'
```

If the result is " (two single quotes), then the cluster mode is disabled. Otherwise, the list of servers is returned, enclosed in single quote. Example: 'server1:9191,server2:9191'.

It is also possible to get the list of servers by using `Connection.getClientInfo()`.

The property list returned from `getClientInfo()` contains a `numServers` property that returns the number of servers that are in the connection list. To get the actual servers, `getClientInfo()` also has properties `server0..serverX`, where `serverX` is the number of servers minus 1.

Example: To get the 2nd server in the connection list one uses `getClientInfo('server1')`. **Note:** The `serverX` property only returns IP addresses and ports and not hostnames.

## Clustering Algorithm and Limitations

Read-only queries are only executed against the first cluster node, but all other statements are executed against all nodes. There is currently no load balancing made to avoid problems with transactions. The following functions may yield different results on different cluster nodes and must be executed with care: `UUID()`, `RANDOM_UUID()`, `SECURE RAND()`, `SESSION_ID()`, `MEMORY_FREE()`, `MEMORY_USED()`, `CSVREAD()`, `CSVWRITE()`, `RAND()` [when not using a seed]. Those functions should not be used directly in modifying statements (for example `INSERT`, `UPDATE`, `MERGE`). However, they can be used in read-only statements and the result can then be used for modifying statements. Using auto-increment and identity columns is currently not supported. Instead, sequence values need to be manually requested and then used to insert data (using two statements).

When using the cluster modes, result sets are read fully in memory by the client, so that there is no problem if the server dies that executed the query. Result sets must fit in memory on the client side.

The SQL statement SET AUTOCOMMIT FALSE is not supported in the cluster mode. To disable autocommit, the method Connection.setAutoCommit(false) needs to be called.

It is possible that a transaction from one connection overtakes a transaction from a different connection. Depending on the operations, this might result in different results, for example when conditionally incrementing a value in a row.

## Two Phase Commit

The two phase commit protocol is supported. 2-phase-commit works as follows:

- Autocommit needs to be switched off
- A transaction is started, for example by inserting a row
- The transaction is marked 'prepared' by executing the SQL statement PREPARE COMMIT transactionName
- The transaction can now be committed or rolled back
- If a problem occurs before the transaction was successfully committed or rolled back (for example because a network problem occurred), the transaction is in the state 'in-doubt'
- When re-connecting to the database, the in-doubt transactions can be listed with SELECT \* FROM INFORMATION\_SCHEMA.IN\_DOUBT
- Each transaction in this list must now be committed or rolled back by executing COMMIT TRANSACTION transactionName or ROLLBACK TRANSACTION transactionName
- The database needs to be closed and re-opened to apply the changes

## Compatibility

This database is (up to a certain point) compatible to other databases such as HSQLDB, MySQL and PostgreSQL. There are certain areas where H2 is incompatible.

## Transaction Commit when Autocommit is On

At this time, this database engine commits a transaction (if autocommit is switched on) just before returning the result. For a query, this means the transaction is committed even before the application scans through the result set, and before the result set is closed. Other database engines may commit the transaction in this case when the result set is closed.

## Keywords / Reserved Words

There is a list of keywords that can't be used as identifiers (table names, column names and so on), unless they are quoted (surrounded with double quotes). The list is currently:

ALL, ARRAY, CASE, CHECK, CONSTRAINT, CROSS, CURRENT\_DATE, CURRENT\_TIME, CURRENT\_TIMESTAMP, CURRENT\_USER, DISTINCT, EXCEPT, EXISTS, FALSE, FETCH, FOR, FOREIGN, FROM, FULL, GROUP, HAVING, IF, INNER, INTERSECT, INTERSECTS, INTERVAL, IS, JOIN, LIKE, LIMIT, LOCALTIME, LOCALTIMESTAMP, MINUS, NATURAL, NOT, NULL, OFFSET, ON, ORDER, PRIMARY, QUALIFY, ROW, \_ROWID\_, ROWNUM, SELECT, SYSDATE, SYSTIME, SYSTIMESTAMP, TABLE, TODAY, TOP, TRUE, UNION, UNIQUE, VALUES, WHERE, WINDOW, WITH

Certain words of this list are keywords because they are functions that can be used without '()', for example CURRENT\_TIMESTAMP.

## Standards Compliance

This database tries to be as much standard compliant as possible. For the SQL language, ANSI/ISO is the main standard. There are several versions that refer to the release date: SQL-92, SQL:1999, and SQL:2003. Unfortunately, the standard documentation is not freely available. Another problem is that important features are not standardized. Whenever this is the case, this database tries to be compatible to other databases.

## Supported Character Sets, Character Encoding, and Unicode

H2 internally uses Unicode, and supports all character encoding systems and character sets supported by the virtual machine you use.

## Run as Windows Service

Using a native wrapper / adapter, Java applications can be run as a Windows Service. There are various tools available to do that. The Java Service Wrapper from [Tanuki Software, Inc.](#) is included in the installation. Batch files are provided to install, start, stop and uninstall the H2 Database Engine Service. This service contains the TCP Server and the H2 Console web application. The batch files are located in the directory h2/service.

The service wrapper bundled with H2 is a 32-bit version. To use a 64-bit version of Windows (x64), you need to use a 64-bit version of the wrapper, for example the one from [Simon Krenger](#).

When running the database as a service, absolute path should be used. Using ~ in the database URL is problematic in this case, because it means to use the home directory of the current user. The service might run without or with the wrong user, so that the database files might end up in an unexpected place.

## **Install the Service**

The service needs to be registered as a Windows Service first. To do that, double click on 1\_install\_service.bat. If successful, a command prompt window will pop up and disappear immediately. If not, a message will appear.

## **Start the Service**

You can start the H2 Database Engine Service using the service manager of Windows, or by double clicking on 2\_start\_service.bat. Please note that the batch file does not print an error message if the service is not installed.

## **Connect to the H2 Console**

After installing and starting the service, you can connect to the H2 Console application using a browser. Double clicking on 3\_start\_browser.bat to do that. The default port (8082) is hard coded in the batch file.

## **Stop the Service**

To stop the service, double click on 4\_stop\_service.bat. Please note that the batch file does not print an error message if the service is not installed or started.

## **Uninstall the Service**

To uninstall the service, double click on 5\_uninstall\_service.bat. If successful, a command prompt window will pop up and disappear immediately. If not, a message will appear.

## **Additional JDBC drivers**

To use other databases (for example MySQL), the location of the JDBC drivers of those databases need to be added to the environment variables H2DRIVERS or CLASSPATH before installing the service. Multiple drivers can be set; each entry needs to be separated with a ; (Windows) or : (other operating systems). Spaces in the path names are supported. The settings must not be quoted.

## ODBC Driver

This database does not come with its own ODBC driver at this time, but it supports the PostgreSQL network protocol. Therefore, the PostgreSQL ODBC driver can be used. Support for the PostgreSQL network protocol is quite new and should be viewed as experimental. It should not be used for production applications.

To use the PostgreSQL ODBC driver on 64 bit versions of Windows, first run `c:/windows/syswow64/odbcad32.exe`. At this point you set up your DSN just like you would on any other system. See also: [Re: ODBC Driver on Windows 64 bit](#)

## ODBC Installation

First, the ODBC driver must be installed. Any recent PostgreSQL ODBC driver should work, however version 8.2 (psqlodbc-08\_02\*) or newer is recommended. The Windows version of the PostgreSQL ODBC driver is available at <https://www.postgresql.org/ftp/odbc/versions/msi/>.

## Starting the Server

After installing the ODBC driver, start the H2 Server using the command line:

```
java -cp h2*.jar org.h2.tools.Server
```

The PG Server (PG for PostgreSQL protocol) is started as well. By default, databases are stored in the current working directory where the server is started. Use `-baseDir` to save databases in another directory, for example the user home directory:

```
java -cp h2*.jar org.h2.tools.Server -baseDir ~
```

The PG server can be started and stopped from within a Java application as follows:

```
Server server = Server.createPgServer("-baseDir", "~");
server.start();
...
server.stop();
```

By default, only connections from localhost are allowed. To allow remote connections, use `-pgAllowOthers` when starting the server.

To map an ODBC database name to a different JDBC database name, use the option `-key` when starting the server. Please note only one mapping is allowed. The following will map the ODBC database named TEST to the database URL `jdbc:h2:~/data/test;cipher=aes`:

```
java org.h2.tools.Server -pg -key TEST "~/data/test;cipher=aes"
```

## ODBC Configuration

After installing the driver, a new Data Source must be added. In Windows, run `odbcad32.exe` to open the Data Source Administrator. Then click on 'Add...' and select the PostgreSQL Unicode driver. Then click 'Finish'. You will be able to change the connection properties. The property column represents the property key in the `odbc.ini` file (which may be different from the GUI).

Property	Example	Remarks
Data Source	H2 Test	The name of the ODBC Data Source
Database	~/test;ifexists=true	The database name. This can include connections settings. By default, the database is stored in the current working directory where the Server is started except when the <code>-baseDir</code> setting is used. The name must be at least 3 characters.
Servename	localhost	The server name or IP address. By default, only remote connections are allowed
Username	sa	The database user name.
SSL	false (disabled)	At this time, SSL is not supported.
Port	5435	The port where the PG Server is listening.
Password	sa	The database password.

To improve performance, please enable 'server side prepare' under Options / Datasource / Page 2 / Server side prepare.

Afterwards, you may use this data source.

## PG Protocol Support Limitations

At this time, only a subset of the PostgreSQL network protocol is implemented. Also, there may be compatibility problems on the SQL level, with the catalog, or with text encoding. Problems are fixed as they are found. Currently, statements

can not be canceled when using the PG protocol. Also, H2 does not provide index meta over ODBC.

PostgreSQL ODBC Driver Setup requires a database password; that means it is not possible to connect to H2 databases without password. This is a limitation of the ODBC driver.

## Security Considerations

Currently, the PG Server does not support challenge response or encrypt passwords. This may be a problem if an attacker can listen to the data transferred between the ODBC driver and the server, because the password is readable to the attacker. Also, it is currently not possible to use encrypted SSL connections. Therefore the ODBC driver should not be used where security is important.

The first connection that opens a database using the PostgreSQL server needs to be an administrator user. Subsequent connections don't need to be opened by an administrator.

## Using Microsoft Access

When using Microsoft Access to edit data in a linked H2 table, you may need to enable the following option: Tools - Options - Edit/Find - ODBC fields.

## Using H2 in Microsoft .NET

The database can be used from Microsoft .NET even without using Java, by using IKVM.NET. You can access a H2 database on .NET using the JDBC API, or using the ADO.NET interface.

## Using the ADO.NET API on .NET

An implementation of the ADO.NET interface is available in the open source project [H2Sharp](#).

## Using the JDBC API on .NET

- Install the .NET Framework from [Microsoft](#). Mono has not yet been tested.
- Install [IKVM.NET](#).
- Copy the h2\*.jar file to ikvm/bin
- Run the H2 Console using: ikvm -jar h2\*.jar
- Convert the H2 Console to an .exe file using: ikvmc -target:winexe h2\*.jar.  
You may ignore the warnings.



- Create a .dll file using (change the version accordingly): ikvmc.exe -target:library -version:1.0.69.0 h2\*.jar

If you want your C# application use H2, you need to add the h2.dll and the IKVM.OpenJDK.ClassLibrary.dll to your C# solution. Here some sample code:

```
using System;
using java.sql;

class Test
{
    static public void Main()
    {
        org.h2.Driver.load();
        Connection conn = DriverManager.getConnection("jdbc:h2:~/test", "sa",
"sa");
        Statement stat = conn.createStatement();
        ResultSet rs = stat.executeQuery("SELECT 'Hello World'");
        while (rs.next())
        {
            Console.WriteLine(rs.getString(1));
        }
    }
}
```

## ACID

In the database world, ACID stands for:

- Atomicity: transactions must be atomic, meaning either all tasks are performed or none.
- Consistency: all operations must comply with the defined constraints.
- Isolation: transactions must be isolated from each other.
- Durability: committed transaction will not be lost.

### Atomicity

Transactions in this database are always atomic.

### Consistency

By default, this database is always in a consistent state. Referential integrity rules are enforced except when explicitly disabled.

## Isolation

For H2, as with most other database systems, the default isolation level is 'read committed'. This provides better performance, but also means that transactions are not completely isolated. H2 supports the transaction isolation levels 'serializable', 'read committed', and 'read uncommitted'.

## Durability

This database does not guarantee that all committed transactions survive a power failure. Tests show that all databases sometimes lose transactions on power failure (for details, see below). Where losing transactions is not acceptable, a laptop or UPS (uninterruptible power supply) should be used. If durability is required for all possible cases of hardware failure, clustering should be used, such as the H2 clustering mode.

## Durability Problems

Complete durability means all committed transaction survive a power failure. Some databases claim they can guarantee durability, but such claims are wrong. A durability test was run against H2, HSQLDB, PostgreSQL, and Derby. All of those databases sometimes lose committed transactions. The test is included in the H2 download, see `org.h2.test.poweroff.Test`.

## Ways to (Not) Achieve Durability

Making sure that committed transactions are not lost is more complicated than it seems first. To guarantee complete durability, a database must ensure that the log record is on the hard drive before the commit call returns. To do that, databases use different methods. One is to use the 'synchronous write' file access mode. In Java, `RandomAccessFile` supports the modes `rws` and `rwd`:

- `rwd`: every update to the file's content is written synchronously to the underlying storage device.
- `rws`: in addition to `rwd`, every update to the metadata is written synchronously.

A test (`org.h2.test.poweroff.TestWrite`) with one of those modes achieves around 50 thousand write operations per second. Even when the operating system write buffer is disabled, the write rate is around 50 thousand operations per second. This feature does not force changes to disk because it does not flush all buffers. The test updates the same byte in the file again and again. If the hard drive was able to write at this rate, then the disk would need to make at least 50 thousand revolutions per second, or 3 million RPM (revolutions per minute). There are no

such hard drives. The hard drive used for the test is about 7200 RPM, or about 120 revolutions per second. There is an overhead, so the maximum write rate must be lower than that.

Calling `fsync` flushes the buffers. There are two ways to do that in Java:

- `FileDescriptor.sync()`. The documentation says that this forces all system buffers to synchronize with the underlying device. This method is supposed to return after all in-memory modified copies of buffers associated with this file descriptor have been written to the physical medium.
- `FileChannel.force()`. This method is supposed to force any updates to this channel's file to be written to the storage device that contains it.

By default, MySQL calls `fsync` for each commit. When using one of those methods, only around 60 write operations per second can be achieved, which is consistent with the RPM rate of the hard drive used. Unfortunately, even when calling `FileDescriptor.sync()` or `FileChannel.force()`, data is not always persisted to the hard drive, because most hard drives do not obey `fsync()`: see [Your Hard Drive Lies to You](#). In Mac OS X, `fsync` does not flush hard drive buffers. See [Bad fsync?](#). So the situation is confusing, and tests prove there is a problem.

Trying to flush hard drive buffers is hard, and if you do the performance is very bad. First you need to make sure that the hard drive actually flushes all buffers. Tests show that this can not be done in a reliable way. Then the maximum number of transactions is around 60 per second. Because of those reasons, the default behavior of H2 is to delay writing committed transactions.

In H2, after a power failure, a bit more than one second of committed transactions may be lost. To change the behavior, use `SET WRITE_DELAY` and `CHECKPOINT SYNC`. Most other databases support commit delay as well. In the performance comparison, commit delay was used for all databases that support it.

## Running the Durability Test

To test the durability / non-durability of this and other databases, you can use the test application in the package `org.h2.test.poweroff`. Two computers with network connection are required to run this test. One computer just listens, while the test application is run (and power is cut) on the other computer. The computer with the listener application opens a TCP/IP port and listens for an incoming connection. The second computer first connects to the listener, and then created the databases and starts inserting records. The connection is set to 'autocommit', which means after each inserted record a commit is performed automatically. Afterwards, the test computer notifies the listener that this record was inserted successfully. The listener computer displays the last inserted record number every 10 seconds. Now, switch off the power manually, then restart the computer, and

run the application again. You will find out that in most cases, none of the databases contains all the records that the listener computer knows about. For details, please consult the source code of the listener and test application.

## Using the Recover Tool

The Recover tool can be used to extract the contents of a database file, even if the database is corrupted. It also extracts the content of the transaction log and large objects (CLOB or BLOB). To run the tool, type on the command line:

```
java -cp h2*.jar org.h2.tools.Recover
```

For each database in the current directory, a text file will be created. This file contains raw insert statements (for the data) and data definition (DDL) statements to recreate the schema of the database. This file can be executed using the RunScript tool or a RUNSCRIPT FROM SQL statement. The script includes at least one CREATE USER statement. If you run the script against a database that was created with the same user, or if there are conflicting users, running the script will fail. Consider running the script against a database that was created with a user name that is not in the script.

The Recover tool creates a SQL script from database file. It also processes the transaction log.

To verify the database can recover at any time, append ;RECOVER\_TEST=64 to the database URL in your test environment. This will simulate an application crash after each 64 writes to the database file. A log file named `databaseName.h2.db.log` is created that lists the operations. The recovery is tested using an in-memory file system, that means it may require a larger heap setting.

## File Locking Protocols

Multiple concurrent connections to the same database are supported, however a database file can only be open for reading and writing (in embedded mode) by one process at the same time. Otherwise, the processes would overwrite each others data and corrupt the database file. To protect against this problem, whenever a database is opened, a lock file is created to signal other processes that the database is in use. If the database is closed, or if the process that opened the database stops normally, this lock file is deleted.

In special cases (if the process did not terminate normally, for example because there was a power failure), the lock file is not deleted by the process that created it. That means the existence of the lock file is not a safe protocol for file locking.

However, this software uses a challenge-response protocol to protect the database files. There are two methods (algorithms) implemented to provide both security (that is, the same database files cannot be opened by two processes at the same time) and simplicity (that is, the lock file does not need to be deleted manually by the user). The two methods are 'file method' and 'socket methods'.

The file locking protocols (except the file locking method 'FS') have the following limitation: if a shared file system is used, and the machine with the lock owner is sent to sleep (standby or hibernate), another machine may take over. If the machine that originally held the lock wakes up, the database may become corrupt. If this situation can occur, the application must ensure the database is closed when the application is put to sleep.

## **File Locking Method 'File'**

The default method for database file locking for version 1.3 and older is the 'File Method'. The algorithm is:

- If the lock file does not exist, it is created (using the atomic operation `File.createNewFile`). Then, the process waits a little bit (20 ms) and checks the file again. If the file was changed during this time, the operation is aborted. This protects against a race condition when one process deletes the lock file just after another one create it, and a third process creates the file again. It does not occur if there are only two writers.
- If the file can be created, a random number is inserted together with the locking method ('file'). Afterwards, a watchdog thread is started that checks regularly (every second once by default) if the file was deleted or modified by another (challenger) thread / process. Whenever that occurs, the file is overwritten with the old data. The watchdog thread runs with high priority so that a change to the lock file does not get through undetected even if the system is very busy. However, the watchdog thread does use very little resources (CPU time), because it waits most of the time. Also, the watchdog only reads from the hard disk and does not write to it.
- If the lock file exists and was recently modified, the process waits for some time (up to two seconds). If it was still changed, an exception is thrown (database is locked). This is done to eliminate race conditions with many concurrent writers. Afterwards, the file is overwritten with a new version (challenge). After that, the thread waits for 2 seconds. If there is a watchdog thread protecting the file, he will overwrite the change and this process will fail to lock the database. However, if there is no watchdog thread, the lock file will still be as written by this thread. In this case, the file is deleted and atomically created again. The watchdog thread is started in this case and the file is locked.

This algorithm is tested with over 100 concurrent threads. In some cases, when there are many concurrent threads trying to lock the database, they block each other (meaning the file cannot be locked by any of them) for some time. However, the file never gets locked by two threads at the same time. However using that many concurrent threads / processes is not the common use case. Generally, an application should throw an error to the user if it cannot open a database, and not try again in a (fast) loop.

## **File Locking Method 'Socket'**

There is a second locking mechanism implemented, but disabled by default. To use it, append ;FILE\_LOCK=SOCKET to the database URL. The algorithm is:

- If the lock file does not exist, it is created. Then a server socket is opened on a defined port, and kept open. The port and IP address of the process that opened the database is written into the lock file.
- If the lock file exists, and the lock method is 'file', then the software switches to the 'file' method.
- If the lock file exists, and the lock method is 'socket', then the process checks if the port is in use. If the original process is still running, the port is in use and this process throws an exception (database is in use). If the original process died (for example due to a power failure, or abnormal termination of the virtual machine), then the port was released. The new process deletes the lock file and starts again.

This method does not require a watchdog thread actively polling (reading) the same file every second. The problem with this method is, if the file is stored on a network share, two processes (running on different computers) could still open the same database files, if they do not have a direct TCP/IP connection.

## **File Locking Method 'FS'**

This is the default mode for version 1.4 and newer. This database file locking mechanism uses native file system lock on the database file. No \*.lock.db file is created in this case, and no background thread is started. This mechanism may not work on all systems as expected. Some systems allow to lock the same file multiple times within the same virtual machine, and on some system native file locking is not supported or files are not unlocked after a power failure.

To enable this feature, append ;FILE\_LOCK=FS to the database URL.

This feature is relatively new. When using it for production, please ensure your system does in fact lock files as expected.

# Using Passwords

## Using Secure Passwords

Remember that weak passwords can be broken regardless of the encryption and security protocols. Don't use passwords that can be found in a dictionary. Appending numbers does not make passwords secure. A way to create good passwords that can be remembered is: take the first letters of a sentence, use upper and lower case characters, and creatively include special characters (but it's more important to use a long password than to use special characters). Example:

i'sE2rtPiUKtT from the sentence it's easy to remember this password if you know the trick.

## Passwords: Using Char Arrays instead of Strings

Java strings are immutable objects and cannot be safely 'destroyed' by the application. After creating a string, it will remain in the main memory of the computer at least until it is garbage collected. The garbage collection cannot be controlled by the application, and even if it is garbage collected the data may still remain in memory. It might also be possible that the part of memory containing the password is swapped to disk (if not enough main memory is available), which is a problem if the attacker has access to the swap file of the operating system.

It is a good idea to use char arrays instead of strings for passwords. Char arrays can be cleared (filled with zeros) after use, and therefore the password will not be stored in the swap file.

This database supports using char arrays instead of string to pass user and file passwords. The following code can be used to do that:

```
import java.sql.*;
import java.util.*;
public class Test {
    public static void main(String[] args) throws Exception {
        String url = "jdbc:h2:~/test";
        Properties prop = new Properties();
        prop.setProperty("user", "sa");
        System.out.print("Password?");
        char[] password = System.console().readPassword();
        prop.put("password", password);
        Connection conn = null;
        try {
            conn = DriverManager.getConnection(url, prop);
```

```
    } finally {  
        Arrays.fill(password, (char) 0);  
    }  
    conn.close();  
}  
}
```

When using Swing, use `javax.swing.JPasswordField`.

## Passing the User Name and/or Password in the URL

Instead of passing the user name as a separate parameter as in `Connection conn = DriverManager.getConnection("jdbc:h2:~/test", "sa", "123");` the user name (and/or password) can be supplied in the URL itself: `Connection conn = DriverManager.getConnection("jdbc:h2:~/test;USER=sa;PASSWORD=123");` The settings in the URL override the settings passed as a separate parameter.

## Password Hash

Sometimes the database password needs to be stored in a configuration file (for example in the `web.xml` file). In addition to connecting with the plain text password, this database supports connecting with the password hash. This means that only the hash of the password (and not the plain text password) needs to be stored in the configuration file. This will only protect others from reading or re-constructing the plain text password (even if they have access to the configuration file); it does not protect others from accessing the database using the password hash.

To connect using the password hash instead of plain text password, append `;PASSWORD_HASH=TRUE` to the database URL, and replace the password with the password hash. To calculate the password hash from a plain text password, run the following command within the H2 Console tool: `@password_hash <upperCaseUserName> <password>`. As an example, if the user name is `sa` and the password is `test`, run the command `@password_hash SA test`. Then use the resulting password hash as you would use the plain text password. When using an encrypted database, then the user password and file password need to be hashed separately. To calculate the hash of the file password, run: `@password_hash file <filePassword>`.



# Protection against SQL Injection

## What is SQL Injection

This database engine provides a solution for the security vulnerability known as 'SQL Injection'. Here is a short description of what SQL injection means. Some applications build SQL statements with embedded user input such as:

```
String sql = "SELECT * FROM USERS WHERE PASSWORD='"+pwd+"'";  
ResultSet rs = conn.createStatement().executeQuery(sql);
```

If this mechanism is used anywhere in the application, and user input is not correctly filtered or encoded, it is possible for a user to inject SQL functionality or statements by using specially built input such as (in this example) this password: ' OR '='. In this case the statement becomes:

```
SELECT * FROM USERS WHERE PASSWORD=" OR "=";
```

Which is always true no matter what the password stored in the database is. For more information about SQL Injection, see [Glossary and Links](#).

## Disabling Literals

SQL Injection is not possible if user input is not directly embedded in SQL statements. A simple solution for the problem above is to use a prepared statement:

```
String sql = "SELECT * FROM USERS WHERE PASSWORD=?";  
PreparedStatement prep = conn.prepareStatement(sql);  
prep.setString(1, pwd);  
ResultSet rs = prep.executeQuery();
```

This database provides a way to enforce usage of parameters when passing user input to the database. This is done by disabling embedded literals in SQL statements. To do this, execute the statement:

```
SET ALLOW_LITERALS NONE;
```

Afterwards, SQL statements with text and number literals are not allowed any more. That means, SQL statement of the form WHERE NAME='abc' or WHERE CustomerId=10 will fail. It is still possible to use prepared statements and parameters as described above. Also, it is still possible to generate SQL statements dynamically, and use the Statement API, as long as the SQL statements do not include literals. There is also a second mode where number

literals are allowed: SET ALLOW\_LITERALS NUMBERS. To allow all literals, execute SET ALLOW\_LITERALS ALL (this is the default setting). Literals can only be enabled or disabled by an administrator.

## Using Constants

Disabling literals also means disabling hard-coded 'constant' literals. This database supports defining constants using the CREATE CONSTANT command. Constants can be defined only when literals are enabled, but used even when literals are disabled. To avoid name clashes with column names, constants can be defined in other schemas:

```
CREATE SCHEMA CONST AUTHORIZATION SA;  
CREATE CONSTANT CONST.ACTIVE VALUE 'Active';  
CREATE CONSTANT CONST.INACTIVE VALUE 'Inactive';  
SELECT * FROM USERS WHERE TYPE=CONST.ACTIVE;
```

Even when literals are enabled, it is better to use constants instead of hard-coded number or text literals in queries or views. With constants, typos are found at compile time, the source code is easier to understand and change.

## Using the ZERO() Function

It is not required to create a constant for the number 0 as there is already a built-in function ZERO():

```
SELECT * FROM USERS WHERE LENGTH(PASSWORD)=ZERO();
```

## Protection against Remote Access

By default this database does not allow connections from other machines when starting the H2 Console, the TCP server, or the PG server. Remote access can be enabled using the command line options -webAllowOthers, -tcpAllowOthers, -pgAllowOthers.

If you enable remote access using -tcpAllowOthers or -pgAllowOthers, please also consider using the options -baseDir, so that remote users can not create new databases or access existing databases with weak passwords. When using the option -baseDir, only databases within that directory may be accessed. Ensure the existing accessible databases are protected using strong passwords.

If you enable remote access using -webAllowOthers, please ensure the web server can only be accessed from trusted networks. The options -baseDir don't

protect access to the tools section, prevent remote shutdown of the web server, changes to the preferences, the saved connection settings, or access to other databases accessible from the system.

## Restricting Class Loading and Usage

By default there is no restriction on loading classes and executing Java code for admins. That means an admin may call system functions such as `System.setProperty` by executing:

```
CREATE ALIAS SET_PROPERTY FOR "java.lang.System.setProperty";  
CALL SET_PROPERTY('abc', '1');  
CREATE ALIAS GET_PROPERTY FOR "java.lang.System.getProperty";  
CALL GET_PROPERTY('abc');
```

To restrict users (including admins) from loading classes and executing code, the list of allowed classes can be set in the system property `h2.allowedClasses` in the form of a comma separated list of classes or patterns (items ending with `*`). By default all classes are allowed. Example:

```
java -Dh2.allowedClasses=java.lang.Math,com.acme.*
```

This mechanism is used for all user classes, including database event listeners, trigger classes, user-defined functions, user-defined aggregate functions, and JDBC driver classes (with the exception of the H2 driver) when using the H2 Console.

## Security Protocols

The following paragraphs document the security protocols used in this database. These descriptions are very technical and only intended for security experts that already know the underlying security primitives.

### User Password Encryption

When a user tries to connect to a database, the combination of user name, @, and password are hashed using SHA-256, and this hash value is transmitted to the database. This step does not protect against an attacker that re-uses the value if he is able to listen to the (unencrypted) transmission between the client and the server. But, the passwords are never transmitted as plain text, even when using an unencrypted connection between client and server. That means if a user reuses the same password for different things, this password is still protected up

to some point. See also 'RFC 2617 - HTTP Authentication: Basic and Digest Access Authentication' for more information.

When a new database or user is created, a new random salt value is generated. The size of the salt is 64 bits. Using the random salt reduces the risk of an attacker pre-calculating hash values for many different (commonly used) passwords.

The combination of user-password hash value (see above) and salt is hashed using SHA-256. The resulting value is stored in the database. When a user tries to connect to the database, the database combines user-password hash value with the stored salt value and calculates the hash value. Other products use multiple iterations (hash the hash value again and again), but this is not done in this product to reduce the risk of denial of service attacks (where the attacker tries to connect with bogus passwords, and the server spends a lot of time calculating the hash value for each password). The reasoning is: if the attacker has access to the hashed passwords, he also has access to the data in plain text, and therefore does not need the password any more. If the data is protected by storing it on another computer and only accessible remotely, then the iteration count is not required at all.

## **File Encryption**

The database files can be encrypted using the AES-128 algorithm.

When a user tries to connect to an encrypted database, the combination of file@ and the file password is hashed using SHA-256. This hash value is transmitted to the server.

When a new database file is created, a new cryptographically secure random salt value is generated. The size of the salt is 64 bits. The combination of the file password hash and the salt value is hashed 1024 times using SHA-256. The reason for the iteration is to make it harder for an attacker to calculate hash values for common passwords.

The resulting hash value is used as the key for the block cipher algorithm. Then, an initialization vector (IV) key is calculated by hashing the key again using SHA-256. This is to make sure the IV is unknown to the attacker. The reason for using a secret IV is to protect against watermark attacks.

Before saving a block of data (each block is 8 bytes long), the following operations are executed: first, the IV is calculated by encrypting the block number with the IV key (using the same block cipher algorithm). This IV is combined with the plain text using XOR. The resulting data is encrypted using the AES-128 algorithm.

When decrypting, the operation is done in reverse. First, the block is decrypted using the key, and then the IV is calculated combined with the decrypted text using XOR.

Therefore, the block cipher mode of operation is CBC (cipher-block chaining), but each chain is only one block long. The advantage over the ECB (electronic codebook) mode is that patterns in the data are not revealed, and the advantage over multi block CBC is that flipped cipher text bits are not propagated to flipped plaintext bits in the next block.

Database encryption is meant for securing the database while it is not in use (stolen laptop and so on). It is not meant for cases where the attacker has access to files while the database is in use. When he has write access, he can for example replace pieces of files with pieces of older versions and manipulate data like this.

File encryption slows down the performance of the database engine. Compared to unencrypted mode, database operations take about 2.5 times longer using AES (embedded mode).

## **Wrong Password / User Name Delay**

To protect against remote brute force password attacks, the delay after each unsuccessful login gets double as long. Use the system properties `h2.delayWrongPasswordMin` and `h2.delayWrongPasswordMax` to change the minimum (the default is 250 milliseconds) or maximum delay (the default is 4000 milliseconds, or 4 seconds). The delay only applies for those using the wrong password. Normally there is no delay for a user that knows the correct password, with one exception: after using the wrong password, there is a delay of up to (randomly distributed) the same delay as for a wrong password. This is to protect against parallel brute force attacks, so that an attacker needs to wait for the whole delay. Delays are synchronized. This is also required to protect against parallel attacks.

There is only one exception message for both wrong user and for wrong password, to make it harder to get the list of user names. It is not possible from the stack trace to see if the user name was wrong or the password.

## **HTTPS Connections**

The web server supports HTTP and HTTPS connections using `SSLServerSocket`. There is a default self-certified certificate to support an easy starting point, but custom certificates are supported as well.

## TLS Connections

Remote TLS connections are supported using the Java Secure Socket Extension (SSLServerSocket, SSLSocket). By default, anonymous TLS is enabled.

To use your own keystore, set the system properties `javax.net.ssl.keyStore` and `javax.net.ssl.keyStorePassword` before starting the H2 server and client. See also [Customizing the Default Key and Trust Stores, Store Types, and Store Passwords](#) for more information.

To disable anonymous TLS, set the system property `h2.enableAnonymousTLS` to `false`.

## Universally Unique Identifiers (UUID)

This database supports UUIDs. Also supported is a function to create new UUIDs using a cryptographically strong pseudo random number generator. With random UUIDs, the chance of two having the same value can be calculated using the probability theory. See also 'Birthday Paradox'. Standardized randomly generated UUIDs have 122 random bits. 4 bits are used for the version (Randomly generated UUID), and 2 bits for the variant (Leach-Salz). This database supports generating such UUIDs using the built-in function `RANDOM_UUID()` or `UUID()`. Here is a small program to estimate the probability of having two identical UUIDs after generating a number of values:

```
public class Test {
    public static void main(String[] args) throws Exception {
        double x = Math.pow(2, 122);
        for (int i = 35; i < 62; i++) {
            double n = Math.pow(2, i);
            double p = 1 - Math.exp(-(n * n) / 2 / x);
            System.out.println("2^" + i + "=" + (1L << i) +
                " probability: 0" +
                String.valueOf(1 + p).substring(1));
        }
    }
}
```

Some values are:

Number of UIIs	Probability of Duplicates
$2^{36}=68'719'476'736$	0.000'000'000'000'000'4
$2^{41}=2'199'023'255'552$	0.000'000'000'000'4

$2^{46}=70'368'744'177'664$	$0.000'000'000'4$
-----------------------------	-------------------

To help non-mathematicians understand what those numbers mean, here a comparison: one's annual risk of being hit by a meteorite is estimated to be one chance in 17 billion, that means the probability is about 0.000'000'000'06.

## Spatial Features

H2 supports the geometry data type and spatial indexes. Here is an example SQL script to create a table with a spatial column and index:

```
CREATE TABLE GEO_TABLE(GID SERIAL, THE_GEOM GEOMETRY);
INSERT INTO GEO_TABLE(THE_GEOM) VALUES
  ('POINT(500 505)'),
  ('LINESTRING(550 551, 525 512, 565 566)'),
  ('POLYGON ((550 521, 580 540, 570 564, 512 566, 550 521))');
CREATE SPATIAL INDEX GEO_TABLE_SPATIAL_INDEX ON
GEO_TABLE(THE_GEOM);
```

To query the table using geometry envelope intersection, use the operation &&, as in PostGIS:

```
SELECT * FROM GEO_TABLE
WHERE THE_GEOM && 'POLYGON ((490 490, 536 490, 536 515, 490 515, 490 490))';
```

You can verify that the spatial index is used using the "explain plan" feature:

```
EXPLAIN SELECT * FROM GEO_TABLE
WHERE THE_GEOM && 'POLYGON ((490 490, 536 490, 536 515, 490 515, 490 490))';
-- Result
SELECT
  GEO_TABLE.GID,
  GEO_TABLE.THE_GEOM
FROM PUBLIC.GEO_TABLE
/* PUBLIC.GEO_TABLE_SPATIAL_INDEX:
  THE_GEOM && 'POLYGON ((490 490, 536 490, 536 515, 490 515, 490 490))'
*/
WHERE INTERSECTS(THE_GEOM,
  'POLYGON ((490 490, 536 490, 536 515, 490 515, 490 490))')
```

For persistent databases, the spatial index is stored on disk; for in-memory databases, the index is kept in memory.

## Recursive Queries

H2 has experimental support for recursive queries using so called "common table expressions" (CTE). Examples:

```
WITH RECURSIVE T(N) AS (  
  SELECT 1  
  UNION ALL  
  SELECT N+1 FROM T WHERE N<10  
)  
SELECT * FROM T;  
-- returns the values 1 .. 10
```

```
WITH RECURSIVE T(N) AS (  
  SELECT 1  
  UNION ALL  
  SELECT N*2 FROM T WHERE N<10  
)  
SELECT * FROM T;  
-- returns the values 1, 2, 4, 8, 16
```

```
CREATE TABLE FOLDER(ID INT PRIMARY KEY, NAME VARCHAR(255), PARENT  
INT);
```

```
INSERT INTO FOLDER VALUES(1, null, null), (2, 'src', 1),  
(3, 'main', 2), (4, 'org', 3), (5, 'test', 2);
```

```
WITH LINK(ID, NAME, LEVEL) AS (  
  SELECT ID, NAME, 0 FROM FOLDER WHERE PARENT IS NULL  
  UNION ALL  
  SELECT FOLDER.ID, IFNULL(LINK.NAME || '/', '') || FOLDER.NAME, LEVEL + 1  
  FROM LINK INNER JOIN FOLDER ON LINK.ID = FOLDER.PARENT  
)  
SELECT NAME FROM LINK WHERE NAME IS NOT NULL ORDER BY ID;  
-- src  
-- src/main  
-- src/main/org  
-- src/test
```

Limitations: Recursive queries need to be of the type UNION ALL, and the recursion needs to be on the second part of the query. No tables or views with the name of the table expression may exist. Different table expression names need to be used when using multiple distinct table expressions within the same transaction and for the same session. All columns of the table expression are of type VARCHAR, and may need to be cast to the required data type. Views with



recursive queries are not supported. Subqueries and INSERT INTO ... FROM with recursive queries are not supported. Parameters are only supported within the last SELECT statement (a workaround is to use session variables like @start within the table expression). The syntax is:

```
WITH RECURSIVE recursiveQueryName(columnName, ...) AS (  
    nonRecursiveSelect  
    UNION ALL  
    recursiveSelect  
)  
select
```

## Settings Read from System Properties

Some settings of the database can be set on the command line using -DpropertyName=value. It is usually not required to change those settings manually. The settings are case sensitive. Example:

```
java -Dh2.serverCachedObjects=256 org.h2.tools.Server
```

The current value of the settings can be read in the table INFORMATION\_SCHEMA.SETTINGS.

For a complete list of settings, see [SysProperties](#).

## Setting the Server Bind Address

Usually server sockets accept connections on any/all local addresses. This may be a problem on multi-homed hosts. To bind only to one address, use the system property h2.bindAddress. This setting is used for both regular server sockets and for TLS server sockets. IPv4 and IPv6 address formats are supported.

## Pluggable File System

This database supports a pluggable file system API. The file system implementation is selected using a file name prefix. Internally, the interfaces are very similar to the Java 7 NIO2 API, but do not (yet) use or require Java 7. The following file systems are included:

- zip: read-only zip-file based file system. Format: zip:/zipFileName!/fileName.
- split: file system that splits files in 1 GB files (stackable with other file systems).

- nio: file system that uses FileChannel instead of RandomAccessFile (faster in some operating systems).
- nioMapped: file system that uses memory mapped files (faster in some operating systems). Please note that there currently is a file size limitation of 2 GB when using this file system. To work around this limitation, combine it with the split file system: split:nioMapped:test.
- async: experimental file system that uses AsynchronousFileChannel instead of RandomAccessFile (faster in some operating systems).
- memFS: in-memory file system (slower than mem; experimental; mainly used for testing the database engine itself).
- memLZF: compressing in-memory file system (slower than memFS but uses less memory; experimental; mainly used for testing the database engine itself).
- nioMemFS: stores data outside of the VM's heap - useful for large memory DBs without incurring GC costs.
- nioMemLZF: stores compressed data outside of the VM's heap - useful for large memory DBs without incurring GC costs. Use "nioMemLZF:12:" to tweak the % of blocks that are stored uncompressed. If you size this to your working set correctly, compressed storage is roughly the same performance as uncompressed. The default value is 1%.

As an example, to use the nio file system with PageStore storage engine, use the following database URL: jdbc:h2:nio:~/test;MV\_STORE=FALSE. With MVStore storage engine nio file system is used by default.

To register a new file system, extend the classes org.h2.store.fs.FilePath, FileBase, and call the method FilePath.register before using it.

For input streams (but not for random access files), URLs may be used in addition to the registered file systems. Example:  
 jar:file:///c:/temp/example.zip!/org/example/nested.csv. To read a stream from the classpath, use the prefix classpath:, as in  
 classpath:/org/h2/samples/newsfeed.sql.

## Split File System

The file system prefix split: is used to split logical files into multiple physical files, for example so that a database can get larger than the maximum file system size of the operating system. If the logical file is larger than the maximum file size, then the file is split as follows:

- <fileName> (first block, is always created)
- <fileName>.1.part (second block)

More physical files (\*.2.part, \*.3.part) are automatically created / deleted if needed. The maximum physical file size of a block is  $2^{30}$  bytes, which is also called 1 GiB or 1 GB. However this can be changed if required, by specifying the block size in the file name. The file name format is: split:<x>:<fileName> where the file size per block is  $2^x$ . For 1 MiB block sizes, use  $x = 20$  (because  $2^{20}$  is 1 MiB). The following file name means the logical file is split into 1 MiB blocks: split:20:test.h2.db. An example database URL for this case is jdbc:h2:split:20:~/test.

## Database Upgrade

In version 1.2, H2 introduced a new file store implementation which is incompatible to the one used in versions < 1.2. To automatically convert databases to the new file store, it is necessary to include an additional jar file. The file can be found at [http://h2database.com/h2mig\\_pagestore\\_addon.jar](http://h2database.com/h2mig_pagestore_addon.jar). If this file is in the classpath, every connect to an older database will result in a conversion process.

The conversion itself is done internally via 'script to' and 'runscript from'. After the conversion process, the files will be renamed from

- dbName.data.db to dbName.data.db.backup
- dbName.index.db to dbName.index.db.backup

by default. Also, the temporary script will be written to the database directory instead of a temporary directory. Both defaults can be customized via

- org.h2.upgrade.DbUpgrade.setDeleteOldDb(boolean)
- org.h2.upgrade.DbUpgrade.setScriptInTmpDir(boolean)

prior opening a database connection.

Since version 1.2.140 it is possible to let the old h2 classes (v 1.2.128) connect to the database. The automatic upgrade .jar file must be present, and the URL must start with jdbc:h2v1\_1: (the JDBC driver class is org.h2.upgrade.v1\_1.Driver). If the database should automatically connect using the old version if a database with the old format exists (without upgrade), and use the new version otherwise, then append ;NO\_UPGRADE=TRUE to the database URL. Please note the old driver did not process the system property "h2.baseDir" correctly, so that using this setting is not supported when upgrading.

## Java Objects Serialization

Java objects serialization is enabled by default for columns of type OTHER, using standard Java serialization/deserialization semantics.

To disable this feature set the system property `h2.serializeJavaObject=false` (default: true).

Serialization and deserialization of java objects is customizable both at system level and at database level providing a [JavaObjectSerializer](#) implementation:

- At system level set the system property `h2.javaObjectSerializer` with the Fully Qualified Name of the `JavaObjectSerializer` interface implementation. It will be used over the entire JVM session to (de)serialize java objects being stored in column of type OTHER. Example  
`h2.javaObjectSerializer=com.acme.SerializerClassName.`
- At database level execute the SQL statement `SET JAVA_OBJECT_SERIALIZER 'com.acme.SerializerClassName'` or append `;JAVA_OBJECT_SERIALIZER='com.acme.SerializerClassName'` to the database URL:  
`jdbc:h2:~/test;JAVA_OBJECT_SERIALIZER='com.acme.SerializerClassName'.`

Please note that this SQL statement can only be executed before any tables are defined.

## Custom Data Types Handler API

It is possible to extend the type system of the database by providing your own implementation of minimal required API basically consisting of type identification and conversion routines.

In order to enable this feature, set the system property `h2.customDataTypesHandler` (default: null) to the fully qualified name of the class providing [CustomDataTypesHandler](#) interface implementation. The instance of that class will be created by H2 and used to:

- resolve the names and identifiers of extrinsic data types.
- convert values of extrinsic data types to and from values of built-in types.
- provide order of the data types.

This is a system-level setting, i.e. affects all the databases.

**Note:** Please keep in mind that this feature may not possibly provide the same ABI stability level as other features as it exposes many of the H2 internals. You may be required to update your code occasionally due to internal changes in H2 if you are going to use this feature.

## Limits and Limitations

This database has the following known limitations:

- Database file size limit: 4 TB (using the default page size of 2 KB) or higher (when using a larger page size). This limit is including CLOB and BLOB data.
- The maximum file size for FAT or FAT32 file systems is 4 GB. That means when using FAT or FAT32, the limit is 4 GB for the data. This is the limitation of the file system. The database does provide a workaround for this problem, it is to use the file name prefix split:. In that case files are split into files of 1 GB by default. An example database URL is:  
jdbc:h2:split:~/test.
- The maximum number of rows per table is  $2^{64}$ .
- The maximum number of open transactions is 65535.
- Main memory requirements: The larger the database, the more main memory is required. With the current storage mechanism (the page store), the minimum main memory required is around 1 MB for each 8 GB database file size.
- Limit on the complexity of SQL statements. Statements of the following form will result in a stack overflow exception:

```
SELECT * FROM DUAL WHERE X = 1
OR X = 2 OR X = 2 OR X = 2 OR X = 2 OR X = 2
-- repeat previous line 500 times --
```

- There is no limit for the following entities, except the memory and storage capacity: maximum identifier length (table name, column name, and so on); maximum number of tables, columns, indexes, triggers, and other database objects; maximum statement length, number of parameters per statement, tables per statement, expressions in order by, group by, having, and so on; maximum rows per query; maximum columns per table, columns per index, indexes per table, lob columns per table, and so on; maximum row length, index row length, select row length; maximum length of a varchar column, decimal column, literal in a statement.
- Querying from the metadata tables is slow if there are many tables (thousands).
- For limitations on data types, see the documentation of the respective Java data type or the data type documentation of this database.

## Glossary and Links

Term	Description
AES-128	A block encryption algorithm. See also: <a href="#">Wikipedia: Advanced Encryption Standard</a>
Birthday Paradox	Describes the higher than expected probability that two persons in a room have the same birthday. Also valid for randomly generated UUIDs. See also: <a href="#">Wikipedia: Birthday problem</a>

Digest	Protocol to protect a password (but not to protect data). See also: <a href="#">RFC 2617: HTTP Digest Access Authentication</a>
HTTPS	A protocol to provide security to HTTP connections. See also: <a href="#">RFC 2818: HTTP Over TLS</a>
Modes of Operation	<a href="#">Wikipedia: Block cipher mode of operation</a>
Salt	Random number to increase the security of passwords. See also: <a href="#">Wikipedia: Key derivation function</a>
SHA-256	A cryptographic one-way hash function. See also: <a href="#">Wikipedia: Secure Hash Algorithms</a>
SQL Injection	A security vulnerability where an application embeds SQL statements or expressions in user input. See also: <a href="#">Wikipedia: SQL injection</a>
Watermark Attack	Security problem of certain encryption programs where the existence of certain data can be proven without decrypting. For more information, search in the internet for 'watermark attack cryptoloop'
SSL/TLS	Secure Sockets Layer / Transport Layer Security. See also: <a href="#">Java Secure Socket Extension (JSSE)</a>

## Index

### Commands (Data Manipulation)

SELECT  
INSERT  
UPDATE  
DELETE  
BACKUP  
CALL  
EXPLAIN  
MERGE  
MERGE USING  
RUNSCRIPT  
SCRIPT  
SHOW  
Explicit table  
Table value  
WITH

### Commands (Data Definition)

ALTER INDEX RENAME  
ALTER SCHEMA RENAME  
ALTER SEQUENCE  
ALTER TABLE ADD  
ALTER TABLE ADD CONSTRAINT  
ALTER TABLE RENAME CONSTRAINT  
ALTER TABLE ALTER COLUMN  
ALTER TABLE DROP COLUMN  
ALTER TABLE DROP CONSTRAINT  
ALTER TABLE SET  
ALTER TABLE RENAME  
ALTER USER ADMIN  
ALTER USER RENAME  
ALTER USER SET PASSWORD  
ALTER VIEW RECOMPILE  
ALTER VIEW RENAME  
ANALYZE  
COMMENT

CREATE AGGREGATE  
CREATE ALIAS  
CREATE CONSTANT  
CREATE DOMAIN  
CREATE INDEX  
CREATE LINKED TABLE  
CREATE ROLE  
CREATE SCHEMA  
CREATE SEQUENCE  
CREATE TABLE  
CREATE TRIGGER  
CREATE USER  
CREATE VIEW  
DROP AGGREGATE  
DROP ALIAS  
DROP ALL OBJECTS  
DROP CONSTANT  
DROP DOMAIN  
DROP INDEX  
DROP ROLE  
DROP SCHEMA  
DROP SEQUENCE  
DROP TABLE  
DROP TRIGGER  
DROP USER  
DROP VIEW  
TRUNCATE TABLE

## **Commands (Other)**

CHECKPOINT  
CHECKPOINT SYNC  
COMMIT  
COMMIT TRANSACTION  
GRANT RIGHT  
GRANT ALTER ANY SCHEMA  
GRANT ROLE  
HELP  
PREPARE COMMIT  
REVOKE RIGHT  
REVOKE ROLE  
ROLLBACK  
ROLLBACK TRANSACTION  
SAVEPOINT



SET @  
SET ALLOW\_LITERALS  
SET AUTOCOMMIT  
SET CACHE\_SIZE  
SET CLUSTER  
SET BINARY\_COLLATION  
SET UUID\_COLLATION  
SET BUILTIN\_ALIAS\_OVERRIDE  
SET COLLATION  
SET COMPRESS\_LOB  
SET DATABASE\_EVENT\_LISTENER  
SET DB\_CLOSE\_DELAY  
SET DEFAULT\_LOCK\_TIMEOUT  
SET DEFAULT\_TABLE\_TYPE  
SET EXCLUSIVE  
SET IGNORECASE  
SET JAVA\_OBJECT\_SERIALIZER  
SET LAZY\_QUERY\_EXECUTION  
SET LOG  
SET LOCK\_MODE  
SET LOCK\_TIMEOUT  
SET MAX\_LENGTH\_INPLACE\_LOB  
SET MAX\_LOG\_SIZE  
SET MAX\_MEMORY\_ROWS  
SET MAX\_MEMORY\_UNDO  
SET MAX\_OPERATION\_MEMORY  
SET MODE  
SET MULTI\_THREADED  
SET OPTIMIZE\_REUSE\_RESULTS  
SET PASSWORD  
SET QUERY\_STATISTICS  
SET QUERY\_STATISTICS\_MAX\_ENTRIES  
SET QUERY\_TIMEOUT  
SET REFERENTIAL\_INTEGRITY  
SET RETENTION\_TIME  
SET SALT\_HASH  
SET SCHEMA  
SET SCHEMA\_SEARCH\_PATH  
SET THROTTLE  
SET TRACE\_LEVEL  
SET TRACE\_MAX\_FILE\_SIZE  
SET UNDO\_LOG  
SET WRITE\_DELAY  
SHUTDOWN

## Datetime fields

Datetime field  
Year field  
Month field  
Day of month field  
Hour field  
Minute field  
Second field  
Millisecond field  
Microsecond field  
Nanosecond field  
Timezone hour field  
Timezone minute field  
Day of week field  
ISO week year field  
ISO day of week field  
Week of year field  
ISO week of year field  
Quarter field  
Day of year field  
Epoch field

## Other Grammar

Alias  
And Condition  
Array  
Boolean  
Bytes  
Case  
Case When  
Cipher  
Column Definition  
Comments  
Compare  
Condition  
Condition Right Hand Side  
Constraint  
Constraint Name Definition  
Csv Options  
Data Type  
Date  
Decimal  
Digit

Dollar Quoted String  
Expression  
Factor  
Hex  
Hex Number  
Index Column  
Insert columns and source  
Insert values  
Int  
Long  
Merge when clause  
Merge when matched clause  
Merge when not matched clause  
Name  
Null  
Number  
Numeric  
Operand  
Order  
Quoted Name  
Referential Constraint  
Referential Action  
Script Compression Encryption  
Row value expression  
Select Expression  
Sequence options  
Set clause list  
String  
Summand  
Table Expression  
Within group specification  
Wildcard expression  
Window name or specification  
Window specification  
Window frame  
Window frame preceding  
Window frame bound  
Term  
Column  
Time  
Timestamp  
Timestamp with time zone  
Date and time  
INTERVAL YEAR  
INTERVAL MONTH

INTERVAL DAY  
INTERVAL HOUR  
INTERVAL MINUTE  
INTERVAL SECOND  
INTERVAL YEAR TO MONTH  
INTERVAL DAY TO HOUR  
INTERVAL DAY TO MINUTE  
INTERVAL DAY TO SECOND  
INTERVAL HOUR TO MINUTE  
INTERVAL HOUR TO SECOND  
INTERVAL MINUTE TO SECOND  
Interval  
Value

## Commands (Data Manipulation)

### SELECT

```
SELECT [ TOP term [ PERCENT ] [ WITH TIES ] ]  
[ DISTINCT [ ON ( expression [,...] ) ] | ALL ]  
selectExpression [,...]  
[ FROM tableExpression [,...] ]  
[ WHERE expression ]  
[ GROUP BY expression [,...] ] [ HAVING expression ]  
[ WINDOW { { windowName AS windowSpecification } [,...] } ]  
[ QUALIFY expression ]  
[ { UNION [ ALL ] | EXCEPT | MINUS | INTERSECT } select ]  
[ ORDER BY order [,...] ]  
[ LIMIT expression [ OFFSET expression ] [ SAMPLE_SIZE rowCountInt ] ]  
[ [ OFFSET expression { ROW | ROWS } ]  
[ FETCH { FIRST | NEXT } [ expression [ PERCENT ] ] { ROW | ROWS }  
  { ONLY | WITH TIES } ] [ SAMPLE_SIZE rowCountInt ] ]  
[ FOR UPDATE ]
```

Selects data from a table or multiple tables.

Command is executed in the following logical order:

1. Data is taken from table value expressions that are specified in the FROM clause, joins are executed. If FROM clause is not specified a single row is constructed.
2. WHERE filters rows. Aggregate or window functions are not allowed in this clause.

3. GROUP BY groups the result by the given expression(s). If GROUP BY clause is not specified, but non-window aggregate functions are used or HAVING is specified all rows are grouped together.
4. Aggregate functions are evaluated, SAMPLE\_SIZE limits the number of rows read.
5. HAVING filters rows after grouping and evaluation of aggregate functions. Non-window aggregate functions are allowed in this clause.
6. Window functions are evaluated.
7. QUALIFY filters rows after evaluation of window functions. Aggregate and window functions are allowed in this clause.
8. DISTINCT removes duplicates. If DISTINCT ON is used only the specified expressions are checked for duplicates.
9. UNION, EXCEPT (MINUS), and INTERSECT combine the result of this query with the results of another query. Multiple set operators (UNION, INTERSECT, MINUS, EXCEPT) are evaluated from left to right. For compatibility with other databases and future versions of H2 please use parentheses.
10. ORDER BY sorts the result by the given column(s) or expression(s).
11. Number of rows in output can be limited either with standard OFFSET / FETCH, with non-standard LIMIT / OFFSET, or with non-standard TOP clauses. Different clauses cannot be used together. OFFSET specifies how many rows to skip. Please note that queries with high offset values can be slow. FETCH FIRST/NEXT, LIMIT or TOP limits the number of rows returned by the query (no limit if null or smaller than zero). If PERCENT is specified number of rows is specified as a percent of the total number of rows and should be an integer value between 0 and 100 inclusive. WITH TIES can be used only together with ORDER BY and means that all additional rows that have the same sorting position as the last row will be also returned.

WINDOW clause specifies window definitions for window functions and window aggregate functions. This clause can be used to reuse the same definition in multiple functions.

If FOR UPDATE is specified, the tables or rows are locked for writing. This clause is not allowed in DISTINCT queries and in queries with non-window aggregates, GROUP BY, or HAVING clauses. When using default MVStore engine only the selected rows are locked as in an UPDATE statement; locking behavior for rows that were excluded from result using OFFSET / FETCH is undefined. With PageStore engine the whole tables are locked.

Example:

```
SELECT * FROM TEST;  
SELECT * FROM TEST ORDER BY NAME;  
SELECT ID, COUNT(*) FROM TEST GROUP BY ID;  
SELECT NAME, COUNT(*) FROM TEST GROUP BY NAME HAVING COUNT(*) > 2;  
SELECT 'ID' COL, MAX(ID) AS MAX FROM TEST UNION SELECT 'NAME',  
MAX(NAME) FROM TEST;  
SELECT * FROM TEST OFFSET 1000 ROWS FETCH FIRST 1000 ROWS ONLY;  
SELECT A, B FROM TEST ORDER BY A FETCH FIRST 10 ROWS WITH TIES;  
SELECT * FROM (SELECT ID, COUNT(*) FROM TEST  
    GROUP BY ID UNION SELECT NULL, COUNT(*) FROM TEST)  
    ORDER BY 1 NULLS LAST;  
SELECT DISTINCT C1, C2 FROM TEST;  
SELECT DISTINCT ON(C1) C1, C2 FROM TEST ORDER BY C1;
```

## INSERT

```
INSERT INTO tableName insertColumnsAndSource
```

Inserts a new row / new rows into a table.

When using DIRECT, then the results from the query are directly applied in the target table without any intermediate step.

When using SORTED, b-tree pages are split at the insertion point. This can improve performance and reduce disk usage.

Example:

```
INSERT INTO TEST VALUES(1, 'Hello')
```

## UPDATE

```
UPDATE tableName [ [ AS ] newTableAlias ] SET setClauseList  
[ WHERE expression ] [ ORDER BY order [,...] ] [ LIMIT expression ]
```

Updates data in a table. ORDER BY is supported for MySQL compatibility, but it is ignored.

Example:

```
UPDATE TEST SET NAME='Hi' WHERE ID=1;  
UPDATE PERSON P SET NAME=(SELECT A.NAME FROM ADDRESS A WHERE  
A.ID=P.ID);
```

## DELETE

```
DELETE [ TOP term ] FROM tableName [ WHERE expression ] [ LIMIT term ]
```

Deletes rows from a table. If TOP or LIMIT is specified, at most the specified number of rows are deleted (no limit if null or smaller than zero).

Example:

```
DELETE FROM TEST WHERE ID=2
```

## BACKUP

```
BACKUP TO fileNameString
```

Backs up the database files to a .zip file. Objects are not locked, but the backup is transactionally consistent because the transaction log is also copied. Admin rights are required to execute this command.

Example:

```
BACKUP TO 'backup.zip'
```

## CALL

```
CALL expression
```

Calculates a simple expression. This statement returns a result set with one row, except if the called function returns a result set itself. If the called function returns an array, then each element in this array is returned as a column.

Example:

```
CALL 15*25
```

## EXPLAIN

```
EXPLAIN { [ PLAN FOR ] | ANALYZE }  
{ select | insert | update | delete | merge }
```

Shows the execution plan for a statement. When using EXPLAIN ANALYZE, the statement is actually executed, and the query plan will include the actual row scan count for each table.

Example:

```
EXPLAIN SELECT * FROM TEST WHERE ID=1
```

## MERGE

```
MERGE INTO tableName [ ( columnName [,...] ) ]  
[ KEY ( columnName [,...] ) ]  
{ insertValues | select }
```

Updates existing rows, and insert rows that don't exist. If no key column is specified, the primary key columns are used to find the row. If more than one row per new row is affected, an exception is thrown.

Example:

```
MERGE INTO TEST KEY(ID) VALUES(2, 'World')
```

## MERGE USING

```
MERGE INTO targetTableName [ [AS] targetAlias]  
USING { ( select ) | sourceTableName } [ [AS] sourceAlias ]  
ON expression  
mergeWhenClause [,...]
```

Updates or deletes existing rows, and insert rows that don't exist.

The ON clause specifies the matching column expression. Different rows from a source table may not match with the same target row, but one source row may be matched with multiple target rows.

If statement doesn't need a source table a DUAL table can be substituted.

Example:

```
MERGE INTO TARGET_TABLE AS T USING SOURCE_TABLE AS S  
  ON T.ID = S.ID  
  WHEN MATCHED AND T.COL2 <> 'FINAL' THEN  
    UPDATE SET T.COL1 = S.COL1  
  WHEN MATCHED AND T.COL2 = 'FINAL' THEN  
    DELETE  
  WHEN NOT MATCHED THEN  
    INSERT (ID, COL1, COL2) VALUES(S.ID, S.COL1, S.COL2)  
MERGE INTO TARGET_TABLE AS T USING (SELECT * FROM SOURCE_TABLE) AS  
S
```



```

ON T.ID = S.ID
WHEN MATCHED AND T.COL2 <> 'FINAL' THEN
    UPDATE SET T.COL1 = S.COL1
WHEN MATCHED AND T.COL2 = 'FINAL' THEN
    DELETE
WHEN NOT MATCHED THEN
    INSERT VALUES (S.ID, S.COL1, S.COL2)
MERGE INTO TARGET_TABLE USING DUAL ON ID = 1
    WHEN NOT MATCHED THEN INSERT VALUES (1, 'Test')
    WHEN MATCHED THEN UPDATE SET NAME = 'Test'

```

## RUNSCRIPT

```

RUNSCRIPT FROM fileNameString scriptCompressionEncryption
[ CHARSET charsetString ]

```

Runs a SQL script from a file. The script is a text file containing SQL statements; each statement must end with ';'. This command can be used to restore a database from a backup. The password must be in single quotes; it is case sensitive and can contain spaces.

Instead of a file name, an URL may be used. To read a stream from the classpath, use the prefix 'classpath:'. See the Pluggable File System section on the Advanced page.

The compression algorithm must match the one used when creating the script. Instead of a file, an URL may be used.

Admin rights are required to execute this command.

Example:

```

RUNSCRIPT FROM 'backup.sql'
RUNSCRIPT FROM 'classpath:/com/acme/test.sql'

```

## SCRIPT

```

SCRIPT [ SIMPLE ] [ NODATA ] [ NOPASSWORDS ] [ NOSETTINGS ]
[ DROP ] [ BLOCKSIZE blockSizeInt ]
[ TO fileNameString scriptCompressionEncryption
[ CHARSET charsetString ] ]
[ TABLE tableName [, ...] ]
[ SCHEMA schemaName [, ...] ]

```

Creates a SQL script from the database.

SIMPLE does not use multi-row insert statements. NODATA will not emit INSERT statements. If the DROP option is specified, drop statements are created for tables, views, and sequences. If the block size is set, CLOB and BLOB values larger than this size are split into separate blocks. BLOCKSIZE is used when writing out LOB data, and specifies the point at the values transition from being inserted as inline values, to be inserted using out-of-line commands. NOSETTINGS turns off dumping the database settings (the SET XXX commands)

If no 'TO fileName' clause is specified, the script is returned as a result set. This command can be used to create a backup of the database. For long term storage, it is more portable than copying the database files.

If a 'TO fileName' clause is specified, then the whole script (including insert statements) is written to this file, and a result set without the insert statements is returned.

The password must be in single quotes; it is case sensitive and can contain spaces.

This command locks objects while it is running. Admin rights are required to execute this command.

When using the TABLE or SCHEMA option, only the selected table(s) / schema(s) are included.

Example:

```
SCRIPT NODATA
```

## SHOW

```
SHOW { SCHEMAS | TABLES [ FROM schemaName ] |  
COLUMNS FROM tableName [ FROM schemaName ] }
```

Lists the schemas, tables, or the columns of a table.

Example:

```
SHOW TABLES
```

## Explicit table

```
TABLE [schemaName.]tableName
```

```
[ ORDER BY order [,...] ]  
[ OFFSET expression { ROW | ROWS } ]  
[ FETCH { FIRST | NEXT } [ expression [ PERCENT ] ] { ROW | ROWS }  
{ ONLY | WITH TIES } ]
```

Selects data from a table.

This command is an equivalent to `SELECT * FROM tableName`. See `SELECT` command for description of `ORDER BY`, `OFFSET`, and `FETCH`.

Example:

```
TABLE TEST;  
TABLE TEST ORDER BY ID FETCH FIRST ROW ONLY;
```

## Table value

```
VALUES rowValueExpression [,...]  
[ ORDER BY order [,...] ]  
[ OFFSET expression { ROW | ROWS } ]  
[ FETCH { FIRST | NEXT } [ expression [ PERCENT ] ] { ROW | ROWS }  
{ ONLY | WITH TIES } ]
```

A list of rows that can be used like a table. See `SELECT` command for description of `ORDER BY`, `OFFSET`, and `FETCH`. The column list of the resulting table is `C1`, `C2`, and so on.

Example:

```
VALUES (1, 'Hello'), (2, 'World');
```

## WITH

```
WITH [ RECURSIVE ] { name [( columnName [,...] )] AS ( select ) [,...] }  
{ select | insert | update | merge | delete | createTable }
```

Can be used to create a recursive or non-recursive query (common table expression). For recursive queries the first select has to be a `UNION`. One or more common table entries can be referred to by name. Column name declarations are now optional - the column names will be inferred from the named select queries. The final action in a `WITH` statement can be a select, insert, update, merge, delete or create table.

Example:

```
WITH RECURSIVE cte(n) AS (  
    SELECT 1  
    UNION ALL  
    SELECT n + 1  
    FROM cte  
    WHERE n < 100  
)  
SELECT sum(n) FROM cte;
```

Example 2:

```
WITH cte1 AS (  
    SELECT 1 AS FIRST_COLUMN  
) , cte2 AS (  
    SELECT FIRST_COLUMN+1 AS FIRST_COLUMN FROM cte1  
)  
SELECT sum(FIRST_COLUMN) FROM cte2;
```

## Commands (Data Definition)

### ALTER INDEX RENAME

```
ALTER INDEX [ IF EXISTS ] indexName RENAME TO newIndexName
```

Renames an index. This command commits an open transaction in this connection.

Example:

```
ALTER INDEX IDXNAME RENAME TO IDX_TEST_NAME
```

### ALTER SCHEMA RENAME

```
ALTER SCHEMA [ IF EXISTS ] schema RENAME TO newSchemaName
```

Renames a schema. This command commits an open transaction in this connection.

Example:

```
ALTER SCHEMA TEST RENAME TO PRODUCTION
```

## ALTER SEQUENCE

```
ALTER SEQUENCE [ IF EXISTS ] sequenceName
[ RESTART WITH long ]
[ INCREMENT BY long ]
[ MINVALUE long | NOMINVALUE | NO MINVALUE ]
[ MAXVALUE long | NOMAXVALUE | NO MAXVALUE ]
[ CYCLE long | NOCYCLE | NO CYCLE ]
[ CACHE long | NOCACHE | NO CACHE ]
```

Changes the parameters of a sequence. This command does not commit the current transaction; however the new value is used by other transactions immediately, and rolling back this command has no effect.

Example:

```
ALTER SEQUENCE SEQ_ID RESTART WITH 1000
```

## ALTER TABLE ADD

```
ALTER TABLE [ IF EXISTS ] tableName ADD [ COLUMN ]
{ [ IF NOT EXISTS ] columnName columnDefinition
| ( { columnName columnDefinition | constraint } [,...] ) }
[ { { BEFORE | AFTER } columnName } | FIRST ]
```

Adds a new column to a table. This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST ADD CREATEDATE TIMESTAMP
```

## ALTER TABLE ADD CONSTRAINT

```
ALTER TABLE [ IF EXISTS ] tableName ADD constraint [ CHECK | NOCHECK ]
```

Adds a constraint to a table. If NOCHECK is specified, existing rows are not checked for consistency (the default is to check consistency for existing rows). The required indexes are automatically created if they don't exist yet. It is not possible to disable checking for unique constraints. This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST ADD CONSTRAINT NAME_UNIQUE UNIQUE(NAME)
```

## ALTER TABLE RENAME CONSTRAINT

```
ALTER TABLE [ IF EXISTS ] tableName RENAME oldConstraintName  
TO newConstraintName
```

Renames a constraint. This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST RENAME CONSTRAINT FOO TO BAR
```

## ALTER TABLE ALTER COLUMN

```
ALTER TABLE [ IF EXISTS ] tableName ALTER COLUMN columnName  
{ { columnDefinition }  
| { RENAME TO name }  
| { RESTART WITH long }  
| { SELECTIVITY int }  
| { SET DEFAULT expression }  
| { DROP DEFAULT }  
| { SET ON UPDATE expression }  
| { DROP ON UPDATE }  
| { SET NOT NULL }  
| { DROP NOT NULL } | { SET NULL }  
| { SET DATA TYPE dataType }  
| { SET { VISIBLE | INVISIBLE } } }
```

Changes the data type of a column, rename a column, change the identity value, or change the selectivity.

Changing the data type fails if the data can not be converted.

RESTART changes the next value of an auto increment column. The column must already be an auto increment column. For RESTART, the same transactional rules as for ALTER SEQUENCE apply.

SELECTIVITY sets the selectivity (1-100) for a column. Setting the selectivity to 0 means the default value. Selectivity is used by the cost based optimizer to calculate the estimated cost of an index. Selectivity 100 means values are unique, 10 means every distinct value appears 10 times on average.

SET DEFAULT changes the default value of a column.

DROP DEFAULT removes the default value of a column.

SET ON UPDATE changes the value that is set on update if value for this column is not specified in update statement.

DROP ON UPDATE removes the value that is set on update of a column.

SET NOT NULL sets a column to not allow NULL. Rows may not contains NULL in this column.

DROP NOT NULL and SET NULL set a column to allow NULL. The row may not be part of a primary key.

SET DATA TYPE changes the data type of a column.

SET INVISIBLE makes the column hidden, i.e. it will not appear in SELECT \* results. SET VISIBLE has the reverse effect.

This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST ALTER COLUMN NAME CLOB;  
ALTER TABLE TEST ALTER COLUMN NAME RENAME TO TEXT;  
ALTER TABLE TEST ALTER COLUMN ID RESTART WITH 10000;  
ALTER TABLE TEST ALTER COLUMN NAME SELECTIVITY 100;  
ALTER TABLE TEST ALTER COLUMN NAME SET DEFAULT "";  
ALTER TABLE TEST ALTER COLUMN NAME SET NOT NULL;  
ALTER TABLE TEST ALTER COLUMN NAME SET NULL;  
ALTER TABLE TEST ALTER COLUMN NAME SET VISIBLE;  
ALTER TABLE TEST ALTER COLUMN NAME SET INVISIBLE;
```

## ALTER TABLE DROP COLUMN

```
ALTER TABLE [ IF EXISTS ] tableName DROP COLUMN [ IF EXISTS ]  
columnName [,...] | ( columnName [,...] )
```

Removes column(s) from a table. This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST DROP COLUMN NAME  
ALTER TABLE TEST DROP COLUMN NAME1, NAME2  
ALTER TABLE TEST DROP COLUMN (NAME1, NAME2)
```

## ALTER TABLE DROP CONSTRAINT

```
ALTER TABLE [ IF EXISTS ] tableName DROP  
{ CONSTRAINT [ IF EXISTS ] constraintName | PRIMARY KEY }
```

Removes a constraint or a primary key from a table. This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST DROP CONSTRAINT UNIQUE_NAME
```

## ALTER TABLE SET

```
ALTER TABLE [ IF EXISTS ] tableName SET REFERENTIAL_INTEGRITY  
{ FALSE | TRUE } [ CHECK | NOCHECK ]
```

Disables or enables referential integrity checking for a table. This command can be used inside a transaction. Enabling referential integrity does not check existing data, except if CHECK is specified. Use SET REFERENTIAL\_INTEGRITY to disable it for all tables; the global flag and the flag for each table are independent.

This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST SET REFERENTIAL_INTEGRITY FALSE
```

## ALTER TABLE RENAME

```
ALTER TABLE [ IF EXISTS ] tableName RENAME TO newName
```

Renames a table. This command commits an open transaction in this connection.

Example:

```
ALTER TABLE TEST RENAME TO MY_DATA
```

## ALTER USER ADMIN

```
ALTER USER userName ADMIN { TRUE | FALSE }
```

Switches the admin flag of a user on or off.



Only unquoted or uppercase user names are allowed. Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
ALTER USER TOM ADMIN TRUE
```

## ALTER USER RENAME

```
ALTER USER userName RENAME TO newUserName
```

Renames a user. After renaming a user, the password becomes invalid and needs to be changed as well.

Only unquoted or uppercase user names are allowed. Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
ALTER USER TOM RENAME TO THOMAS
```

## ALTER USER SET PASSWORD

```
ALTER USER userName SET { PASSWORD string | SALT bytes HASH bytes }
```

Changes the password of a user. Only unquoted or uppercase user names are allowed. The password must be enclosed in single quotes. It is case sensitive and can contain spaces. The salt and hash values are hex strings.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
ALTER USER SA SET PASSWORD 'rioyxlt'
```

## ALTER VIEW RECOMPILE

```
ALTER VIEW [ IF EXISTS ] viewName RECOMPILE
```

Recompiles a view after the underlying tables have been changed or created. This command is used for views created using CREATE FORCE VIEW. This command commits an open transaction in this connection.

Example:

```
ALTER VIEW ADDRESS_VIEW RECOMPILE
```

## ALTER VIEW RENAME

```
ALTER VIEW [ IF EXISTS ] viewName RENAME TO newName
```

Renames a view. This command commits an open transaction in this connection.

Example:

```
ALTER VIEW TEST RENAME TO MY_VIEW
```

## ANALYZE

```
ANALYZE [ TABLE tableName ] [ SAMPLE_SIZE rowCountInt ]
```

Updates the selectivity statistics of tables. If no table name is given, all tables are analyzed. The selectivity is used by the cost based optimizer to select the best index for a given query. If no sample size is set, up to 10000 rows per table are read. The value 0 means all rows are read. The selectivity can be set manually using ALTER TABLE ALTER COLUMN SELECTIVITY. Manual values are overwritten by this statement. The selectivity is available in the INFORMATION\_SCHEMA.COLUMNS table.

This command commits an open transaction in this connection.

Example:

```
ANALYZE SAMPLE_SIZE 1000
```

## COMMENT

```
COMMENT ON  
{ { COLUMN [ schemaName. ] tableName.columnName }  
| { { TABLE | VIEW | CONSTANT | CONSTRAINT | ALIAS | INDEX | ROLE  
| SCHEMA | SEQUENCE | TRIGGER | USER | DOMAIN } [ schemaName. ]  
objectName } }  
IS expression
```

Sets the comment of a database object. Use NULL to remove the comment.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
COMMENT ON TABLE TEST IS 'Table used for testing'
```

## CREATE AGGREGATE

```
CREATE AGGREGATE [ IF NOT EXISTS ] newAggregateName FOR className
```

Creates a new user-defined aggregate function. The method name must be the full qualified class name. The class must implement the interface `org.h2.api.Aggregate` or `org.h2.api.AggregateFunction`.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
CREATE AGGREGATE SIMPLE_MEDIAN FOR "com.acme.db.Median"
```

## CREATE ALIAS

```
CREATE ALIAS [ IF NOT EXISTS ] newFunctionAliasName [ DETERMINISTIC ]  
[ NOBUFFER ] { FOR classAndMethodName | AS sourceCodeString }
```

Creates a new function alias. If this is a `ResultSet` returning function, by default the return value is cached in a local temporary file.

NOBUFFER - disables caching of `ResultSet` return value to temporary file.

DETERMINISTIC - Deterministic functions must always return the same value for the same parameters.

The method name must be the full qualified class and method name, and may optionally include the parameter classes as in `java.lang.Integer.parseInt(java.lang.String, int)`. The class and the method must both be public, and the method must be static. The class must be available in the classpath of the database engine (when using the server mode, it must be in the classpath of the server).

When defining a function alias with source code, the Sun javac is compiler is used if the file tools.jar is in the classpath. If not, javac is run as a separate process. Only the source code is stored in the database; the class is compiled each time the database is re-opened. Source code is usually passed as dollar quoted text to avoid escaping problems. If import statements are used, then the tag @CODE must be added before the method.

If the method throws an SQLException, it is directly re-thrown to the calling application; all other exceptions are first converted to a SQLException.

If the first parameter of the Java function is a java.sql.Connection, then a connection to the database is provided. This connection must not be closed. If the class contains multiple methods with the given name but different parameter count, all methods are mapped.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

If you have the Groovy jar in your classpath, it is also possible to write methods using Groovy.

Example:

```
CREATE ALIAS MY_SQRT FOR "java.lang.Math.sqrt";
CREATE ALIAS GET_SYSTEM_PROPERTY FOR "java.lang.System.getProperty";
CALL GET_SYSTEM_PROPERTY('java.class.path');
CALL GET_SYSTEM_PROPERTY('com.acme.test', 'true');
CREATE ALIAS REVERSE AS $$ String reverse(String s) { return new
StringBuilder(s).reverse().toString(); } $$;
CALL REVERSE('Test');
CREATE ALIAS tr AS $$@groovy.transform.CompileStatic
    static String tr(String str, String sourceSet, String replacementSet){
        return str.tr(sourceSet, replacementSet);
    }
    $$
```

## CREATE CONSTANT

```
CREATE CONSTANT [ IF NOT EXISTS ] newConstantName VALUE expression
```

Creates a new constant. This command commits an open transaction in this connection.

Example:

```
CREATE CONSTANT ONE VALUE 1
```

## CREATE DOMAIN

```
CREATE DOMAIN [ IF NOT EXISTS ] newDomainName AS dataType  
[ DEFAULT expression ] [ [ NOT ] NULL ] [ SELECTIVITY selectivity ]  
[ CHECK condition ]
```

Creates a new data type (domain). The check condition must evaluate to true or to NULL (to prevent NULL, use NOT NULL). In the condition, the term VALUE refers to the value being tested.

Domains are usable within the whole database. They can not be created in a specific schema.

This command commits an open transaction in this connection.

Example:

```
CREATE DOMAIN EMAIL AS VARCHAR(255) CHECK (POSITION('@', VALUE) > 1)
```

## CREATE INDEX

```
CREATE  
{ [ UNIQUE ] [ HASH | SPATIAL ] INDEX [ [ IF NOT EXISTS ] newIndexName ]  
| PRIMARY KEY [ HASH ] }  
ON tableName ( indexColumn [,...] )
```

Creates a new index. This command commits an open transaction in this connection.

Hash indexes are meant for in-memory databases and memory tables (CREATE MEMORY TABLE) when PageStore engine is used. For other tables, or if the index contains multiple columns, the HASH keyword is ignored. Hash indexes can only test for equality, do not support range queries (similar to a hash table), use more memory, but can perform lookups faster. Non-unique keys are supported.

Spatial indexes are supported only on Geometry columns.

Example:

```
CREATE INDEX IDXNAME ON TEST(NAME)
```

## CREATE LINKED TABLE

```
CREATE [ FORCE ] [ [ GLOBAL | LOCAL ] TEMPORARY ]  
LINKED TABLE [ IF NOT EXISTS ]
```

```
name ( driverString, urlString, urlString, passwordString,  
[ originalSchemaString, ] originalTableString ) [ EMIT UPDATES | READONLY ]
```

Creates a table link to an external table. The driver name may be empty if the driver is already loaded. If the schema name is not set, only one table with that name may exist in the target database.

FORCE - Create the LINKED TABLE even if the remote database/table does not exist.

EMIT UPDATES - Usually, for update statements, the old rows are deleted first and then the new rows are inserted. It is possible to emit update statements (except on rollback), however in this case multi-row unique key updates may not always work. Linked tables to the same database share one connection.

READONLY - is set, the remote table may not be updated. This is enforced by H2.

If the connection to the source database is lost, the connection is re-opened (this is a workaround for MySQL that disconnects after 8 hours of inactivity by default).

If a query is used instead of the original table name, the table is read only. Queries must be enclosed in parenthesis: (SELECT \* FROM ORDERS).

To use JNDI to get the connection, the driver class must be a javax.naming.Context (for example javax.naming.InitialContext), and the URL must be the resource name (for example java:comp/env/jdbc/Test).

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
CREATE LINKED TABLE LINK('org.h2.Driver', 'jdbc:h2:test2',  
    'sa', 'sa', 'TEST');  
CREATE LINKED TABLE LINK("", 'jdbc:h2:test2', 'sa', 'sa',  
    '(SELECT * FROM TEST WHERE ID>0)');  
CREATE LINKED TABLE LINK('javax.naming.InitialContext',  
    'java:comp/env/jdbc/Test', NULL, NULL,  
    '(SELECT * FROM TEST WHERE ID>0)');
```

## CREATE ROLE

```
CREATE ROLE [ IF NOT EXISTS ] newRoleName
```

Creates a new role. This command commits an open transaction in this connection.

Example:

```
CREATE ROLE READONLY
```

## CREATE SCHEMA

```
CREATE SCHEMA [ IF NOT EXISTS ] name  
[ AUTHORIZATION ownerUserName ]  
[ WITH tableEngineParamName [...]]
```

Creates a new schema. If no owner is specified, the current user is used. The user that executes the command must have admin rights, as well as the owner. Specifying the owner currently has no effect. Optional table engine parameters are used when CREATE TABLE command is run on this schema without having its engine params set.

This command commits an open transaction in this connection.

Example:

```
CREATE SCHEMA TEST_SCHEMA AUTHORIZATION SA
```

## CREATE SEQUENCE

```
CREATE SEQUENCE [ IF NOT EXISTS ] newSequenceName  
sequenceOptions
```

Creates a new sequence. The data type of a sequence is BIGINT. Used values are never re-used, even when the transaction is rolled back.

The cache is the number of pre-allocated numbers. If the system crashes without closing the database, at most this many numbers are lost. The default cache size is 32. To disable caching, use the cache size 1 or lower.

This command commits an open transaction in this connection.

Example:

```
CREATE SEQUENCE SEQ_ID
```

## CREATE TABLE

```
CREATE [ CACHED | MEMORY ] [ TEMP | [ GLOBAL | LOCAL ] TEMPORARY ]  
TABLE [ IF NOT EXISTS ] name
```

```
[ ( { columnName [columnDefinition] | constraint } [,...] ) ]  
[ ENGINE tableEngineName ]  
[ WITH tableEngineParamName [,...] ]  
[ NOT PERSISTENT ] [ TRANSACTIONAL ]  
[ AS select [ WITH [ NO ] DATA ] ]
```

Creates a new table.

Cached tables (the default for regular tables) are persistent, and the number of rows is not limited by the main memory. Memory tables (the default for temporary tables) are persistent, but the index data is kept in main memory, that means memory tables should not get too large.

Temporary tables are deleted when closing or opening a database. Temporary tables can be global (accessible by all connections) or local (only accessible by the current connection). The default for temporary tables is global. Indexes of temporary tables are kept fully in main memory, unless the temporary table is created using CREATE CACHED TABLE.

The ENGINE option is only required when custom table implementations are used. The table engine class must implement the interface org.h2.api.TableEngine. Any table engine parameters are passed down in the tableEngineParams field of the CreateTableData object.

Either ENGINE, or WITH (table engine params), or both may be specified. If ENGINE is not specified in CREATE TABLE, then the engine specified by DEFAULT\_TABLE\_ENGINE option of database params is used.

Tables with the NOT PERSISTENT modifier are kept fully in memory, and all rows are lost when the database is closed.

The column definitions are optional if a query is specified. In that case the column list of the query is used. If the query is specified its results are inserted into created table unless WITH NO DATA is specified.

This command commits an open transaction, except when using TRANSACTIONAL (only supported for temporary tables).

Example:

```
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255))
```

## CREATE TRIGGER

```
CREATE TRIGGER [ IF NOT EXISTS ] newTriggerName
```



```
{ BEFORE | AFTER | INSTEAD OF }  
{ INSERT | UPDATE | DELETE | SELECT | ROLLBACK }  
[,...] ON tableName [ FOR EACH ROW ]  
[ QUEUE int ] [ NOWAIT ]  
{ CALL triggeredClassName | AS sourceCodeString }
```

Creates a new trigger. The trigger class must be public and implement `org.h2.api.Trigger`. Inner classes are not supported. The class must be available in the classpath of the database engine (when using the server mode, it must be in the classpath of the server).

The `sourceCodeString` must define a single method with no parameters that returns `org.h2.api.Trigger`. See `CREATE ALIAS` for requirements regarding the compilation. Alternatively, `javax.script.ScriptEngineManager` can be used to create an instance of `org.h2.api.Trigger`. Currently javascript (included in every JRE) and ruby (with JRuby) are supported. In that case the source must begin respectively with `//javascript` or `#ruby`.

BEFORE triggers are called after data conversion is made, default values are set, null and length constraint checks have been made; but before other constraints have been checked. If there are multiple triggers, the order in which they are called is undefined.

ROLLBACK can be specified in combination with INSERT, UPDATE, and DELETE. Only row based AFTER trigger can be called on ROLLBACK. Exceptions that occur within such triggers are ignored. As the operations that occur within a trigger are part of the transaction, ROLLBACK triggers are only required if an operation communicates outside of the database.

INSTEAD OF triggers are implicitly row based and behave like BEFORE triggers. Only the first such trigger is called. Such triggers on views are supported. They can be used to make views updatable.

A BEFORE SELECT trigger is fired just before the database engine tries to read from the table. The trigger can be used to update a table on demand. The trigger is called with both 'old' and 'new' set to null.

The MERGE statement will call both INSERT and UPDATE triggers. Not supported are SELECT triggers with the option FOR EACH ROW, and AFTER SELECT triggers.

Committing or rolling back a transaction within a trigger is not allowed, except for SELECT triggers.

By default a trigger is called once for each statement, without the old and new rows. FOR EACH ROW triggers are called once for each inserted, updated, or deleted row.

QUEUE is implemented for syntax compatibility with HSQL and has no effect.

The trigger need to be created in the same schema as the table. The schema name does not need to be specified when creating the trigger.

This command commits an open transaction in this connection.

Example:

```
CREATE TRIGGER TRIG_INS BEFORE INSERT ON TEST FOR EACH ROW CALL
"MyTrigger";
CREATE TRIGGER TRIG_SRC BEFORE INSERT ON TEST AS $$org.h2.api.Trigger
create() { return new MyTrigger("constructorParam"); } $$;
CREATE TRIGGER TRIG_JS BEFORE INSERT ON TEST AS $$//javascript\nreturn
new Packages.MyTrigger("constructorParam"); $$;
CREATE TRIGGER TRIG_RUBY BEFORE INSERT ON TEST AS $
$#ruby\nJava::MyPackage::MyTrigger.new("constructorParam") $$;
```

## CREATE USER

```
CREATE USER [ IF NOT EXISTS ] newUserName
{ PASSWORD string | SALT bytes HASH bytes } [ ADMIN ]
```

Creates a new user. For compatibility, only unquoted or uppercase user names are allowed. The password must be in single quotes. It is case sensitive and can contain spaces. The salt and hash values are hex strings.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
CREATE USER GUEST PASSWORD 'abc'
```

## CREATE VIEW

```
CREATE [ OR REPLACE ] [ FORCE ] VIEW [ IF NOT EXISTS ] newViewName
[ ( columnName [,...] ) ] AS select
```

Creates a new view. If the force option is used, then the view is created even if the underlying table(s) don't exist.

If the OR REPLACE clause is used an existing view will be replaced, and any dependent views will not need to be recreated. If dependent views will become

invalid as a result of the change an error will be generated, but this error can be ignored if the FORCE clause is also used.

Views are not updatable except when using 'instead of' triggers.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
CREATE VIEW TEST_VIEW AS SELECT * FROM TEST WHERE ID < 100
```

## DROP AGGREGATE

```
DROP AGGREGATE [ IF EXISTS ] aggregateName
```

Drops an existing user-defined aggregate function.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
DROP AGGREGATE SIMPLE_MEDIAN
```

## DROP ALIAS

```
DROP ALIAS [ IF EXISTS ] existingFunctionAliasName
```

Drops an existing function alias.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
DROP ALIAS MY_SQRT
```

## DROP ALL OBJECTS

```
DROP ALL OBJECTS [ DELETE FILES ]
```

Drops all existing views, tables, sequences, schemas, function aliases, roles, user-defined aggregate functions, domains, and users (except the current user). If

DELETE FILES is specified, the database files will be removed when the last user disconnects from the database. Warning: this command can not be rolled back.

Admin rights are required to execute this command.

Example:

DROP ALL OBJECTS

## DROP CONSTANT

```
DROP CONSTANT [ IF EXISTS ] constantName
```

Drops a constant. This command commits an open transaction in this connection.

Example:

DROP CONSTANT ONE

## DROP DOMAIN

```
DROP DOMAIN [ IF EXISTS ] domainName [ RESTRICT | CASCADE ]
```

Drops a data type (domain). The command will fail if it is referenced by a column (the default). Column descriptors are replaced with original definition of specified domain if the CASCADE clause is used. This command commits an open transaction in this connection.

Example:

DROP DOMAIN EMAIL

## DROP INDEX

```
DROP INDEX [ IF EXISTS ] indexName
```

Drops an index. This command commits an open transaction in this connection.

Example:

DROP INDEX IF EXISTS IDXNAME

## DROP ROLE

```
DROP ROLE [ IF EXISTS ] roleName
```

Drops a role. This command commits an open transaction in this connection.

Example:

```
DROP ROLE READONLY
```

## DROP SCHEMA

```
DROP SCHEMA [ IF EXISTS ] schemaName [ RESTRICT | CASCADE ]
```

Drops a schema. The command will fail if objects in this schema exist and the RESTRICT clause is used (the default). All objects in this schema are dropped as well if the CASCADE clause is used. This command commits an open transaction in this connection.

Example:

```
DROP SCHEMA TEST_SCHEMA
```

## DROP SEQUENCE

```
DROP SEQUENCE [ IF EXISTS ] sequenceName
```

Drops a sequence. This command commits an open transaction in this connection.

Example:

```
DROP SEQUENCE SEQ_ID
```

## DROP TABLE

```
DROP TABLE [ IF EXISTS ] tableName [,...] [ RESTRICT | CASCADE ]
```

Drops an existing table, or a list of tables. The command will fail if dependent objects exist and the RESTRICT clause is used (the default). All dependent views and constraints are dropped as well if the CASCADE clause is used. This command commits an open transaction in this connection.

Example:

## DROP TABLE TEST

## DROP TRIGGER

```
DROP TRIGGER [ IF EXISTS ] triggerName
```

Drops an existing trigger. This command commits an open transaction in this connection.

Example:

```
DROP TRIGGER TRIG_INS
```

## DROP USER

```
DROP USER [ IF EXISTS ] userName
```

Drops a user. The current user cannot be dropped. For compatibility, only unquoted or uppercase user names are allowed.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
DROP USER TOM
```

## DROP VIEW

```
DROP VIEW [ IF EXISTS ] viewName [ RESTRICT | CASCADE ]
```

Drops an existing view. All dependent views are dropped as well if the CASCADE clause is used (the default). The command will fail if dependent views exist and the RESTRICT clause is used. This command commits an open transaction in this connection.

Example:

```
DROP VIEW TEST_VIEW
```

## TRUNCATE TABLE

```
TRUNCATE TABLE tableName [ [ CONTINUE | RESTART ] IDENTITY ]
```

Removes all rows from a table. Unlike DELETE FROM without where clause, this command can not be rolled back. This command is faster than DELETE without where clause. Only regular data tables without foreign key constraints can be truncated (except if referential integrity is disabled for this database or for this table). Linked tables can't be truncated. If RESTART IDENTITY is specified next values for auto-incremented columns are restarted.

This command commits an open transaction in this connection.

Example:

```
TRUNCATE TABLE TEST
```

## Commands (Other)

### CHECKPOINT

```
CHECKPOINT
```

Flushes the data to disk.

Admin rights are required to execute this command.

Example:

```
CHECKPOINT
```

### CHECKPOINT SYNC

```
CHECKPOINT SYNC
```

Flushes the data to disk and forces all system buffers be written to the underlying device.

Admin rights are required to execute this command.

Example:

```
CHECKPOINT SYNC
```

### COMMIT

```
COMMIT [ WORK ]
```

Commits a transaction.

Example:

```
COMMIT
```

## COMMIT TRANSACTION

```
COMMIT TRANSACTION transactionName
```

Sets the resolution of an in-doubt transaction to 'commit'.

Admin rights are required to execute this command. This command is part of the 2-phase-commit protocol.

Example:

```
COMMIT TRANSACTION XID_TEST
```

## GRANT RIGHT

```
GRANT { SELECT | INSERT | UPDATE | DELETE | ALL } [,...] ON  
{ { SCHEMA schemaName } | { tableName [,...] } }  
TO { PUBLIC | userName | roleName }
```

Grants rights for a table to a user or role.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
GRANT SELECT ON TEST TO READONLY
```

## GRANT ALTER ANY SCHEMA

```
GRANT ALTER ANY SCHEMA TO userName
```

Grant schema altering rights to a user.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:



GRANT ALTER ANY SCHEMA TO Bob

## GRANT ROLE

```
GRANT roleName TO { PUBLIC | userName | roleName }
```

Grants a role to a user or role.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
GRANT READONLY TO PUBLIC
```

## HELP

```
HELP [ anything [...] ]
```

Displays the help pages of SQL commands or keywords.

Example:

```
HELP SELECT
```

## PREPARE COMMIT

```
PREPARE COMMIT newTransactionName
```

Prepares committing a transaction. This command is part of the 2-phase-commit protocol.

Example:

```
PREPARE COMMIT XID_TEST
```

## REVOKE RIGHT

```
REVOKE { SELECT | INSERT | UPDATE | DELETE | ALL } [,...] ON  
{ { SCHEMA schemaName } | { tableName [,...] } }  
FROM { PUBLIC | userName | roleName }
```

Removes rights for a table from a user or role.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
REVOKE SELECT ON TEST FROM READONLY
```

## REVOKE ROLE

```
REVOKE roleName FROM { PUBLIC | userName | roleName }
```

Removes a role from a user or role.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
REVOKE READONLY FROM TOM
```

## ROLLBACK

```
ROLLBACK [ TO SAVEPOINT savepointName ]
```

Rolls back a transaction. If a savepoint name is used, the transaction is only rolled back to the specified savepoint.

Example:

```
ROLLBACK
```

## ROLLBACK TRANSACTION

```
ROLLBACK TRANSACTION transactionName
```

Sets the resolution of an in-doubt transaction to 'rollback'.

Admin rights are required to execute this command. This command is part of the 2-phase-commit protocol.

Example:

```
ROLLBACK TRANSACTION XID_TEST
```

## SAVEPOINT

```
SAVEPOINT savepointName
```

Create a new savepoint. See also ROLLBACK. Savepoints are only valid until the transaction is committed or rolled back.

Example:

```
SAVEPOINT HALF_DONE
```

## SET @

```
SET @variableName [ = ] expression
```

Updates a user-defined variable. Variables are not persisted and session scoped, that means only visible from within the session in which they are defined. This command does not commit a transaction, and rollback does not affect it.

Example:

```
SET @TOTAL=0
```

## SET ALLOW\_LITERALS

```
SET ALLOW_LITERALS { NONE | ALL | NUMBERS }
```

This setting can help solve the SQL injection problem. By default, text and number literals are allowed in SQL statements. However, this enables SQL injection if the application dynamically builds SQL statements. SQL injection is not possible if user data is set using parameters ('?').

NONE means literals of any kind are not allowed, only parameters and constants are allowed. NUMBERS mean only numerical and boolean literals are allowed. ALL means all literals are allowed (default).

See also CREATE CONSTANT.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:  
`jdbc:h2:test;ALLOW_LITERALS=NONE`

Example:

SET ALLOW\_LITERALS NONE

## SET AUTOCOMMIT

```
SET AUTOCOMMIT { TRUE | ON | FALSE | OFF }
```

Switches auto commit on or off. This setting can be appended to the database URL: jdbc:h2:test;AUTOCOMMIT=OFF - however this will not work as expected when using a connection pool (the connection pool manager will re-enable autocommit when returning the connection to the pool, so autocommit will only be disabled the first time the connection is used).

Example:

```
SET AUTOCOMMIT OFF
```

## SET CACHE\_SIZE

```
SET CACHE_SIZE int
```

Sets the size of the cache in KB (each KB being 1024 bytes) for the current database. The default is 65536 per available GB of RAM, i.e. 64 MB per GB. The value is rounded to the next higher power of two. Depending on the virtual machine, the actual memory required may be higher.

This setting is persistent and affects all connections as there is only one cache per database. Using a very small value (specially 0) will reduce performance a lot. This setting only affects the database engine (the server in a client/server environment; in embedded mode, the database engine is in the same process as the application). It has no effect for in-memory databases.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:  
jdbc:h2:test;CACHE\_SIZE=8192

Example:

```
SET CACHE_SIZE 8192
```

## SET CLUSTER

```
SET CLUSTER serverListString
```

This command should not be used directly by an application, the statement is executed automatically by the system. The behavior may change in future releases. Sets the cluster server list. An empty string switches off the cluster mode. Switching on the cluster mode requires admin rights, but any user can switch it off (this is automatically done when the client detects the other server is not responding).

This command is effective immediately, but does not commit an open transaction.

Example:

```
SET CLUSTER "
```

## **SET BINARY\_COLLATION**

```
SET BINARY_COLLATION { UNSIGNED | SIGNED }
```

Sets the collation used for comparing BINARY columns, the default is SIGNED for version 1.3 and older, and UNSIGNED for version 1.4 and newer. This command can only be executed if there are no tables defined.

Admin rights are required to execute this command. This command commits an open transaction in this connection. This setting is persistent.

Example:

```
SET BINARY_COLLATION SIGNED
```

## **SET UUID\_COLLATION**

```
SET UUID_COLLATION { UNSIGNED | SIGNED }
```

Sets the collation used for comparing UUID columns, the default is SIGNED. This command can only be executed if there are no tables defined.

SIGNED means signed comparison between first 64 bits of compared values treated as long values and if they are equal a signed comparison of the last 64 bits of compared values treated as long values. See also `Java UUID.compareTo()`. UNSIGNED means RFC 4122 compatible unsigned comparison.

Admin rights are required to execute this command. This command commits an open transaction in this connection. This setting is persistent.

Example:

SET UUID\_COLLATION UNSIGNED

## SET BUILTIN\_ALIAS\_OVERRIDE

```
SET BUILTIN_ALIAS_OVERRIDE { TRUE | FALSE }
```

Allows the overriding of the builtin system date/time functions for unit testing purposes.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
SET BUILTIN_ALIAS_OVERRIDE TRUE
```

## SET COLLATION

```
SET [ DATABASE ] COLLATION  
{ OFF | collationName  
[ STRENGTH { PRIMARY | SECONDARY | TERTIARY | IDENTICAL } ] }
```

Sets the collation used for comparing strings. This command can only be executed if there are no tables defined. See `java.text.Collator` for details about the supported collations and the STRENGTH (PRIMARY is usually case- and umlaut-insensitive; SECONDARY is case-insensitive but umlaut-sensitive; TERTIARY is both case- and umlaut-sensitive; IDENTICAL is sensitive to all differences and only affects ordering).

The ICU4J collator is used if it is in the classpath. It is also used if the collation name starts with ICU4J\_ (in that case, the ICU4J must be in the classpath, otherwise an exception is thrown). The default collator is used if the collation name starts with DEFAULT\_ (even if ICU4J is in the classpath). The charset collator is used if the collation name starts with CHARSET\_ (e.g. CHARSET\_CP500). This collator sorts strings according to the binary representation in the given charset.

Admin rights are required to execute this command. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL: `jdbc:h2:test;COLLATION='ENGLISH'`

Example:

```
SET COLLATION ENGLISH  
SET COLLATION CHARSET_CP500
```

## SET COMPRESS\_LOB

```
SET COMPRESS_LOB { NO | LZF | DEFLATE }
```

This feature is only available for the PageStore storage engine. For the MVStore engine (the default for H2 version 1.4.x), append ;COMPRESS=TRUE to the database URL instead.

Sets the compression algorithm for BLOB and CLOB data. Compression is usually slower, but needs less disk space. LZF is faster but uses more space.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent.

Example:

```
SET COMPRESS_LOB LZF
```

## SET DATABASE\_EVENT\_LISTENER

```
SET DATABASE_EVENT_LISTENER classNameString
```

Sets the event listener class. An empty string (") means no listener should be used. This setting is not persistent.

Admin rights are required to execute this command, except if it is set when opening the database (in this case it is reset just after opening the database). This setting can be appended to the database URL:  
jdbc:h2:test;DATABASE\_EVENT\_LISTENER='sample.MyListener'

Example:

```
SET DATABASE_EVENT_LISTENER 'sample.MyListener'
```

## SET DB\_CLOSE\_DELAY

```
SET DB_CLOSE_DELAY int
```

Sets the delay for closing a database if all connections are closed. The value -1 means the database is never closed until the close delay is set to some other value or SHUTDOWN is called. The value 0 means no delay (default; the database is closed if the last connection to it is closed). Values 1 and larger mean the number of seconds the database is left open after closing the last connection.

If the application exits normally or `System.exit` is called, the database is closed immediately, even if a delay is set.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:  
`jdbc:h2:test;DB_CLOSE_DELAY=-1`

Example:

```
SET DB_CLOSE_DELAY -1
```

## SET DEFAULT\_LOCK\_TIMEOUT

```
SET DEFAULT_LOCK_TIMEOUT int
```

Sets the default lock timeout (in milliseconds) in this database that is used for the new sessions. The default value for this setting is 1000 (one second).

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent.

Example:

```
SET DEFAULT_LOCK_TIMEOUT 5000
```

## SET DEFAULT\_TABLE\_TYPE

```
SET DEFAULT_TABLE_TYPE { MEMORY | CACHED }
```

Sets the default table storage type that is used when creating new tables. Memory tables are kept fully in the main memory (including indexes), however the data is still stored in the database file. The size of memory tables is limited by the memory. The default is CACHED.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. It has no effect for in-memory databases.

Example:

```
SET DEFAULT_TABLE_TYPE MEMORY
```



## SET EXCLUSIVE

```
SET EXCLUSIVE { 0 | 1 | 2 }
```

Switched the database to exclusive mode (1, 2) and back to normal mode (0).

In exclusive mode, new connections are rejected, and operations by other connections are paused until the exclusive mode is disabled. When using the value 1, existing connections stay open. When using the value 2, all existing connections are closed (and current transactions are rolled back) except the connection that executes SET EXCLUSIVE. Only the connection that set the exclusive mode can disable it. When the connection is closed, it is automatically disabled.

Admin rights are required to execute this command. This command commits an open transaction in this connection.

Example:

```
SET EXCLUSIVE 1
```

## SET IGNORECASE

```
SET IGNORECASE { TRUE | FALSE }
```

If IGNORECASE is enabled, text columns in newly created tables will be case-insensitive. Already existing tables are not affected. The effect of case-insensitive columns is similar to using a collation with strength PRIMARY. Case-insensitive columns are compared faster than when using a collation. String literals and parameters are however still considered case sensitive even if this option is set.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:  
jdbc:h2:test;IGNORECASE=TRUE

Example:

```
SET IGNORECASE TRUE
```

## SET JAVA\_OBJECT\_SERIALIZER

```
SET JAVA_OBJECT_SERIALIZER  
{ null | className }
```

Sets the object used to serialize and deserialize java objects being stored in column of type OTHER. The serializer class must be public and implement `org.h2.api.JavaObjectSerializer`. Inner classes are not supported. The class must be available in the classpath of the database engine (when using the server mode, it must be both in the classpath of the server and the client). This command can only be executed if there are no tables defined.

Admin rights are required to execute this command. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:

```
jdbc:h2:test;JAVA_OBJECT_SERIALIZER='com.acme.SerializerClassName'
```

Example:

```
SET JAVA_OBJECT_SERIALIZER 'com.acme.SerializerClassName'
```

## SET LAZY\_QUERY\_EXECUTION

```
SET LAZY_QUERY_EXECUTION int
```

Sets the lazy query execution mode. The values 0, 1 are supported.

If true, then large results are retrieved in chunks.

Note that not all queries support this feature, queries which do not are processed normally.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is not persistent. This setting can be appended to the database URL:

```
jdbc:h2:test;LAZY_QUERY_EXECUTION=1
```

Example:

```
SET LAZY_QUERY_EXECUTION 1
```

## SET LOG

```
SET LOG int
```

Sets the transaction log mode. The values 0, 1, and 2 are supported, the default is 2. This setting affects all connections.

LOG 0 means the transaction log is disabled completely. It is the fastest mode, but also the most dangerous: if the process is killed while the database is open in

this mode, the data might be lost. It must only be used if this is not a problem, for example when initially loading a database, or when running tests.

LOG 1 means the transaction log is enabled, but FileDescriptor.sync is disabled. This setting is about half as fast as with LOG 0. This setting is useful if no protection against power failure is required, but the data must be protected against killing the process.

LOG 2 (the default) means the transaction log is enabled, and FileDescriptor.sync is called for each checkpoint. This setting is about half as fast as LOG 1. Depending on the file system, this will also protect against power failure in the majority of cases.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is not persistent. This setting can be appended to the database URL:  
jdbc:h2:test;LOG=0

Example:

SET LOG 1

## SET LOCK\_MODE

SET LOCK\_MODE [int](#)

Sets the lock mode. The values 0, 1, 2, and 3 are supported. The default is 3 (READ\_COMMITTED). This setting affects all connections.

The value 0 means no locking (should only be used for testing; also known as READ\_UNCOMMITTED). Please note that using SET LOCK\_MODE 0 while at the same time using multiple connections may result in inconsistent transactions.

The value 1 means table level locking (also known as SERIALIZABLE).

The value 2 means table level locking with garbage collection (if the application does not close all connections).

The value 3 means table level locking, but read locks are released immediately (default; also known as READ\_COMMITTED).

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:  
jdbc:h2:test;LOCK\_MODE=3

Example:

```
SET LOCK_MODE 1
```

## SET LOCK\_TIMEOUT

```
SET LOCK_TIMEOUT int
```

Sets the lock timeout (in milliseconds) for the current session. The default value for this setting is 1000 (one second).

This command does not commit a transaction, and rollback does not affect it. This setting can be appended to the database URL:

```
jdbc:h2:test;LOCK_TIMEOUT=10000
```

Example:

```
SET LOCK_TIMEOUT 1000
```

## SET MAX\_LENGTH\_INPLACE\_LOB

```
SET MAX_LENGTH_INPLACE_LOB int
```

Sets the maximum size of an in-place LOB object.

This is the maximum length of an LOB that is stored with the record itself, and the default value is 128.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent.

Example:

```
SET MAX_LENGTH_INPLACE_LOB 128
```

## SET MAX\_LOG\_SIZE

```
SET MAX_LOG_SIZE int
```

Sets the maximum size of the transaction log, in megabytes. If the log is larger, and if there is no open transaction, the transaction log is truncated. If there is an open transaction, the transaction log will continue to grow however. The default max size is 16 MB. This setting has no effect for in-memory databases.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent.

Example:

```
SET MAX_LOG_SIZE 2
```

## SET MAX\_MEMORY\_ROWS

```
SET MAX_MEMORY_ROWS int
```

The maximum number of rows in a result set that are kept in-memory. If more rows are read, then the rows are buffered to disk. The default is 40000 per GB of available RAM.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. It has no effect for in-memory databases.

Example:

```
SET MAX_MEMORY_ROWS 1000
```

## SET MAX\_MEMORY\_UNDO

```
SET MAX_MEMORY_UNDO int
```

The maximum number of undo records per a session that are kept in-memory. If a transaction is larger, the records are buffered to disk. The default value is 50000. Changes to tables without a primary key can not be buffered to disk. This setting is not supported when using multi-version concurrency.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. It has no effect for in-memory databases.

Example:

```
SET MAX_MEMORY_UNDO 1000
```

## SET MAX\_OPERATION\_MEMORY

```
SET MAX_OPERATION_MEMORY int
```

Sets the maximum memory used for large operations (delete and insert), in bytes. Operations that use more memory are buffered to disk, slowing down the operation. The default max size is 100000. 0 means no limit.

This setting is not persistent. Admin rights are required to execute this command, as it affects all connections. It has no effect for in-memory databases. This setting can be appended to the database URL:

```
jdbc:h2:test;MAX_OPERATION_MEMORY=10000
```

Example:

```
SET MAX_OPERATION_MEMORY 0
```

## SET MODE

```
SET MODE { REGULAR | DB2 | DERBY | HSQLDB | MSSQLSERVER | MYSQL |  
ORACLE | POSTGRESQL }
```

Changes to another database compatibility mode. For details, see Compatibility Modes in the feature section.

This setting is not persistent. Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting can be appended to the database URL:

```
jdbc:h2:test;MODE=MYSQL
```

Example:

```
SET MODE HSQLDB
```

## SET MULTI\_THREADED

```
SET MULTI_THREADED { 0 | 1 }
```

Enabled (1) or disabled (0) multi-threading inside the database engine. MULTI\_THREADED is enabled by default with default MVStore storage engine. MULTI\_THREADED is disabled by default when using PageStore storage engine, enabling this with PageStore is experimental only.

This is a global setting, which means it is not possible to open multiple databases with different modes at the same time in the same virtual machine. This setting is not persistent, however the value is kept until the virtual machine exits or it is changed.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting can be appended to the database URL: jdbc:h2:test;MULTI\_THREADED=1

Example:

```
SET MULTI_THREADED 1
```

## SET OPTIMIZE\_REUSE\_RESULTS

```
SET OPTIMIZE_REUSE_RESULTS { 0 | 1 }
```

Enabled (1) or disabled (0) the result reuse optimization. If enabled, subqueries and views used as subqueries are only re-run if the data in one of the tables was changed. This option is enabled by default.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting can be appended to the database URL: jdbc:h2:test;OPTIMIZE\_REUSE\_RESULTS=0

Example:

```
SET OPTIMIZE_REUSE_RESULTS 0
```

## SET PASSWORD

```
SET PASSWORD string
```

Changes the password of the current user. The password must be in single quotes. It is case sensitive and can contain spaces.

This command commits an open transaction in this connection.

Example:

```
SET PASSWORD 'abcstzri!.5'
```

## SET QUERY\_STATISTICS

```
SET QUERY_STATISTICS { TRUE | FALSE }
```

Disabled or enables query statistics gathering for the whole database. The statistics are reflected in the INFORMATION\_SCHEMA.QUERY\_STATISTICS meta-table.

This setting is not persistent. This command commits an open transaction in this connection. Admin rights are required to execute this command, as it affects all connections.

Example:

```
SET QUERY_STATISTICS FALSE
```

## SET QUERY\_STATISTICS\_MAX\_ENTRIES

```
SET QUERY_STATISTICS int
```

Set the maximum number of entries in query statistics meta-table. Default value is 100.

This setting is not persistent. This command commits an open transaction in this connection. Admin rights are required to execute this command, as it affects all connections.

Example:

```
SET QUERY_STATISTICS_MAX_ENTRIES 500
```

## SET QUERY\_TIMEOUT

```
SET QUERY_TIMEOUT int
```

Set the query timeout of the current session to the given value. The timeout is in milliseconds. All kinds of statements will throw an exception if they take longer than the given value. The default timeout is 0, meaning no timeout.

This command does not commit a transaction, and rollback does not affect it.

Example:

```
SET QUERY_TIMEOUT 10000
```

## SET REFERENTIAL\_INTEGRITY

```
SET REFERENTIAL_INTEGRITY { TRUE | FALSE }
```

Disabled or enables referential integrity checking for the whole database. Enabling it does not check existing data. Use ALTER TABLE SET to disable it only for one table.



This setting is not persistent. This command commits an open transaction in this connection. Admin rights are required to execute this command, as it affects all connections.

Example:

```
SET REFERENTIAL_INTEGRITY FALSE
```

## SET RETENTION\_TIME

```
SET RETENTION_TIME int
```

This property is only used when using the MVStore storage engine. How long to retain old, persisted data, in milliseconds. The default is 45000 (45 seconds), 0 means overwrite data as early as possible. It is assumed that a file system and hard disk will flush all write buffers within this time. Using a lower value might be dangerous, unless the file system and hard disk flush the buffers earlier. To manually flush the buffers, use CHECKPOINT SYNC, however please note that according to various tests this does not always work as expected depending on the operating system and hardware.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting is persistent. This setting can be appended to the database URL:  
jdbc:h2:test;RETENTION\_TIME=0

Example:

```
SET RETENTION_TIME 0
```

## SET SALT HASH

```
SET SALT bytes HASH bytes
```

Sets the password salt and hash for the current user. The password must be in single quotes. It is case sensitive and can contain spaces.

This command commits an open transaction in this connection.

Example:

```
SET SALT '00' HASH '1122'
```

## SET SCHEMA

```
SET SCHEMA schemaName
```

Changes the default schema of the current connection. The default schema is used in statements where no schema is set explicitly. The default schema for new connections is PUBLIC.

This command does not commit a transaction, and rollback does not affect it. This setting can be appended to the database URL: jdbc:h2:test;SCHEMA=ABC

Example:

```
SET SCHEMA INFORMATION_SCHEMA
```

## SET SCHEMA\_SEARCH\_PATH

```
SET SCHEMA_SEARCH_PATH schemaName [...]
```

Changes the schema search path of the current connection. The default schema is used in statements where no schema is set explicitly. The default schema for new connections is PUBLIC.

This command does not commit a transaction, and rollback does not affect it. This setting can be appended to the database URL:  
jdbc:h2:test;SCHEMA\_SEARCH\_PATH=ABC,DEF

Example:

```
SET SCHEMA_SEARCH_PATH INFORMATION_SCHEMA, PUBLIC
```

## SET THROTTLE

```
SET THROTTLE int
```

Sets the throttle for the current connection. The value is the number of milliseconds delay after each 50 ms. The default value is 0 (throttling disabled).

This command does not commit a transaction, and rollback does not affect it. This setting can be appended to the database URL: jdbc:h2:test;THROTTLE=50

Example:

```
SET THROTTLE 200
```

## SET TRACE\_LEVEL

```
SET { TRACE_LEVEL_FILE | TRACE_LEVEL_SYSTEM_OUT } int
```

Sets the trace level for file the file or system out stream. Levels are: 0=off, 1=error, 2=info, 3=debug. The default level is 1 for file and 0 for system out. To use SLF4J, append ;TRACE\_LEVEL\_FILE=4 to the database URL when opening the database.

This setting is not persistent. Admin rights are required to execute this command, as it affects all connections. This command does not commit a transaction, and rollback does not affect it. This setting can be appended to the database URL: jdbc:h2:test;TRACE\_LEVEL\_SYSTEM\_OUT=3

Example:

```
SET TRACE_LEVEL_SYSTEM_OUT 3
```

## SET TRACE\_MAX\_FILE\_SIZE

```
SET TRACE_MAX_FILE_SIZE int
```

Sets the maximum trace file size. If the file exceeds the limit, the file is renamed to .old and a new file is created. If another .old file exists, it is deleted. The default max size is 16 MB.

This setting is persistent. Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting can be appended to the database URL: jdbc:h2:test;TRACE\_MAX\_FILE\_SIZE=3

Example:

```
SET TRACE_MAX_FILE_SIZE 10
```

## SET UNDO\_LOG

```
SET UNDO_LOG int
```

Enables (1) or disables (0) the per session undo log. The undo log is enabled by default. When disabled, transactions can not be rolled back. This setting should only be used for bulk operations that don't need to be atomic.

This command commits an open transaction in this connection.

Example:

```
SET UNDO_LOG 0
```

## SET WRITE\_DELAY

```
SET WRITE_DELAY int
```

Set the maximum delay between a commit and flushing the log, in milliseconds. This setting is persistent. The default is 500 ms.

Admin rights are required to execute this command, as it affects all connections. This command commits an open transaction in this connection. This setting can be appended to the database URL: jdbc:h2:test;WRITE\_DELAY=0

Example:

```
SET WRITE_DELAY 2000
```

## SHUTDOWN

```
SHUTDOWN [ IMMEDIATELY | COMPACT | DEFRAG ]
```

This statement closes all open connections to the database and closes the database. This command is usually not required, as the database is closed automatically when the last connection to it is closed.

If no option is used, then the database is closed normally. All connections are closed, open transactions are rolled back.

SHUTDOWN COMPACT fully compacts the database (re-creating the database may further reduce the database size). If the database is closed normally (using SHUTDOWN or by closing all connections), then the database is also compacted, but only for at most the time defined by the database setting `h2.maxCompactTime` in milliseconds (see there).

SHUTDOWN IMMEDIATELY closes the database files without any cleanup and without compacting.

SHUTDOWN DEFRAG re-orders the pages when closing the database so that table scans are faster. In case of MVStore it is currently equivalent to COMPACT.

Admin rights are required to execute this command.

Example:

## Datetime fields

### Datetime field

yearField | monthField | dayOfMonthField  
| hourField | minuteField | secondField  
| millisecondField | microsecondField | nanosecondField  
| timezoneHourField | timezoneMinuteField  
| dayOfWeekField | isoWeekYearField | isoDayOfWeekField  
| weekOfYearField | isoWeekOfYearField  
| quarterField | dayOfYearField | epochField

Fields for EXTRACT, DATEADD, and DATEDIFF functions.

Example:

YEAR

### Year field

YEAR | YYYY | YY | SQL\_TSI\_YEAR

Year.

Example:

YEAR

### Month field

MONTH | MM | M | SQL\_TSI\_MONTH

Month (1-12).

Example:

MONTH

## Day of month field

DAY | DD | D | SQL\_TSI\_DAY

Day of month (1-31).

Example:

DAY

## Hour field

HOUR | HH | SQL\_TSI\_HOUR

Hour (0-23).

Example:

HOUR

## Minute field

MINUTE | MI | N | SQL\_TSI\_MINUTE

Minute (0-59).

Example:

MINUTE

## Second field

SECOND | SS | S | SQL\_TSI\_SECOND

Second (0-59).

Example:

SECOND

## Millisecond field

MILLISECOND | MS

Millisecond (0-999).

Example:

MILLISECOND

### **Microsecond field**

MICROSECOND | MCS

Microsecond (0-999999).

Example:

MICROSECOND

### **Nanosecond field**

NANOSECOND | NS

Nanosecond (0-999999999).

Example:

NANOSECOND

### **Timezone hour field**

TIMEZONE\_HOUR

Timezone hour (from -18 to +18).

Example:

TIMEZONE\_HOUR

### **Timezone minute field**

TIMEZONE\_MINUTE

Timezone minute (from -59 to +59).

Example:

TIMEZONE\_MINUTE

## Day of week field

DAY\_OF\_WEEK | DAYOFWEEK | DOW

Day of week (1-7). Sunday is 1.

Example:

DAY\_OF\_WEEK

## ISO week year field

ISO\_YEAR | ISOYEAR

Returns the ISO week year from a date/time value.

Example:

ISO\_YEAR

## ISO day of week field

ISO\_DAY\_OF\_WEEK | ISODOW

ISO day of week (1-7). Monday is 1.

Example:

ISO\_DAY\_OF\_WEEK

## Week of year field

WEEK | WW | W | SQL\_TSI\_WEEK

Week of year (1-53). EXTRACT function uses local rules to get number of week in year. DATEDIFF function uses Sunday as a first day of week.

Example:

WEEK



## ISO week of year field

ISO\_WEEK

ISO week of year (1-53). ISO definition is used when first week of year should have at least four days and week is started with Monday.

Example:

ISO\_WEEK

## Quarter field

QUARTER

Quarter (1-4).

Example:

QUARTER

## Day of year field

DAYOFYEAR | DAY\_OF\_YEAR | DOY | DY

Day of year (1-366).

Example:

DAYOFYEAR

## Epoch field

EPOCH

For TIMESTAMP values number of seconds since 1970-01-01 00:00:00 in local time zone. For TIMESTAMP WITH TIME ZONE values number of seconds since 1970-01-01 00:00:00 in UTC time zone. For DATE values number of seconds since 1970-01-01. For TIME values number of seconds since midnight.

Example:

EPOCH

## Other Grammar

### Alias

`name`

An alias is a name that is only valid in the context of the statement.

Example:

A

### And Condition

`condition [ { AND condition } [...] ]`

Value or condition.

Example:

ID=1 AND NAME='Hi'

### Array

`ARRAY '[' [ expression, [...] ] '['`

An array of values.

Example:

ARRAY[1, 2]

ARRAY[1]

ARRAY[]

### Boolean

`TRUE | FALSE`

A boolean value.

Example:

TRUE

## Bytes

```
X'hex'
```

A binary value. The hex value is not case sensitive.

Example:

```
X'01FF'
```

## Case

```
CASE expression { WHEN expression THEN expression } [...]  
[ ELSE expression ] END
```

Returns the first expression where the value is equal to the test expression. If no else part is specified, return NULL.

Example:

```
CASE CNT WHEN 0 THEN 'No' WHEN 1 THEN 'One' ELSE 'Some' END
```

## Case When

```
CASE { WHEN expression THEN expression } [...]  
[ ELSE expression ] END
```

Returns the first expression where the condition is true. If no else part is specified, return NULL.

Example:

```
CASE WHEN CNT<10 THEN 'Low' ELSE 'High' END
```

## Cipher

```
AES
```

Only the algorithm AES (AES-128) is supported currently.

Example:

```
AES
```

## Column Definition

```
dataType [ VISIBLE | INVISIBLE ]  
[ { DEFAULT expression  
| AS computedColumnExpression  
| GENERATED {ALWAYS | BY DEFAULT} AS IDENTITY [(sequenceOptions)] } ]  
[ ON UPDATE expression ] [ [ NOT ] NULL ]  
[ { AUTO_INCREMENT | IDENTITY } [ ( startInt [, incrementInt ] ) ] ]  
[ SELECTIVITY selectivity ] [ COMMENT expression ]  
[ PRIMARY KEY [ HASH ] | UNIQUE ] [ CHECK condition ]
```

Default expressions are used if no explicit value was used when adding a row. The computed column expression is evaluated and assigned whenever the row changes. On update column expression is used if row is updated, at least one column have a new value that is different from its previous value and value for this column is not set explicitly in update statement.

Identity, auto-increment, or generated as identity columns are columns with a sequence as the default. The column declared as the identity columns with IDENTITY data type or with IDENTITY () clause is implicitly the primary key column of this table. AUTO\_INCREMENT and GENERATED clauses do not create the primary key constraint. GENERATED ALWAYS is accepted by treated in the same way as GENERATED BY DEFAULT.

The invisible column will not be displayed as a result of SELECT \* query. Otherwise, it works as normal column.

The options PRIMARY KEY, UNIQUE, and CHECK are not supported for ALTER statements.

Check constraints can reference columns of the table, and they can reference objects that exist while the statement is executed. Conditions are only checked when a row is added or modified in the table where the constraint exists.

Example:

```
CREATE TABLE TEST(ID INT PRIMARY KEY, NAME VARCHAR(255) DEFAULT "");  
CREATE TABLE TEST(ID BIGINT IDENTITY);  
CREATE TABLE TEST(QUANTITY INT, PRICE DECIMAL, AMOUNT DECIMAL AS  
QUANTITY*PRICE);
```

## Comments

```
-- anything | // anything | /* anything */
```

Comments can be used anywhere in a command and are ignored by the database. Line comments end with a newline. Block comments cannot be nested, but can be multiple lines long.

Example:

```
// This is a comment
```

## Compare

```
<> | <= | >= | = | < | > | != | &&
```

Comparison operator. The operator != is the same as <>. The operator && means overlapping; it can only be used with geometry types.

Example:

```
<>
```

## Condition

```
operand [ conditionRightHandSide ] | NOT condition | EXISTS ( select )
```

Boolean value or condition.

Example:

```
ID<>2
```

## Condition Right Hand Side

```
compare { { { ALL | ANY | SOME } ( select ) } | operand }  
| IS [ NOT ] NULL  
| IS [ NOT ] [ DISTINCT FROM ] operand  
| BETWEEN operand AND operand  
| IN ( { select | expression [,...] } )  
| [ NOT ] [ LIKE | ILIKE ] operand [ ESCAPE string ]  
| [ NOT ] REGEXP operand
```

The right hand side of a condition.

Quantified comparison predicate ALL returns TRUE if specified comparison operation between left side of condition and each row from a subquery returns

TRUE, including case when there are no rows. ALL predicate returns FALSE if at least one such comparison returns FALSE. Otherwise it returns NULL.

Quantified comparison predicates ANY and SOME return TRUE if specified comparison operation between left side of condition and at least one row from a subquery returns TRUE. ANY and SOME predicates return FALSE if all such comparisons return FALSE. Otherwise it returns NULL. Note that these predicates have priority over ANY and SOME aggregate functions with subquery on the right side. Use parentheses around aggregate function.

The conditions IS [ NOT ] and IS [ NOT ] DISTINCT FROM are null-safe, meaning NULL is considered the same as NULL, and the condition never evaluates to NULL.

When comparing with LIKE, the wildcard characters are \_ (any one character) and % (any characters). The database uses an index when comparing with LIKE except if the operand starts with a wildcard. To search for the characters % and \_, the characters need to be escaped. The default escape character is \ (backslash). To select no escape character, use ESCAPE '' (empty string). At most one escape character is allowed. Each character that follows the escape character in the pattern needs to match exactly. Patterns that end with an escape character are invalid and the expression returns NULL.

ILIKE does a case-insensitive compare.

When comparing with REGEXP, regular expression matching is used. See Java Matcher.find for details.

Example:

```
VALUE > 10
A IS NOT DISTINCT FROM B
LIKE 'Jo%'
```

## Constraint

```
[ constraintNameDefinition ]
{ CHECK expression
| UNIQUE ( columnName [,...] )
| referentialConstraint
| PRIMARY KEY [ HASH ] ( columnName [,...] ) }
```

Defines a constraint. The check condition must evaluate to TRUE, FALSE or NULL. TRUE and NULL mean the operation is to be permitted, and FALSE means the operation is to be rejected. To prevent NULL in a column, use NOT NULL instead of a check constraint.

Example:

PRIMARY KEY(ID, NAME)

## Constraint Name Definition

```
CONSTRAINT [ IF NOT EXISTS ] newConstraintName
```

Defines a constraint name.

Example:

CONSTRAINT CONST\_ID

## Csv Options

```
charsetString [, fieldSepString [, fieldDelimString [, escString [, nullString]]]]  
| optionString
```

Optional parameters for CSVREAD and CSVWRITE. Instead of setting the options one by one, all options can be combined into a space separated key-value pairs, as follows: STRINGDECODE('charset=UTF-8 escape=\" fieldDelimiter=\" fieldSeparator=, ' || 'lineComment=# lineSeparator=\n null= rowSeparator='). The following options are supported:

caseSensitiveColumnNames (true or false; disabled by default),

charset (for example 'UTF-8'),

escape (the character that escapes the field delimiter),

fieldDelimiter (a double quote by default),

fieldSeparator (a comma by default),

lineComment (disabled by default),

lineSeparator (the line separator used for writing; ignored for reading),

null, Support reading existing CSV files that contain explicit null delimiters. Note that an empty, unquoted values are also treated as null.

preserveWhitespace (true or false; disabled by default),

writeColumnHeader (true or false; enabled by default).

For a newline or other special character, use STRINGDECODE as in the example above. A space needs to be escaped with a backslash ('\ '), and a backslash needs to be escaped with another backslash ('\\'). All other characters are not to be escaped, that means newline and tab characters are written as such.

Example:

```
CALL CSVWRITE('test2.csv', 'SELECT * FROM TEST', 'charset=UTF-8  
fieldSeparator=|');
```

## Data Type

[intType](#) | [booleanType](#) | [tinyintType](#) | [smallintType](#) | [bigintType](#) | [identityType](#)  
| [decimalType](#) | [doubleType](#) | [realType](#) | [dateType](#) | [timeType](#) | [timestampType](#)  
| [timestampWithTimeZoneType](#) | [binaryType](#) | [otherType](#) | [varcharType](#)  
| [varcharIgnorecaseType](#) | [charType](#) | [blobType](#) | [clobType](#) | [uuidType](#)  
| [arrayType](#) | [enumType](#) | [intervalType](#)

A data type definition.

Example:

INT

## Date

DATE 'yyyy-MM-dd'

A date literal. The limitations are the same as for the Java data type `java.sql.Date`, but for compatibility with other databases the suggested minimum and maximum years are 0001 and 9999.

Example:

DATE '2004-12-31'

## Decimal

[ + | - ] { { [number](#) [ . [number](#) ] } | { . [number](#) } }  
[ E [ + | - ] [expNumber](#) [...] ] ]

A decimal number with fixed precision and scale. Internally, `java.lang.BigDecimal` is used. To ensure the floating point representation is used, use `CAST(X AS DOUBLE)`. There are some special decimal values: to represent positive infinity,



use POWER(0, -1); for negative infinity, use (-POWER(0, -1)); for -0.0, use (-CAST(0 AS DOUBLE)); for NaN (not a number), use SQRT(-1).

Example:

```
SELECT -1600.05
SELECT CAST(0 AS DOUBLE)
SELECT -1.4e-10
```

## Digit

0-9

A digit.

Example:

0

## Dollar Quoted String

\$\$anything\$\$

A string starts and ends with two dollar signs. Two dollar signs are not allowed within the text. A whitespace is required before the first set of dollar signs. No escaping is required within the text.

Example:

\$\$John's car\$\$

## Expression

`andCondition` [ { OR `andCondition` } [...] ]

Value or condition.

Example:

ID=1 OR NAME='Hi'

## Factor

```
term [ { { * | / | % } term } [...] ]
```

A value or a numeric factor.

Example:

ID \* 10

## Hex

```
{ { digit | a-f | A-F } { digit | a-f | A-F } } [...]
```

The hexadecimal representation of a number or of bytes. Two characters are one byte.

Example:

cafe

## Hex Number

```
[ + | - ] 0x hex
```

A number written in hexadecimal notation.

Example:

0xff

## Index Column

```
columnName [ ASC | DESC ] [ NULLS { FIRST | LAST } ]
```

Indexes this column in ascending or descending order. Usually it is not required to specify the order; however doing so will speed up large queries that order the column in the same way.

Example:

NAME

## Insert columns and source

```
{ [ ( columnName [,...] ) ]  
{ insertValues | [ DIRECT ] [ SORTED ] select | DEFAULT VALUES } }  
| { SET { columnName = { DEFAULT | expression } } [,...] }
```

Names of columns and their values for INSERT statement.

Example:

(ID, NAME) VALUES (1, 'Test')

## Insert values

```
VALUES { DEFAULT|expression | [ROW] ({DEFAULT|expression} [,...]) }, [,...]
```

Values for INSERT statement.

Example:

VALUES (1, 'Test')

## Int

```
[ + | - ] number
```

The maximum integer number is 2147483647, the minimum is -2147483648.

Example:

10

## Long

```
[ + | - ] number
```

Long numbers are between -9223372036854775808 and 9223372036854775807.

Example:

100000

## Merge when clause

```
mergeWhenMatchedClause | mergeWhenNotMatchedClause
```

WHEN MATCHED or WHEN NOT MATCHED clause for MERGE USING command.

Example:

WHEN MATCHED THEN DELETE

## Merge when matched clause

```
WHEN MATCHED [ AND expression ] THEN  
UPDATE SET setClauseList  
| DELETE  
| {UPDATE SET setClauseList [ WHERE expression ] DELETE [ WHERE expression  
]}
```

WHEN MATCHED clause for MERGE USING command.

If both UPDATE and DELETE are specified, DELETE can delete only rows that were updated, WHERE condition in DELETE clause can be used to specify which updated rows should be deleted. This condition checks values in updated row.

Example:

WHEN MATCHED THEN UPDATE SET VALUE = S.VALUE  
WHEN MATCHED THEN DELETE

## Merge when not matched clause

```
WHEN NOT MATCHED [ AND expression ] THEN INSERT  
insertColumnsAndSource
```

WHEN NOT MATCHED clause for MERGE USING command.

Example:

WHEN NOT MATCHED THEN INSERT (ID, NAME) VALUES (S.ID, S.NAME)

## Name

```
{ { A-Z|_ } [ { A-Z|_|0-9 } [...] ] } | quotedName
```

Names are not case sensitive. There is no maximum name length.

Example:

TEST

## Null

NULL

NULL is a value without data type and means 'unknown value'.

Example:

NULL

## Number

digit [...]

The maximum length of the number depends on the data type used.

Example:

100

## Numeric

decimal | int | long | hexNumber

The data type of a numeric value is always the lowest possible for the given value. If the number contains a dot this is decimal; otherwise it is int, long, or decimal (depending on the value).

Example:

```
SELECT -1600.05
SELECT CAST(0 AS DOUBLE)
SELECT -1.4e-10
```

## Operand

summand [ { || summand } [...] ]

A value or a concatenation of values. In the default mode, the result is NULL if either parameter is NULL.

Example:

'Hi' || ' Eva'

## Order

```
{ int | expression } [ ASC | DESC ] [ NULLS { FIRST | LAST } ]
```

Sorts the result by the given column number, or by an expression. If the expression is a single parameter, then the value is interpreted as a column number. Negative column numbers reverse the sort order.

Example:

NAME DESC NULLS LAST

## Quoted Name

```
"anything"
```

Quoted names are case sensitive, and can contain spaces. There is no maximum name length. Two double quotes can be used to create a single double quote inside an identifier.

Example:

"FirstName"

## Referential Constraint

```
FOREIGN KEY ( columnName [,...] )  
REFERENCES [ refTableName ] [ ( refColumnName [,...] ) ]  
[ ON DELETE referentialAction ] [ ON UPDATE referentialAction ]
```

Defines a referential constraint. If the table name is not specified, then the same table is referenced. RESTRICT is the default action. If the referenced columns are not specified, then the primary key columns are used. The required indexes are automatically created if required. Some tables may not be referenced, such as metadata tables.

Example:

FOREIGN KEY(ID) REFERENCES TEST(ID)

## Referential Action

```
CASCADE | RESTRICT | NO ACTION | SET { DEFAULT | NULL }
```

The action CASCADE will cause conflicting rows in the referencing (child) table to be deleted or updated. RESTRICT is the default action. As this database does not support deferred checking, RESTRICT and NO ACTION will both throw an exception if the constraint is violated. The action SET DEFAULT will set the column in the referencing (child) table to the default value, while SET NULL will set it to NULL.

Example:

```
FOREIGN KEY(ID) REFERENCES TEST(ID) ON UPDATE CASCADE
```

## Script Compression Encryption

```
[ COMPRESSION { DEFLATE | LZF | ZIP | GZIP } ]  
[ CIPHER cipher PASSWORD string ]
```

The compression and encryption algorithm to use for script files. When using encryption, only DEFLATE and LZF are supported. LZF is faster but uses more space.

Example:

```
COMPRESSION LZF
```

## Row value expression

```
ROW (expression, [...])  
| ( [ expression, expression [...] ] )  
| expression
```

A row value expression.

Example:

```
ROW (1)  
(1, 2)  
1
```

## Select Expression

```
wildcardExpression | expression [ [ AS ] columnAlias ]
```

An expression in a SELECT statement.

Example:

ID AS VALUE

## Sequence options

```
[ START WITH long ]  
[ INCREMENT BY long ]  
[ MINVALUE long | NOMINVALUE | NO MINVALUE ]  
[ MAXVALUE long | NOMAXVALUE | NO MAXVALUE ]  
[ CYCLE long | NOCYCLE | NO CYCLE ]  
[ CACHE long | NOCACHE | NO CACHE ]
```

Options of a sequence.

Example:

START WITH 1

## Set clause list

```
{ { columnName = { DEFAULT | expression } } [,...] }  
| { ( columnName [,...] ) = { rowValueExpression | (select) } }
```

List of SET clauses.

Example:

NAME = 'Test', VALUE = 2

(A, B) = (1, 2)

(A, B) = (SELECT X, Y FROM OTHER T2 WHERE T1.ID = T2.ID)

## String

```
'anything'
```

A string starts and ends with a single quote. Two single quotes can be used to create a single quote inside a string.



Example:

'John's car'

## Summand

```
factor [ { { + | - } factor } [...] ]
```

A value or a numeric sum.

Please note the text concatenation operator is ||.

Example:

ID + 20

## Table Expression

```
{ [ schemaName. ] tableName | ( select ) | valuesExpression }  
[ [ AS ] newTableAlias [ ( columnName [,...] ) ] ]  
[ USE INDEX ( [ indexName [,...] ] ) ]  
[ { { LEFT | RIGHT } [ OUTER ] | [ INNER ] | CROSS | NATURAL }  
JOIN tableExpression [ ON expression ] ]
```

Joins a table. The join expression is not supported for cross and natural joins. A natural join is an inner join, where the condition is automatically on the columns with the same name.

Example:

TEST AS T LEFT JOIN TEST AS T1 ON T.ID = T1.ID

## Within group specification

```
WITHIN GROUP (ORDER BY {expression [ASC|DESC]} [,...])
```

Group specification for ordered set functions.

Example:

WITHIN GROUP (ORDER BY ID DESC)

## Wildcard expression

```
{* | tableAlias.*} [EXCEPT ([tableAlias.]columnName, [...])]
```

A wildcard expression in a SELECT statement. A wildcard expression represents all visible columns. Some columns can be excluded with optional EXCEPT clause.

Example:

```
*  
* EXCEPT (DATA)
```

## Window name or specification

```
windowName | windowSpecification
```

A window name or inline specification for a window function or aggregate.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
W1  
(ORDER BY ID)
```

## Window specification

```
([existingWindowName]  
[PARTITION BY expression [...]] [ORDER BY order [...]]  
[windowFrame])
```

A window specification for a window, window function or aggregate.

If name of an existing window is specified its clauses are used by default.

Optional window partition clause separates rows into independent partitions. Each partition is processed separately. If this clause is not present there is one implicit partition with all rows.

Optional window order clause specifies order of rows in the partition. If some rows have the same order position they are considered as a group of rows in optional window frame clause.

Optional window frame clause specifies which rows are processed by a window function, see its documentation for a more details.

Example:

```
()  
(W1 ORDER BY ID)  
(PARTITION BY CATEGORY)  
(PARTITION BY CATEGORY ORDER BY NAME, ID)  
(ORDER BY Y RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW  
EXCLUDE TIES)
```

## Window frame

```
ROWS|RANGE|GROUP  
{windowFramePreceding|BETWEEN windowFrameBound AND  
windowFrameBound}  
[EXCLUDE {CURRENT ROW|GROUP|TIES|NO OTHERS}]
```

A window frame clause. May be specified only for aggregates and FIRST\_VALUE(), LAST\_VALUE(), and NTH\_VALUE() window functions.

If this clause is not specified for an aggregate or window function that supports this clause the default window frame depends on window order clause. If window order clause is also not specified the default window frame contains all the rows in the partition. If window order clause is specified the default window frame contains all preceding rows and all rows from the current group.

Window frame unit determines how rows or groups of rows are selected and counted. If ROWS is specified rows are not grouped in any way and relative numbers of rows are used in bounds. If RANGE is specified rows are grouped according window order clause, preceding and following values mean the difference between value in the current row and in the target rows, and CURRENT ROW in bound specification means current group of rows. If GROUPS is specified rows are grouped according window order clause, preceding and following values means relative number of groups of rows, and CURRENT ROW in bound specification means current group of rows.

If only window frame preceding clause is specified it is treated as BETWEEN [windowFramePreceding](#) AND CURRENT ROW.

Optional window frame exclusion clause specifies rows that should be excluded from the frame. EXCLUDE CURRENT ROW excludes only the current row regardless the window frame unit. EXCLUDE GROUP excludes the whole current group of rows, including the current row. EXCLUDE TIES excludes the current

group of rows, but not the current row. EXCLUDE NO OTHERS is default and it does not exclude anything.

Example:

ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW EXCLUDE TIES

## Window frame preceding

```
UNBOUNDED PRECEDING|value PRECEDING|CURRENT ROW
```

A window frame preceding clause. If value is specified it should not be negative.

Example:

```
UNBOUNDED PRECEDING  
1 PRECEDING  
CURRENT ROW
```

## Window frame bound

```
UNBOUNDED PRECEDING|value PRECEDING|CURRENT ROW  
|value FOLLOWING|UNBOUNDED FOLLOWING
```

A window frame bound clause. If value is specified it should not be negative.

Example:

```
UNBOUNDED PRECEDING  
UNBOUNDED FOLLOWING  
1 FOLLOWING  
CURRENT ROW
```

## Term

```
value  
| column  
| ?[ int ]  
| NEXT VALUE FOR sequenceName  
| function  
| { - | + } term  
| ( expression )  
| select  
| case
```

```
| caseWhen  
| userDefinedFunctionName
```

A value. Parameters can be indexed, for example ?1 meaning the first parameter.

Example:

'Hello'

## Column

```
[[schemaName.]tableAlias.] { columnName | _ROWID_ }
```

A column name with optional table alias and schema. \_ROWID\_ can be used to access unique row identifier.

Example:

ID

## Time

```
TIME [ WITHOUT TIME ZONE ] 'hh:mm:ss[.nnnnnnnnnn]'
```

A time literal. A value is between 0:00:00 and 23:59:59.999999999 and has nanosecond resolution.

Example:

TIME '23:59:59'

## Timestamp

```
TIMESTAMP [ WITHOUT TIME ZONE ] 'yyyy-MM-dd hh:mm:ss[.nnnnnnnnnn]'
```

A timestamp literal. The limitations are the same as for the Java data type `java.sql.Timestamp`, but for compatibility with other databases the suggested minimum and maximum years are 0001 and 9999.

Example:

TIMESTAMP '2005-12-31 23:59:59'

## Timestamp with time zone

```
TIMESTAMP WITH TIME ZONE 'yyyy-MM-dd hh:mm:ss[.nnnnnnnnnn]  
[Z | { - | + } timeZoneOffsetString | timeZoneNameString ]'
```

A timestamp with time zone literal. If name of time zone is specified it will be converted to time zone offset.

Example:

```
TIMESTAMP WITH TIME ZONE '2005-12-31 23:59:59Z'  
TIMESTAMP WITH TIME ZONE '2005-12-31 23:59:59-10:00'  
TIMESTAMP WITH TIME ZONE '2005-12-31 23:59:59.123+05'  
TIMESTAMP WITH TIME ZONE '2005-12-31 23:59:59.123456789 Europe/London'
```

## Date and time

```
date | time | timestamp | timestampWithTimeZone
```

A literal value of any date-time data type.

Example:

```
TIMESTAMP '1999-01-31 10:00:00'
```

## INTERVAL YEAR

```
INTERVAL [-|+] '[-|+]yearInt' YEAR
```

An INTERVAL YEAR literal.

Example:

```
INTERVAL '10' YEAR
```

## INTERVAL MONTH

```
INTERVAL [-|+] '[-|+]monthInt' MONTH
```

An INTERVAL MONTH literal.

Example:

```
INTERVAL '10' MONTH
```

## INTERVAL DAY

```
INTERVAL [-|+] '[-|+]dayInt' DAY
```

An INTERVAL DAY literal.

Example:

```
INTERVAL '10' DAY
```

## INTERVAL HOUR

```
INTERVAL [-|+] '[-|+]hourInt' HOUR
```

An INTERVAL HOUR literal.

Example:

```
INTERVAL '10' HOUR
```

## INTERVAL MINUTE

```
INTERVAL [-|+] '[-|+]minuteInt' MINUTE
```

An INTERVAL MINUTE literal.

Example:

```
INTERVAL '10' MINUTE
```

## INTERVAL SECOND

```
INTERVAL [-|+] '[-|+]secondInt[.nnnnnnnnnn]' SECOND
```

An INTERVAL SECOND literal.

Example:

```
INTERVAL '10.123' SECOND
```

## INTERVAL YEAR TO MONTH

```
INTERVAL [-|+] '[-|+]yearInt-monthInt' YEAR TO MONTH
```

An INTERVAL YEAR TO MONTH literal.

Example:

INTERVAL '1-6' YEAR TO MONTH

## INTERVAL DAY TO HOUR

```
INTERVAL [-|+] '[-|+]dayInt hoursInt' DAY TO HOUR
```

An INTERVAL DAY TO HOUR literal.

Example:

INTERVAL '10 11' DAY TO HOUR

## INTERVAL DAY TO MINUTE

```
INTERVAL [-|+] '[-|+]dayInt hh:mm' DAY TO MINUTE
```

An INTERVAL DAY TO MINUTE literal.

Example:

INTERVAL '10 11:12' DAY TO MINUTE

## INTERVAL DAY TO SECOND

```
INTERVAL [-|+] '[-|+]dayInt hh:mm:ss[.nnnnnnnnnn]' DAY TO SECOND
```

An INTERVAL DAY TO SECOND literal.

Example:

INTERVAL '10 11:12:13.123' DAY TO SECOND

## INTERVAL HOUR TO MINUTE

```
INTERVAL [-|+] '[-|+]hh:mm' HOUR TO MINUTE
```

An INTERVAL HOUR TO MINUTE literal.

Example:



INTERVAL '10:11' HOUR TO MINUTE

## INTERVAL HOUR TO SECOND

INTERVAL [-|+] '[-|+]hh:mm:ss[.nnnnnnnnnn]' HOUR TO SECOND

An INTERVAL HOUR TO SECOND literal.

Example:

INTERVAL '10:11:12.123' HOUR TO SECOND

## INTERVAL MINUTE TO SECOND

INTERVAL [-|+] '[-|+]mm:ss[.nnnnnnnnnn]' MINUTE TO SECOND

An INTERVAL MINUTE TO SECOND literal.

Example:

INTERVAL '11:12.123' MINUTE TO SECOND

## Interval

[intervalYear](#) | [intervalMonth](#) | [intervalDay](#) | [intervalHour](#) | [intervalMinute](#)  
| [intervalSecond](#) | [intervalYearToMonth](#) | [intervalDayToHour](#)  
| [intervalDayToMinute](#) | [intervalDayToSecond](#) | [intervalHourToMinute](#)  
| [intervalHourToSecond](#) | [intervalMinuteToSecond](#)

An interval literal.

Example:

INTERVAL '1-2' YEAR TO MONTH

## Value

[string](#) | [dollarQuotedString](#) | [numeric](#) | [dateAndTime](#) | [boolean](#) | [bytes](#)  
| [interval](#) | [array](#) | [null](#)

A literal value of any data type, or null.

Example:



# Functions

## Index

### Aggregate Functions

AVG  
BIT\_AND  
BIT\_OR  
EVERY  
ANY  
COUNT  
LISTAGG  
ARRAY\_AGG  
MAX  
MIN  
SUM  
SELECTIVITY  
STDDEV\_POP  
STDDEV\_SAMP  
VAR\_POP  
VAR\_SAMP  
RANK aggregate  
DENSE\_RANK aggregate  
PERCENT\_RANK aggregate  
CUME\_DIST aggregate  
PERCENTILE\_CONT  
PERCENTILE\_DISC  
MEDIAN  
MODE  
ENVELOPE

### Numeric Functions

ABS  
ACOS  
ASIN  
ATAN  
COS  
COSH  
COT  
SIN

SINH  
TAN  
TANH  
ATAN2  
BITAND  
BITGET  
BITOR  
BITXOR  
MOD  
CEILING  
DEGREES  
EXP  
FLOOR  
LN  
LOG10  
ORA\_HASH  
RADIANS  
SQRT  
PI  
POWER  
RAND  
RANDOM\_UUID  
ROUND  
ROUNDMAGIC  
SECURE\_RAND  
SIGN  
ENCRYPT  
DECRYPT  
HASH  
TRUNCATE  
COMPRESS  
EXPAND  
ZERO

## **String Functions**

ASCII  
BIT\_LENGTH  
LENGTH  
OCTET\_LENGTH  
CHAR  
CONCAT  
CONCAT\_WS  
DIFFERENCE

HEXTORAW  
RAWTOHEX  
INSTR  
INSERT Function  
LOWER  
UPPER  
LEFT  
RIGHT  
LOCATE  
POSITION  
LPAD  
RPAD  
LTRIM  
RTRIM  
TRIM  
REGEXP\_REPLACE  
REGEXP\_LIKE  
REPEAT  
REPLACE  
SOUNDEX  
SPACE  
STRINGDECODE  
STRINGENCODE  
STRINGTOUTF8  
SUBSTRING  
UTF8TOSTRING  
XMLATTR  
XMLNODE  
XMLCOMMENT  
XMLCDATA  
XMLSTARTDOC  
XMLTEXT  
TO\_CHAR  
TRANSLATE

## **Time and Date Functions**

CURRENT\_DATE  
CURRENT\_TIME  
CURRENT\_TIMESTAMP  
LOCALTIME  
LOCALTIMESTAMP  
DATEADD  
DATEDIFF

DAYNAME  
DAY\_OF\_MONTH  
DAY\_OF\_WEEK  
ISO\_DAY\_OF\_WEEK  
DAY\_OF\_YEAR  
EXTRACT  
FORMATDATETIME  
HOUR  
MINUTE  
MONTH  
MONTHNAME  
PARSEDATETIME  
QUARTER  
SECOND  
WEEK  
ISO\_WEEK  
YEAR  
ISO\_YEAR

## **System Functions**

ARRAY\_GET  
ARRAY\_LENGTH  
ARRAY\_CONTAINS  
ARRAY\_CAT  
ARRAY\_APPEND  
ARRAY\_SLICE  
AUTOCOMMIT  
CANCEL\_SESSION  
CASEWHEN Function  
CAST  
COALESCE  
CONVERT  
CURRVAL  
CSVREAD  
CSVWRITE  
DATABASE  
DATABASE\_PATH  
DECODE  
DISK\_SPACE\_USED  
SIGNAL  
ESTIMATED\_ENVELOPE  
FILE\_READ  
FILE\_WRITE

GREATEST  
IDENTITY  
IFNULL  
LEAST  
LOCK\_MODE  
LOCK\_TIMEOUT  
LINK\_SCHEMA  
MEMORY\_FREE  
MEMORY\_USED  
NEXTVAL  
NULLIF  
NVL2  
READONLY  
ROWNUM  
SCHEMA  
SCOPE\_IDENTITY  
SESSION\_ID  
SET  
TABLE  
TRANSACTION\_ID  
TRUNCATE\_VALUE  
UNNEST  
USER  
H2VERSION

## **Window Functions**

ROW\_NUMBER  
RANK  
DENSE\_RANK  
PERCENT\_RANK  
CUME\_DIST  
NTILE  
LEAD  
LAG  
FIRST\_VALUE  
LAST\_VALUE  
NTH\_VALUE  
RATIO\_TO\_REPORT

# Aggregate Functions

## AVG

```
AVG ( [ DISTINCT|ALL ] { numeric } )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The average (mean) value. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements. The returned value is of the same data type as the parameter.

Example:

AVG(X)

## BIT\_AND

```
BIT_AND(expression)  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The bitwise AND of all non-null values. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

BIT\_AND(ID)

## BIT\_OR

```
BIT_OR(expression)  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The bitwise OR of all non-null values. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

BIT\_OR(ID)

## EVERY

```
{EVERY|BOOL_AND}(boolean)  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```



Returns true if all expressions are true. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
EVERY(ID>10)
```

## ANY

```
{ANY|SOME|BOOL_OR}(boolean)
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns true if any expression is true. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Note that if ANY or SOME aggregate function is placed on the right side of comparison operation and argument of this function is a subquery additional parentheses around aggregate function are required, otherwise it will be parsed as quantified comparison predicate.

Example:

```
ANY(NAME LIKE 'W%')
A = (ANY((SELECT B FROM T)))
```

## COUNT

```
COUNT( { * | { [ DISTINCT|ALL ] expression } } )
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The count of all row, or of the non-null values. This method returns a long. If no rows are selected, the result is 0. Aggregates are only allowed in select statements.

Example:

```
COUNT(*)
```

## LISTAGG

```
{ LISTAGG ( [ DISTINCT|ALL ] string [, separatorString] [ ON OVERFLOW
ERROR ] )
withinGroupSpecification }
| { GROUP_CONCAT ( [ DISTINCT|ALL ] string
```

```
[ ORDER BY { expression [ ASC | DESC ] } [,...] ]  
[ SEPARATOR separatorString ] ) }  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Concatenates strings with a separator. Separator must be the same for all rows in the same group. The default separator is a ',' (without space). This method returns a string. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
LISTAGG(NAME, ' ' ) WITHIN GROUP (ORDER BY ID)  
LISTAGG(ID, ' ' ) WITHIN GROUP (ORDER BY ID) OVER (ORDER BY ID)
```

## ARRAY\_AGG

```
ARRAY_AGG ( [ DISTINCT|ALL ] string  
[ ORDER BY { expression [ ASC | DESC ] } [,...] ] )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Aggregate the value into an array. This method returns an array. If no rows are selected, the result is NULL. If ORDER BY is not specified order of values is not determined. When this aggregate is used with OVER clause that contains ORDER BY subclause it does not enforce exact order of values. This aggregate needs additional own ORDER BY clause to make it deterministic. Aggregates are only allowed in select statements.

Example:

```
ARRAY_AGG(NAME ORDER BY ID)  
ARRAY_AGG(ID ORDER BY ID) OVER (ORDER BY ID)
```

## MAX

```
MAX(value)  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The highest value. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements. The returned value is of the same data type as the parameter.

Example:

```
MAX(NAME)
```

## MIN

```
MIN(value)  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The lowest value. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements. The returned value is of the same data type as the parameter.

Example:

```
MIN(NAME)
```

## SUM

```
SUM( [ DISTINCT|ALL ] { numeric } )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The sum of all values. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements. The data type of the returned value depends on the parameter data type like this: BOOLEAN, TINYINT, SMALLINT, INT -> BIGINT, BIGINT -> DECIMAL, REAL -> DOUBLE

Example:

```
SUM(X)
```

## SELECTIVITY

```
SELECTIVITY(value)  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Estimates the selectivity (0-100) of a value. The value is defined as  $(100 * \text{distinctCount} / \text{rowCount})$ . The selectivity of 0 rows is 0 (unknown). Up to 10000 values are kept in memory. Aggregates are only allowed in select statements.

Example:

```
SELECT SELECTIVITY(FIRSTNAME), SELECTIVITY(NAME) FROM TEST WHERE  
ROWNUM()<20000
```

## STDDEV\_POP

```
STDDEV_POP( [ DISTINCT|ALL ] numeric )
```

```
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The population standard deviation. This method returns a double. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
STDDEV_POP(X)
```

## STDDEV\_SAMP

```
STDDEV_SAMP( [ DISTINCT|ALL ] numeric )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The sample standard deviation. This method returns a double. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
STDDEV(X)
```

## VAR\_POP

```
VAR_POP( [ DISTINCT|ALL ] numeric )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The population variance (square of the population standard deviation). This method returns a double. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
VAR_POP(X)
```

## VAR\_SAMP

```
VAR_SAMP( [ DISTINCT|ALL ] numeric )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The sample variance (square of the sample standard deviation). This method returns a double. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

VAR\_SAMP(X)

## RANK aggregate

```
RANK(value [,...])  
withinGroupSpecification  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns the rank of the hypothetical row in specified collection of rows. The rank of a row is the number of rows that precede this row plus 1. If two or more rows have the same values in ORDER BY columns, these rows get the same rank from the first row with the same values. It means that gaps in ranks are possible.

Example:

```
SELECT RANK(5) WITHIN GROUP (ORDER BY V) FROM TEST;
```

## DENSE\_RANK aggregate

```
DENSE_RANK(value [,...])  
withinGroupSpecification  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns the dense rank of the hypothetical row in specified collection of rows. The rank of a row is the number of groups of rows with the same values in ORDER BY columns that precede group with this row plus 1. If two or more rows have the same values in ORDER BY columns, these rows get the same rank. Gaps in ranks are not possible.

Example:

```
SELECT DENSE_RANK(5) WITHIN GROUP (ORDER BY V) FROM TEST;
```

## PERCENT\_RANK aggregate

```
PERCENT_RANK(value [,...])  
withinGroupSpecification  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns the relative rank of the hypothetical row in specified collection of rows. The relative rank is calculated as  $(\text{RANK} - 1) / (\text{NR} - 1)$ , where RANK is a rank of the row and NR is a total number of rows in the collection including hypothetical row.

Example:

```
SELECT PERCENT_RANK(5) WITHIN GROUP (ORDER BY V) FROM TEST;
```

## CUME\_DIST aggregate

```
CUME_DIST(value [...])  
withinGroupSpecification  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns the relative rank of the hypothetical row in specified collection of rows. The relative rank is calculated as  $NP / NR$  where NP is a number of rows that precede the current row or have the same values in ORDER BY columns and NR is a total number of rows in the collection including hypothetical row.

Example:

```
SELECT CUME_DIST(5) WITHIN GROUP (ORDER BY V) FROM TEST;
```

## PERCENTILE\_CONT

```
PERCENTILE_CONT(numeric) WITHIN GROUP (ORDER BY value [ASC|DESC])  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Return percentile of values from the group with interpolation. Interpolation is only supported for numeric, date-time, and interval data types. Argument must be between 0 and 1 inclusive. Argument must be the same for all rows in the same group. If argument is NULL, the result is NULL. NULL values are ignored in the calculation. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY V)
```

## PERCENTILE\_DISC

```
PERCENTILE_DISC(numeric) WITHIN GROUP (ORDER BY value [ASC|DESC])  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Return percentile of values from the group. Interpolation is not performed. Argument must be between 0 and 1 inclusive. Argument must be the same for all rows in the same group. If argument is NULL, the result is NULL. NULL values are

ignored in the calculation. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY V)
```

## MEDIAN

```
MEDIAN( [ DISTINCT|ALL ] value )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

The value separating the higher half of a values from the lower half. Returns the middle value or an interpolated value between two middle values if number of values is even. Interpolation is only supported for numeric, date-time, and interval data types. NULL values are ignored in the calculation. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
MEDIAN(X)
```

## MODE

```
{ MODE( value ) [ ORDER BY expression [ ASC | DESC ] ] }  
| { MODE() WITHIN GROUP (ORDER BY expression [ ASC | DESC ]) }  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns the value that occurs with the greatest frequency. If there are multiple values with the same frequency only one value will be returned. In this situation value will be chosen based on optional ORDER BY clause that should specify exactly the same expression as argument of this function. Use ascending order to get smallest value or descending order to get largest value from multiple values with the same frequency. If this clause is not specified the exact chosen value is not determined in this situation. NULL values are ignored in the calculation. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
MODE(X)  
MODE(X ORDER BY X)  
MODE() WITHIN GROUP (ORDER BY X)
```

## ENVELOPE

```
ENVELOPE( value )  
[FILTER (WHERE expression)] [OVER windowNameOrSpecification]
```

Returns the minimum bounding box that encloses all specified GEOMETRY values. Only 2D coordinate plane is supported. NULL values are ignored in the calculation. If no rows are selected, the result is NULL. Aggregates are only allowed in select statements.

Example:

```
ENVELOPE(X)
```

## Numeric Functions

### ABS

```
ABS(numeric|interval)
```

Returns the absolute value of a specified value. The returned value is of the same data type as the parameter.

Note that TINYINT, SMALLINT, INT, and BIGINT data types cannot represent absolute values of their minimum negative values, because they have more negative values than positive. For example, for INT data type allowed values are from -2147483648 to 2147483647. ABS(-2147483648) should be 2147483648, but this value is not allowed for this data type. It leads to an exception. To avoid it cast argument of this function to a higher data type.

Example:

```
ABS(VALUE)  
ABS(CAST(VALUE AS BIGINT))
```

### ACOS

```
ACOS(numeric)
```

Calculate the arc cosine. See also Java Math.acos. This method returns a double.

Example:

```
ACOS(D)
```



## ASIN

```
ASIN(numeric)
```

Calculate the arc sine. See also Java Math.asin. This method returns a double.

Example:

```
ASIN(D)
```

## ATAN

```
ATAN(numeric)
```

Calculate the arc tangent. See also Java Math.atan. This method returns a double.

Example:

```
ATAN(D)
```

## COS

```
COS(numeric)
```

Calculate the trigonometric cosine. See also Java Math.cos. This method returns a double.

Example:

```
COS(ANGLE)
```

## COSH

```
COSH(numeric)
```

Calculate the hyperbolic cosine. See also Java Math.cosh. This method returns a double.

Example:

```
COSH(X)
```

## COT

COT([numeric](#))

Calculate the trigonometric cotangent ( $1/\text{TAN}(\text{ANGLE})$ ). See also Java Math.\* functions. This method returns a double.

Example:

COT(ANGLE)

## SIN

SIN([numeric](#))

Calculate the trigonometric sine. See also Java Math.sin. This method returns a double.

Example:

SIN(ANGLE)

## SINH

SINH([numeric](#))

Calculate the hyperbolic sine. See also Java Math.sinh. This method returns a double.

Example:

SINH(ANGLE)

## TAN

TAN([numeric](#))

Calculate the trigonometric tangent. See also Java Math.tan. This method returns a double.

Example:

TAN(ANGLE)

## TANH

```
TANH(numeric)
```

Calculate the hyperbolic tangent. See also Java Math.tanh. This method returns a double.

Example:

```
TANH(X)
```

## ATAN2

```
ATAN2(numeric, numeric)
```

Calculate the angle when converting the rectangular coordinates to polar coordinates. See also Java Math.atan2. This method returns a double.

Example:

```
ATAN2(X, Y)
```

## BITAND

```
BITAND(long, long)
```

The bitwise AND operation. This method returns a long. See also Java operator &.

Example:

```
BITAND(A, B)
```

## BITGET

```
BITGET(long, int)
```

Returns true if and only if the first parameter has a bit set in the position specified by the second parameter. This method returns a boolean. The second parameter is zero-indexed; the least significant bit has position 0.

Example:

```
BITGET(A, 1)
```

## BITOR

```
BITOR(long, long)
```

The bitwise OR operation. This method returns a long. See also Java operator |.

Example:

```
BITOR(A, B)
```

## BITXOR

```
BITXOR(long, long)
```

The bitwise XOR operation. This method returns a long. See also Java operator ^.

Example:

```
BITXOR(A, B)
```

## MOD

```
MOD(long, long)
```

The modulo operation. This method returns a long. See also Java operator %.

Example:

```
MOD(A, B)
```

## CEILING

```
{ CEILING | CEIL } (numeric)
```

See also Java Math.ceil. This method returns a double.

Example:

```
CEIL(A)
```

## DEGREES

```
DEGREES(numeric)
```

See also Java Math.toDegrees. This method returns a double.

Example:

DEGREES(A)

## EXP

```
EXP(numeric)
```

See also Java Math.exp. This method returns a double.

Example:

EXP(A)

## FLOOR

```
FLOOR(numeric)
```

See also Java Math.floor. This method returns a double.

Example:

FLOOR(A)

## LN

```
{LN|LOG}(numeric)
```

Calculates the natural (base e) logarithm as a double value. In the PostgreSQL mode, LOG(x) is base 10. See also Java Math.log.

Example:

LN(A)

## LOG10

```
LOG10(numeric)
```

Calculates the base 10 logarithm as a double value. See also Java Math.log10.

Example:

LOG10(A)

## ORA\_HASH

```
ORA_HASH(expression [, bucketLong [, seedLong]])
```

Computes a hash value. Optional bucket argument determines the maximum returned value. This argument should be between 0 and 4294967295, default is 4294967295. Optional seed argument is combined with the given expression to return the different values for the same expression. This argument should be between 0 and 4294967295, default is 0. This method returns a long value between 0 and the specified or default bucket value inclusive.

Example:

```
ORA_HASH(A)
```

## RADIANS

```
RADIANS(numeric)
```

See also Java Math.toRadians. This method returns a double.

Example:

```
RADIANS(A)
```

## SQRT

```
SQRT(numeric)
```

See also Java Math.sqrt. This method returns a double.

Example:

```
SQRT(A)
```

## PI

```
PI()
```

See also Java Math.PI. This method returns a double.

Example:

PI()

## POWER

```
POWER(numeric, numeric)
```

See also Java Math.pow. This method returns a double.

Example:

POWER(A, B)

## RAND

```
{ RAND | RANDOM } ( [ int ] )
```

Calling the function without parameter returns the next a pseudo random number. Calling it with an parameter seeds the session's random number generator. This method returns a double between 0 (including) and 1 (excluding).

Example:

RAND()

## RANDOM\_UUID

```
{ RANDOM_UUID | UUID } ()
```

Returns a new UUID with 122 pseudo random bits.

Please note that using an index on randomly generated data will result on poor performance once there are millions of rows in a table. The reason is that the cache behavior is very bad with randomly distributed data. This is a problem for any database system.

Example:

RANDOM\_UUID()

## ROUND

```
ROUND(numeric [, digitsInt])
```

Rounds to a number of digits, or to the nearest long if the number of digits is not set. This method returns a numeric (the same type as the input).

Example:

```
ROUND(VALUE, 2)
```

## ROUNDMAGIC

```
ROUNDMAGIC(numeric)
```

This function rounds numbers in a good way, but it is slow. It has a special handling for numbers around 0. Only numbers smaller or equal +/- 1000000000000 are supported. The value is converted to a String internally, and then the last 4 characters are checked. '000x' becomes '0000' and '999x' becomes '999999', which is rounded automatically. This method returns a double.

Example:

```
ROUNDMAGIC(VALUE/3*3)
```

## SECURE\_RAND

```
SECURE_RAND(int)
```

Generates a number of cryptographically secure random numbers. This method returns bytes.

Example:

```
CALL SECURE_RAND(16)
```

## SIGN

```
SIGN(numeric|interval)
```

Returns -1 if the value is smaller than 0, 0 if zero, and otherwise 1.

Example:



SIGN(VALUE)

## ENCRYPT

```
ENCRYPT(algorithmString, keyBytes, dataBytes)
```

Encrypts data using a key. The supported algorithm is AES. The block size is 16 bytes. This method returns bytes.

Example:

```
CALL ENCRYPT('AES', '00', STRINGTOUTF8('Test'))
```

## DECRYPT

```
DECRYPT(algorithmString, keyBytes, dataBytes)
```

Decrypts data using a key. The supported algorithm is AES. The block size is 16 bytes. This method returns bytes.

Example:

```
CALL TRIM(CHAR(0) FROM UTF8TOSTRING(  
    DECRYPT('AES', '00', '3fabb4de8f1ee2e97d7793bab2db1116')))
```

## HASH

```
HASH(algorithmString, expression [, iterationInt])
```

Calculate the hash value using an algorithm, and repeat this process for a number of iterations. Currently, the only algorithm supported is SHA256. This method returns bytes.

Example:

```
CALL HASH('SHA256', STRINGTOUTF8('Password'), 1000)
```

## TRUNCATE

```
{ TRUNC | TRUNCATE } ( { { numeric, digitsInt }  
| timestamp | timestampWithTimeZone | date | timestampString } )
```

Truncates to a number of digits (to the next value closer to 0). This method returns a double. When used with a timestamp, truncates a timestamp to a date (day) value. When used with a date, truncates a date to a date (day) value less time part. When used with a timestamp as string, truncates a timestamp to a date (day) value.

Example:

```
TRUNCATE(VALUE, 2)
```

## COMPRESS

```
COMPRESS(dataBytes [, algorithmString])
```

Compresses the data using the specified compression algorithm. Supported algorithms are: LZF (faster but lower compression; default), and DEFLATE (higher compression). Compression does not always reduce size. Very small objects and objects with little redundancy may get larger. This method returns bytes.

Example:

```
COMPRESS(STRINGTOUTF8('Test'))
```

## EXPAND

```
EXPAND(bytes)
```

Expands data that was compressed using the COMPRESS function. This method returns bytes.

Example:

```
UTF8TOSTRING(EXPAND(COMPRESS(STRINGTOUTF8('Test'))))
```

## ZERO

```
ZERO()
```

Returns the value 0. This function can be used even if numeric literals are disabled.

Example:

```
ZERO()
```

# String Functions

## ASCII

```
ASCII(string)
```

Returns the ASCII value of the first character in the string. This method returns an int.

Example:

```
ASCII('Hi')
```

## BIT\_LENGTH

```
BIT_LENGTH(string)
```

Returns the number of bits in a string. This method returns a long. For BLOB, CLOB, BYTES and JAVA\_OBJECT, the precision is used. Each character needs 16 bits.

Example:

```
BIT_LENGTH(NAME)
```

## LENGTH

```
{ LENGTH | CHAR_LENGTH | CHARACTER_LENGTH } ( string )
```

Returns the number of characters in a string. This method returns a long. For BLOB, CLOB, BYTES and JAVA\_OBJECT, the precision is used.

Example:

```
LENGTH(NAME)
```

## OCTET\_LENGTH

```
OCTET_LENGTH(string)
```

Returns the number of bytes in a string. This method returns a long. For BLOB, CLOB, BYTES and JAVA\_OBJECT, the precision is used. Each character needs 2 bytes.

Example:

OCTET\_LENGTH(NAME)

## CHAR

```
{ CHAR | CHR } ( int )
```

Returns the character that represents the ASCII value. This method returns a string.

Example:

CHAR(65)

## CONCAT

```
CONCAT(string, string [...])
```

Combines strings. Unlike with the operator ||, NULL parameters are ignored, and do not cause the result to become NULL. This method returns a string.

Example:

CONCAT(NAME, '!')

## CONCAT\_WS

```
CONCAT_WS(separatorString, string, string [...])
```

Combines strings with separator. Unlike with the operator ||, NULL parameters are ignored, and do not cause the result to become NULL. This method returns a string.

Example:

CONCAT\_WS(',', NAME, '!')

## DIFFERENCE

```
DIFFERENCE(string, string)
```

Returns the difference between the sounds of two strings. The difference is calculated as a number of matched characters in the same positions in SOUNDEX representations of arguments. This method returns an int between 0 and 4 inclusive, or null if any of its parameters is null. Note that value of 0 means that strings are not similar to each other. Value of 4 means that strings are fully similar to each other (have the same SOUNDEX representation).

Example:

```
DIFFERENCE(T1.NAME, T2.NAME)
```

## HEXTORAW

```
HEXTORAW(string)
```

Converts a hex representation of a string to a string. 4 hex characters per string character are used.

Example:

```
HEXTORAW(DATA)
```

## RAWTOHEX

```
RAWTOHEX(string)
```

Converts a string to the hex representation. 4 hex characters per string character are used. This method returns a string.

Example:

```
RAWTOHEX(DATA)
```

## INSTR

```
INSTR(string, searchString, [, startInt])
```

Returns the location of a search string in a string. If a start position is used, the characters before it are ignored. If position is negative, the rightmost location is returned. 0 is returned if the search string is not found. Please note this function is case sensitive, even if the parameters are not.

Example:

INSTR(EMAIL,'@')

## INSERT Function

```
INSERT(originalString, startInt, lengthInt, addString)
```

Inserts a additional string into the original string at a specified start position. The length specifies the number of characters that are removed at the start position in the original string. This method returns a string.

Example:

```
INSERT(NAME, 1, 1, ' ')
```

## LOWER

```
{ LOWER | LCASE } ( string )
```

Converts a string to lowercase.

Example:

```
LOWER(NAME)
```

## UPPER

```
{ UPPER | UCASE } ( string )
```

Converts a string to uppercase.

Example:

```
UPPER(NAME)
```

## LEFT

```
LEFT(string, int)
```

Returns the leftmost number of characters.

Example:

```
LEFT(NAME, 3)
```

## RIGHT

```
RIGHT(string, int)
```

Returns the rightmost number of characters.

Example:

```
RIGHT(NAME, 3)
```

## LOCATE

```
LOCATE(searchString, string [, startInt])
```

Returns the location of a search string in a string. If a start position is used, the characters before it are ignored. If position is negative, the rightmost location is returned. 0 is returned if the search string is not found.

Example:

```
LOCATE('.', NAME)
```

## POSITION

```
POSITION(searchString, string)
```

Returns the location of a search string in a string. See also LOCATE.

Example:

```
POSITION('.', NAME)
```

## LPAD

```
LPAD(string, int[, paddingString])
```

Left pad the string to the specified length. If the length is shorter than the string, it will be truncated at the end. If the padding string is not set, spaces will be used.

Example:

```
LPAD(AMOUNT, 10, '*')
```

## RPAD

```
RPAD(string, int[, paddingString])
```

Right pad the string to the specified length. If the length is shorter than the string, it will be truncated. If the padding string is not set, spaces will be used.

Example:

```
RPAD(TEXT, 10, '-')
```

## LTRIM

```
LTRIM(string)
```

Removes all leading spaces from a string.

Example:

```
LTRIM(NAME)
```

## RTRIM

```
RTRIM(string)
```

Removes all trailing spaces from a string.

Example:

```
RTRIM(NAME)
```

## TRIM

```
TRIM ( [ { LEADING | TRAILING | BOTH } [ string ] FROM ] string )
```

Removes all leading spaces, trailing spaces, or spaces at both ends, from a string. Other characters can be removed as well.

Example:

```
TRIM(BOTH '_' FROM NAME)
```



## REGEXP\_REPLACE

```
REGEXP_REPLACE(inputString, regexString, replacementString [, flagsString])
```

Replaces each substring that matches a regular expression. For details, see the Java String.replaceAll() method. If any parameter is null (except optional flagsString parameter), the result is null.

Flags values limited to 'i', 'c', 'n', 'm'. Other symbols causes exception. Multiple symbols could be uses in one flagsString parameter (like 'im'). Later flags overrides first ones, for example 'ic' equivalent to case sensitive matching 'c'.

'i' enables case insensitive matching (Pattern.CASE\_INSENSITIVE)

'c' disables case insensitive matching (Pattern.CASE\_INSENSITIVE)

'n' allows the period to match the newline character (Pattern.DOTALL)

'm' enables multiline mode (Pattern.MULTILINE)

Example:

```
REGEXP_REPLACE('Hello  World', ' +', ' ')  
REGEXP_REPLACE('Hello WWWWorld', 'w+', 'W', 'i')
```

## REGEXP\_LIKE

```
REGEXP_LIKE(inputString, regexString [, flagsString])
```

Matches string to a regular expression. For details, see the Java Matcher.find() method. If any parameter is null (except optional flagsString parameter), the result is null.

Flags values limited to 'i', 'c', 'n', 'm'. Other symbols causes exception. Multiple symbols could be uses in one flagsString parameter (like 'im'). Later flags overrides first ones, for example 'ic' equivalent to case sensitive matching 'c'.

'i' enables case insensitive matching (Pattern.CASE\_INSENSITIVE)

'c' disables case insensitive matching (Pattern.CASE\_INSENSITIVE)

'n' allows the period to match the newline character (Pattern.DOTALL)

'm' enables multiline mode (Pattern.MULTILINE)

Example:

REGEXP\_LIKE('Hello World', '[A-Z ]\*', 'i')

## REPEAT

REPEAT([string](#), [int](#))

Returns a string repeated some number of times.

Example:

REPEAT(NAME || ' ', 10)

## REPLACE

REPLACE([string](#), [searchString](#) [, [replacementString](#)])

Replaces all occurrences of a search string in a text with another string. If no replacement is specified, the search string is removed from the original string. If any parameter is null, the result is null.

Example:

REPLACE(NAME, ' ')

## SOUNDEX

SOUNDEX([string](#))

Returns a four character code representing the sound of a string. This method returns a string, or null if parameter is null. See <https://en.wikipedia.org/wiki/Soundex> for more information.

Example:

SOUNDEX(NAME)

## SPACE

SPACE([int](#))

Returns a string consisting of a number of spaces.

Example:

SPACE(80)

## STRINGDECODE

```
STRINGDECODE(string)
```

Converts a encoded string using the Java string literal encoding format. Special characters are \b, \t, \n, \f, \r, \", \\, \<octal>, \u<unicode>. This method returns a string.

Example:

```
CALL STRINGENCODE(STRINGDECODE('Lines 1\nLine 2'))
```

## STRINGENCODE

```
STRINGENCODE(string)
```

Encodes special characters in a string using the Java string literal encoding format. Special characters are \b, \t, \n, \f, \r, \", \\, \<octal>, \u<unicode>. This method returns a string.

Example:

```
CALL STRINGENCODE(STRINGDECODE('Lines 1\nLine 2'))
```

## STRINGTOUTF8

```
STRINGTOUTF8(string)
```

Encodes a string to a byte array using the UTF8 encoding format. This method returns bytes.

Example:

```
CALL UTF8TOSTRING(STRINGTOUTF8('This is a test'))
```

## SUBSTRING

```
{ SUBSTRING | SUBSTR } ( string, startInt [, lengthInt ] )
```

Returns a substring of a string starting at a position. If the start index is negative, then the start index is relative to the end of the string. The length is optional. Also supported is: SUBSTRING(string [FROM start] [FOR length]).

Example:

```
CALL SUBSTR('[Hello]', 2, 5);  
CALL SUBSTR('Hello World', -5);
```

## UTF8TOSTRING

```
UTF8TOSTRING(bytes)
```

Decodes a byte array in the UTF8 format to a string.

Example:

```
CALL UTF8TOSTRING(STRINGTOUTF8('This is a test'))
```

## XMLATTR

```
XMLATTR(nameString, valueString)
```

Creates an XML attribute element of the form name=value. The value is encoded as XML text. This method returns a string.

Example:

```
CALL XMLNODE('a', XMLATTR('href', 'http://h2database.com'))
```

## XMLNODE

```
XMLNODE(elementString [, attributesString [, contentString [, indentBoolean]]])
```

Create an XML node element. An empty or null attribute string means no attributes are set. An empty or null content string means the node is empty. The content is indented by default if it contains a newline. This method returns a string.

Example:

```
CALL XMLNODE('a', XMLATTR('href', 'http://h2database.com'), 'H2')
```

## XMLCOMMENT

```
XMLCOMMENT(commentString)
```

Creates an XML comment. Two dashes (--) are converted to - -. This method returns a string.

Example:

```
CALL XMLCOMMENT('Test')
```

## XMLCDATA

```
XMLCDATA(valueString)
```

Creates an XML CDATA element. If the value contains ]]>, an XML text element is created instead. This method returns a string.

Example:

```
CALL XMLCDATA('data')
```

## XMLSTARTDOC

```
XMLSTARTDOC()
```

Returns the XML declaration. The result is always <?xml version=1.0?>.

Example:

```
CALL XMLSTARTDOC()
```

## XMLTEXT

```
XMLTEXT(valueString [, escapeNewlineBoolean])
```

Creates an XML text element. If enabled, newline and linefeed is converted to an XML entity (&#). This method returns a string.

Example:

```
CALL XMLTEXT('test')
```

## TO\_CHAR

```
TO_CHAR(value [, formatString[, nlsParamString]])
```

Oracle-compatible TO\_CHAR function that can format a timestamp, a number, or text.

Example:

```
CALL TO_CHAR(TIMESTAMP '2010-01-01 00:00:00', 'DD MON, YYYY')
```

## TRANSLATE

```
TRANSLATE(value, searchString, replacementString)
```

Oracle-compatible TRANSLATE function that replaces a sequence of characters in a string with another set of characters.

Example:

```
CALL TRANSLATE('Hello world', 'eo', 'EO')
```

## Time and Date Functions

### CURRENT\_DATE

```
{ CURRENT_DATE | CURDATE() | SYSDATE | TODAY }
```

Returns the current date. These methods always return the same value within a transaction (default) or within a command depending on database mode.

Example:

```
CURRENT_DATE
```

### CURRENT\_TIME

```
CURRENT_TIME [ (int) ]
```

Returns the current time. The returned value does not have time zone information, because TIME WITH TIME ZONE data type is not supported in H2. If fractional seconds precision is specified it should be from 0 to 9, 0 is default. The specified value can be used only to limit precision of a result. The actual

maximum available precision depends on operating system and JVM and can be 3 (milliseconds) or higher. Higher precision is not available before Java 9. These methods always return the same value within a transaction (default) or within a command depending on database mode.

Example:

```
CURRENT_TIME  
CURRENT_TIME(9)
```

## CURRENT\_TIMESTAMP

```
CURRENT_TIMESTAMP [ (int) ]
```

Returns the current timestamp with time zone. Time zone offset is set to a current time zone offset. If fractional seconds precision is specified it should be from 0 to 9, 6 is default. The specified value can be used only to limit precision of a result. The actual maximum available precision depends on operating system and JVM and can be 3 (milliseconds) or higher. Higher precision is not available before Java 9. This method always returns the same value within a transaction (default) or within a command depending on database mode.

Example:

```
CURRENT_TIMESTAMP  
CURRENT_TIMESTAMP(9)
```

## LOCALTIME

```
{ LOCALTIME [ (int) ] | CURTIME([ int ]) }
```

Returns the current time. If fractional seconds precision is specified it should be from 0 to 9, 0 is default. The specified value can be used only to limit precision of a result. The actual maximum available precision depends on operating system and JVM and can be 3 (milliseconds) or higher. Higher precision is not available before Java 9. These methods always return the same value within a transaction (default) or within a command depending on database mode.

Example:

```
LOCALTIME  
LOCALTIME(9)
```

## LOCALTIMESTAMP

```
{ LOCALTIMESTAMP [ (int) ] | NOW( [ int ] ) }
```

Returns the current timestamp without time zone. If fractional seconds precision is specified it should be from 0 to 9, 6 is default. The specified value can be used only to limit precision of a result. The actual maximum available precision depends on operating system and JVM and can be 3 (milliseconds) or higher. Higher precision is not available before Java 9. These methods always return the same value within a transaction (default) or within a command depending on database mode.

Example:

```
LOCALTIMESTAMP
```

```
LOCALTIMESTAMP(9)
```

## DATEADD

```
{ DATEADD| TIMESTAMPADD } (datetimeField, addIntLong, dateAndTime)
```

Adds units to a date-time value. The datetimeField indicates the unit. Use negative values to subtract units. addIntLong may be a long value when manipulating milliseconds, microseconds, or nanoseconds otherwise its range is restricted to int. This method returns a value with the same type as specified value if unit is compatible with this value. If specified field is a HOUR, MINUTE, SECOND, MILLISECOND, etc and value is a DATE value DATEADD returns combined TIMESTAMP. Fields DAY, MONTH, YEAR, WEEK, etc are not allowed for TIME values. Fields TIMEZONE\_HOUR and TIMEZONE\_MINUTE are only allowed for TIMESTAMP WITH TIME ZONE values.

Example:

```
DATEADD('MONTH', 1, DATE '2001-01-31')
```

## DATEDIFF

```
{ DATEDIFF | TIMESTAMPDIFF } (datetimeField, aDateAndTime, bDateAndTime)
```

Returns the number of crossed unit boundaries between two date/time values. This method returns a long. The datetimeField indicates the unit. Only TIMEZONE\_HOUR and TIMEZONE\_MINUTE fields use the time zone offset component. With all other fields if date/time values have time zone offset component it is ignored.



Example:

DATEDIFF(YEAR, T1.CREATED, T2.CREATED)

## DAYNAME

```
DAYNAME(dateAndTime)
```

Returns the name of the day (in English).

Example:

DAYNAME(CREATED)

## DAY\_OF\_MONTH

```
DAY_OF_MONTH(dateAndTime|interval)
```

Returns the day of the month (1-31).

Example:

DAY\_OF\_MONTH(CREATED)

## DAY\_OF\_WEEK

```
DAY_OF_WEEK(dateAndTime)
```

Returns the day of the week (1 means Sunday).

Example:

DAY\_OF\_WEEK(CREATED)

## ISO\_DAY\_OF\_WEEK

```
ISO_DAY_OF_WEEK(dateAndTime)
```

Returns the ISO day of the week (1 means Monday).

Example:

ISO\_DAY\_OF\_WEEK(CREATED)

## DAY\_OF\_YEAR

```
DAY_OF_YEAR(dateAndTime|interval)
```

Returns the day of the year (1-366).

Example:

```
DAY_OF_YEAR(CREATED)
```

## EXTRACT

```
EXTRACT ( datetimeField FROM { dateAndTime | interval } )
```

Returns a value of the specific time unit from a date/time value. This method returns a numeric value with EPOCH field and an int for all other fields.

Example:

```
EXTRACT(SECOND FROM CURRENT_TIMESTAMP)
```

## FORMATDATETIME

```
FORMATDATETIME ( dateAndTime, formatString  
[ , localeString [ , timeZoneString ] ] )
```

Formats a date, time or timestamp as a string. The most important format characters are: y year, M month, d day, H hour, m minute, s second. For details of the format, see `java.text.SimpleDateFormat`. `timeZoneString` may be specified if `dateAndTime` is a DATE, TIME or TIMESTAMP. `timeZoneString` is ignored if `dateAndTime` is TIMESTAMP WITH TIME ZONE. This method returns a string.

Example:

```
CALL FORMATDATETIME(TIMESTAMP '2001-02-03 04:05:06',  
    'EEE, d MMM yyyy HH:mm:ss z', 'en', 'GMT')
```

## HOUR

```
HOUR(dateAndTime|interval)
```

Returns the hour (0-23) from a date/time value.

Example:

hour(created)

## MINUTE

minute([dateAndTime](#)|[interval](#))

Returns the minute (0-59) from a date/time value.

Example:

minute(created)

## MONTH

month([dateAndTime](#)|[interval](#))

Returns the month (1-12) from a date/time value.

Example:

month(created)

## MONTHNAME

monthname([dateAndTime](#))

Returns the name of the month (in English).

Example:

monthname(created)

## PARSEDATETIME

parsedatetime([string](#), [formatString](#)  
[, [localeString](#) [, [timeZoneString](#)]])

Parses a string and returns a timestamp. The most important format characters are: y year, M month, d day, H hour, m minute, s second. For details of the format, see `java.text.SimpleDateFormat`.

Example:

```
CALL PARSEDATETIME('Sat, 3 Feb 2001 03:05:06 GMT',  
    'EEE, d MMM yyyy HH:mm:ss z', 'en', 'GMT')
```

## QUARTER

```
QUARTER(dateAndTime)
```

Returns the quarter (1-4) from a date/time value.

Example:

```
QUARTER(CREATED)
```

## SECOND

```
SECOND(dateAndTime)
```

Returns the second (0-59) from a date/time value.

Example:

```
SECOND(CREATED|interval)
```

## WEEK

```
WEEK(dateAndTime)
```

Returns the week (1-53) from a date/time value. This method uses the current system locale.

Example:

```
WEEK(CREATED)
```

## ISO\_WEEK

```
ISO_WEEK(dateAndTime)
```

Returns the ISO week (1-53) from a date/time value. This function uses the ISO definition when first week of year should have at least four days and week is started with Monday.

Example:

ISO\_WEEK(CREATED)

## YEAR

```
YEAR(dateAndTime|interval)
```

Returns the year from a date/time value.

Example:

```
YEAR(CREATED)
```

## ISO\_YEAR

```
ISO_YEAR(dateAndTime)
```

Returns the ISO week year from a date/time value.

Example:

```
ISO_YEAR(CREATED)
```

## System Functions

### ARRAY\_GET

```
ARRAY_GET(arrayExpression, indexExpression)
```

Returns element at the specified 1-based index from an array. Returns NULL if there is no such element or array is NULL.

Example:

```
CALL ARRAY_GET(ARRAY['Hello', 'World'], 2)
```

### ARRAY\_LENGTH

```
ARRAY_LENGTH(arrayExpression)
```

Returns the length of an array. Returns NULL if the specified array is NULL.

Example:

```
CALL ARRAY_LENGTH(ARRAY['Hello', 'World'])
```

## ARRAY\_CONTAINS

```
ARRAY_CONTAINS(arrayExpression, value)
```

Returns a boolean TRUE if the array contains the value or FALSE if it does not contain it. Returns NULL if the specified array is NULL.

Example:

```
CALL ARRAY_CONTAINS(ARRAY['Hello', 'World'], 'Hello')
```

## ARRAY\_CAT

```
ARRAY_CAT(arrayExpression, arrayExpression)
```

Returns the concatenation of two arrays. Returns NULL if any parameter is NULL.

Example:

```
CALL ARRAY_CAT(ARRAY[1, 2], ARRAY[3, 4])
```

## ARRAY\_APPEND

```
ARRAY_APPEND(arrayExpression, value)
```

Append an element to the end of an array. Returns NULL if any parameter is NULL.

Example:

```
CALL ARRAY_APPEND(ARRAY[1, 2], 3)
```

## ARRAY\_SLICE

```
ARRAY_SLICE(arrayExpression, lowerBoundInt, upperBoundInt)
```

Returns elements from the array as specified by the lower and upper bound parameters. Both parameters are inclusive and the first element has index 1, i.e. ARRAY\_SLICE(a, 2, 2) has only the second element. Returns NULL if any parameter is NULL or if an index is out of bounds.

Example:

```
CALL ARRAY_SLICE(ARRAY[1, 2, 3, 4], 1, 3)
```

## AUTOCOMMIT

```
AUTOCOMMIT()
```

Returns true if auto commit is switched on for this session.

Example:

```
AUTOCOMMIT()
```

## CANCEL\_SESSION

```
CANCEL_SESSION(sessionInt)
```

Cancels the currently executing statement of another session. The method only works if the multithreaded kernel is enabled (see SET MULTI\_THREADED). Returns true if the statement was canceled, false if the session is closed or no statement is currently executing.

Admin rights are required to execute this command.

Example:

```
CANCEL_SESSION(3)
```

## CASEWHEN Function

```
CASEWHEN(boolean, aValue, bValue)
```

Returns 'a' if the boolean expression is true, otherwise 'b'. Returns the same data type as the parameter.

Example:

```
CASEWHEN(ID=1, 'A', 'B')
```

## CAST

```
CAST(value AS dataType)
```

Converts a value to another data type. The following conversion rules are used: When converting a number to a boolean, 0 is false and every other value is true. When converting a boolean to a number, false is 0 and true is 1. When converting a number to a number of another type, the value is checked for overflow. When converting a number to binary, the number of bytes matches the precision. When converting a string to binary, it is hex encoded (every byte two characters); a hex string can be converted to a number by first converting it to binary. If a direct conversion is not possible, the value is first converted to a string. Note that some data types may need explicitly specified precision to avoid overflow or rounding.

Example:

```
CAST(NAME AS INT);  
CAST(65535 AS BINARY);  
CAST(CAST('FFFF' AS BINARY) AS INT);  
CAST(TIMESTAMP '2010-01-01 10:40:00.123456' AS TIME(6))
```

## COALESCE

```
{ COALESCE | NVL } (aValue, bValue [,...])
```

Returns the first value that is not null.

Example:

```
COALESCE(A, B, C)
```

## CONVERT

```
CONVERT(value, dataType)
```

Converts a value to another data type.

Example:

```
CONVERT(NAME, INT)
```

## CURRVAL

```
CURRVAL( [ schemaName, ] sequenceString )
```

Returns the current (last) value of the sequence, independent of the session. If the sequence was just created, the method returns (start - interval). If the schema name is not set, the current schema is used. If the schema name is not



set, the sequence name is converted to uppercase (for compatibility). This method returns a long.

Example:

```
CURRVAL('TEST_SEQ')
```

## CSVREAD

```
CSVREAD(fileNameString [, columnsString [, csvOptions ] ] )
```

Returns the result set of reading the CSV (comma separated values) file. For each parameter, NULL means the default value should be used.

If the column names are specified (a list of column names separated with the fieldSeparator), those are used, otherwise (or if they are set to NULL) the first line of the file is interpreted as the column names. In that case, column names that contain no special characters (only letters, '\_', and digits; similar to the rule for Java identifiers) are considered case insensitive. Other column names are case sensitive, that means you need to use quoted identifiers (see below).

The default charset is the default value for this system, and the default field separator is a comma. Missing unquoted values as well as data that matches nullString is parsed as NULL. All columns of type VARCHAR.

The BOM (the byte-order-mark) character 0xfeff at the beginning of the file is ignored.

This function can be used like a table: SELECT \* FROM CSVREAD(...).

Instead of a file, an URL may be used, for example jar:file:///c:/temp/example.zip!/org/example/nested.csv. To read a stream from the classpath, use the prefix classpath:. To read from HTTP, use the prefix http: (as in a browser).

For performance reason, CSVREAD should not be used inside a join. Instead, import the data first (possibly into a temporary table) and then use the table.

Admin rights are required to execute this command.

Example:

```
CALL CSVREAD('test.csv');  
-- Read a file containing the columns ID, NAME with  
CALL CSVREAD('test2.csv', 'ID|NAME', 'charset=UTF-8 fieldSeparator=|');  
SELECT * FROM CSVREAD('data/test.csv', null, 'rowSeparator=;');
```

```
-- Read a tab-separated file
SELECT * FROM CSVREAD('data/test.tsv', null, 'rowSeparator=' || CHAR(9));
SELECT "Last Name" FROM CSVREAD('address.csv');
SELECT "Last Name" FROM CSVREAD('classpath:/org/acme/data/address.csv');
```

## CSVWRITE

```
CSVWRITE ( fileNameString, queryString [, csvOptions [, lineSepString] ] )
```

Writes a CSV (comma separated values). The file is overwritten if it exists. If only a file name is specified, it will be written to the current working directory. For each parameter, NULL means the default value should be used. The default charset is the default value for this system, and the default field separator is a comma.

The values are converted to text using the default string representation; if another conversion is required you need to change the select statement accordingly. The parameter nullString is used when writing NULL (by default nothing is written when NULL appears). The default line separator is the default value for this system (system property line.separator).

The returned value is the number of rows written. Admin rights are required to execute this command.

Example:

```
CALL CSVWRITE('data/test.csv', 'SELECT * FROM TEST');
CALL CSVWRITE('data/test2.csv', 'SELECT * FROM TEST', 'charset=UTF-8
fieldSeparator=|');
-- Write a tab-separated file
CALL CSVWRITE('data/test.tsv', 'SELECT * FROM TEST', 'charset=UTF-8
fieldSeparator=' || CHAR(9));
```

## DATABASE

```
DATABASE()
```

Returns the name of the database.

Example:

```
CALL DATABASE();
```

## DATABASE\_PATH

```
DATABASE_PATH()
```

Returns the directory of the database files and the database name, if it is file based. Returns NULL otherwise.

Example:

```
CALL DATABASE_PATH();
```

## DECODE

```
DECODE(value, whenValue, thenValue [,...])
```

Returns the first matching value. NULL is considered to match NULL. If no match was found, then NULL or the last parameter (if the parameter count is even) is returned. This function is provided for Oracle compatibility (see there for details).

Example:

```
CALL DECODE(RAND()>0.5, 0, 'Red', 1, 'Black');
```

## DISK\_SPACE\_USED

```
DISK_SPACE_USED(tableNameString)
```

Returns the approximate amount of space used by the table specified. Does not currently take into account indexes or LOB's. This function may be expensive since it has to load every page in the table.

Example:

```
CALL DISK_SPACE_USED('my_table');
```

## SIGNAL

```
SIGNAL(sqlStateString, messageString)
```

Throw an SQLException with the passed SQLState and reason.

Example:

```
CALL SIGNAL('23505', 'Duplicate user ID: ' || user_id);
```

## ESTIMATED\_ENVELOPE

```
ESTIMATED_ENVELOPE(tableNameString, columnNameString)
```

Returns the estimated minimum bounding box that encloses all specified GEOMETRY values. Only 2D coordinate plane is supported. NULL values are ignored. This function is only supported by MVStore engine. Column must have a spatial index. This function is fast, but estimation may include uncommitted data (including data from other transactions), may return approximate bounds, or be different with actual value due to other reasons. Use with caution. If estimation is not available this function returns NULL. For accurate and reliable result use ESTIMATE aggregate function instead.

Example:

```
CALL ESTIMATED_ENVELOPE('MY_TABLE', 'GEOMETRY_COLUMN');
```

## FILE\_READ

```
FILE_READ(fileNameString [,encodingString])
```

Returns the contents of a file. If only one parameter is supplied, the data are returned as a BLOB. If two parameters are used, the data is returned as a CLOB (text). The second parameter is the character set to use, NULL meaning the default character set for this system.

File names and URLs are supported. To read a stream from the classpath, use the prefix classpath:..

Admin rights are required to execute this command.

Example:

```
SELECT LENGTH(FILE_READ('~/.h2.server.properties')) LEN;  
SELECT FILE_READ('http://localhost:8182/styleSheet.css', NULL) CSS;
```

## FILE\_WRITE

```
FILE_WRITE(blobValue, fileNameString)
```

Write the supplied parameter into a file. Return the number of bytes written.

Write access to folder, and admin rights are required to execute this command.

Example:

```
SELECT FILE_WRITE('Hello world', '/tmp/hello.txt')) LEN;
```

## GREATEST

```
GREATEST(aValue, bValue [,...])
```

Returns the largest value that is not NULL, or NULL if all values are NULL.

Example:

```
CALL GREATEST(1, 2, 3);
```

## IDENTITY

```
IDENTITY()
```

Returns the last inserted identity value for this session. This value changes whenever a new sequence number was generated, even within a trigger or Java function. See also SCOPE\_IDENTITY. This method returns a long.

Example:

```
CALL IDENTITY();
```

## IFNULL

```
IFNULL(aValue, bValue)
```

Returns the value of 'a' if it is not null, otherwise 'b'.

Example:

```
CALL IFNULL(NULL, "");
```

## LEAST

```
LEAST(aValue, bValue [,...])
```

Returns the smallest value that is not NULL, or NULL if all values are NULL.

Example:

```
CALL LEAST(1, 2, 3);
```

## LOCK\_MODE

```
LOCK_MODE()
```

Returns the current lock mode. See SET LOCK\_MODE. This method returns an int.

Example:

```
CALL LOCK_MODE();
```

## LOCK\_TIMEOUT

```
LOCK_TIMEOUT()
```

Returns the lock timeout of the current session (in milliseconds).

Example:

```
LOCK_TIMEOUT()
```

## LINK\_SCHEMA

```
LINK_SCHEMA(targetSchemaString, driverString, urlString,  
userString, passwordString, sourceSchemaString)
```

Creates table links for all tables in a schema. If tables with the same name already exist, they are dropped first. The target schema is created automatically if it does not yet exist. The driver name may be empty if the driver is already loaded. The list of tables linked is returned in the form of a result set. Admin rights are required to execute this command.

Example:

```
CALL LINK_SCHEMA('TEST2', '', 'jdbc:h2:test2', 'sa', 'sa', 'PUBLIC');
```

## MEMORY\_FREE

```
MEMORY_FREE()
```

Returns the free memory in KB (where 1024 bytes is a KB). This method returns an int. The garbage is run before returning the value. Admin rights are required to execute this command.

Example:

MEMORY\_FREE()

## MEMORY\_USED

MEMORY\_USED()

Returns the used memory in KB (where 1024 bytes is a KB). This method returns an int. The garbage is run before returning the value. Admin rights are required to execute this command.

Example:

MEMORY\_USED()

## NEXTVAL

NEXTVAL ( [ [schemaName](#), ] [sequenceString](#) )

Returns the next value of the sequence. Used values are never re-used, even when the transaction is rolled back. If the schema name is not set, the current schema is used, and the sequence name is converted to uppercase (for compatibility). This method returns a long.

Example:

NEXTVAL('TEST\_SEQ')

## NULLIF

NULLIF([aValue](#), [bValue](#))

Returns NULL if 'a' is equals to 'b', otherwise 'a'.

Example:

NULLIF(A, B)

## NVL2

NVL2([testValue](#), [aValue](#), [bValue](#))

If the test value is null, then 'b' is returned. Otherwise, 'a' is returned. The data type of the returned value is the data type of 'a' if this is a text type.

Example:

```
NVL2(X, 'not null', 'null')
```

## READONLY

READONLY()

Returns true if the database is read-only.

Example:

```
READONLY()
```

## ROWNUM

ROWNUM()

Returns the number of the current row. This method returns a long value. It is supported for SELECT statements, as well as for DELETE and UPDATE. The first row has the row number 1, and is calculated before ordering and grouping the result set, but after evaluating index conditions (even when the index conditions are specified in an outer query). Use the ROW\_NUMBER() OVER () function to get row numbers after grouping or in specified order.

Example:

```
SELECT ROWNUM(), * FROM TEST;  
SELECT ROWNUM(), * FROM (SELECT * FROM TEST ORDER BY NAME);  
SELECT ID FROM (SELECT T.*, ROWNUM AS R FROM TEST T) WHERE R  
BETWEEN 2 AND 3;
```

## SCHEMA

SCHEMA()

Returns the name of the default schema for this session.

Example:

```
CALL SCHEMA()
```



## SCOPE\_IDENTITY

```
SCOPE_IDENTITY()
```

Returns the last inserted identity value for this session for the current scope (the current statement). Changes within triggers and Java functions are ignored. See also `IDENTITY()`. This method returns a long.

Example:

```
CALL SCOPE_IDENTITY();
```

## SESSION\_ID

```
SESSION_ID()
```

Returns the unique session id number for the current database connection. This id stays the same while the connection is open. This method returns an int. The database engine may re-use a session id after the connection is closed.

Example:

```
CALL SESSION_ID()
```

## SET

```
SET(@variableName, value)
```

Updates a variable with the given value. The new value is returned. When used in a query, the value is updated in the order the rows are read. When used in a subquery, not all rows might be read depending on the query plan. This can be used to implement running totals / cumulative sums.

Example:

```
SELECT X, SET(@I, IFNULL(@I, 0)+X) RUNNING_TOTAL FROM  
SYSTEM_RANGE(1, 10)
```

## TABLE

```
{ TABLE | TABLE_DISTINCT }  
( { name dataType = array|rowValueExpression } [,...] )
```

Returns the result set. `TABLE_DISTINCT` removes duplicate rows.

Example:

```
SELECT * FROM TABLE(VALUE INT = ARRAY[1, 2]);  
SELECT * FROM TABLE(ID INT=(1, 2), NAME VARCHAR=('Hello', 'World'));
```

## TRANSACTION\_ID

TRANSACTION\_ID()

Returns the current transaction id for this session. This method returns NULL if there is no uncommitted change, or if the database is not persisted. Otherwise a value of the following form is returned: logFileId-position-sessionId. This method returns a string. The value is unique across database restarts (values are not re-used).

Example:

```
CALL TRANSACTION_ID()
```

## TRUNCATE\_VALUE

TRUNCATE\_VALUE([value](#), [precisionInt](#), [forceBoolean](#))

Truncate a value to the required precision. The precision of the returned value may be a bit larger than requested, because fixed precision values are not truncated (unlike the numeric TRUNCATE method). Unlike CAST, the truncating a decimal value may lose precision if the force flag is set to true. The method returns a value with the same data type as the first parameter.

Example:

```
CALL TRUNCATE_VALUE(X, 10, TRUE);
```

## UNNEST

UNNEST([array](#), [...]) [WITH ORDINALITY]

Returns the result set. Number of columns is equal to number of arguments, plus one additional column with row number if WITH ORDINALITY is specified. Number of rows is equal to length of longest specified array. If multiple arguments are specified and they have different length, cells with missing values will contain null values.

Example:

```
SELECT * FROM UNNEST(ARRAY['a', 'b', 'c']);
```

## USER

```
{ USER | CURRENT_USER } ()
```

Returns the name of the current user of this session.

Example:

```
CURRENT_USER()
```

## H2VERSION

```
H2VERSION()
```

Returns the H2 version as a String.

Example:

```
H2VERSION()
```

## Window Functions

### ROW\_NUMBER

```
ROW_NUMBER() OVER windowNameOrSpecification
```

Returns the number of the current row starting with 1. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT ROW_NUMBER() OVER (), * FROM TEST;  
SELECT ROW_NUMBER() OVER (ORDER BY ID), * FROM TEST;  
SELECT ROW_NUMBER() OVER (PARTITION BY CATEGORY ORDER BY ID), *  
FROM TEST;
```

## RANK

**RANK()** OVER [windowNameOrSpecification](#)

Returns the rank of the current row. The rank of a row is the number of rows that precede this row plus 1. If two or more rows have the same values in ORDER BY columns, these rows get the same rank from the first row with the same values. It means that gaps in ranks are possible. This function requires window order clause. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT RANK() OVER (ORDER BY ID), * FROM TEST;  
SELECT RANK() OVER (PARTITION BY CATEGORY ORDER BY ID), * FROM TEST;
```

## DENSE\_RANK

**DENSE\_RANK()** OVER [windowNameOrSpecification](#)

Returns the dense rank of the current row. The rank of a row is the number of groups of rows with the same values in ORDER BY columns that precede group with this row plus 1. If two or more rows have the same values in ORDER BY columns, these rows get the same rank. Gaps in ranks are not possible. This function requires window order clause. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT DENSE_RANK() OVER (ORDER BY ID), * FROM TEST;  
SELECT DENSE_RANK() OVER (PARTITION BY CATEGORY ORDER BY ID), *  
FROM TEST;
```

## PERCENT\_RANK

**PERCENT\_RANK()** OVER [windowNameOrSpecification](#)

Returns the relative rank of the current row. The relative rank is calculated as  $(\text{RANK} - 1) / (\text{NR} - 1)$ , where RANK is a rank of the row and NR is a number of

rows in window partition with this row. Note that result is always 0 if window order clause is not specified. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT PERCENT_RANK() OVER (ORDER BY ID), * FROM TEST;  
SELECT PERCENT_RANK() OVER (PARTITION BY CATEGORY ORDER BY ID), *  
FROM TEST;
```

## CUME\_DIST

**CUME\_DIST()** OVER [windowNameOrSpecification](#)

Returns the relative rank of the current row. The relative rank is calculated as  $NP / NR$  where NP is a number of rows that precede the current row or have the same values in ORDER BY columns and NR is a number of rows in window partition with this row. Note that result is always 1 if window order clause is not specified. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT CUME_DIST() OVER (ORDER BY ID), * FROM TEST;  
SELECT CUME_DIST() OVER (PARTITION BY CATEGORY ORDER BY ID), * FROM  
TEST;
```

## NTILE

**NTILE(long)** OVER [windowNameOrSpecification](#)

Distributes the rows into a specified number of groups. Number of groups should be a positive long value. NTILE returns the 1-based number of the group to which the current row belongs. First groups will have more rows if number of rows is not divisible by number of groups. For example, if 5 rows are distributed into 2 groups this function returns 1 for the first 3 row and 2 for the last 2 rows. This function requires window order clause. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT NTILE(10) OVER (ORDER BY ID), * FROM TEST;  
SELECT NTILE(5) OVER (PARTITION BY CATEGORY ORDER BY ID), * FROM  
TEST;
```

## LEAD

```
LEAD(value [, offsetInt [, defaultValue]]) [{RESPECT|IGNORE} NULLS]  
OVER windowNameOrSpecification
```

Returns the value in a next row with specified offset relative to the current row. Offset must be non-negative. If IGNORE NULLS is specified rows with null values in selected expression are skipped. If number of considered rows is less than specified relative number this function returns NULL or the specified default value, if any. If offset is 0 the value from the current row is returned unconditionally. This function requires window order clause. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT LEAD(X) OVER (ORDER BY ID), * FROM TEST;  
SELECT LEAD(X, 2, 0) IGNORE NULLS OVER (  
    PARTITION BY CATEGORY ORDER BY ID  
) , * FROM TEST;
```

## LAG

```
LAG(value [, offsetInt [, defaultValue]]) [{RESPECT|IGNORE} NULLS]  
OVER windowNameOrSpecification
```

Returns the value in a previous row with specified offset relative to the current row. Offset must be non-negative. If IGNORE NULLS is specified rows with null values in selected expression are skipped. If number of considered rows is less than specified relative number this function returns NULL or the specified default value, if any. If offset is 0 the value from the current row is returned unconditionally. This function requires window order clause. Window frame clause is not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT LAG(X) OVER (ORDER BY ID), * FROM TEST;  
SELECT LAG(X, 2, 0) IGNORE NULLS OVER (  
    PARTITION BY CATEGORY ORDER BY ID  
) , * FROM TEST;
```

## FIRST\_VALUE

**FIRST\_VALUE**([value](#)) [{RESPECT|IGNORE} NULLS]  
OVER [windowNameOrSpecification](#)

Returns the first value in a window. If IGNORE NULLS is specified null values are skipped and the function returns first non-null value, if any.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT FIRST_VALUE(X) OVER (ORDER BY ID), * FROM TEST;  
SELECT FIRST_VALUE(X) IGNORE NULLS OVER (PARTITION BY CATEGORY  
ORDER BY ID), * FROM TEST;
```

## LAST\_VALUE

**LAST\_VALUE**([value](#)) [{RESPECT|IGNORE} NULLS]  
OVER [windowNameOrSpecification](#)

Returns the last value in a window. If IGNORE NULLS is specified null values are skipped and the function returns last non-null value before them, if any; if there is no non-null value it returns NULL. Note that the last value is actually a value in the current group of rows if window order clause is specified and window frame clause is not specified.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT LAST_VALUE(X) OVER (ORDER BY ID), * FROM TEST;  
SELECT LAST_VALUE(X) IGNORE NULLS OVER (  
    PARTITION BY CATEGORY ORDER BY ID  
    RANGE BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING  
) , * FROM TEST;
```

## NTH\_VALUE

```
NTH_VALUE(value, nInt) [FROM {FIRST|LAST}] [{RESPECT|IGNORE} NULLS]  
OVER windowNameOrSpecification
```

Returns the value in a row with a specified relative number in a window. Relative row number must be positive. If FROM LAST is specified rows are counted backwards from the last row. If IGNORE NULLS is specified rows with null values in selected expression are skipped. If number of considered rows is less than specified relative number this function returns NULL. Note that the last row is actually a last row in the current group of rows if window order clause is specified and window frame clause is not specified.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT NTH_VALUE(X) OVER (ORDER BY ID), * FROM TEST;  
SELECT NTH_VALUE(X) IGNORE NULLS OVER (  
    PARTITION BY CATEGORY ORDER BY ID  
    RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING  
) , * FROM TEST;
```

## RATIO\_TO\_REPORT

```
RATIO_TO_REPORT(value)  
OVER windowNameOrSpecification
```

Returns the ratio of a value to the sum of all values. If argument is NULL or sum of all values is 0, then the value of function is NULL. Window ordering and window frame clauses are not allowed for this function.

Window functions are currently experimental in H2 and should be used with caution. They also may require a lot of memory for large queries.

Example:

```
SELECT X, RATIO_TO_REPORT(X) OVER (PARTITION BY CATEGORY), CATEGORY  
FROM TEST;
```



# Data Types

## Index

INT  
BOOLEAN  
TINYINT  
SMALLINT  
BIGINT  
IDENTITY  
DECIMAL  
DOUBLE  
REAL  
TIME  
DATE  
TIMESTAMP  
TIMESTAMP WITH TIME ZONE  
BINARY  
OTHER  
VARCHAR  
VARCHAR\_IGNORECASE  
CHAR  
BLOB  
CLOB  
UUID  
ARRAY  
ENUM  
GEOMETRY  
INTERVAL

## INT

INT | INTEGER | MEDIUMINT | INT4 | SIGNED

Possible values: -2147483648 to 2147483647.

Mapped to java.lang.Integer.

Example:

INT

## BOOLEAN

BOOLEAN | BIT | BOOL

Possible values: TRUE and FALSE.

Mapped to java.lang.Boolean.

Example:

BOOLEAN

## TINYINT

TINYINT

Possible values are: -128 to 127.

Mapped to java.lang.Byte.

Example:

TINYINT

## SMALLINT

SMALLINT | INT2 | YEAR

Possible values: -32768 to 32767.

Mapped to java.lang.Short.

Example:

SMALLINT

## BIGINT

BIGINT | INT8

Possible values: -9223372036854775808 to 9223372036854775807.

Mapped to java.lang.Long.

Example:

BIGINT

## IDENTITY

IDENTITY

Auto-Increment value. Possible values: -9223372036854775808 to 9223372036854775807. Used values are never re-used, even when the transaction is rolled back.

Mapped to java.lang.Long.

Example:

IDENTITY

## DECIMAL

{ DECIMAL | NUMBER | DEC | NUMERIC } ( [precisionInt](#) [ , [scaleInt](#) ] )

Data type with fixed precision and scale. This data type is recommended for storing currency values.

Mapped to java.math.BigDecimal.

Example:

DECIMAL(20, 2)

## DOUBLE

{ DOUBLE [ PRECISION ] | FLOAT [ ( [precisionInt](#) ) ] | FLOAT8 }

A floating point number. Should not be used to represent currency values, because of rounding problems. If precision value is specified for FLOAT type name, it should be from 25 to 53.

Mapped to java.lang.Double.

Example:

DOUBLE

## REAL

```
{ REAL | FLOAT ( precisionInt ) | FLOAT4 }
```

A single precision floating point number. Should not be used to represent currency values, because of rounding problems. Precision value for FLOAT type name should be from 0 to 24.

Mapped to `java.lang.Float`.

Example:

REAL

## TIME

```
TIME [ ( precisionInt ) ] [ WITHOUT TIME ZONE ]
```

The time data type. The format is `hh:mm:ss[.nnnnnnnnnn]`. If fractional seconds precision is specified it should be from 0 to 9, 0 is default.

Mapped to `java.sql.Time`. When converted to a `java.sql.Date`, the date is set to 1970-01-01. `java.time.LocalTime` is also supported on Java 8 and later versions. Use `java.time.LocalTime` or `String` instead of `java.sql.Time` when non-zero precision is needed. Cast from higher fractional seconds precision to lower fractional seconds precision performs round half up; if result of rounding is higher than maximum supported value 23:59:59.999999999 it is saturated to 23:59:59.999999999.

Example:

TIME

TIME(9)

## DATE

```
DATE
```

The date data type. The format is `yyyy-MM-dd`.

Mapped to `java.sql.Date`, with the time set to 00:00:00 (or to the next possible time if midnight doesn't exist for the given date and timezone due to a daylight saving change). `java.time.LocalDate` is also supported on Java 8 and later versions.

Example:

DATE

## TIMESTAMP

```
{ TIMESTAMP [ ( precisionInt ) ] [ WITHOUT TIME ZONE ]  
| DATETIME [ ( precisionInt ) ] | SMALLDATETIME }
```

The timestamp data type. The format is yyyy-MM-dd hh:mm:ss[.nnnnnnnnnn]. Stored internally as a BCD-encoded date, and nanoseconds since midnight. If fractional seconds precision is specified it should be from 0 to 9, 6 is default. Fractional seconds precision of SMALLDATETIME is always 0 and cannot be specified.

Mapped to java.sql.Timestamp (java.util.Date may be used too). java.time.LocalDateTime is also supported on Java 8 and later versions. Cast from higher fractional seconds precision to lower fractional seconds precision performs round half up.

Example:

TIMESTAMP

TIMESTAMP(9)

## TIMESTAMP WITH TIME ZONE

```
TIMESTAMP [ ( precisionInt ) ] WITH TIME ZONE
```

The timestamp with time zone data type. Stored internally as a BCD-encoded date, nanoseconds since midnight, and time zone offset in minutes. If fractional seconds precision is specified it should be from 0 to 9, 6 is default.

Mapped to org.h2.api.TimestampWithTimeZone. java.time.OffsetDateTime and java.time.Instant are also supported on Java 8 and later versions.

Values of this data type are compared by UTC values. It means that 2010-01-01 10:00:00+01 is greater than 2010-01-01 11:00:00+03.

Conversion to TIMESTAMP uses time zone offset to get UTC time and converts it to local time using the system time zone. Conversion from TIMESTAMP does the same operations in reverse and sets time zone offset to offset of the system time zone. Cast from higher fractional seconds precision to lower fractional seconds precision performs round half up.

Example:

TIMESTAMP WITH TIME ZONE  
TIMESTAMP(9) WITH TIME ZONE

## BINARY

```
{ BINARY | VARBINARY | BINARY VARYING  
| LONGVARBINARY | RAW | BYTEA }  
[ ( precisionInt ) ]
```

Represents a byte array. For very long arrays, use BLOB. The maximum size is 2 GB, but the whole object is kept in memory when using this data type. The precision is a size constraint; only the actual data is persisted. For large text data BLOB or CLOB should be used.

Mapped to `byte[]`.

Example:

BINARY(1000)

## OTHER

OTHER

This type allows storing serialized Java objects. Internally, a byte array is used. Serialization and deserialization is done on the client side only. Deserialization is only done when `getObject` is called. Java operations cannot be executed inside the database engine for security reasons. Use `PreparedStatement.setObject` to store values.

Mapped to `java.lang.Object` (or any subclass).

Example:

OTHER

## VARCHAR

```
{ VARCHAR | CHARACTER VARYING | LONGVARCHAR | VARCHAR2 | NVARCHAR  
| NVARCHAR2 | VARCHAR_CASESENSITIVE } [ ( precisionInt ) ]
```

A Unicode String. Use two single quotes (") to create a quote.

The maximum precision is Integer.MAX\_VALUE. The precision is a size constraint; only the actual data is persisted.

The whole text is loaded into memory when using this data type. For large text data CLOB should be used; see there for details.

Mapped to java.lang.String.

Example:

VARCHAR(255)

## VARCHAR\_IGNORECASE

```
VARCHAR_IGNORECASE [ ( precisionInt ) ]
```

Same as VARCHAR, but not case sensitive when comparing. Stored in mixed case.

The maximum precision is Integer.MAX\_VALUE. The precision is a size constraint; only the actual data is persisted.

The whole text is loaded into memory when using this data type. For large text data CLOB should be used; see there for details.

Mapped to java.lang.String.

Example:

VARCHAR\_IGNORECASE

## CHAR

```
{ CHAR | CHARACTER | NCHAR } [ ( precisionInt ) ]
```

A Unicode String. This type is supported for compatibility with other databases and older applications. The difference to VARCHAR is that trailing spaces are ignored and not persisted.

The maximum precision is Integer.MAX\_VALUE. The precision is a size constraint; only the actual data is persisted.

The whole text is kept in memory when using this data type. For large text data CLOB should be used; see there for details.

Mapped to java.lang.String.

Example:

CHAR(10)

## BLOB

```
{ BLOB | BINARY LARGE OBJECT  
| TINYBLOB | MEDIUMBLOB | LONGBLOB | IMAGE | OID }  
[ ( precisionInt [K|M|G|T|P]) ]
```

Like BINARY, but intended for very large values such as files or images. Unlike when using BINARY, large objects are not kept fully in-memory. Use `PreparedStatement.setBinaryStream` to store values. See also CLOB and Advanced / Large Objects.

Mapped to `java.sql.Blob` (`java.io.InputStream` is also supported).

Example:

BLOB  
BLOB(10K)

## CLOB

```
{ CLOB | CHARACTER LARGE OBJECT  
| TINYTEXT | TEXT | MEDIUMTEXT | LONGTEXT | NTEXT | NCLOB }  
[ ( precisionInt [K|M|G|T|P] [CHARACTERS|OCTETS]) ]
```

CLOB is like VARCHAR, but intended for very large values. Unlike when using VARCHAR, large CLOB objects are not kept fully in-memory; instead, they are streamed. CLOB should be used for documents and texts with arbitrary size such as XML or HTML documents, text files, or memo fields of unlimited size. Use `PreparedStatement.setCharacterStream` to store values. See also Advanced / Large Objects.

VARCHAR should be used for text with relatively short average size (for example shorter than 200 characters). Short CLOB values are stored inline, but there is an overhead compared to VARCHAR.

Precision, if any, should be specified in characters, CHARACTERS and OCTETS units have no effect in H2.

Mapped to `java.sql.Clob` (`java.io.Reader` is also supported).

Example:



CLOB  
CLOB(10K)

## UUID

### UUID

Universally unique identifier. This is a 128 bit value. To store values, use `PreparedStatement.setBytes`, `setString`, or `setObject(uuid)` (where `uuid` is a `java.util.UUID`). `ResultSet.getObject` will return a `java.util.UUID`.

Please note that using an index on randomly generated data will result on poor performance once there are millions of rows in a table. The reason is that the cache behavior is very bad with randomly distributed data. This is a problem for any database system.

For details, see the documentation of `java.util.UUID`.

Example:

UUID

## ARRAY

### ARRAY

An array of values. Mapped to `java.lang.Object[]` (arrays of any non-primitive type are also supported).

Use a value list (1, 2) or `PreparedStatement.setObject(.., new Object[] {..})` to store values, and `ResultSet.getObject(..)` or `ResultSet.getArray(..)` to retrieve the values.

Example:

ARRAY

## ENUM

```
{ ENUM (string [, ... ]) }
```

A type with enumerated values. Mapped to `java.lang.Integer`.

The first provided value is mapped to 0, the second mapped to 1, and so on.

Duplicate and empty values are not permitted.

Example:

```
ENUM('clubs', 'diamonds', 'hearts', 'spades')
```

## GEOMETRY

```
GEOMETRY  
[( { GEOMETRY |  
  { POINT  
    | LINESTRING  
    | POLYGON  
    | MULTIPOINT  
    | MULTILINESTRING  
    | MULTIPOLYGON  
    | GEOMETRYCOLLECTION } [Z|M|ZM]}  
[, sridInt ] )]
```

A spatial geometry type. If additional constraints are not specified this type accepts all supported types of geometries. A constraint with required geometry type and dimension system can be set by specifying name of the type and dimension system. A whitespace between them is optional. 2D dimension system does not have a name and assumed if only a geometry type name is specified. POINT means 2D point, POINT Z or POINTZ means 3D point. GEOMETRY constraint means no restrictions on type or dimension system of geometry. A constraint with required spatial reference system identifier (SRID) can be set by specifying this identifier.

Mapped to `org.locationtech.jts.geom.Geometry` if JTS library is in classpath and to `java.lang.String` otherwise. May be represented in textual format using the WKT (well-known text) or EWKT (extended well-known text) format. Values are stored internally in EWKB (extended well-known binary) format. Only a subset of EWKB and EWKT features is supported. Supported objects are POINT, LINESTRING, POLYGON, MULTIPOINT, MULTILINESTRING, MULTIPOLYGON, and GEOMETRYCOLLECTION. Supported dimension systems are 2D (XY), Z (XYZ), M (XYM), and ZM (XYZM). SRID (spatial reference system identifier) is supported.

Use a quoted string containing a WKT/EWKT formatted string or `PreparedStatement.setObject()` to store values, and `ResultSet.getObject(..)` or `ResultSet.getString(..)` to retrieve the values.

Example:

GEOMETRY  
GEOMETRY(POINT)  
GEOMETRY(POINT Z)  
GEOMETRY(POINT Z, 4326)  
GEOMETRY(GEOMETRY, 4326)

## INTERVAL

[intervalYearType](#) | [intervalMonthType](#) | [intervalDayType](#)  
| [intervalHourType](#) | [intervalMinuteType](#) | [intervalSecondType](#)  
| [intervalYearToMonthType](#) | [intervalDayToHourType](#)  
| [intervalDayToMinuteType](#) | [intervalDayToSecondType](#)  
| [intervalHourToMinuteType](#) | [intervalHourToSecondType](#)  
| [intervalMinuteToSecondType](#)

Interval data type. There are two classes of intervals. Year-month intervals can store years and months. Day-time intervals can store days, hours, minutes, and seconds. Year-month intervals are comparable only with another year-month intervals. Day-time intervals are comparable only with another day-time intervals.

Mapped to `org.h2.api.Interval`.

Example:

INTERVAL DAY TO SECOND

## Interval Data Types

### INTERVAL YEAR

INTERVAL YEAR [ ( [precisionInt](#) ) ]

Interval data type. If precision is specified it should be from 1 to 18, 2 is default.

Mapped to `org.h2.api.Interval`. `java.time.Period` is also supported on Java 8 and later versions.

Example:

INTERVAL YEAR

## INTERVAL MONTH

```
INTERVAL MONTH [ ( precisionInt ) ]
```

Interval data type. If precision is specified it should be from 1 to 18, 2 is default.

Mapped to `org.h2.api.Interval`. `java.time.Period` is also supported on Java 8 and later versions.

Example:

INTERVAL MONTH

## INTERVAL DAY

```
INTERVAL DAY [ ( precisionInt ) ]
```

Interval data type. If precision is specified it should be from 1 to 18, 2 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL DAY

## INTERVAL HOUR

```
INTERVAL HOUR [ ( precisionInt ) ]
```

Interval data type. If precision is specified it should be from 1 to 18, 2 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL HOUR

## INTERVAL MINUTE

```
INTERVAL MINUTE [ ( precisionInt ) ]
```

Interval data type. If precision is specified it should be from 1 to 18, 2 is default.

Mapped to org.h2.api.Interval. java.time.Duration is also supported on Java 8 and later versions.

Example:

INTERVAL MINUTE

## INTERVAL SECOND

```
INTERVAL SECOND [ ( precisionInt [, fractionalPrecisionInt ] ) ]
```

Interval data type. If precision is specified it should be from 1 to 18, 2 is default. If fractional seconds precision is specified it should be from 0 to 9, 6 is default.

Mapped to org.h2.api.Interval. java.time.Duration is also supported on Java 8 and later versions.

Example:

INTERVAL SECOND

## INTERVAL YEAR TO MONTH

```
INTERVAL YEAR [ ( precisionInt ) ] TO MONTH
```

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default.

Mapped to org.h2.api.Interval. java.time.Period is also supported on Java 8 and later versions.

Example:

INTERVAL YEAR TO MONTH

## INTERVAL DAY TO HOUR

```
INTERVAL DAY [ ( precisionInt ) ] TO HOUR
```

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default.

Mapped to org.h2.api.Interval. java.time.Duration is also supported on Java 8 and later versions.

Example:

INTERVAL DAY TO HOUR

## INTERVAL DAY TO MINUTE

```
INTERVAL DAY [ ( precisionInt ) ] TO MINUTE
```

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL DAY TO MINUTE

## INTERVAL DAY TO SECOND

```
INTERVAL DAY [ ( precisionInt ) ] TO SECOND [ ( fractionalPrecisionInt ) ]
```

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default. If fractional seconds precision is specified it should be from 0 to 9, 6 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL DAY TO SECOND

## INTERVAL HOUR TO MINUTE

```
INTERVAL HOUR [ ( precisionInt ) ] TO MINUTE
```

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL HOUR TO MINUTE

## INTERVAL HOUR TO SECOND

INTERVAL HOUR [ ( [precisionInt](#) ) ] TO SECOND [ ( [fractionalPrecisionInt](#) ) ]

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default. If fractional seconds precision is specified it should be from 0 to 9, 6 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL HOUR TO SECOND

## INTERVAL MINUTE TO SECOND

INTERVAL MINUTE [ ( [precisionInt](#) ) ] TO SECOND [ ( [fractionalPrecisionInt](#) ) ]

Interval data type. If leading field precision is specified it should be from 1 to 18, 2 is default. If fractional seconds precision is specified it should be from 0 to 9, 6 is default.

Mapped to `org.h2.api.Interval`. `java.time.Duration` is also supported on Java 8 and later versions.

Example:

INTERVAL MINUTE TO SECOND

# System Tables

Information Schema  
Range Table

## Information Schema

The system tables in the schema INFORMATION\_SCHEMA contain the meta data of all tables in the database as well as the current settings.

### CATALOGS

CATALOG\_NAME

### COLLATIONS

NAME, KEY

### COLUMNS

TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, COLUMN\_NAME,  
ORDINAL\_POSITION, DOMAIN\_CATALOG, DOMAIN\_SCHEMA, DOMAIN\_NAME,  
COLUMN\_DEFAULT, IS\_NULLABLE, DATA\_TYPE,  
CHARACTER\_MAXIMUM\_LENGTH, CHARACTER\_OCTET\_LENGTH,  
NUMERIC\_PRECISION, NUMERIC\_PRECISION\_RADIX, NUMERIC\_SCALE,  
DATETIME\_PRECISION, INTERVAL\_TYPE, INTERVAL\_PRECISION,  
CHARACTER\_SET\_NAME, COLLATION\_NAME, TYPE\_NAME, NULLABLE,  
IS\_COMPUTED, SELECTIVITY, CHECK\_CONSTRAINT, SEQUENCE\_NAME,  
REMARKS, SOURCE\_DATA\_TYPE, COLUMN\_TYPE, COLUMN\_ON\_UPDATE,  
IS\_VISIBLE

### COLUMN\_PRIVILEGES

GRANTOR, GRANTEE, TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME,  
COLUMN\_NAME, PRIVILEGE\_TYPE, IS\_GRANTABLE

### CONSTANTS

CONSTANT\_CATALOG, CONSTANT\_SCHEMA, CONSTANT\_NAME, DATA\_TYPE,  
REMARKS, SQL, ID



## CONSTRAINTS

CONSTRAINT\_CATALOG, CONSTRAINT\_SCHEMA, CONSTRAINT\_NAME, CONSTRAINT\_TYPE, TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, UNIQUE\_INDEX\_NAME, CHECK\_EXPRESSION, COLUMN\_LIST, REMARKS, SQL, ID

## CROSS\_REFERENCES

PKTABLE\_CATALOG, PKTABLE\_SCHEMA, PKTABLE\_NAME, PKCOLUMN\_NAME, FKTABLE\_CATALOG, FKTABLE\_SCHEMA, FKTABLE\_NAME, FKCOLUMN\_NAME, ORDINAL\_POSITION, UPDATE\_RULE, DELETE\_RULE, FK\_NAME, PK\_NAME, DEFERRABILITY

## DOMAINS

DOMAIN\_CATALOG, DOMAIN\_SCHEMA, DOMAIN\_NAME, COLUMN\_DEFAULT, IS\_NULLABLE, DATA\_TYPE, PRECISION, SCALE, TYPE\_NAME, SELECTIVITY, CHECK\_CONSTRAINT, REMARKS, SQL, ID

## FUNCTION\_ALIASES

ALIAS\_CATALOG, ALIAS\_SCHEMA, ALIAS\_NAME, JAVA\_CLASS, JAVA\_METHOD, DATA\_TYPE, TYPE\_NAME, COLUMN\_COUNT, RETURNS\_RESULT, REMARKS, ID, SOURCE

## FUNCTION\_COLUMNS

ALIAS\_CATALOG, ALIAS\_SCHEMA, ALIAS\_NAME, JAVA\_CLASS, JAVA\_METHOD, COLUMN\_COUNT, POS, COLUMN\_NAME, DATA\_TYPE, TYPE\_NAME, PRECISION, SCALE, RADIX, NULLABLE, COLUMN\_TYPE, REMARKS, COLUMN\_DEFAULT

## HELP

ID, SECTION, TOPIC, SYNTAX, TEXT

## INDEXES

TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, NON\_UNIQUE, INDEX\_NAME, ORDINAL\_POSITION, COLUMN\_NAME, CARDINALITY, PRIMARY\_KEY, INDEX\_TYPE\_NAME, IS\_GENERATED, INDEX\_TYPE, ASC\_OR\_DESC, PAGES, FILTER\_CONDITION, REMARKS, SQL, ID, SORT\_TYPE, CONSTRAINT\_NAME, INDEX\_CLASS, AFFINITY

## **IN\_DOUBT**

TRANSACTION, STATE

## **KEY\_COLUMN\_USAGE**

CONSTRAINT\_CATALOG, CONSTRAINT\_SCHEMA, CONSTRAINT\_NAME,  
TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, COLUMN\_NAME,  
ORDINAL\_POSITION, POSITION\_IN\_UNIQUE\_CONSTRAINT

## **LOCKS**

TABLE\_SCHEMA, TABLE\_NAME, SESSION\_ID, LOCK\_TYPE

## **QUERY\_STATISTICS**

SQL\_STATEMENT, EXECUTION\_COUNT, MIN\_EXECUTION\_TIME,  
MAX\_EXECUTION\_TIME, CUMULATIVE\_EXECUTION\_TIME,  
AVERAGE\_EXECUTION\_TIME, STD\_DEV\_EXECUTION\_TIME, MIN\_ROW\_COUNT,  
MAX\_ROW\_COUNT, CUMULATIVE\_ROW\_COUNT, AVERAGE\_ROW\_COUNT,  
STD\_DEV\_ROW\_COUNT

## **REFERENTIAL\_CONSTRAINTS**

CONSTRAINT\_CATALOG, CONSTRAINT\_SCHEMA, CONSTRAINT\_NAME,  
UNIQUE\_CONSTRAINT\_CATALOG, UNIQUE\_CONSTRAINT\_SCHEMA,  
UNIQUE\_CONSTRAINT\_NAME, MATCH\_OPTION, UPDATE\_RULE, DELETE\_RULE

## **RIGHTS**

GRANTEE, GRANTEETYPE, GRANTEDROLE, RIGHTS, TABLE\_SCHEMA,  
TABLE\_NAME, ID

## **ROLES**

NAME, REMARKS, ID

## **SCHEMATA**

CATALOG\_NAME, SCHEMA\_NAME, SCHEMA\_OWNER,  
DEFAULT\_CHARACTER\_SET\_NAME, DEFAULT\_COLLATION\_NAME, IS\_DEFAULT,  
REMARKS, ID

## **SEQUENCES**

SEQUENCE\_CATALOG, SEQUENCE\_SCHEMA, SEQUENCE\_NAME,  
CURRENT\_VALUE, INCREMENT, IS\_GENERATED, REMARKS, CACHE, MIN\_VALUE,  
MAX\_VALUE, IS\_CYCLE, ID

## **SESSIONS**

ID, USER\_NAME, SESSION\_START, STATEMENT, STATEMENT\_START,  
CONTAINS\_UNCOMMITTED, STATE, BLOCKER\_ID

## **SESSION\_STATE**

KEY, SQL

## **SETTINGS**

NAME, VALUE

## **SYNONYMS**

SYNONYM\_CATALOG, SYNONYM\_SCHEMA, SYNONYM\_NAME, SYNONYM\_FOR,  
SYNONYM\_FOR\_SCHEMA, TYPE\_NAME, STATUS, REMARKS, ID

## **TABLES**

TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, TABLE\_TYPE, STORAGE\_TYPE,  
SQL, REMARKS, LAST\_MODIFICATION, ID, TYPE\_NAME, TABLE\_CLASS,  
ROW\_COUNT\_ESTIMATE

## **TABLE\_CONSTRAINTS**

CONSTRAINT\_CATALOG, CONSTRAINT\_SCHEMA, CONSTRAINT\_NAME,  
CONSTRAINT\_TYPE, TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME,  
IS\_DEFERRABLE, INITIALLY\_DEFERRED

## **TABLE\_PRIVILEGES**

GRANTOR, GRANTEE, TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME,  
PRIVILEGE\_TYPE, IS\_GRANTABLE

## TABLE\_TYPES

TYPE

## TRIGGERS

TRIGGER\_CATALOG, TRIGGER\_SCHEMA, TRIGGER\_NAME, TRIGGER\_TYPE, TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, BEFORE, JAVA\_CLASS, QUEUE\_SIZE, NO\_WAIT, REMARKS, SQL, ID

## TYPE\_INFO

TYPE\_NAME, DATA\_TYPE, PRECISION, PREFIX, SUFFIX, PARAMS, AUTO\_INCREMENT, MINIMUM\_SCALE, MAXIMUM\_SCALE, RADIX, POS, CASE\_SENSITIVE, NULLABLE, SEARCHABLE

## USERS

NAME, ADMIN, REMARKS, ID

## VIEWS

TABLE\_CATALOG, TABLE\_SCHEMA, TABLE\_NAME, VIEW\_DEFINITION, CHECK\_OPTION, IS\_UPDATABLE, STATUS, REMARKS, ID

## Range Table

The range table is a dynamic system table that contains all values from a start to an end value. Non-zero step value may be also specified, default is 1. Start value, end value, and optional step value are converted to BIGINT data type. The table contains one column called X. If start value is greater than end value and step is positive the result is empty. If start value is less than end value and step is negative the result is empty too. If start value is equal to end value the result contains only start value. Start value, start value plus step, start value plus step multiplied by two and so on are included in result. If step is positive the last value is less than or equal to the specified end value. If step is negative the last value is greater than or equal to the specified end value. The table is used as follows:

Examples:

```
SELECT X FROM SYSTEM_RANGE(1, 10);  
-- 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
```

```
SELECT X FROM SYSTEM_RANGE(1, 10, 2);  
-- 1, 3, 5, 7, 9  
SELECT X FROM SYSTEM_RANGE(1, 10, -1);  
-- No rows  
SELECT X FROM SYSTEM_RANGE(10, 2, -2);  
-- 10, 8, 6, 4, 2
```

# Build

[Portability](#)  
[Environment](#)  
[Building the Software](#)  
[Build Targets](#)  
[Using Maven 2](#)  
[Using Eclipse](#)  
[Translating](#)  
[Submitting Source Code Changes](#)  
[Reporting Problems or Requests](#)  
[Automated Build](#)  
[Generating Railroad Diagrams](#)

## Portability

This database is written in Java and therefore works on many platforms.

## Environment

To run this database, a Java Runtime Environment (JRE) version 7 or higher is required.

To create the database executables, the following software stack was used. To use this database, it is not required to install this software however.

- Mac OS X and Windows
- [Oracle JDK Version 7](#) (version 7 is not available for free download any more)
- [Eclipse](#)
- Eclipse Plugins: [Subclipse](#), [Eclipse Checkstyle Plug-in](#), [EclEmma Java Code Coverage](#)
- [Mozilla Firefox](#)
- [OpenOffice](#)
- [NSIS](#) (Nullsoft Scriptable Install System)
- [Maven](#)

## Building the Software

You need to install a JDK, for example the Oracle JDK version 7 or 8. Ensure that Java binary directory is included in the PATH environment variable, and that the environment variable JAVA\_HOME points to your Java installation. On the command line, go to the directory h2 and execute the following command:

```
build -?
```

For Linux and OS X, use `./build.sh` instead of `build`.

You will get a list of targets. If you want to build the jar file, execute (Windows):

```
build jar
```

To run the build tool in shell mode, use the command line option `-`:

```
./build.sh -
```

## Build Targets

The build system can generate smaller jar files as well. The following targets are currently supported:

- `jarClient` creates the file `h2client.jar`. This only contains the JDBC client.
- `jarSmall` creates the file `h2small.jar`. This only contains the embedded database. Debug information is disabled.
- `javadocImpl` creates the Javadocs of the implementation.

To create the file `h2client.jar`, go to the directory `h2` and execute the following command:

```
build jarClient
```

## Using Apache Lucene

Apache Lucene 5.5.5 is used for testing. Newer versions up to 7.6.0 can also be used.

## Using Maven 2

### Using a Central Repository

You can include the database in your Maven 2 project as a dependency. Example:

```
<dependency>  
  <groupId>com.h2database</groupId>  
  <artifactId>h2</artifactId>  
  <version>1.4.198</version>
```

```
</dependency>
```

New versions of this database are first uploaded to <http://hsqldb.sourceforge.net/m2-repo/> and then automatically synchronized with the main [Maven repository](#); however after a new release it may take a few hours before they are available there.

## Maven Plugin to Start and Stop the TCP Server

A Maven plugin to start and stop the H2 TCP server is available from [Laird Nelson at GitHub](#). To start the H2 server, use:

```
mvn com.edugility.h2-maven-plugin:1.0-SNAPSHOT:spawn
```

To stop the H2 server, use:

```
mvn com.edugility.h2-maven-plugin:1.0-SNAPSHOT:stop
```

## Using Snapshot Version

To build a h2-\*-SNAPSHOT.jar file and upload it to the local Maven 2 repository, execute the following command:

```
build mavenInstallLocal
```

Afterwards, you can include the database in your Maven 2 project as a dependency:

```
<dependency>
  <groupId>com.h2database</groupId>
  <artifactId>h2</artifactId>
  <version>1.0-SNAPSHOT</version>
</dependency>
```

## Using Eclipse

To create an Eclipse project for H2, use the following steps:

- Install Git and [Eclipse](#).
- Get the H2 source code from Github:  
git clone <https://github.com/h2database/h2database>



- Download all dependencies:  
build.bat download(Windows)  
./build.sh download(otherwise)
- In Eclipse, create a new Java project from existing source code: File, New, Project, Java Project, Create project from existing source.
- Select the h2 folder, click Next and Finish.
- To resolve com.sun.javadoc import statements, you may need to manually add the file <java.home>/../lib/tools.jar to the build path.

## Translating

The translation of this software is split into the following parts:

- H2 Console: src/main/org/h2/server/web/res/\_text\_\*.prop
- Error messages: src/main/org/h2/res/\_messages\_\*.prop

To translate the H2 Console, start it and select Preferences / Translate. After you are done, send the translated \*.prop file to the Google Group. The web site is currently translated using Google.

## Submitting Source Code Changes

If you'd like to contribute bug fixes or new features, please consider the following guidelines to simplify merging them:

- Only use Java 7 features (do not use Java 8/9/etc) (see [Environment](#)).
- Follow the coding style used in the project, and use Checkstyle (see above) to verify. For example, do not use tabs (use spaces instead). The checkstyle configuration is in src/installer/checkstyle.xml.
- A template of the Eclipse settings are in src/installer/eclipse.settings/\*. If you want to use them, you need to copy them to the .settings directory. The formatting options (eclipseCodeStyle) are also included.
- Please provide test cases and integrate them into the test suite. For Java level tests, see src/test/org/h2/test/TestAll.java. For SQL level tests, see SQL files in src/test/org/h2/test/scripts.
- The test cases should cover at least 90% of the changed and new code; use a code coverage tool to verify that (see above). or use the build target coverage.
- Verify that you did not break other features: run the test cases by executing build test.
- Provide end user documentation if required (src/docsrc/html/\*).
- Document grammar changes in src/docsrc/help/help.csv
- Provide a change log entry (src/docsrc/html/changelog.html).

- Verify the spelling using build spellcheck. If required add the new words to `src/tools/org/h2/build/doc/dictionary.txt`.
- Run `src/installer/buildRelease` to find and fix formatting errors.
- Verify the formatting using build docs and build javadoc.
- Submit changes using GitHub's "pull requests". You'll require a free [GitHub](#) account. If you are not familiar with pull requests, please read GitHub's [Using pull requests](#) page.

For legal reasons, patches need to be public in the form of an [issue report or attachment](#) or in the form of an email to the [group](#). Significant contributions need to include the following statement:

"I wrote the code, it's mine, and I'm contributing it to H2 for distribution multiple-licensed under the MPL 2.0, and the EPL 1.0 (<http://h2database.com/html/license.html>)."

## Reporting Problems or Requests

Please consider the following checklist if you have a question, want to report a problem, or if you have a feature request:

- For bug reports, please provide a [short, self contained, correct \(compilable\), example](#) of the problem.
- Feature requests are always welcome, even if the feature is already on the [roadmap](#). Your mail will help prioritize feature requests. If you urgently need a feature, consider [providing a patch](#).
- Before posting problems, check the [FAQ](#) and do a [Google search](#).
- When got an unexpected exception, please try the [Error Analyzer tool](#). If this doesn't help, please report the problem, including the complete error message and stack trace, and the root cause stack trace(s).
- When sending source code, please use a public web clipboard such as [Pastebin](#) or [Mystic Paste](#) to avoid formatting problems. Please keep test cases as simple and short as possible, but so that the problem can still be reproduced. As a template, use: [HelloWorld.java](#). Method that simply call other methods should be avoided, as well as unnecessary exception handling. Please use the JDBC API and no external tools or libraries. The test should include all required initialization code, and should be started with the main method.
- For large attachments, use a public storage such as [Google Drive](#).
- Google Group versus issue tracking: Use the [Google Group](#) for questions or if you are not sure it's a bug. If you are sure it's a bug, you can create an [issue](#), but you don't need to (sending an email to the group is enough). Please note that only few people monitor the issue tracking system.
- For out-of-memory problems, please analyze the problem yourself first, for example using the command line option `-XX:`

+HeapDumpOnOutOfMemoryError (to create a heap dump file on out of memory) and a memory analysis tool such as the [Eclipse Memory Analyzer \(MAT\)](#).

- It may take a few days to get an answers. Please do not double post.

## Automated Build

This build process is automated and runs regularly. The build process includes running the tests and code coverage, using the command line `./build.sh clean jar coverage -Dh2.ftpPassword=... uploadBuild`. The last results are available here:

- [Test Output](#)
- [Code Coverage Summary](#)
- [Code Coverage Details \(download, 1.3 MB\)](#)
- [Build Newsfeed](#)

## Generating Railroad Diagrams

The railroad diagrams of the [SQL grammar](#) are HTML, formatted as nested tables. The diagrams are generated as follows:

- The BNF parser (`org.h2.bnf.Bnf`) reads and parses the BNF from the file `help.csv`.
- The page parser (`org.h2.server.web.PageParser`) reads the template HTML file and fills in the diagrams.
- The rail images (one straight, four junctions, two turns) are generated using a simple Java application.

To generate railroad diagrams for other grammars, see the package `org.h2.jcr`. This package is used to generate the SQL-2 railroad diagrams for the JCR 2.0 specification.

# History and Roadmap

[Change Log](#)

[Roadmap](#)

[History of this Database Engine](#)

[Why Java](#)

[Supporters](#)

## Change Log

The up-to-date change log is available at

<http://www.h2database.com/html/changelog.html>

## Roadmap

The current roadmap is available at

<http://www.h2database.com/html/roadmap.html>

## History of this Database Engine

The development of H2 was started in May 2004, but it was first published on December 14th 2005. The original author of H2, Thomas Mueller, is also the original developer of Hypersonic SQL. In 2001, he joined PointBase Inc. where he wrote PointBase Micro, a commercial Java SQL database. At that point, he had to discontinue Hypersonic SQL. The HSQLDB Group was formed to continued to work on the Hypersonic SQL codebase. The name H2 stands for Hypersonic 2, however H2 does not share code with Hypersonic SQL or HSQLDB. H2 is built from scratch.

## Why Java

The main reasons to use a Java database are:

- Very simple to integrate in Java applications
- Support for many different platforms
- More secure than native applications (no buffer overflows)
- User defined functions (or triggers) run very fast
- Unicode support

Some think Java is too slow for low level operations, but this is no longer true. Garbage collection for example is now faster than manual memory management.

Developing Java code is faster than developing C or C++ code. When using Java, most time can be spent on improving the algorithms instead of porting the code to different platforms or doing memory management. Features such as Unicode and network libraries are already built-in. In Java, writing secure code is easier because buffer overflows can not occur. Features such as reflection can be used for randomized testing.

Java is future proof: a lot of companies support Java. Java is now open source.

To increase the portability and ease of use, this software depends on very few libraries. Features that are not available in open source Java implementations (such as Swing) are not used, or only used for optional features.

## Supporters

Many thanks for those who reported bugs, gave valuable feedback, spread the word, and translated this project.

Also many thanks to the donors. To become a donor, use PayPal (at the very bottom of the main web page). Donators are:

- Martin Wildam, Austria
- [tagtraum industries incorporated, USA](#)
- [TimeWriter, Netherlands](#)
- [Cognitect, USA](#)
- [Code 42 Software, Inc., Minneapolis](#)
- [Code Lutin, France](#)
- [NetSuxxess GmbH, Germany](#)
- [Poker Copilot, Steve McLeod, Germany](#)
- [SkyCash, Poland](#)
- [Lumber-mill, Inc., Japan](#)
- [StockMarketEye, USA](#)
- [Eckenfelder GmbH & Co.KG, Germany](#)
- Jun Iyama, Japan
- Steven Branda, USA
- Anthony Goubard, Netherlands
- Richard Hickey, USA
- Alessio Jacopo D'Adamo, Italy
- Ashwin Jayaprakash, USA
- Donald Bleyl, USA
- Frank Berger, Germany
- Florent Ramiere, France
- Antonio Casqueiro, Portugal
- Oliver Computing LLC, USA

- Harpal Grover Consulting Inc., USA
- Elisabetta Berlini, Italy
- William Gilbert, USA
- Antonio Dieguez Rojas, Chile
- [Ontology Works, USA](#)
- Pete Haidinyak, USA
- William Osmond, USA
- Joachim Ansorg, Germany
- Oliver Soerensen, Germany
- Christos Vasilakis, Greece
- Fyodor Kupolov, Denmark
- Jakob Jenkov, Denmark
- Stéphane Chartrand, Switzerland
- Glenn Kidd, USA
- Gustav Trede, Sweden
- Joonas Pulakka, Finland
- Bjorn Darri Sigurdsson, Iceland
- Gray Watson, USA
- Erik Dick, Germany
- Pengxiang Shao, China
- Bilingual Marketing Group, USA
- Philippe Marschall, Switzerland
- Knut Staring, Norway
- Theis Borg, Denmark
- Mark De Mendonca Duske, USA
- Joel A. Garringer, USA
- Olivier Chafik, France
- Rene Schwietzke, Germany
- Jalpesh Patadia, USA
- Takanori Kawashima, Japan
- Terrence JC Huang, China
- [JiaDong Huang, Australia](#)
- Laurent van Roy, Belgium
- Qian Chen, China
- Clinton Hyde, USA
- Kritchai Phromros, Thailand
- Alan Thompson, USA
- Ladislav Jech, Czech Republic
- Dimitrijs Fedotovs, Latvia
- Richard Manley-Reeve, United Kingdom
- Daniel Cyr, ThirdHalf.com, LLC, USA
- Peter Jünger, Germany
- Dan Keegan, USA
- Rafel Israels, Germany

- Fabien Todescato, France
- Cristan Meijer, Netherlands
- Adam McMahon, USA
- Fábio Gomes Lisboa Gomes, Brasil
- Lyderic Landry, England
- Mederp, Morocco
- Joaquim Golay, Switzerland
- Clemens Quoss, Germany
- Kervin Pierre, USA
- Jake Bellotti, Australia
- Arun Chittanoor, USA

# Frequently Asked Questions

I Have a Problem or Feature Request  
Are there Known Bugs? When is the Next Release?  
Is this Database Engine Open Source?  
Is Commercial Support Available?  
How to Create a New Database?  
How to Connect to a Database?  
Where are the Database Files Stored?  
What is the Size Limit (Maximum Size) of a Database?  
Is it Reliable?  
Why is Opening my Database Slow?  
My Query is Slow  
H2 is Very Slow  
Column Names are Incorrect?  
Float is Double?  
How to Translate this Project?  
How to Contribute to this Project?

## I Have a Problem or Feature Request

Please read the [support checklist](#).

## Are there Known Bugs? When is the Next Release?

Usually, bugs get fixes as they are found. There is a release every few weeks. Here is the list of known and confirmed issues:

- When opening a database file in a timezone that has different daylight saving rules: the time part of dates where the daylight saving doesn't match will differ. This is not a problem within regions that use the same rules (such as, within USA, or within Europe), even if the timezone itself is different. As a workaround, export the database to a SQL script using the old timezone, and create a new database in the new timezone.
- Tomcat and Glassfish 3 set most static fields (final or non-final) to null when unloading a web application. This can cause a NullPointerException in H2 versions 1.1.107 and older, and may still not work in newer versions. Please report it if you run into this issue. In Tomcat  $\geq 6.0$  this behavior can be disabled by setting the system property `org.apache.catalina.loader.WebappClassLoader.ENABLE_CLEAR_REFERENCE` `S=false`, however Tomcat may then run out of memory. A known workaround is to put the `h2*.jar` file in a shared lib directory (`common/lib`).



- Some problems have been found with right outer join. Internally, it is converted to left outer join, which does not always produce the same results as other databases when used in combination with other joins. This problem is fixed in H2 version 1.3.

For a complete list, see [Open Issues](#).

## Is this Database Engine Open Source?

Yes. It is free to use and distribute, and the source code is included. See also under license.

## Is Commercial Support Available?

No, currently commercial support is not available.

## How to Create a New Database?

By default, a new database is automatically created if it does not yet exist when [embedded](#) URL is used. See [Creating New Databases](#).

## How to Connect to a Database?

The database driver is `org.h2.Driver`, and the database URL starts with `jdbc:h2:`. To connect to a database using JDBC, use the following code:

```
Connection conn = DriverManager.getConnection("jdbc:h2:~/test", "sa", "");
```

## Where are the Database Files Stored?

When using database URLs like `jdbc:h2:~/test`, the database is stored in the user directory. For Windows, this is usually `C:\Documents and Settings\<userName>` or `C:\Users\<userName>`. If the base directory is not set (as in `jdbc:h2:./test`), the database files are stored in the directory where the application is started (the current working directory). When using the H2 Console application from the start menu, this is `<Installation Directory>/bin`. The base directory can be set in the database URL. A fixed or relative path can be used. When using the URL `jdbc:h2:file:./data/sample`, the database is stored in the directory `data` (relative to the current working directory). The directory is created automatically if it does not yet exist. It is also possible to use the fully qualified directory name (and for Windows, drive name). Example: `jdbc:h2:file:C:/data/test`

## What is the Size Limit (Maximum Size) of a Database?

See [Limits and Limitations](#).

## Is it Reliable?

That is not easy to say. It is still a quite new product. A lot of tests have been written, and the code coverage of these tests is higher than 80% for each package. Randomized stress tests are run regularly. But there are probably still bugs that have not yet been found (as with most software). Some features are known to be dangerous, they are only supported for situations where performance is more important than reliability. Those dangerous features are:

- Disabling the transaction log or `FileDescriptor.sync()` using `LOG=0` or `LOG=1`.
- Using the transaction isolation level `READ_UNCOMMITTED` (`LOCK_MODE 0`) while at the same time using multiple connections.
- Disabling database file protection using (setting `FILE_LOCK` to `NO` in the database URL).
- Disabling referential integrity using `SET REFERENTIAL_INTEGRITY FALSE`.

In addition to that, running out of memory should be avoided. In older versions, `OutOfMemory` errors while using the database could corrupt a databases.

This database is well tested using automated test cases. The tests run every night and run for more than one hour. But not all areas of this database are equally well tested. When using one of the following features for production, please ensure your use case is well tested (if possible with automated test cases). The areas that are not well tested are:

- Platforms other than Windows, Linux, Mac OS X, or runtime environments other than Oracle / OpenJDK 7, 8, 9.
- The features `AUTO_SERVER` and `AUTO_RECONNECT`.
- Cluster mode, 2-phase commit, savepoints.
- Fulltext search.
- Operations on LOBs over 2 GB.
- The optimizer may not always select the best plan.
- Using the ICU4J collator.

Areas considered experimental are:

- The PostgreSQL server
- Clustering (there are cases where transaction isolation can be broken due to timing issues, for example one session overtaking another session).

- Multi-threading within the old PageStore engine using SET MULTI\_THREADED=1. Default MVStore engine is multi-threaded by default.
- Compatibility modes for other databases (only some features are implemented).
- The soft reference cache (CACHE\_TYPE=SOFT\_LRU). It might not improve performance, and out of memory issues have been reported.

Some users have reported that after a power failure, the database cannot be opened sometimes. In this case, use a backup of the database or the Recover tool. Please report such problems. The plan is that the database automatically recovers in all situations.

## Why is Opening my Database Slow?

To find out what the problem is, use the H2 Console and click on "Test Connection" instead of "Login". After the "Login Successful" appears, click on it (it's a link). This will list the top stack traces. Then either analyze this yourself, or post those stack traces in the Google Group.

Other possible reasons are: the database is very big (many GB), or contains linked tables that are slow to open.

## My Query is Slow

Slow SELECT (or DELETE, UPDATE, MERGE) statement can have multiple reasons. Follow this checklist:

- Run ANALYZE (see documentation for details).
- Run the query with EXPLAIN and check if indexes are used (see documentation for details).
- If required, create additional indexes and try again using ANALYZE and EXPLAIN.
- If it doesn't help please report the problem.

## H2 is Very Slow

By default, H2 closes the database when the last connection is closed. If your application closes the only connection after each operation, the database is opened and closed a lot, which is quite slow. There are multiple ways to solve this problem, see [Database Performance Tuning](#).

## Column Names are Incorrect?

For the query `SELECT ID AS X FROM TEST` the method `ResultSetMetaData.getColumnNames()` returns `ID`, I expect it to return `X`. What's wrong?

This is not a bug. According the JDBC specification, the method `ResultSetMetaData.getColumnNames()` should return the name of the column and not the alias name. If you need the alias name, use [ResultSetMetaData.getColumnLabel\(\)](#). Some other database don't work like this yet (they don't follow the JDBC specification). If you need compatibility with those databases, use the [Compatibility Mode](#), or append `;ALIAS_COLUMN_NAME=TRUE` to the database URL.

This also applies to `DatabaseMetaData` calls that return a result set. The columns in the JDBC API are column labels, not column names.

## Float is Double?

For a table defined as `CREATE TABLE TEST(X FLOAT)` the method `ResultSet.getObject()` returns a `java.lang.Double`, I expect it to return a `java.lang.Float`. What's wrong?

This is not a bug. According the JDBC specification, the JDBC data type `FLOAT` is equivalent to `DOUBLE`, and both are mapped to `java.lang.Double`. See also [Mapping SQL and Java Types - 8.3.10 FLOAT](#).

Use `REAL` or `FLOAT(24)` data type for `java.lang.Float` values.

## How to Translate this Project?

For more information, see [Build/Translating](#).

## How to Contribute to this Project?

There are various way to help develop an open source project like H2. The first step could be to [translate](#) the error messages and the GUI to your native language. Then, you could [provide patches](#). Please start with small patches. That could be adding a test case to improve the [code coverage](#) (the target code coverage for this project is 90%, higher is better). You will have to [develop, build and run the tests](#). Once you are familiar with the code, you could implement missing features from the [feature request list](#). I suggest to start with very small features that are easy to implement. Keep in mind to provide test cases as well.